

# Introduction - 1

Last Update: Feb 3, 2023

# Outline and Objectives


## Outline

- What Operating Systems Do
- Computer-System Organization
- OS structure and operation
- Major OS Functions
  - Process Management
  - Memory Management
  - Storage Management
  - Protection and Security
- Computing Environments

## Objectives

- To provide a grand tour of the major operating systems **components**
- To provide coverage of basic **computer system**

# Basic components of a computer system: place of OS




- A computer system can be divided into four components
  - **Hardware** – provides basic computing and storage resources
    - CPU, memory, I/O devices
  - **Operating system**
    - Controls and coordinates use of hardware among various applications and users
  - **Application programs** – solve the problems of the users: use system resources
    - Word processors, compilers, web browsers, database systems, video games
  - **Users**
    - People, machines, other computers

# What is an operating system?

---

- A program that manages hardware and acts as an *intermediary* between a users/applications and the hardware
- Operating system functionalities
  - Start, terminate, **control executing user programs**
  - Make system **convenient** to use
  - **Control and coordinate use of hardware**
    - Perform and manage **I/O**; setup devices
    - Manage and allocate **resources**
    - Use hardware **efficiently**
  - Implement **common services**




**“operating system”**: software responsible for proper “operation” of the computer.

# Operating System Definition (as a software)

---

- No universally accepted definition
  - “*Everything a vendor ships when you order an operating system*” is good approximation
    - But varies wildly.
- **Kernel**: running all the time; having most of the functionality
- Everything else: either a **system program** (associated with the operation of the system) or an **application program**

**System programs**: programs that are **associated** with the operating system.

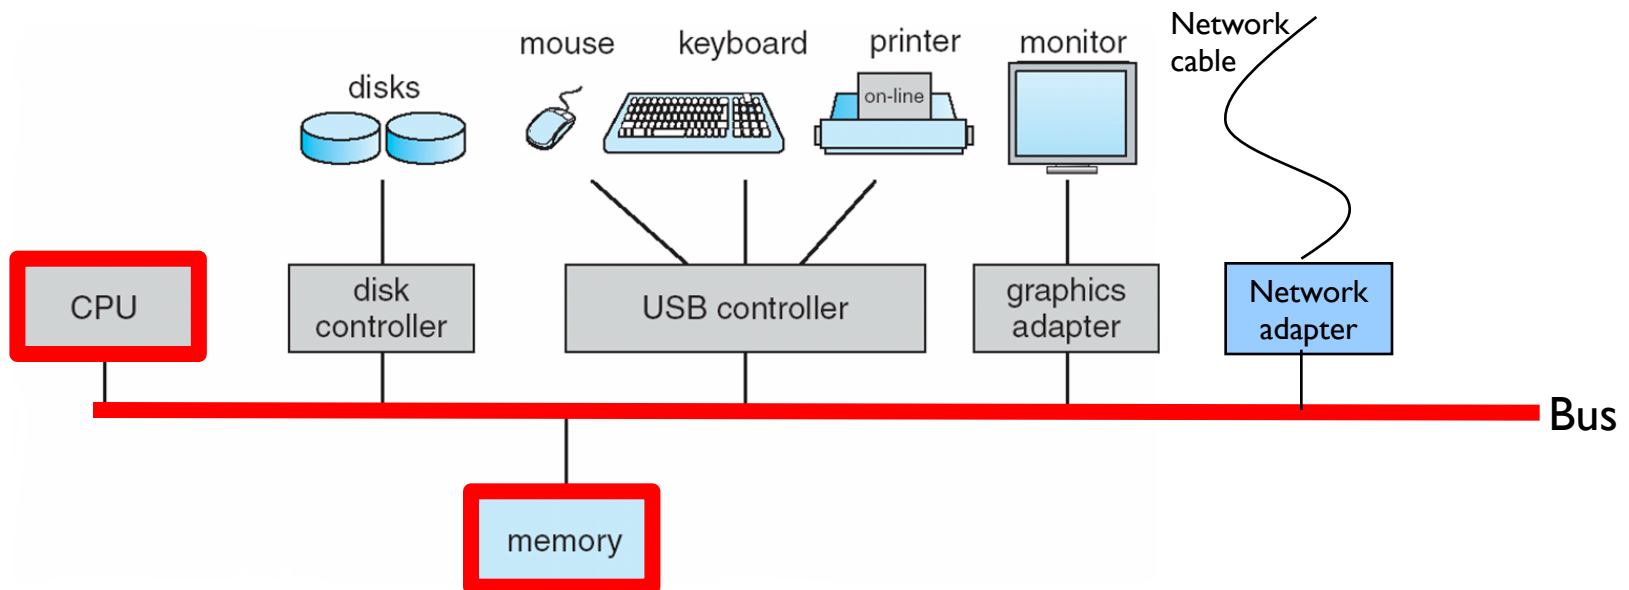


# Computer System Organization and Operation

# Computer System Organization

---

- Computer-system operation
  - One or more **CPUs**, **device controllers** connect through **common bus** providing access to **shared memory**
  - **Concurrent execution** of CPUs and Device Controllers competing for memory cycles



# Computer Startup

---


- **bootstrap program** is loaded at power-up or reboot
  - Typically stored in ROM or EPROM,
  - Generally known as **firmware**
  - Initializes all aspects of the system
  - **Loads** operating system **kernel** into memory and starts its execution.
- **Kernel** runs and make the system ready for running applications.
  - Kernel is always ready to run (always in memory)



# Computer system operation: I/O and device interaction

- I/O devices (and controllers) and the CPU can execute **concurrently**
- Each device controller has a **local buffer**
  - Data movement (I/O) between **device** and **local buffer** (by device)
  - Data movement between **memory** and **local buffer** (by CPU)

Device **controller** informs CPU that it has finished its current output operation or it has some input data by causing an **interrupt**.



# Hardware interrupts

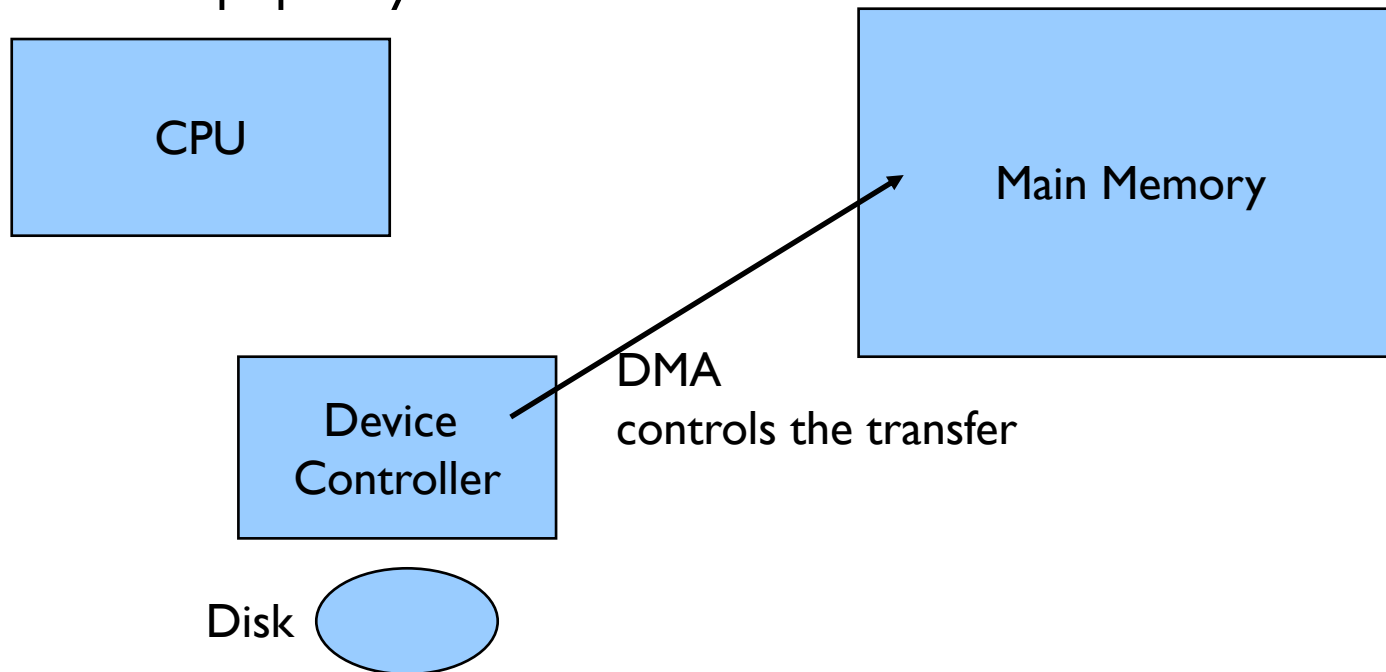
---

- When interrupt occurs, hardware does the following:
  - CPU is interrupted
    - at that time application code or kernel code might be running
    - registers and the program counter **saved** into RAM to preserve CPU state
  - CPU starts running the respective **Interrupt Service Routine (ISR)**
    - ISR is a kernel routine
    - ISR is found through interrupt vector: a table containing addresses of ISRs

# Direct Memory Access Structure

---

- With DMA, device controller **transfers blocks of data** from device buffer **directly** to main memory **without CPU intervention**
  - Only **one interrupt is generated per block**, rather than one interrupt per byte

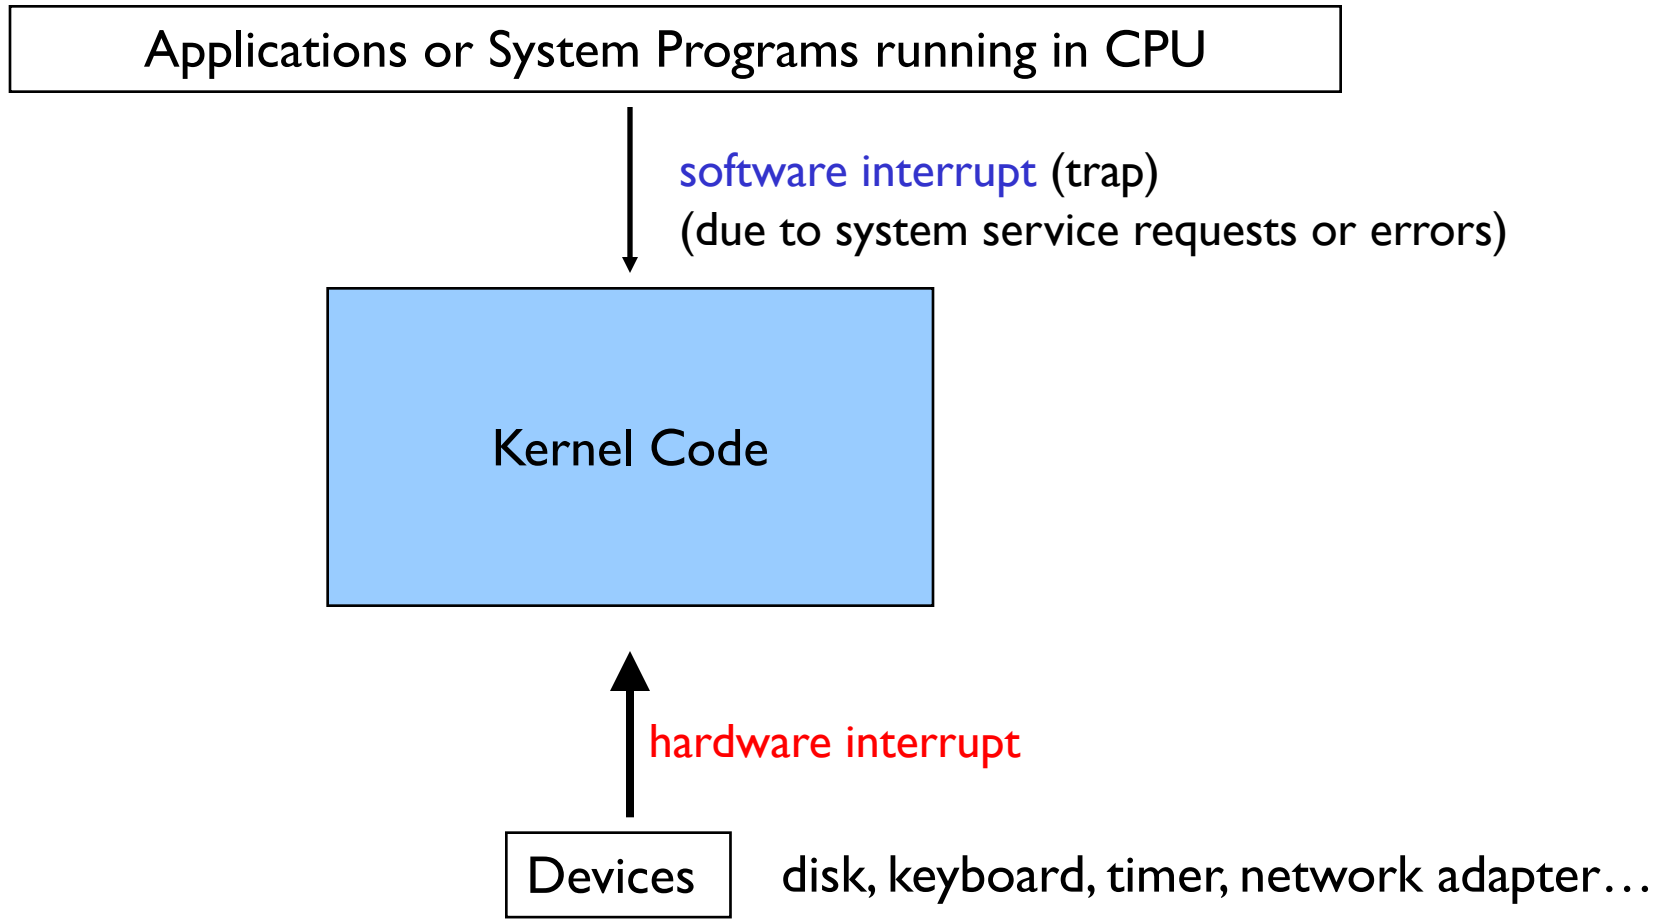


# Software interrupts

---

- Running application software (executing program) may generate interrupts as well.
  - They are called **software interrupts** (trap)
    - 1. Exceptions (caused by errors, such as division by zero)
    - 2. System calls (service request)
      - syscall (or trap) instruction is used
- *An operating system (kernel) is **interrupt-driven** (event driven)*


# Interrupt-Driven OS



# Storage Structure

---

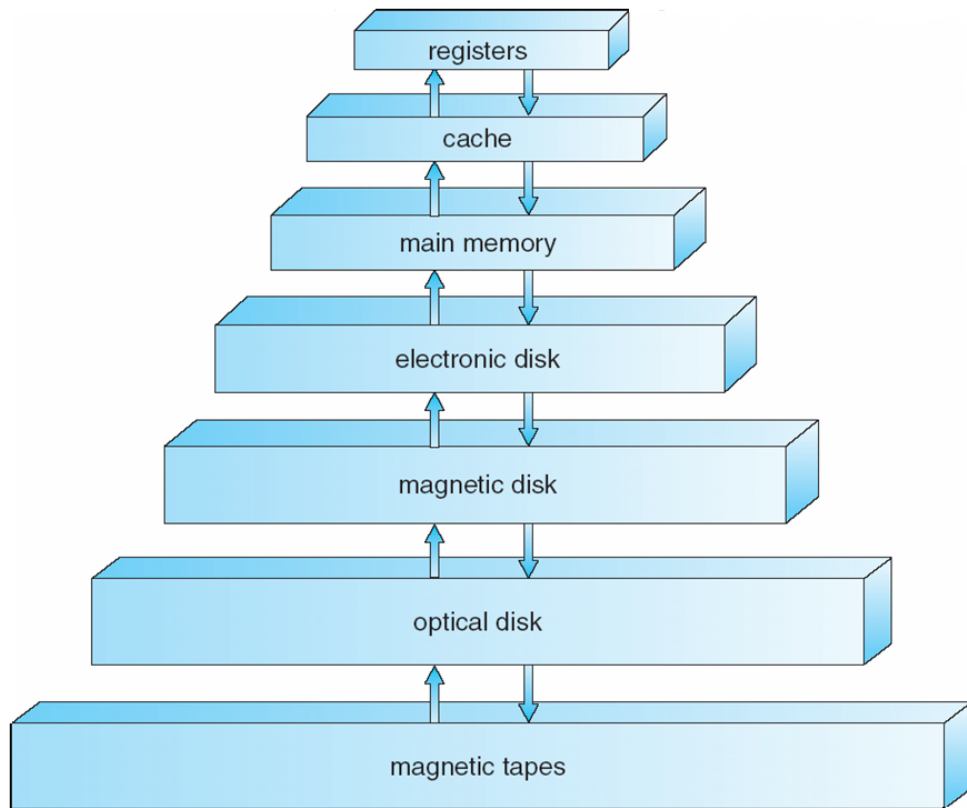
- **Main memory** (primary storage)
  - CPU can access directly
- **Secondary storage** – extension of main memory that provides large **nonvolatile** storage capacity
  - Magnetic disks (HDD disks)
    - Platters
    - The **disk controller** handles the interaction between the device and the computer
  - SSD disks



# Storage Hierarchy

---

- Storage systems organized in hierarchy
  - They differ in
    - Speed (ms)
    - Cost (\$)
    - Capacity (GB)
    - Volatility



# Storage units

---


- Bit: 0 or 1 (one of two value)
- Byte: 8 bits
- Word (architecture's native unit of data – e.g., 1 word = 4 bytes)
- 1 KB = 1024 bytes =  $2^{10}$  bytes
- 1 MB =  $2^{20}$  bytes =  $1024^2$  bytes
- 1 GB =  $2^{30}$  bytes =  $1024^3$  bytes
- 1 TB =  $2^{40}$  bytes =  $1024^4$  bytes
- 1 PB =  $2^{50}$  bytes =  $1024^5$  bytes



# Caching

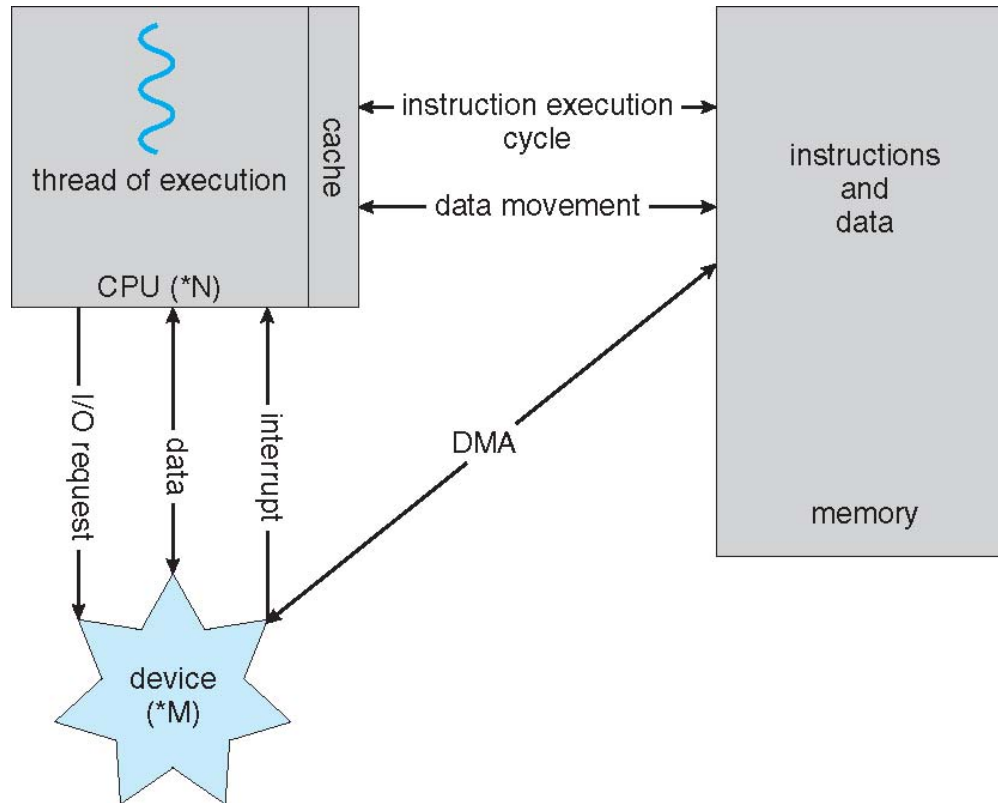
---

- **Caching** – copying information **into faster storage**
  - there is **tradeoff** between **size** and **speed** of storage devices.
- Performed at many levels in a computer  
(at hardware, operating system, or application level)
- Cache is checked first for an item.



# Interplay of all Hardware Components

---




# Computer System Architecture

# Single processor systems

---

- A lot of systems use a **single general-purpose processor** (CPU) or a limited number of CPUs
  - Most systems have special-purpose processors as well
- CPU is capable of executing a general purpose instruction set, including instructions from user programs.
- Computers have **device-specific processors** as well.
  - They don't run user programs.
  - Some may be commanded/managed by the operation system.
    - By the device drivers.




# Multiprocessor Systems

---

- Multiprocessor systems growing in use and importance
  - They are parallel systems
  - Tightly-coupled systems
  - Advantages include
    - Increased throughput
    - Economy of scale (cheaper than using multiple computers)
    - Increased reliability – graceful degradation or fault tolerance
  - Two types of multiprocessor architecture
    1. Asymmetric Multiprocessing
    2. Symmetric Multiprocessing (SMP) (very common)


# Symmetric Multiprocessing Architecture (SMP)

Each CPU has equal role  
(can execute a user program or kernel)




# A Dual Core Design

Single chip



# Clustered Systems: multicomputers

- Like multiprocessor systems, but **multiple systems** working together
  - Usually sharing storage via a storage-area network (SAN)
  - Provides a **high-availability** service which survives failures
    - **Asymmetric clustering** has one machine in **hot-standby** mode
    - **Symmetric clustering** has multiple nodes running applications, monitoring each other
  - Some clusters are for **high-performance computing (HPC)**
    - Applications must be written to use **parallelization**






# Operating System Operations

# Operating Systems: providing multiprogramming

- **Multiprogramming**: multiple programs can be started and loaded.
- A subset of total jobs in system is kept in memory.
- It is **convenient**
- It is **efficient**:
  - Single user cannot keep CPU and I/O devices busy at all times
- One job selected and run via **job scheduling**
  - OS selects which job
  - When the job has to wait (for I/O for example), CPU is given to another job.
- “**job**”, “**process**”, “**running program**” will be used interchangeably.



# Operating Systems: providing time sharing

---


- **Timesharing (Multitasking)** is logical extension in which CPU switches jobs so **frequently** that users can **interact** with each job while it is running, creating **interactive** computing
  - **Response time** should be  $< 1$  second
  - program loaded in memory  $\Rightarrow$  **process**
  - If several processes ready to run at the same time  $\Rightarrow$  **CPU scheduling**

# Operating System execution

---

- OS system kernel is **interrupt driven**
  - Hardware interrupt causes ISR to run (which is a routine of OS)
  - Software error or system request (system call) causes exception handler or system call handler to run
    - example: “division by zero” (**exception**)
    - Example: request for an operating system service (“open a file”) (**system call**)

OS Code  
(Kernel Code)



# Dual mode operation

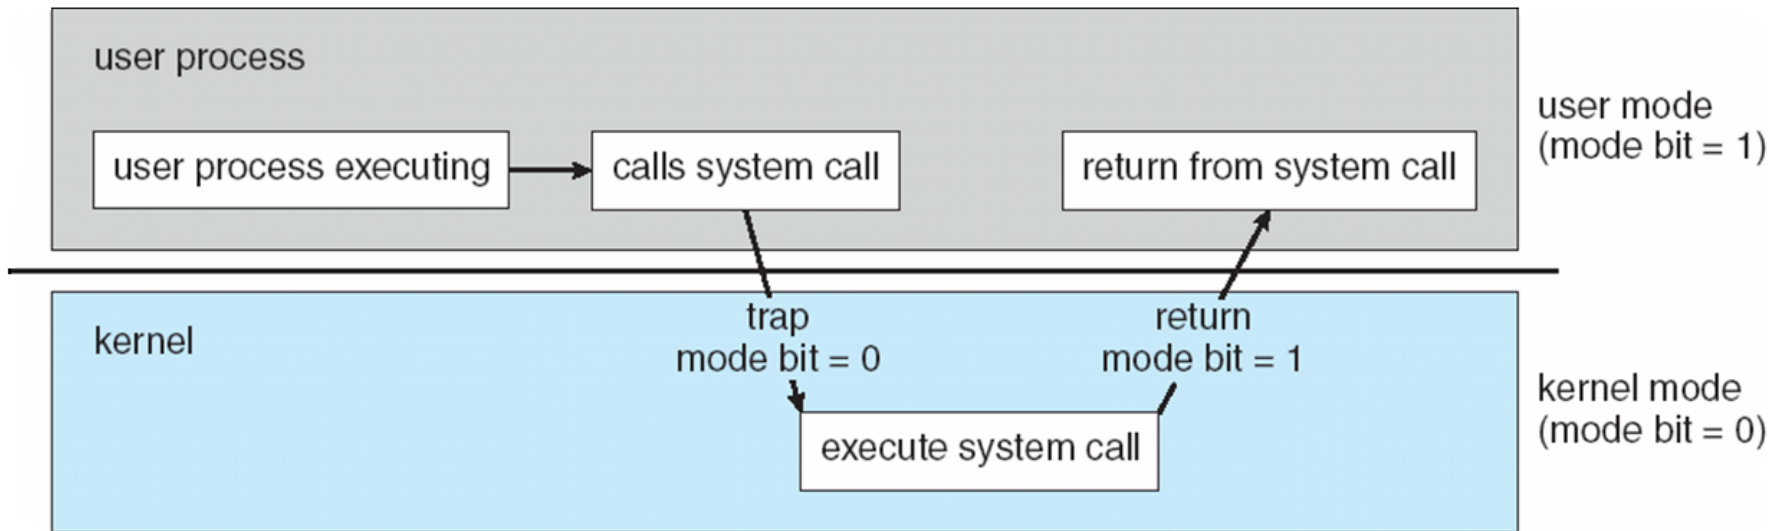
---

- **Dual-mode** operation (a hardware property) allows OS to **protect** itself and programs and other system components
- CPU can run in one of two (at least) modes:
  - **User mode** and **kernel mode**
  - **Mode bit** provided by hardware (CPU)
- User code runs in user mode;
- Kernel code runs in kernel mode.
- Some machine instructions designated as **privileged/special**, only executable in kernel mode; other instructions are **normal** instructions.
- In **kernel mode**, where all instructions (normal + privileged) can be executed.
- In **user mode**, only normal instructions are allowed to execute.

# Dual mode operation

## Dual mode system operation

### Transition from User to Kernel Mode and Vice Versa



*system-call* instruction (executed by user program) changes the mode to **kernel mode**

*return-from-system-call* (executed by the kernel) instruction resets the mode to **user mode**

# Periodic timer interrupts

- **Timer device** to prevent an infinite loop / process hogging resources
  - 1) **Set the timer device** to interrupt after a while later
    - Can be a fixed (for example 10 ms) or variable time period
  - 2) CPU executes a program (a process)
  - **3) Timer device sends an interrupt after that period**
  - 4) CPU starts executing timer handler: OS gains control
  - 5) OS can schedule the same process or other process
  - 6) OS **sets the timer again** before giving the CPU to the scheduled process

# Major OS Functionalities

---

- Process management
- Memory management
- Storage (HDD or SDD) management
  - File concept, file mapping to disk blocks, disk scheduling
- I/O control and management
  - Device drivers (doing I/O), buffering, providing uniform access interface
- Protection and security
  - Controlled access to resources,
  - Preventing processes interfering with each other and OS



# Process Management

- A process is a **program in execution**.  
(unit of work) (active)
- Process executes instructions **sequentially**, one at a time, until completion
- Process needs **resources**
  - CPU, memory, I/O, files
- Typically system has many processes running **concurrently**
  - Some of them may be OS processes
- Upon termination, resources are **released**

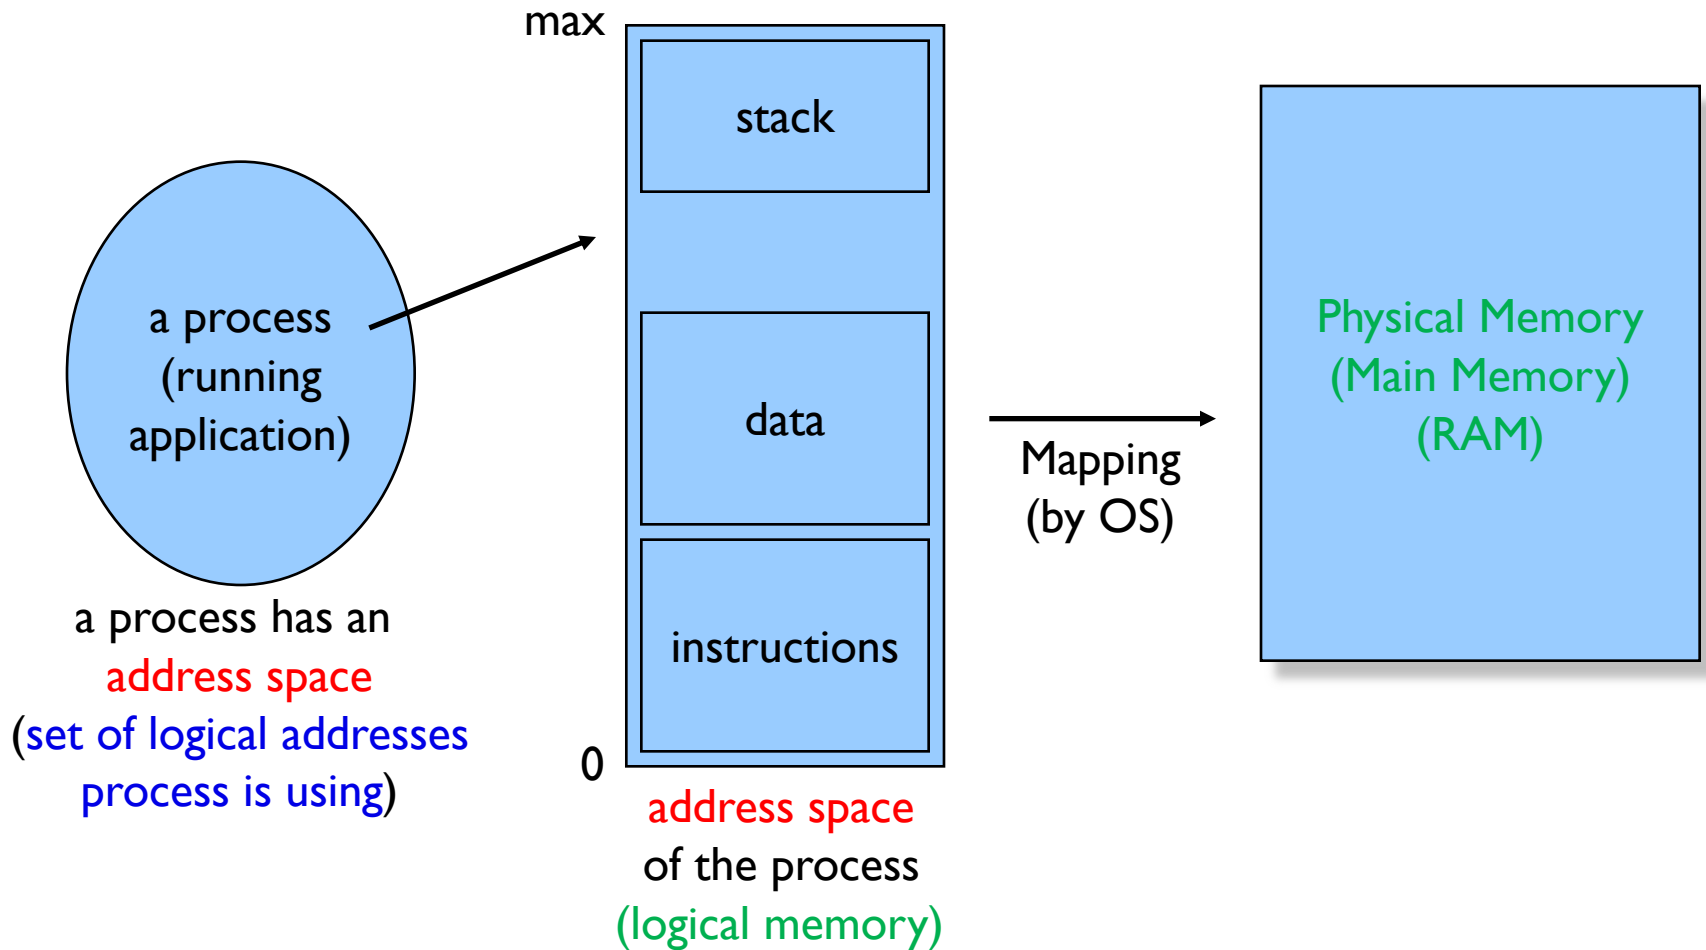
For process management:

- **Creating** and deleting both user and system processes and
- **Suspending**/resuming processes
- Providing mechanisms for process **synchronization**
- Providing mechanisms for process **communication**
- Providing mechanisms for **deadlock** handling

# Memory Management

- All **data in memory** before and after processing
- All **instructions in memory** in order to execute
- Memory management determines what is in memory, where and when
- Memory management activities
  - **Keeping track** of which parts of memory are currently being used and by which program (process)
  - Deciding which processes (or parts of a process) and data to move into and out of memory
  - **Allocating** and **deallocating** memory space as needed

# Process Address Space



# File-System Management

---

- OS provides **uniform, logical view** of information storage
  - Abstracts physical storage to logical storage unit (a file)
  - Various storage device types varying in medium type, access speed, capacity, data-transfer rate, access method
- **File-System management**
  - Files usually organized into **directories**
  - **Access control** on most systems to determine who can access what
- OS activities include
  - Creating and deleting files and directories;
  - Primitives to manipulate files/dirs;
  - Mapping files onto secondary storage

# Mass-Storage Management

---

- **Mass Storage:**  
HDD disks, SDD disks (secondary storage);  
CDs, tapes, etc. (tertiary storage)
- Proper management of mass storage devices is of central importance
  - For improving **performance** of the computer system
  - Since they are **slow devices**.
- OS activities
  - **Free-space management**; Storage allocation
  - Disk **scheduling**
  - **Uniform naming** ....

# Performance of various levels of storage

---

- Movement between **levels of storage hierarchy** can be explicit or implicit.


Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape

# Input/Output Subsystem

---

- One purpose of OS is to **hide peculiarities of hardware devices** from the user
- **I/O subsystem** responsible for
  - Buffering, caching,
  - General device-driver interface
  - Drivers for specific hardware devices  
Interacting with the device and doing I/O


I/O sub-system of Kernel



# I/O Structure

---

- Application programs do I/O via OS.
  - The request is done by calling a system call (OS routine)
  - System call routine in OS performs the I/O via the help of device driver routines in OS.
  - After issuing a system call, an application
    - may wait for the call to finish (blocking call), or
    - may continue to do something else (non-blocking call)





# Protection and Security

---

- **Protection** – any mechanism for **controlling access** of processes or users to resources defined by the OS
- **Security** – defense of the system against internal and external **attacks**
  - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally first distinguish among users to determine who can do what
  - User identities (**user IDs**, security IDs) include name and associated number, one per user
  - **User ID** of the user is then associated with all the **files** and **processes** of the user to do access control

# Kernel Data Structures

- Lists
  - Singly linked lists
  - Double linked lists
  - Circular linked lists
- Queues
- Stacks
- Trees
  - Binary search tree
  - Balanced binary search tree (red-black tree)
- Hash functions and hash tables
- Bitmaps

## *LINUX KERNEL DATA STRUCTURES*


The data structures used in the Linux kernel are available in the kernel source code. The *include* file `<linux/list.h>` provides details of the linked-list data structure used throughout the kernel. A queue in Linux is known as a `kfifo`, and its implementation can be found in the `kfifo.c` file in the `kernel` directory of the source code. Linux also provides a balanced binary search tree implementation using *red-black trees*. Details can be found in the include file `<linux/rbtree.h>`.

# Virtual Machines

---

- Hardware is abstracted into several different execution environments
  - Virtual machines
- Each virtual machine provides an interface that is identical to the bare hardware
- A *guest* kernel (and processes) can run on top of a virtual machine.
  - We can run several operating systems on the same *host*.
  - Each virtual machine can run another operating system.

# Virtual Machines





# Different Types of Computer Systems and Applications

## (Computing Environments)

# Distributing Computing

---

- Earlier systems executed tasks on a single system
- Now we have systems interconnected (networked)  
Operating systems have now **support** for **networking** multiple systems,
  - enabling data communication
  - Enabling distributed computing
  - Enabling resource sharing
  - Enabling distributing file storage
- Therefore, the computing environment is no longer a single system.



# Computing Environments


---

- Traditionally



a single system with a user


Computing and OS  
in a single machine



# Computing Environments

---

- Client-Server Computing
  - Dumb terminals replaced by smart PCs
  - Many systems now are **servers**, responding to requests generated by **clients**
  - A **compute-server** provides an interface to clients to request services (i.e., database) executed in the server
  - **File-server** provides interface for clients to store and access files





# Peer-To-Peer Computing


---

- Another model of distributed system.
- P2P **does not distinguish** clients and servers
  - Instead **all nodes** are considered **peers**
  - Each may **act as a client, as a server, or both**
  - A node must join P2P network
  - A peer **registers its service** with central lookup service on network, or
  - Peer **broadcasts requests** for service and the serving peer(s) responds (resource discovery/lookup protocol)

# Web Based Computing

---

- Web has become ubiquitous
- More devices becoming networked to allow web access
- OSs run web servers and web clients
- **Web based applications** can be developed to **run over web servers** and **clients**.
  - Having a browser at the client is enough to run an application.
  - No special client software required



# Mobile Computing

---

- Computing on smart phones and tablets.
- Potable and lightweight devices: mobile devices
- Many sensors: GPS, accelerometers, gyroscope, etc.
- Small screen, touch screen, no keyboard/mouse
- Wireless interfaces (port): 3G/4G, WiFi, Bluetooth.
- Mobile OS: iOS or Android

## Applications:


- New types:
  - Applications that use sensors
  - Location based applications
- How they are developed and run:
  - Web based applications or
  - Native applications

# Cloud Computing

---

- Type of computing that delivers **computing, storage, applications as a service** across a network.
- Different types of services:
  - *Computing as a service*: remote **virtual machine instances** or **platforms-APIs** (software stacks) - **IaaS** or **PaaS**
  - *Storage as a service*: block storage (remote virtual disks), object storage (blob storage) - **IaaS**
  - *Software as a service*: Internet services, email services, web based services, ... - **SaaS**
- Public Cloud: Can use anyone
- Private Cloud: internal to a company

# Cloud Computing



A public cloud providing IaaS

# Real Time Embedded Systems

---

- Embedded Computing
- Embedded computers in car engines, robots, microwave ovens, ...
- They do **specific tasks**.
- Little or no interface (no monitor)
- Some use general purpose processors (CPUs) and OSs (Linux)
- Some use ASICs – No OS
- OS is **real time OS**
  - Rigid timing requirements for tasks to be performed

# Open-Source Operating Systems

---

- Some operating systems made available in **source-code** format rather than just binary closed-source
- Examples include
  - **GNU/Linux**,
  - BSD Unix (**FreeBSD**, etc.)
  - **Sun Solaris**
- Closed source: Windows
- Hybrid: Mac OS X, iOS

# References

---


- Operating System Concepts, Silberschatz et al.
- Modern Operating Systems, Andrew S. Tanenbaum et al.
- OSTEP, Remzi Arpaci-Dusseau et al.



# Migration of Integer A from Disk to Register

---

- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide cache coherency in hardware such that all CPUs have the most recent value in their cache

