



Bilkent University  
Department of Computer Engineering  
CS342 Operating Systems

---

# Introduction - 2

## Operating System Structures

Last Update: Sep 20, 2022

# Outline

---

## Outline

- Operating System **Services**
- User Operating System **Interface**
- **System Calls**
- System Programs
- Operating System Structure

## Objectives

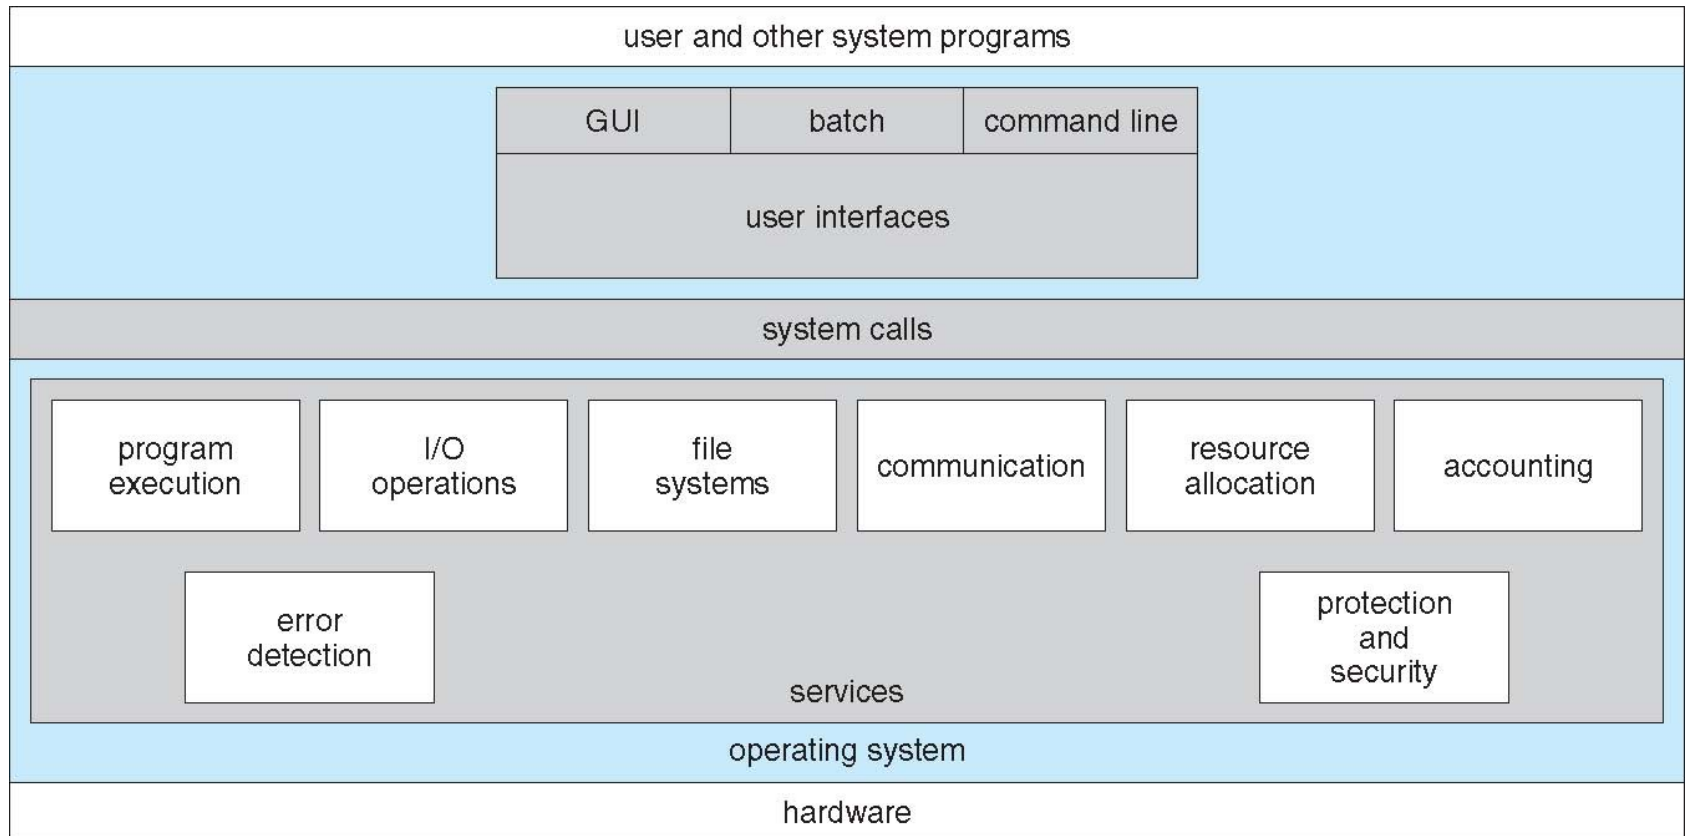
- To describe the services an operating system provides to users, processes, and other systems
- To discuss the various ways of structuring an operating system

# Operating System Services

---

- For user
  - User interface
  - Program execution
  - I/O operations
  - File-system manipulation
  - Process communication
  - Error detection and handling
- For System: efficiency and sharing
  - Resource allocation
  - Accounting
  - Protection and security

# OS Services




# [User - Operating System] Interface - CLI

---

- **CLI**: Command Line Interface (CLI) or

## command interpreter (shell)

- fetches a **command** from user and executes it
- Interpreter is in kernel or is a system program,
- many flavors
  - Some commands may be **built-in to the shell**,
  - Come commands may be **another program**

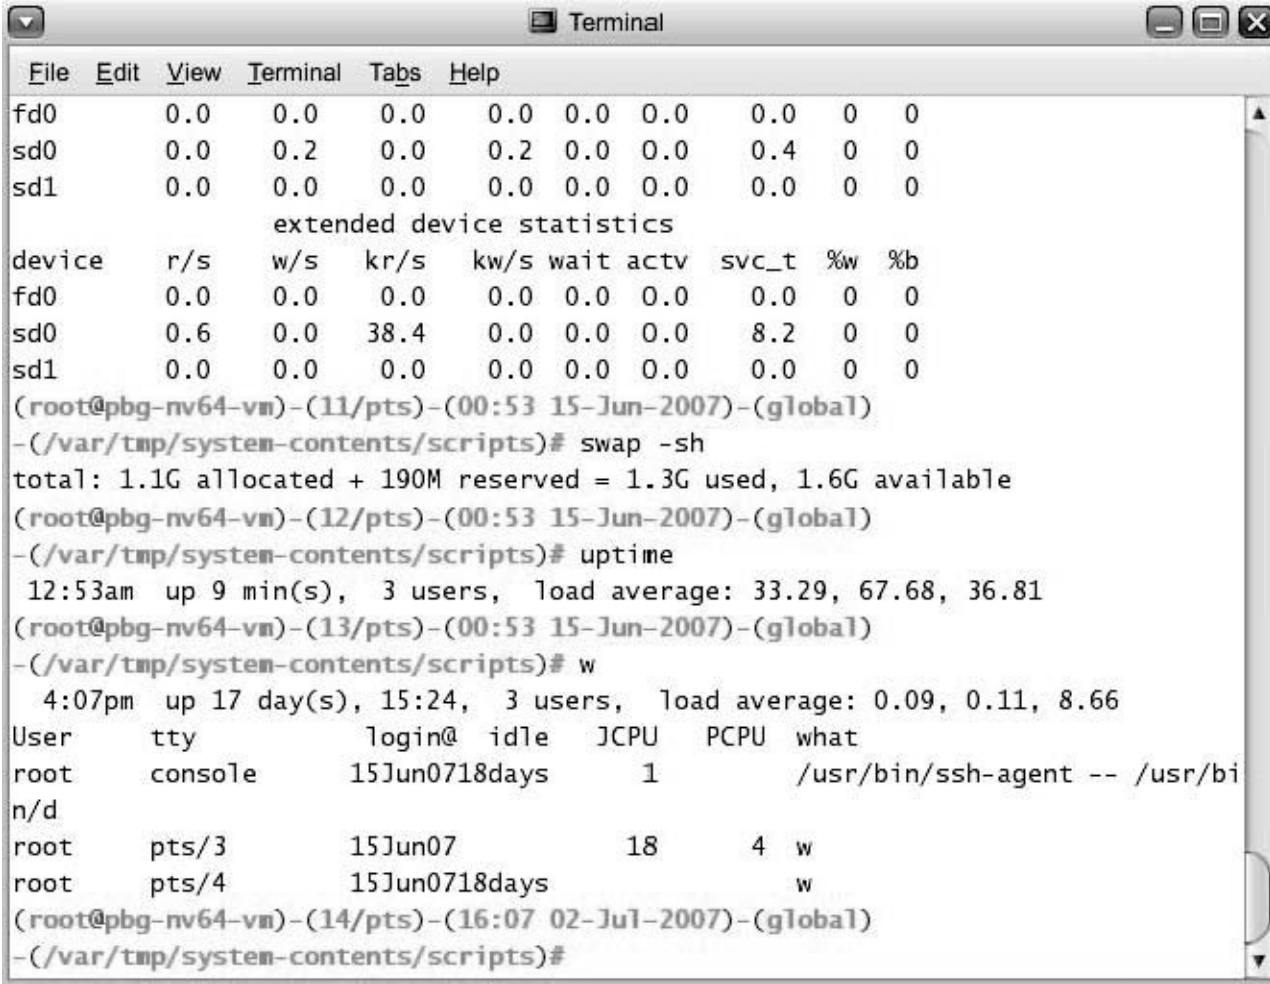


- **GUI**: User-friendly **desktop** interface

- **Icons** represent files, programs, actions, etc.

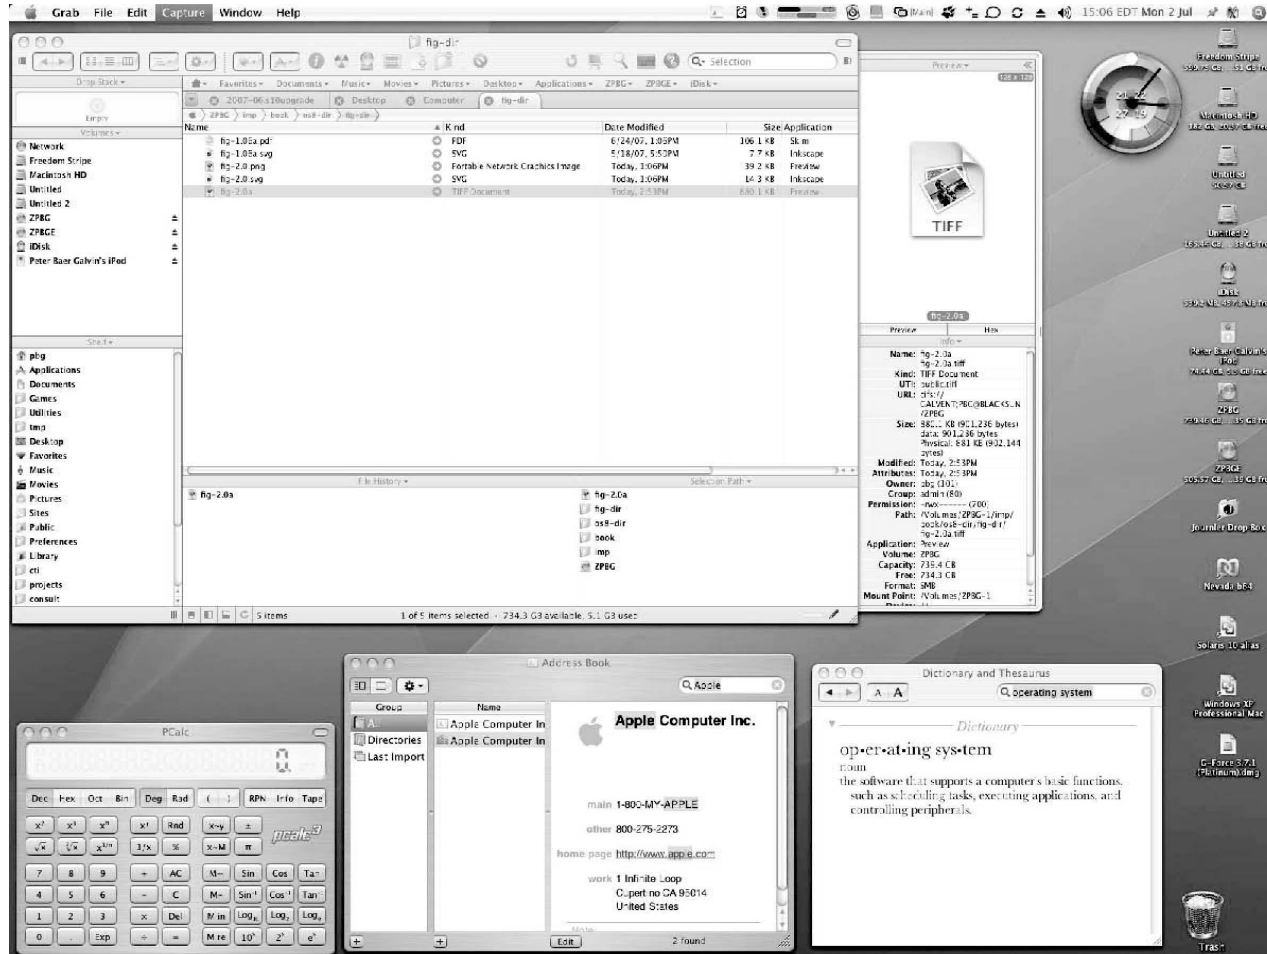
*Many operating systems now include both CLI and GUI interfaces  
Linux: command shells available (CLI); KDE as GUI*

# Bourne Shell Command Interpreter



```
File Edit View Terminal Tabs Help
fd0      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
sd0      0.0    0.2    0.0    0.2  0.0  0.0    0.4  0  0
sd1      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
          extended device statistics
device   r/s    w/s    kr/s    kw/s wait actv  svc_t  %w  %b
fd0      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
sd0      0.6    0.0   38.4    0.0  0.0  0.0    8.2  0  0
sd1      0.0    0.0    0.0    0.0  0.0  0.0    0.0  0  0
(root@pbg-nv64-vm)-(11/pts)-(00:53 15-Jun-2007)-(global)
-(/var/tmp/system-contents/scripts)# swap -sh
total: 1.1G allocated + 190M reserved = 1.3G used, 1.6G available
(root@pbg-nv64-vm)-(12/pts)-(00:53 15-Jun-2007)-(global)
-(/var/tmp/system-contents/scripts)# uptime
12:53am up 9 min(s), 3 users, load average: 33.29, 67.68, 36.81
(root@pbg-nv64-vm)-(13/pts)-(00:53 15-Jun-2007)-(global)
-(/var/tmp/system-contents/scripts)# w
4:07pm up 17 day(s), 15:24, 3 users, load average: 0.09, 0.11, 8.66
User      tty          login@ idle   JCPU   PCPU   what
root      console      15Jun0718days      1      /usr/bin/ssh-agent -- /usr/bi
n/d
root      pts/3        15Jun07      18      4   w
root      pts/4        15Jun0718days      w
(root@pbg-nv64-vm)-(14/pts)-(16:07 02-Jul-2007)-(global)
-(/var/tmp/system-contents/scripts)#
```

# The MacOS X GUI



# System Calls

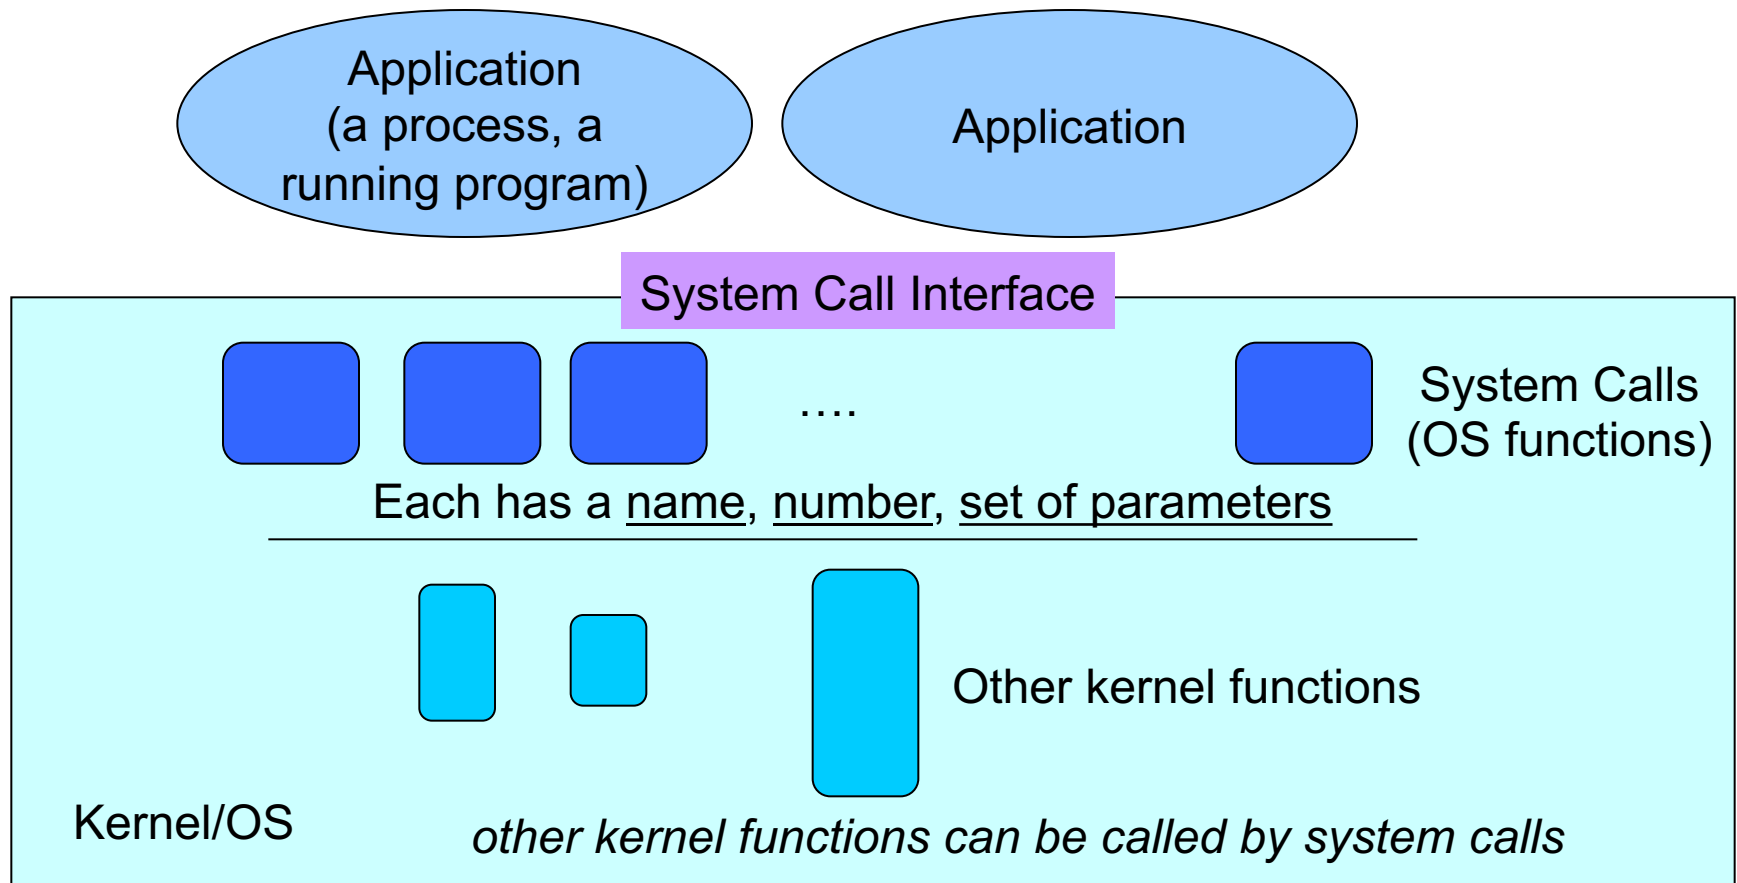
---

- **Programming interface** to the services provided by the OS
  - i.e., interface provided to applications (to commands and programs)
- Typically written in a high-level language (C or C++)
- Are called by a running program to get service from kernel
  - We say “**making a system call**” or “**calling a system call**”
- Even a simple program may make a lot of system calls per second.



# System Calls

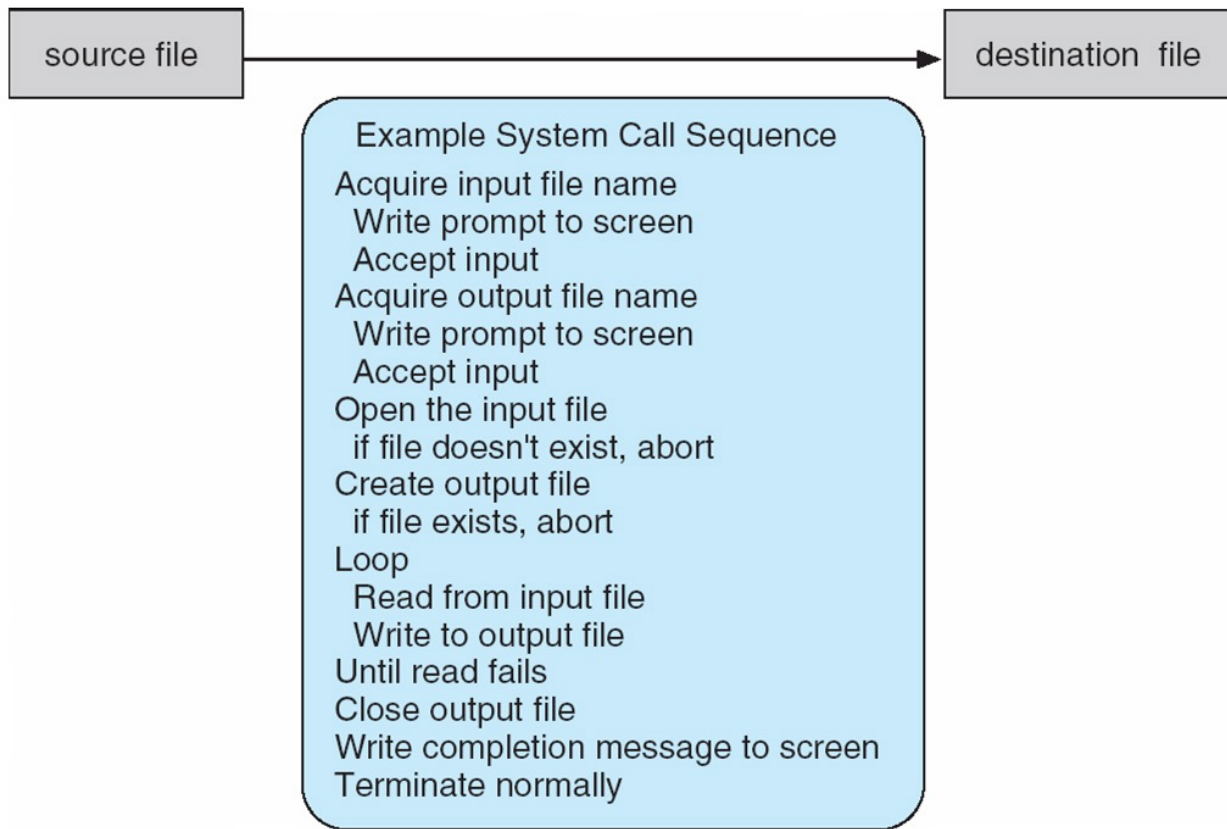
---



# Example of System Calls

---

- System call sequence to **copy** the contents of one file to another file



# Example copy program

```
=====
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include<sys/types.h>
#include<sys/stat.h>
#include <fcntl.h>

#define INSIZE 10

char infilename[128];
char outfilename[128];

char buf[1024];

int fd1;
int fd2;

int main()
{
    printf ("enter input filename\n");
    scanf ("%s", infilename);
    printf ("enter output filename\n");
    scanf ("%s", outfilename);

    printf ("%s %s\n", infilename, outfilename);

    fd1 = open (infilename, O_RDONLY);
    if (fd1 < 0) {
        printf ("error in opening");
        exit (1);
    }
    fd2 = open (outfilename, O_WRONLY | O_CREAT, 0600);

    int n, m;
    while (1) {
        n = read (fd1, buf, INSIZE);
        if (n > 1) {
            m = write (fd2, buf, n);
            if (m < 0)
                printf ("write failed\n");

        } else
            break;
    }

    close (fd1);
    close (fd2);
    printf ("finished copying\n");
}
```

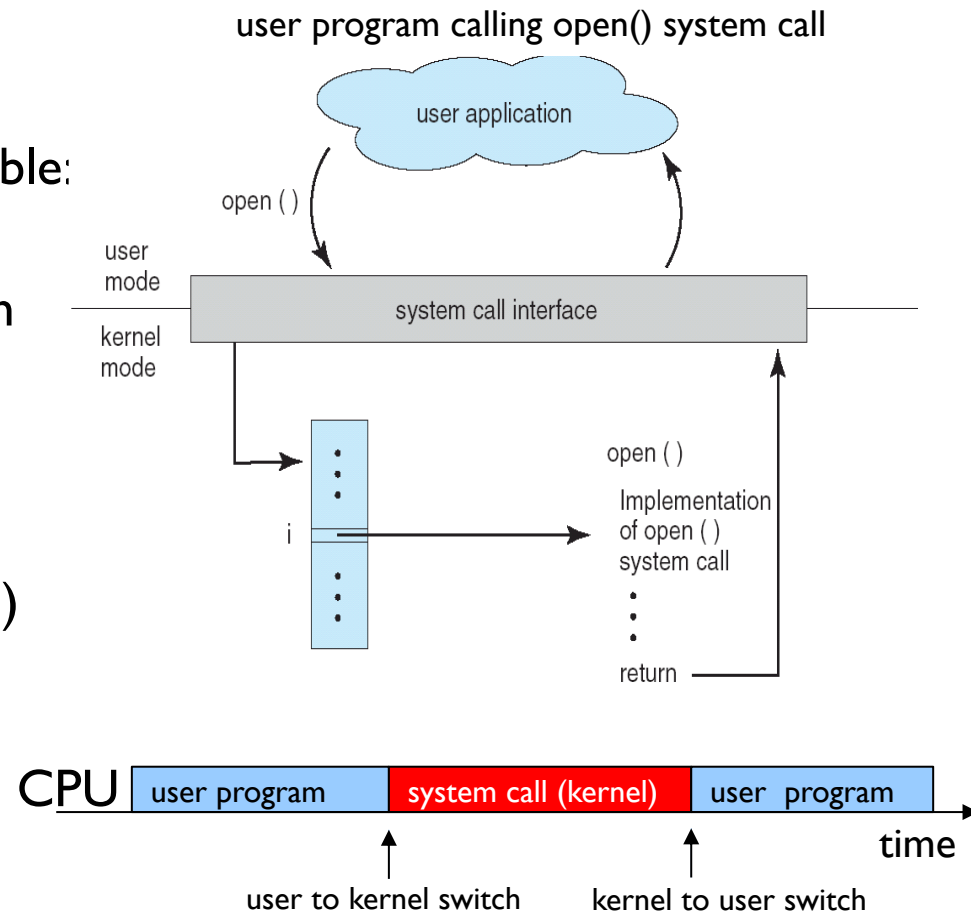
# strace output

---

```
...
read(0, Makefile
"Makefile\n", 1024)          = 9
write(1, "enter output file\n", 18enter output file
)          = 18
read(0, out
"out\n", 1024)                = 4
write(1, "Makefile out\n", 13Makefile out
)          = 13
openat(AT_FDCWD, "Makefile", O_RDONLY) = 3
openat(AT_FDCWD, "out", O_WRONLY|O_CREAT, 0600) = 4
read(3, "\n\nall: myc", 10)      = 10
write(4, "\n\nall: myc", 10)      = 10
read(3, "opy\n\nmycop", 10)      = 10
write(4, "opy\n\nmycop", 10)      = 10
read(3, "y:\tmycopy.", 10)       = 10
write(4, "y:\tmycopy.", 10)       = 10
read(3, "c\n\tgcc -Wa", 10)      = 10
write(4, "c\n\tgcc -Wa", 10)      = 10
read(3, "ll -o myco", 10)        = 10
write(4, "ll -o myco", 10)        = 10
read(3, "py mycopy.", 10)        = 10
write(4, "py mycopy.", 10)        = 10
read(3, "c\n\n\nclean:", 10)     = 10
write(4, "c\n\n\nclean:", 10)     = 10
...
```

# System Call Implementation and Calling

- Typically,
  - a **number** associated with each system call
  - Number is used as index to a table: **System Call Table**
  - Table keeps **addresses** of system calls (system routines)
- When called, system call:
  - **runs**
  - **may block** (e.g., read system call)
  - **returns**
- Caller does not know system call **implementation**
  - Just knows the **interface**



# Some Linux x86-64 system calls (64 bit)

---

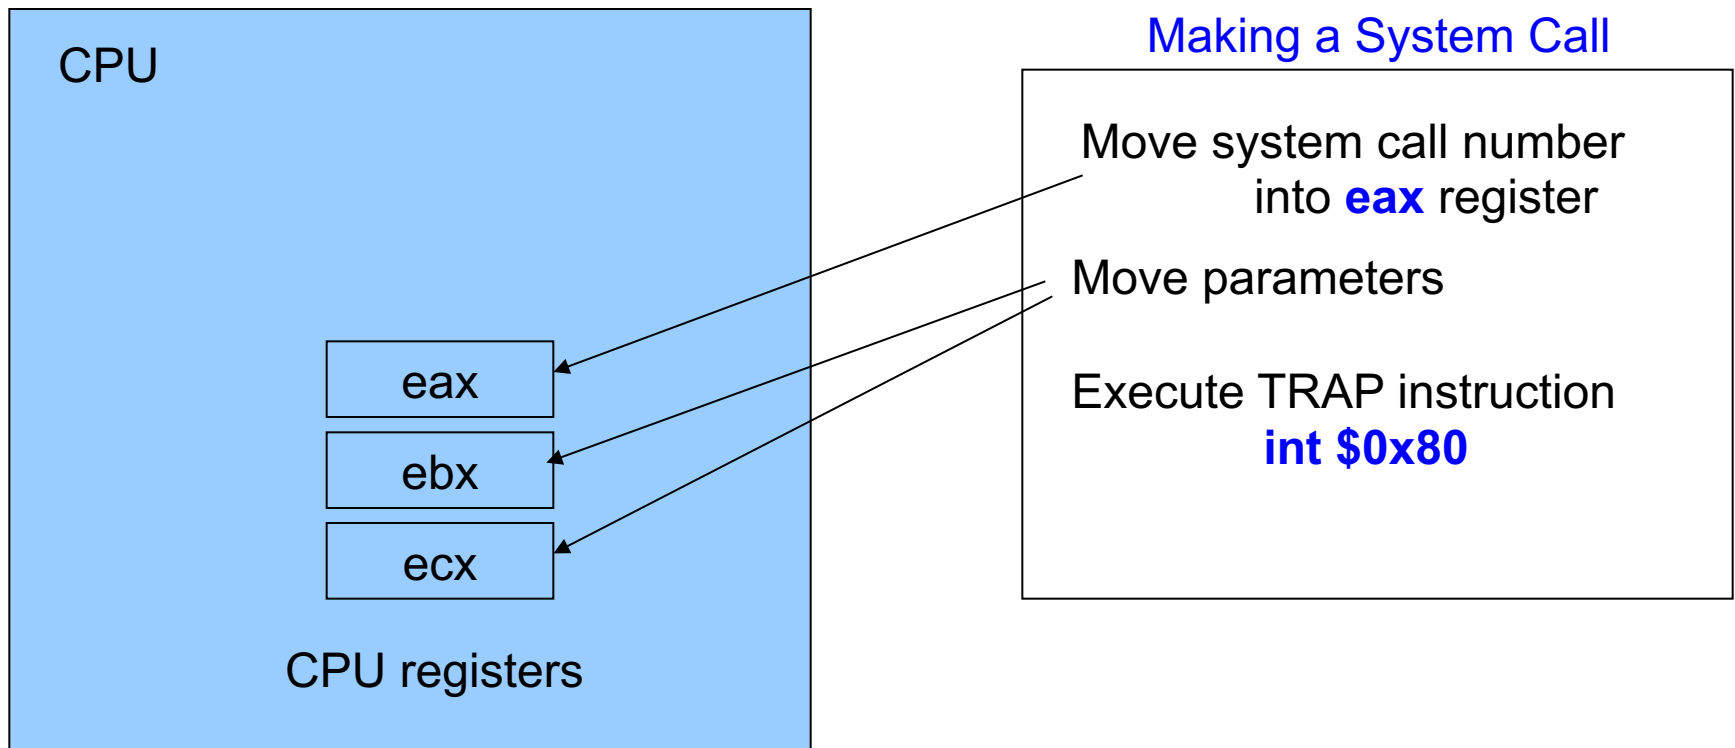
0 read	// read from a file	33 pause	
1 write	// write into a file	37 alarm	// set an alarm
2 open	// open a file	39 getpid	// get process id
3 close	// close a file	57 fork	// create a new process
4 stat	// get file attributes	59 execve	// run a new program
9 mmap	// map a file	60 _exit	// terminate the program
12 brk	// extend heap segment	61 wait4	// wait for child
32 dup2	// dup file desc	62 kill	// send signal

*Around 400 system calls in total in Linux Kernel.*

# Invoking a system:

## Linux and x86 32 bit architecture

---



**syscall** or **int** instruction is used to make system call

# System Call Parameter Passing

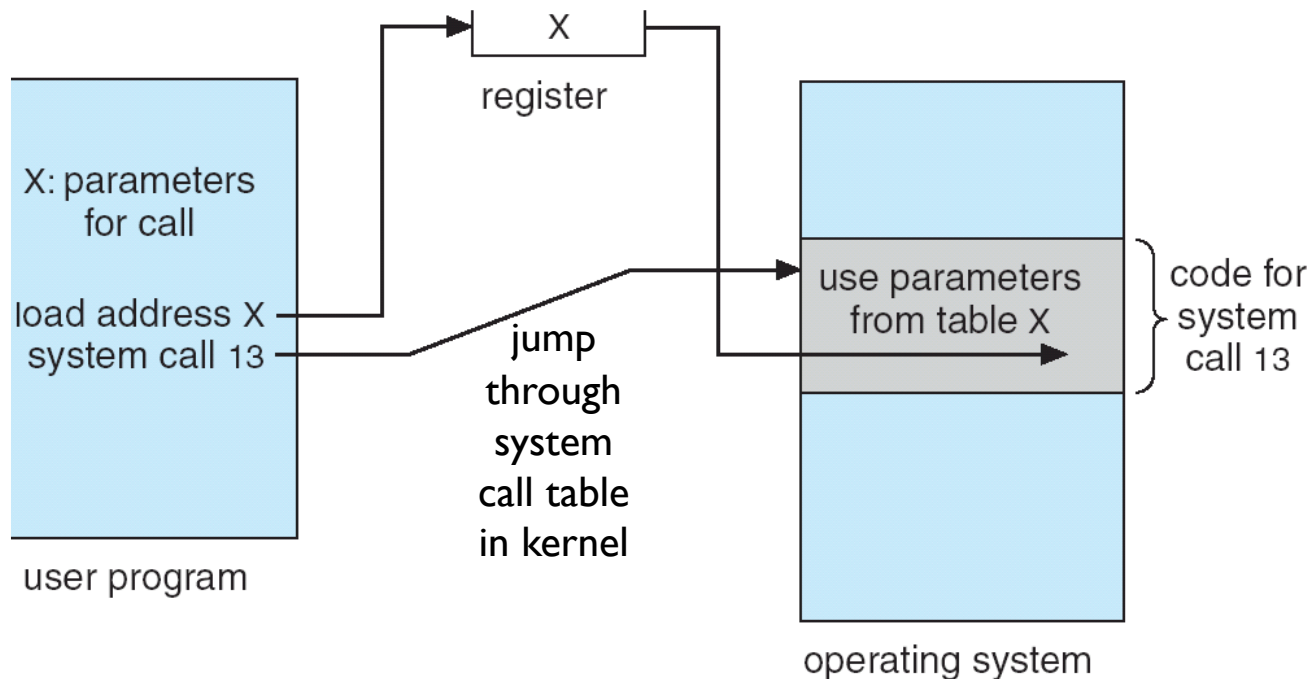
---

- Often, more information is required than the identity (number) of the desired system call
  - Exact type and amount of information vary according to OS and call
- Three general methods used **to pass parameters** to the system call
  - 1) Simplest: pass the parameters in *registers*
    - In some cases, may be more parameters than registers
  - 2) Parameters stored in a *block*, or *table*, in memory, and address of the block passed as a parameter in a register
  - 3) Parameters placed, or *pushed*, onto the *stack* by the program and *popped* off the stack by the operating system

*Last two methods do not limit the number or length of parameters being passed*



# Parameter Passing via Table




kernel can access user memory

# Accessing and executing System Calls

---


- System calls typically **not accessed directly** by programs
- Mostly accessed by programs via a high-level **Application Program Interface (API)** (i.e., a library) rather than direct system call use
- Three most common APIs are :
  - **Win32** API for Windows 32 bit system,
  - **POSIX API** for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X),
  - **Java API** for the Java virtual machine (JVM)



# Example of Standard API - Windows

---

- Consider the `ReadFile()` function in the Win32 API — a function for reading from a file



- A description of the `parameters` passed to `ReadFile()`
  - `HANDLE file`—the file to be read
  - `LPVOID buffer`—a buffer where the data will be read into and written from
  - `DWORD bytesToRead`—the number of bytes to be read into the buffer
  - `LPDWORD bytesRead`—the number of bytes read during the last read
  - `LPOVERLAPPED ovl`—indicates if overlapped I/O is being used

# Example of Standard API - Linux

---

## EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count)
```

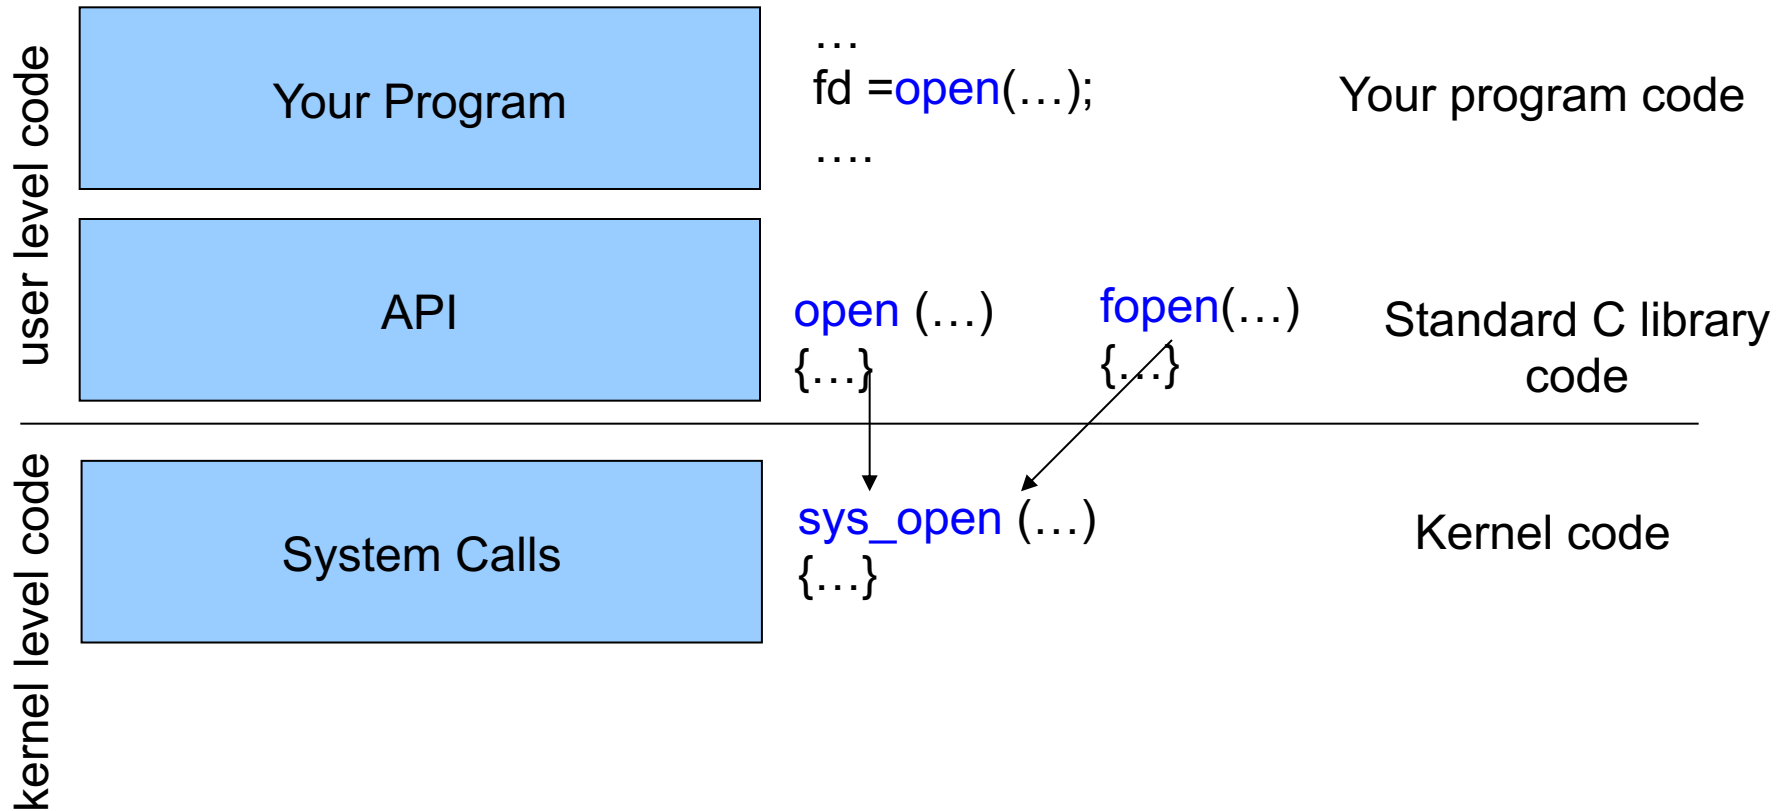
return value	function name	parameters
-----------------	------------------	------------

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.

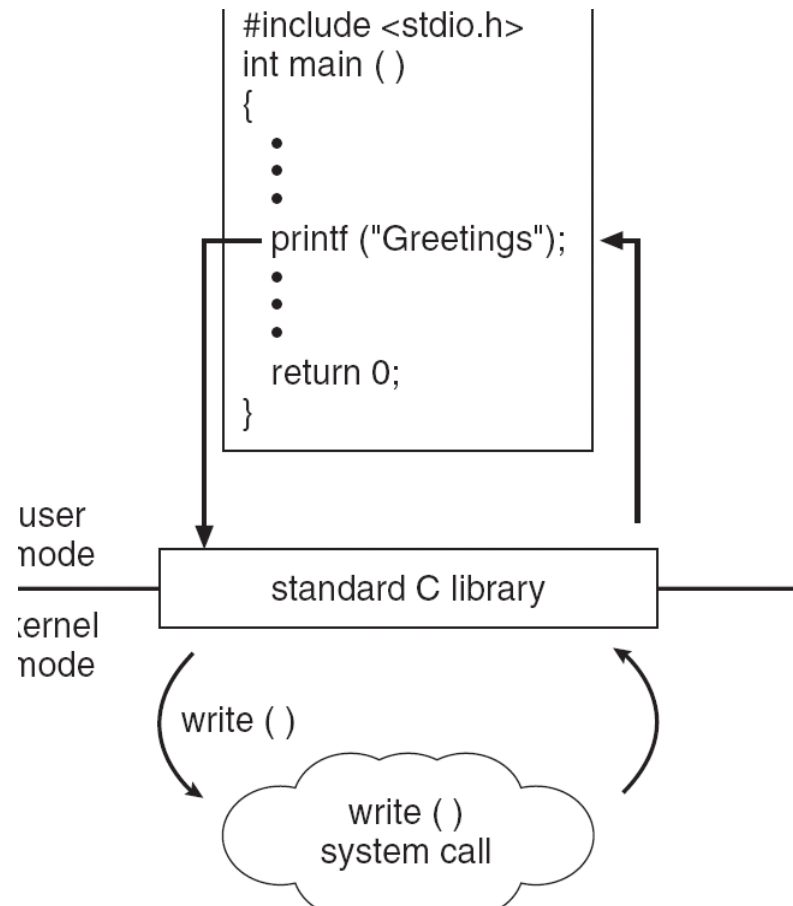
# Why use APIs rather than system calls directly?



# Standard C Library Example

---

- C program invoking `printf()` library call, which calls `write()` system call



# Types of System Calls

---

- We can categorize system calls:
  - Process control
  - File management
  - Device management
  - Information maintenance
  - Communications
  - Protection

# Examples of Windows and Unix System Calls

---

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



# Hello World with Linux System Call

hello.c

```
int main()
{
    write(1,
        "hello, world\n",
        13);
    _exit(0)
}
```

In 64 bit Linux  
for x86\_64 arch

hello.S

```
.section .data
string:
    .ascii "hello, world\n"
string_end:
    .equ len, string_end - string.section
.text.globl
main:

    movq $1, %rax    # write system call number
    movq $1, %rdi    # put file handle for output
    movq $string, %rsi #put the string
    movq $len, %rdx   #put length
    syscall          #make the call
    movq $60, %rax   #exit system call number
    movq $0, %rdi
    syscall.         #make the call
```

# Hello World with direct Linux System Call

---

Compile and run the assembly version as below:

```
$ gcc -c hello.S -o hello.o
```

```
$ gcc -no-pie hello.o -o hello
```

```
$ ./hello
```

# System Services/Programs

---

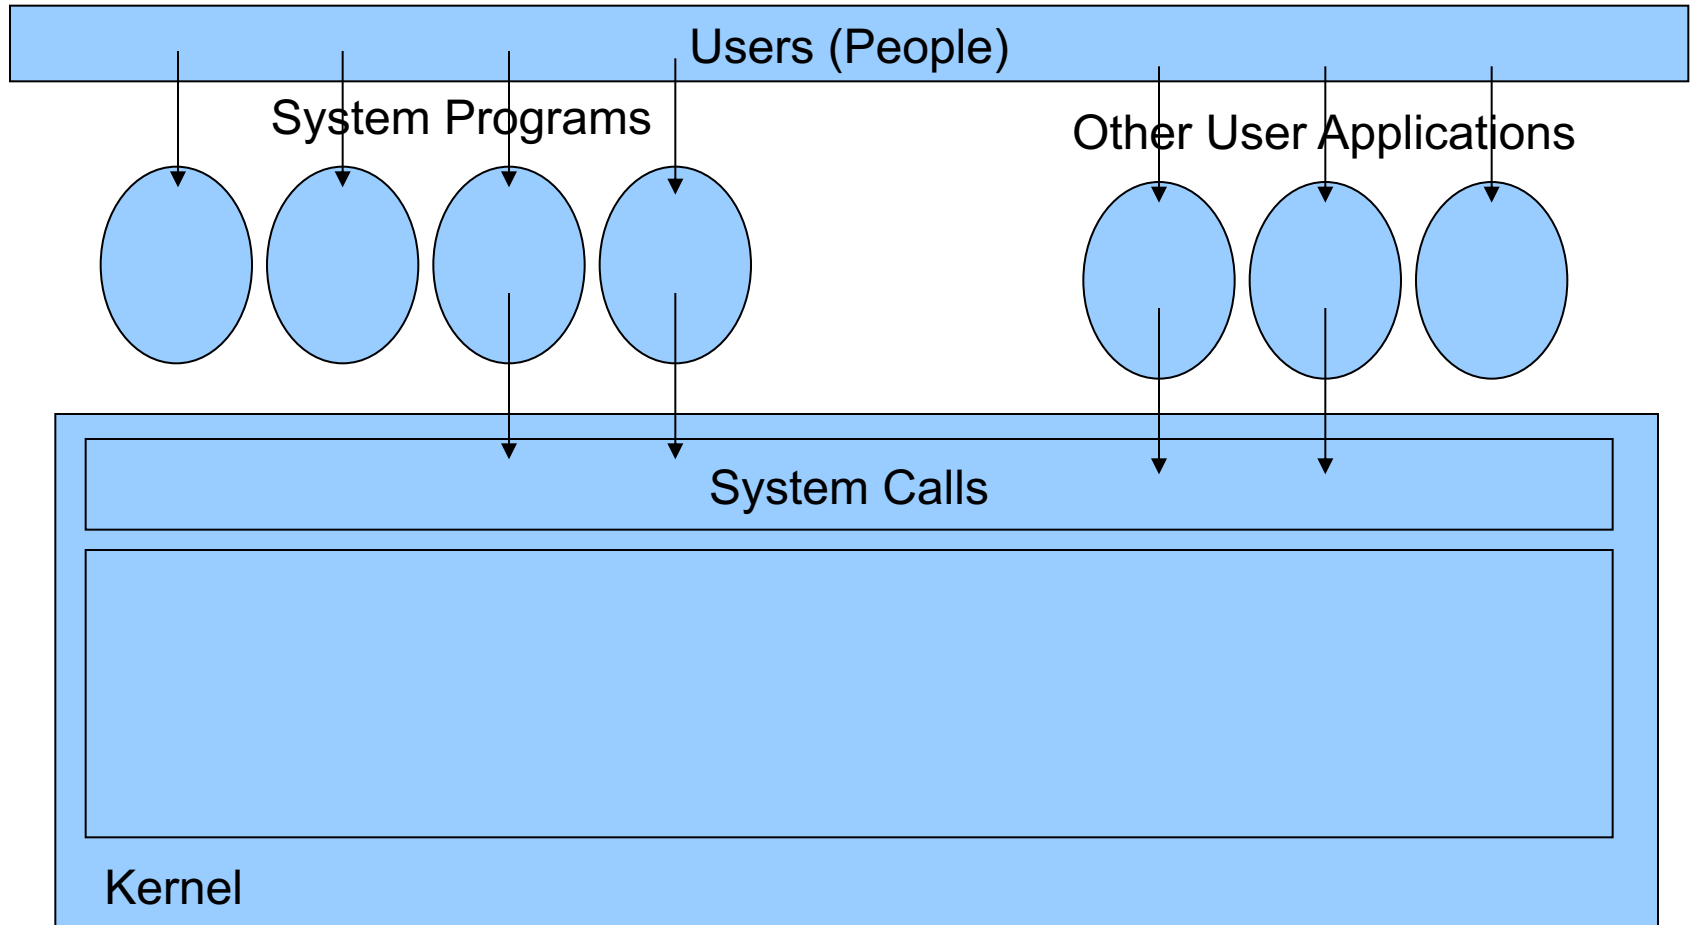
- **System programs** provide a convenient environment for program development and execution. They help you to manage the system. They can be divided into categories:
  - **File** manipulation (create, delete, copy, rename, print, list, ...)
  - **Status** information (date, time, amount of available memory, disk space, who is logged on, ...)
  - File **modification** (text editors, grep, ...)
  - Programming **language** support (compiler, debuggers, ...)
  - Program loading and **execution** (loaders, linkers)
  - **Communications** (ftp, browsers, ssh, ...)
- Other System Utilities/Applications may come with OS CD (games, math solvers, plotting tools, database systems, spreadsheets, word processors, ...).

# System Services/Programs

---

- Most **users' view** of the operation system is defined by system programs, not the actual system calls.
- Some of the system programs are **simply user interfaces** to system calls; others are considerably more **complex**.
  - **create file**: simple system program that can just call “create” system call or something similar (open system call with O\_CREAT flag).
  - **compiler**: complex system program.

# System Programs



From OS' s view: system+user programs are all applications

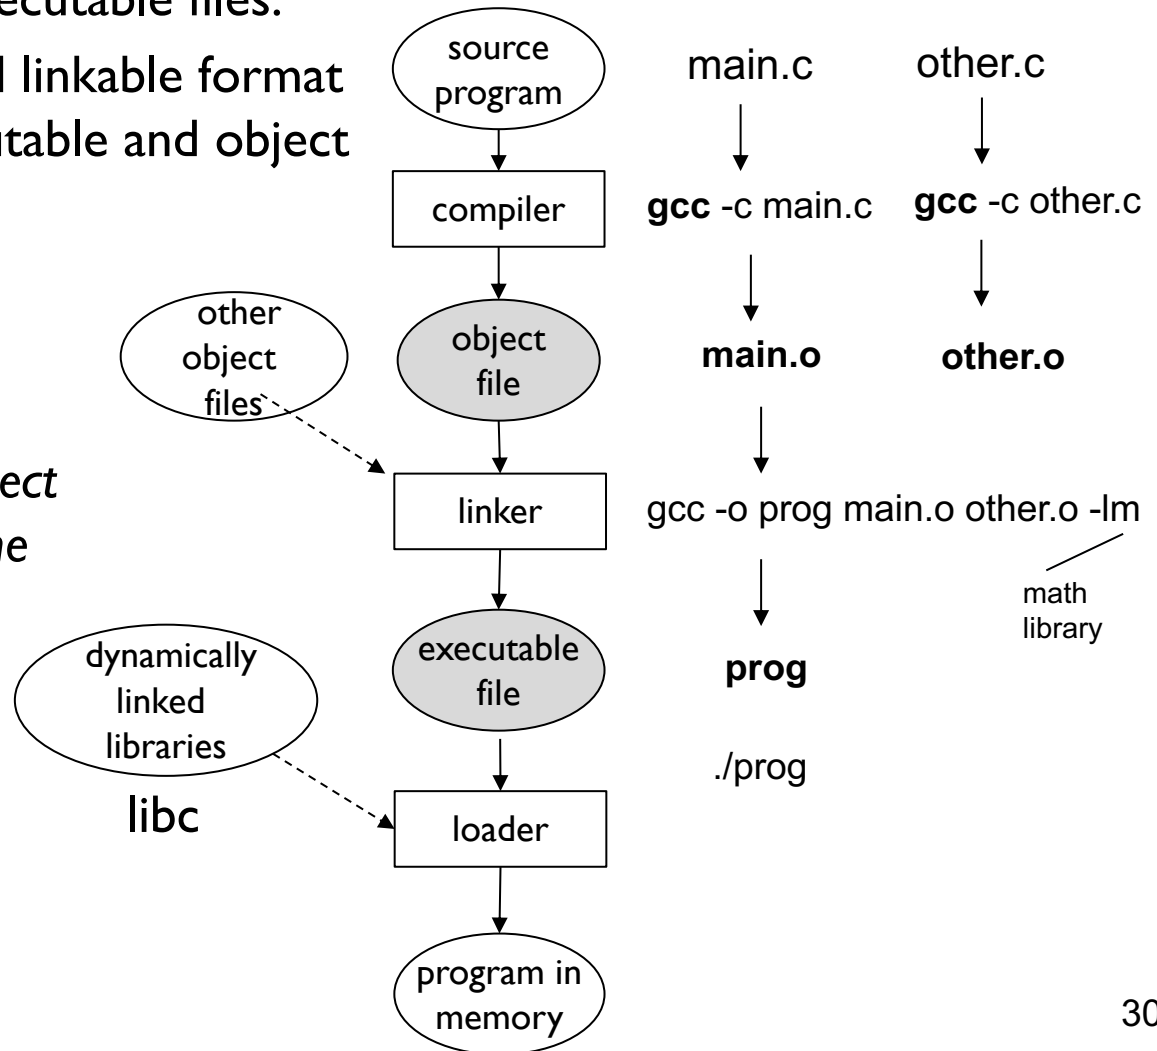
# Linkers and loaders

- They generate and load executable files.
- **ELF** format: executable and linkable format (this is the format of executable and object files in Linux)

*To get information about an object or executable file, we can use the following tools:*

**readelf**

**objdump** (more general)



# OS Design and Implementation Issues

---

- We need to set design **goals**
  - **User goals**: some goals are user oriented.
  - **System goals**: some goals are system related, not directly visible to users.
- Mechanisms and policies
  - Principle: **Separate mechanisms from policies**
  - Mechanism: **how** to do
  - Policy: **what** to do
- **Implementation**
  - C, C++ (high level language)
  - Porting to different architectures

---

# Structuring Operating System



# OS Structure

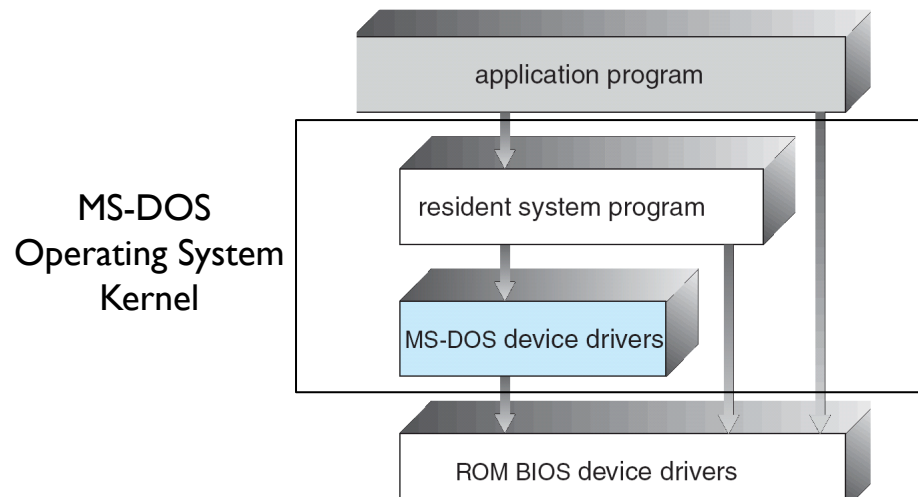
---

- Simple Structure (MSDOS)
- Layered Approach
- Microkernel Approach
- Modules Approach

# Simple Structure

---


- MS-DOS: written to provide the most functionality in the least space
  - Not divided into modules
- Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated.



# Layered Approach


---

- The operating system is divided into a number of layers (levels), each built on each other. The bottom layer (layer 0) is the hardware; the highest (layer N) is the user interface.
- Layers are selected such that each layer uses functions (operations) and services of only the layer below it.



# Layered Operating System

---



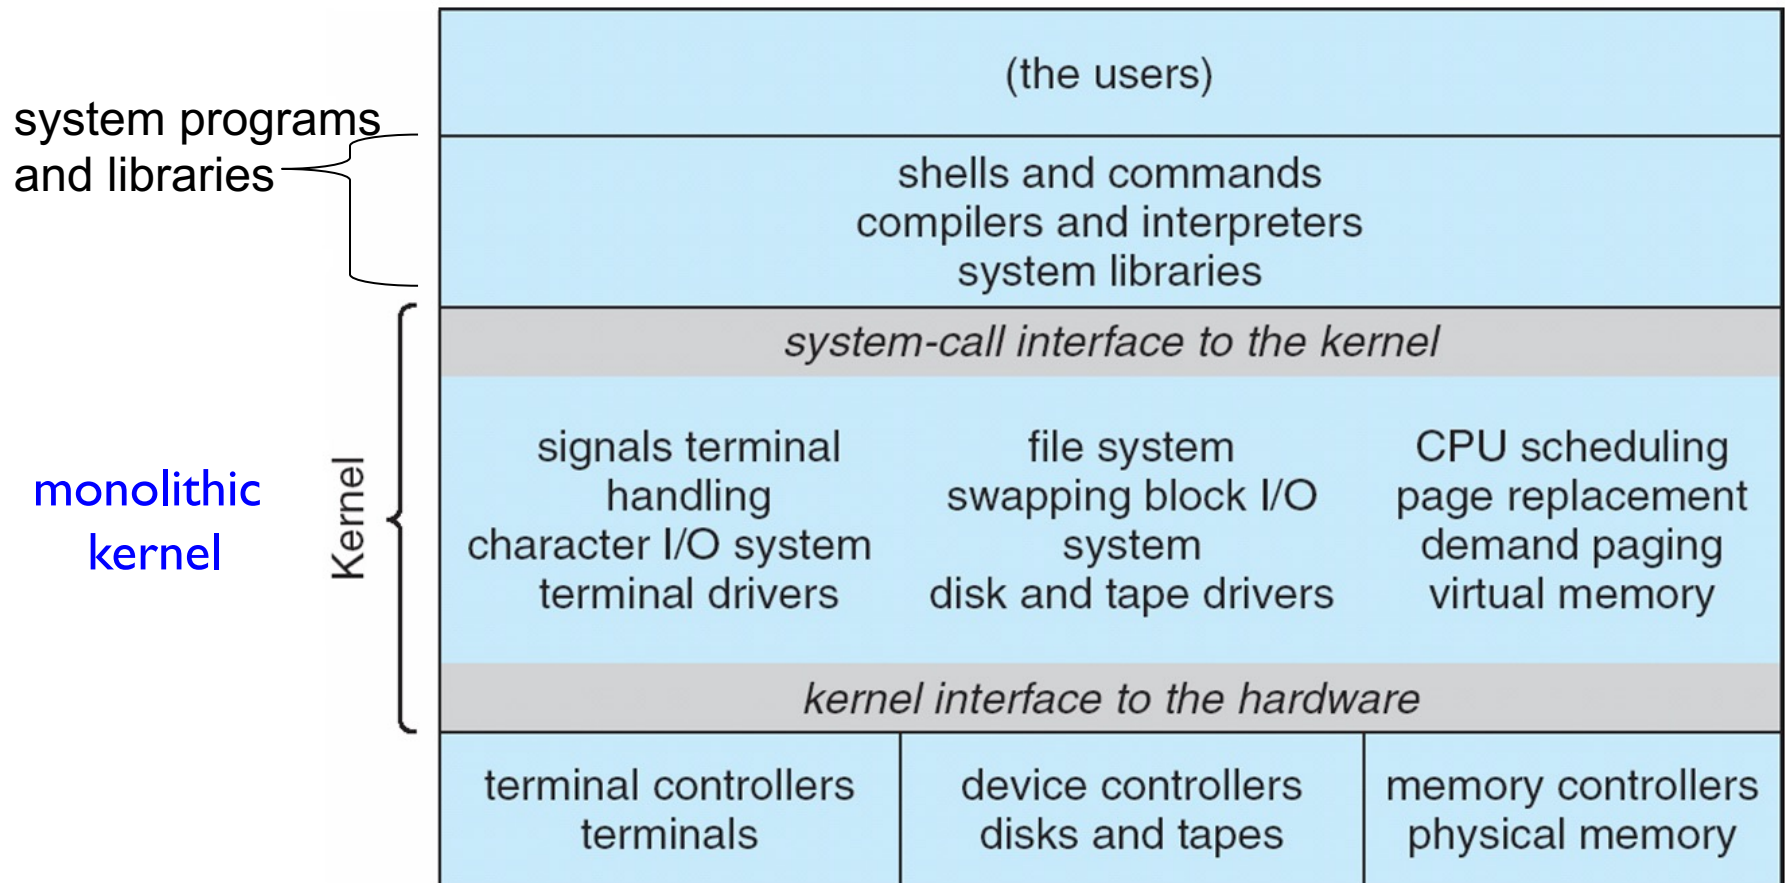
# Unix

---

- UNIX: limited by hardware functionality, the original UNIX operating system had **limited structuring**. The **UNIX OS** consists of **two separable parts**:
  - **Systems programs and libraries** (run in user mode)
  - **The kernel** (runs in kernel mode)
- The kernel consists of everything **below the system-call interface** and above the physical hardware. Provides:
  - **file system**,
  - **CPU scheduling**,
  - **memory** management,
  - and other operating-system **functions**.

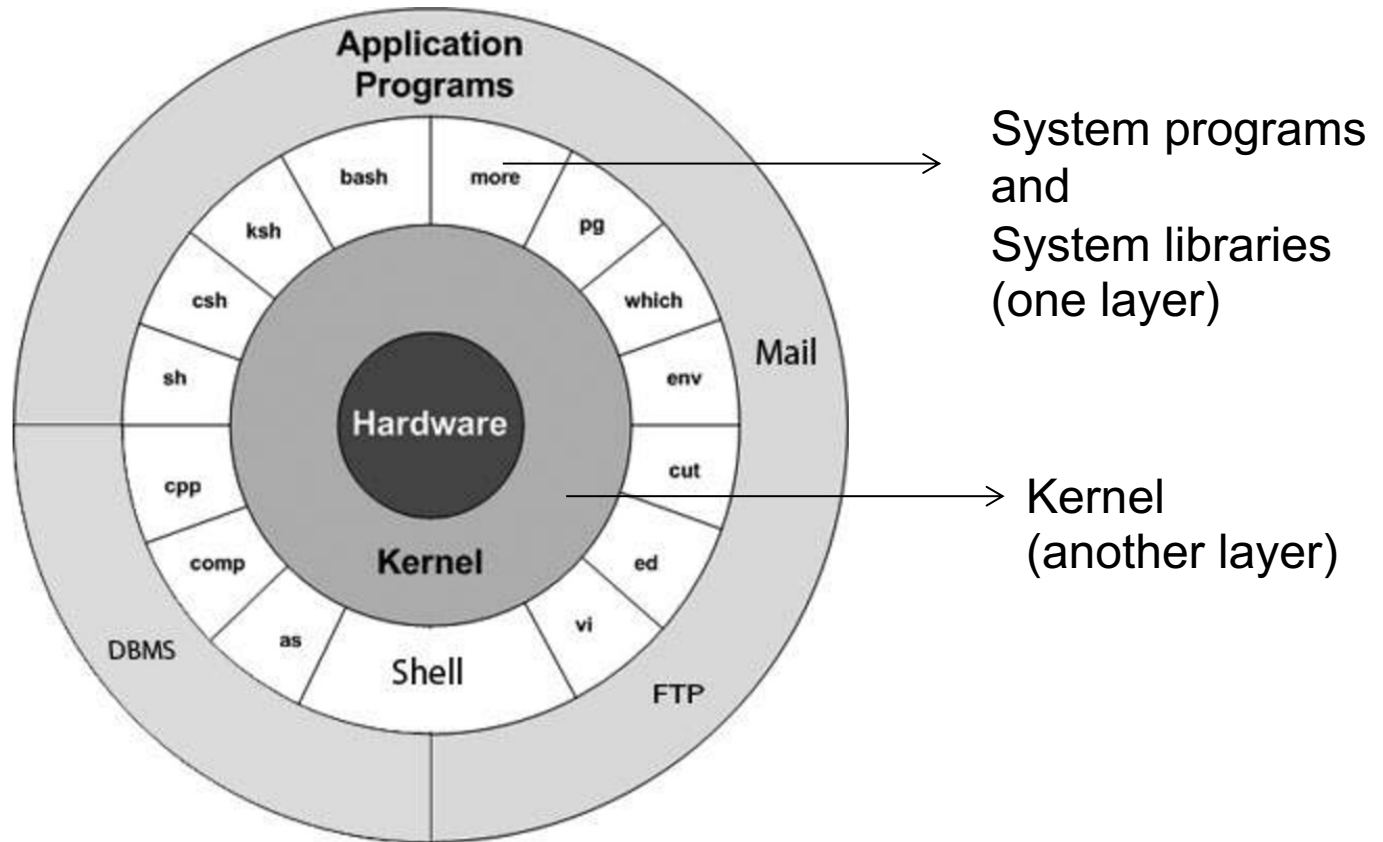
*a large number of functions for one level (monolithic kernel)*

# Traditional UNIX System Structure



# Unix OS architecture

---



From <https://www.tutorialspoint.com>

# Linux OS architecture

Various layers within Linux, also showing separation between the **userland** and **kernel space**

User mode	User applications	For example, <code>bash</code> , LibreOffice, GIMP, Blender, 0 A.D., Mozilla Firefox, etc.				
	Low-level system components:	<b>System daemons:</b> <i>systemd, runit, logind, networkd, PulseAudio, ...</i>	<b>Windowing system:</b> <i>X11, Wayland, Mir, SurfaceFlinger (Android)</i>	<b>Other libraries:</b> <i>GTK+, Qt, EFL, SDL, SFML, FLTK, GNUstep, etc.</i>		<b>Graphics:</b> <i>Mesa, AMD Catalyst, ...</i>
	C standard library	<code>open()</code> , <code>exec()</code> , <code>sbrk()</code> , <code>socket()</code> , <code>fopen()</code> , <code>calloc()</code> , ... (up to 2000 subroutines) <i>glibc</i> aims to be POSIX/SUS-compatible, <i>uClibc</i> targets embedded systems, <i>bionic</i> written for Android, etc.				
Kernel mode	Linux kernel	<code>stat</code> , <code>splice</code> , <code>dup</code> , <code>read</code> , <code>open</code> , <code>ioctl</code> , <code>write</code> , <code>mmap</code> , <code>close</code> , <code>exit</code> , etc. (about 380 system calls) The Linux kernel System Call Interface (SCI, aims to be POSIX/SUS-compatible)				
		Process scheduling subsystem	IPC subsystem	Memory management subsystem	Virtual files subsystem	Network subsystem
		Other components: ALSA, DRI, evdev, LVM, device mapper, Linux Network Scheduler, Netfilter Linux Security Modules: SELinux, TOMOYO, AppArmor, Smack				
Hardware (CPU, main memory, data storage devices, etc.)						

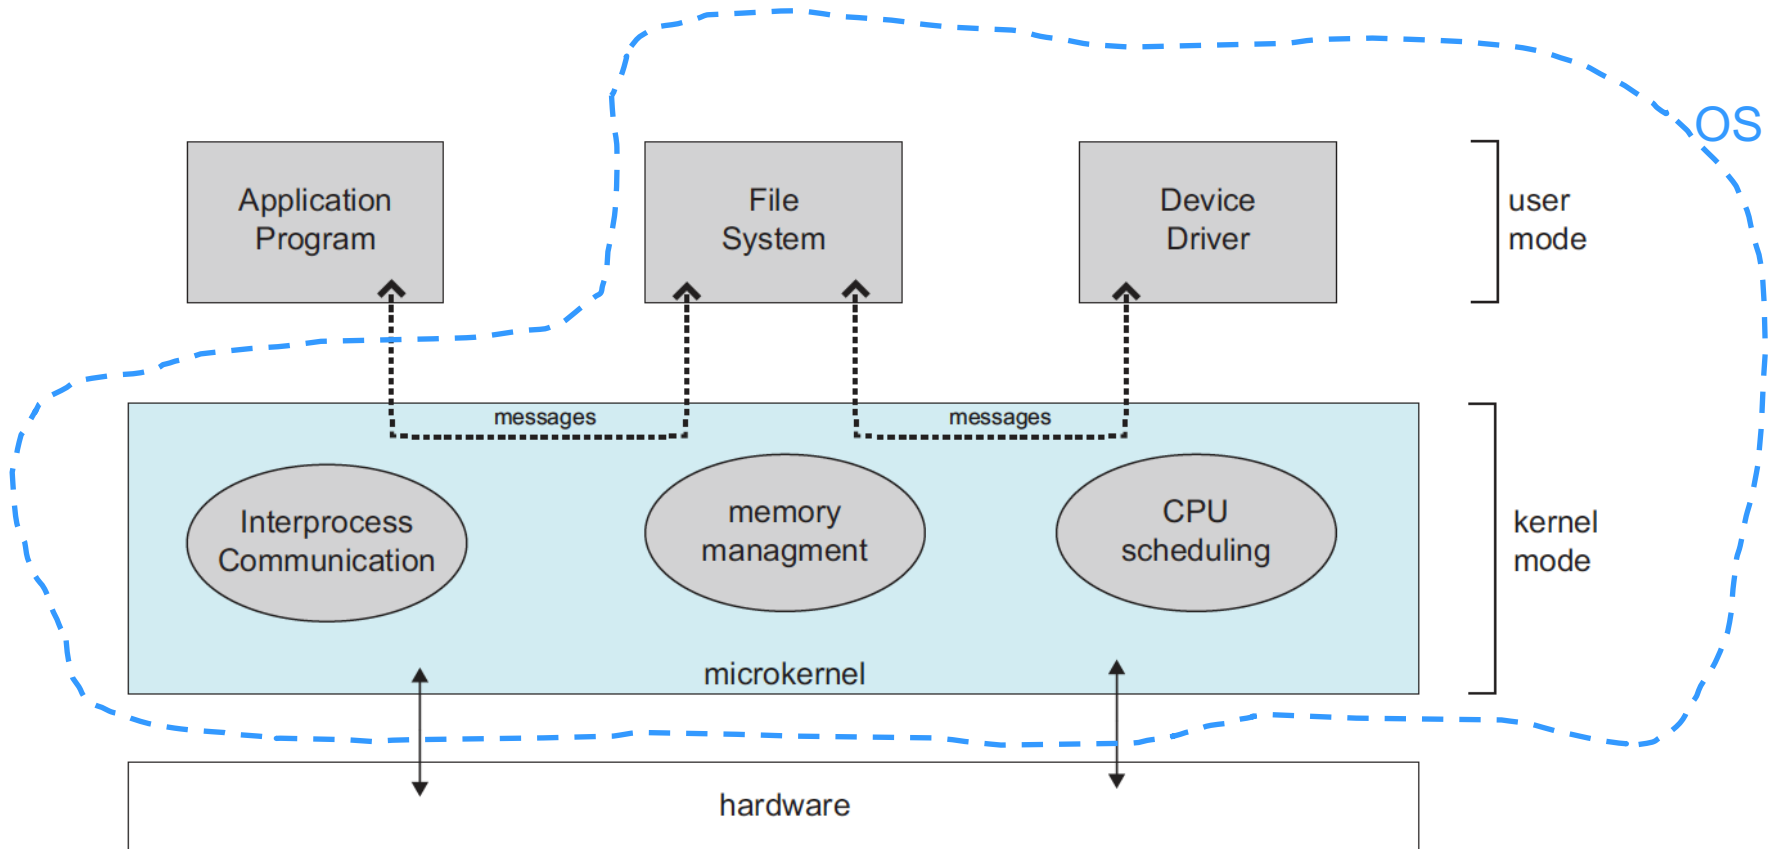


# Microkernel System Structure

---

- Moves as much from the *kernel* into *user* space
- **Communication** takes place between user-mode OS processes using **message passing**.
- Benefits:
  - Easier to **extend** a microkernel
  - Easier to **port** the operating system to new architectures
  - More **reliable** (less code is running in kernel mode)
  - More **secure**
- Detriments:
  - **Performance overhead** of user space to kernel space communication.

# Microkernels

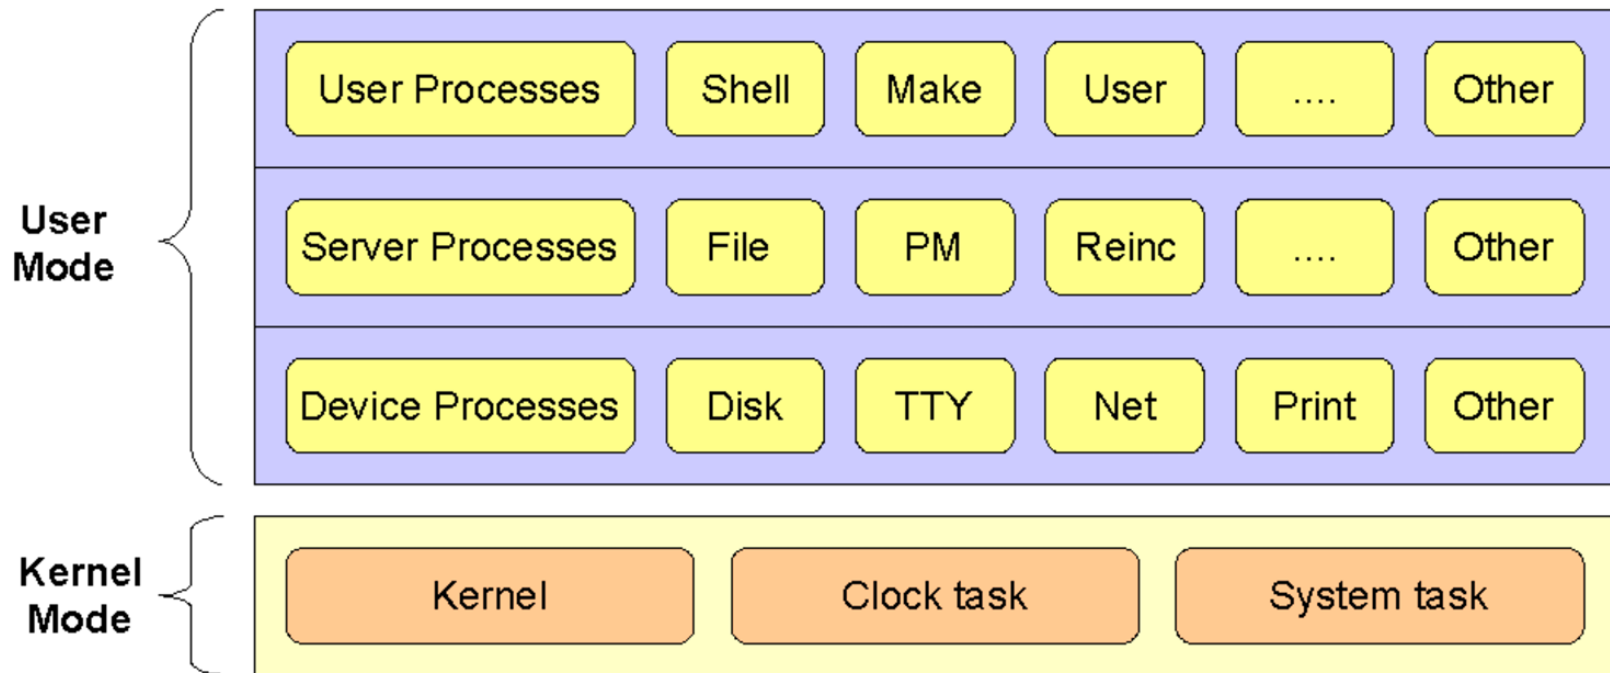


Architecture of a typical microkernel OS

# MINIX3 OS architecture

## a microkernel operating system

---




**The MINIX 3 Microkernel Architecture**

from wikipedia

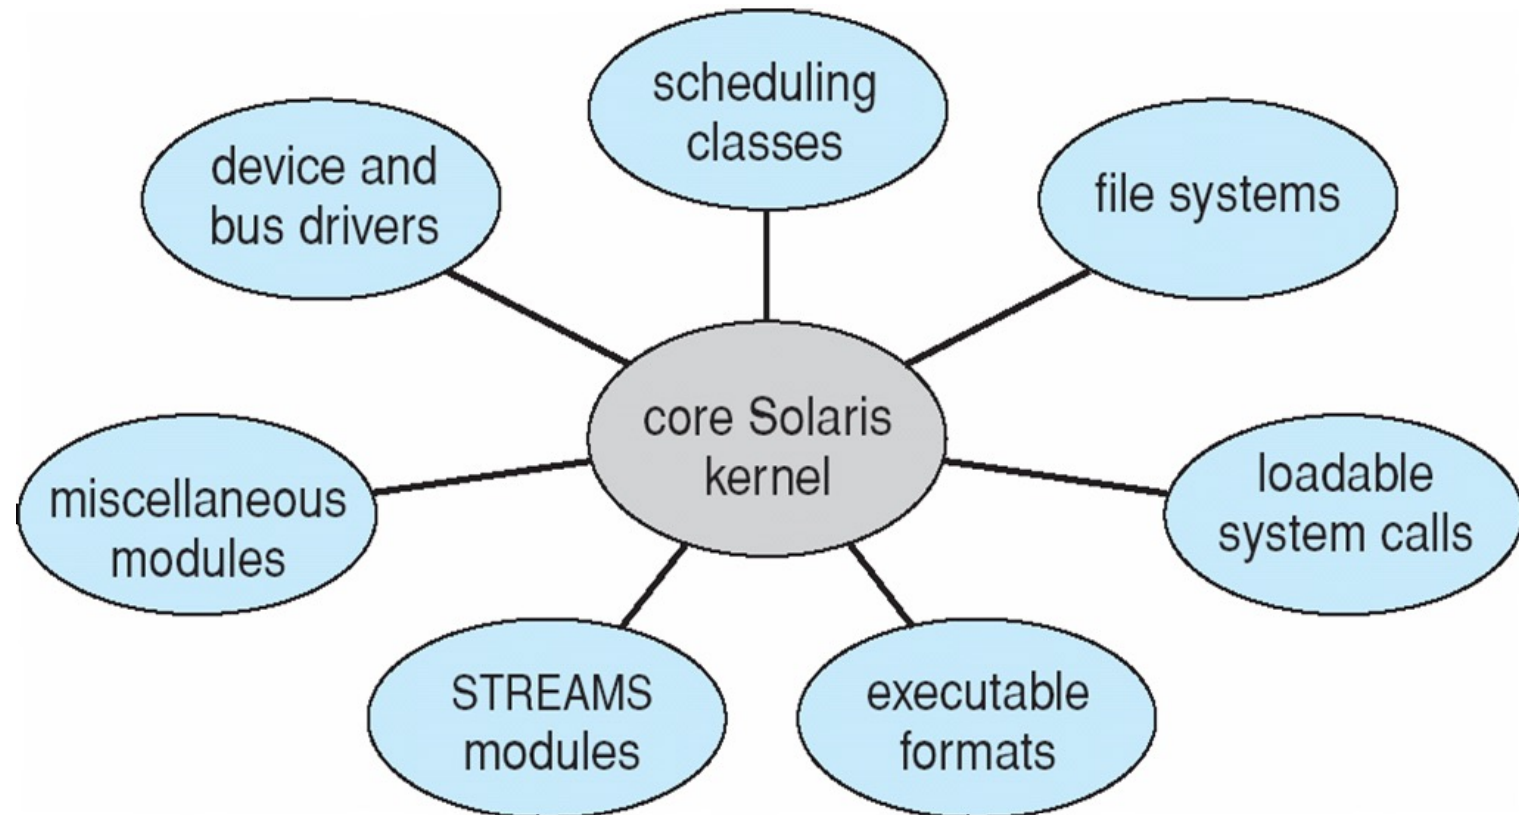
# Modules

- Most modern operating systems implement **kernel modules**. In this way they improve their monolithic nature.
  - Uses object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces (function calls in the same address space)
  - Each is **loadable** as needed within the kernel
- Overall, similar to layers but more flexible.
- Linux supports **modules**.
- You can insert modules.
- You can remove modules.



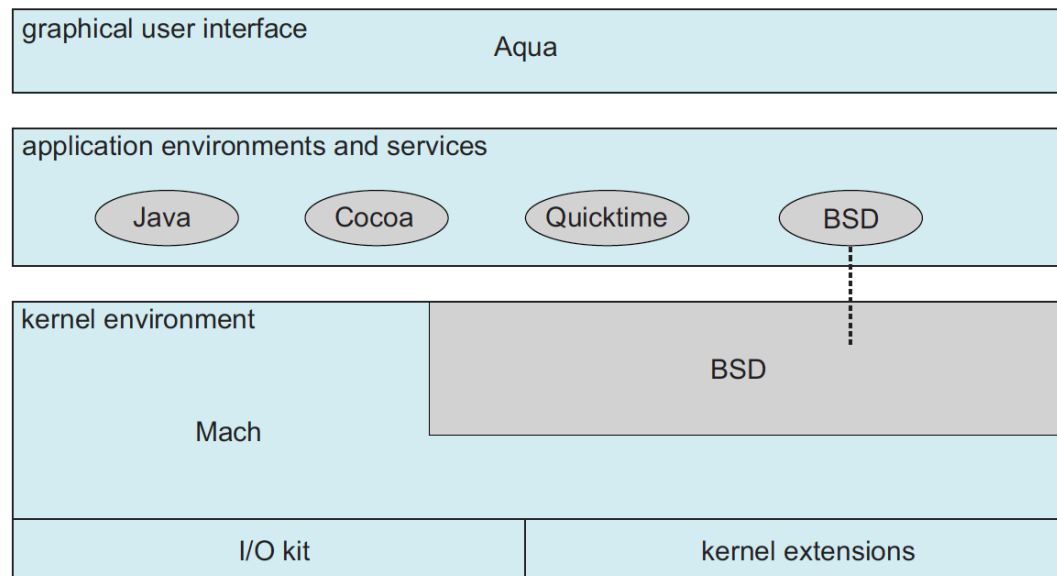
# Solaris Modular Approach

---



# Hybrid Systems

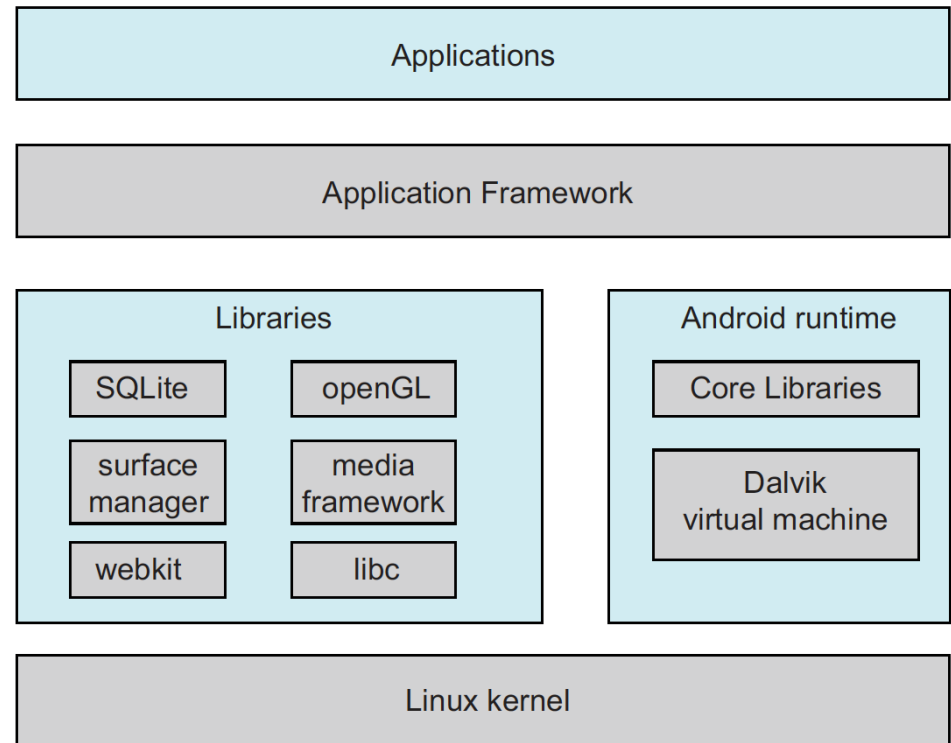
- Mac OS X
  - Uses Mach and BSD kernels
  - Mach: microkernel
  - BSD: monolithic



# Hybrid Systems

---

- Android
  - Uses **Linux kernel**
  - **Layered stack of software** on Linux kernel.
- A rich **set of frameworks** to develop and run Android applications.



# Virtualization

---

- Virtualization abstracts/virtualizes the hardware so that OSs and programs can run on a **virtualized machine** or a **virtual environment**
- Many different **types of virtualizations**:
  - **Virtual machines**: run OS/Apps on a VM that is identical to the **bare hardware**
    - For example: run Linux OS compiled for Intel x86 on a virtual machine duplicating x86 hardware.
  - **Emulation**: run OS or App developed for a machine that is **different** than the bare hardware
    - For example: run app developed for MIPS on Intel x86 machine using an emulator (MIPS emulation on Intel x86).
  - **Abstract machines**: run applications developed/compiled for an **abstract machine** running on bare hardware (interpretation)
    - For example JVM. Run a Java app on JVM running on x86

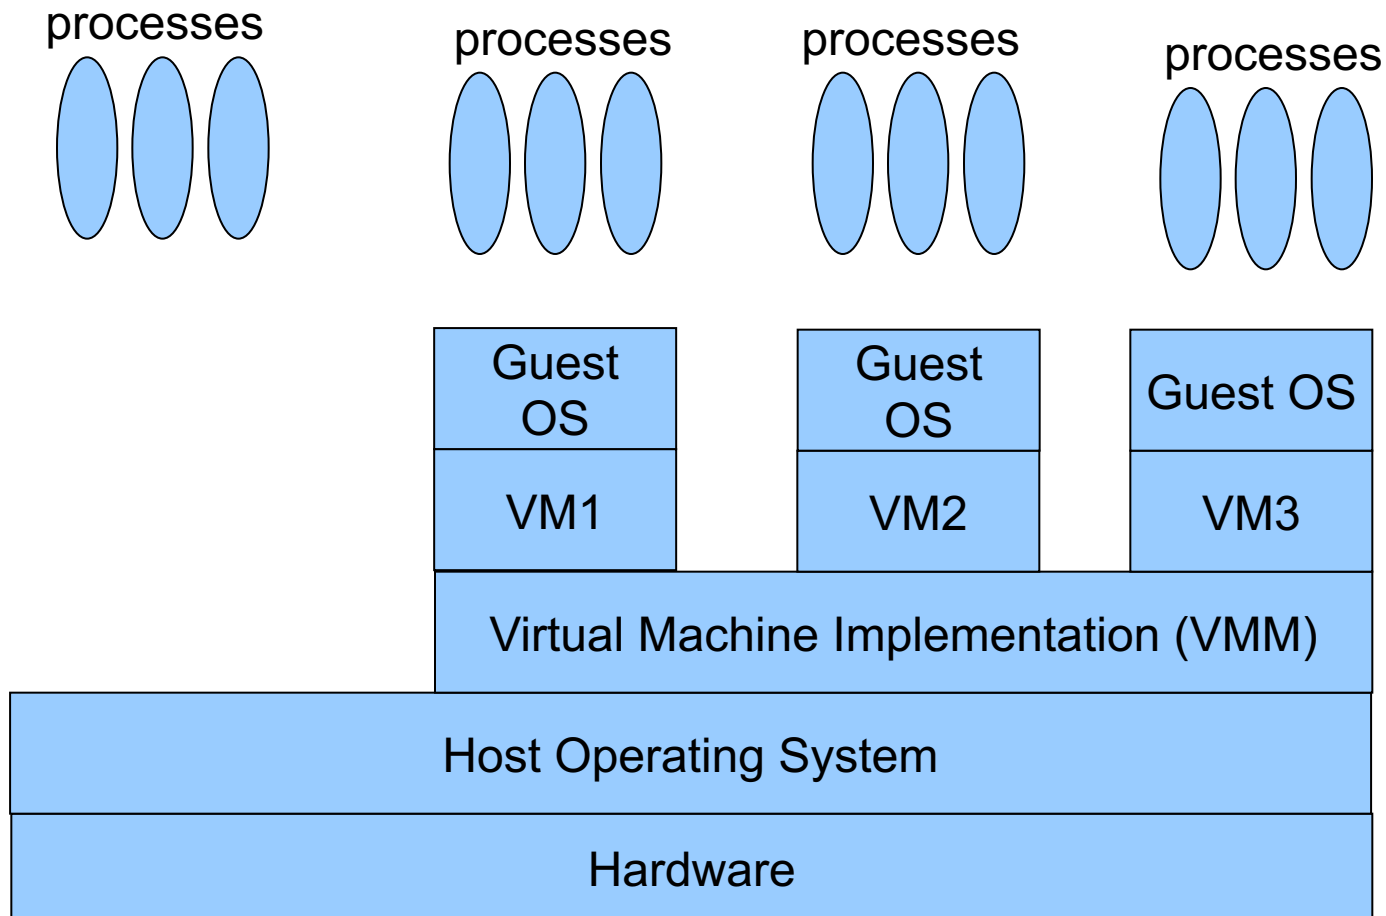


# Virtual Machines

---

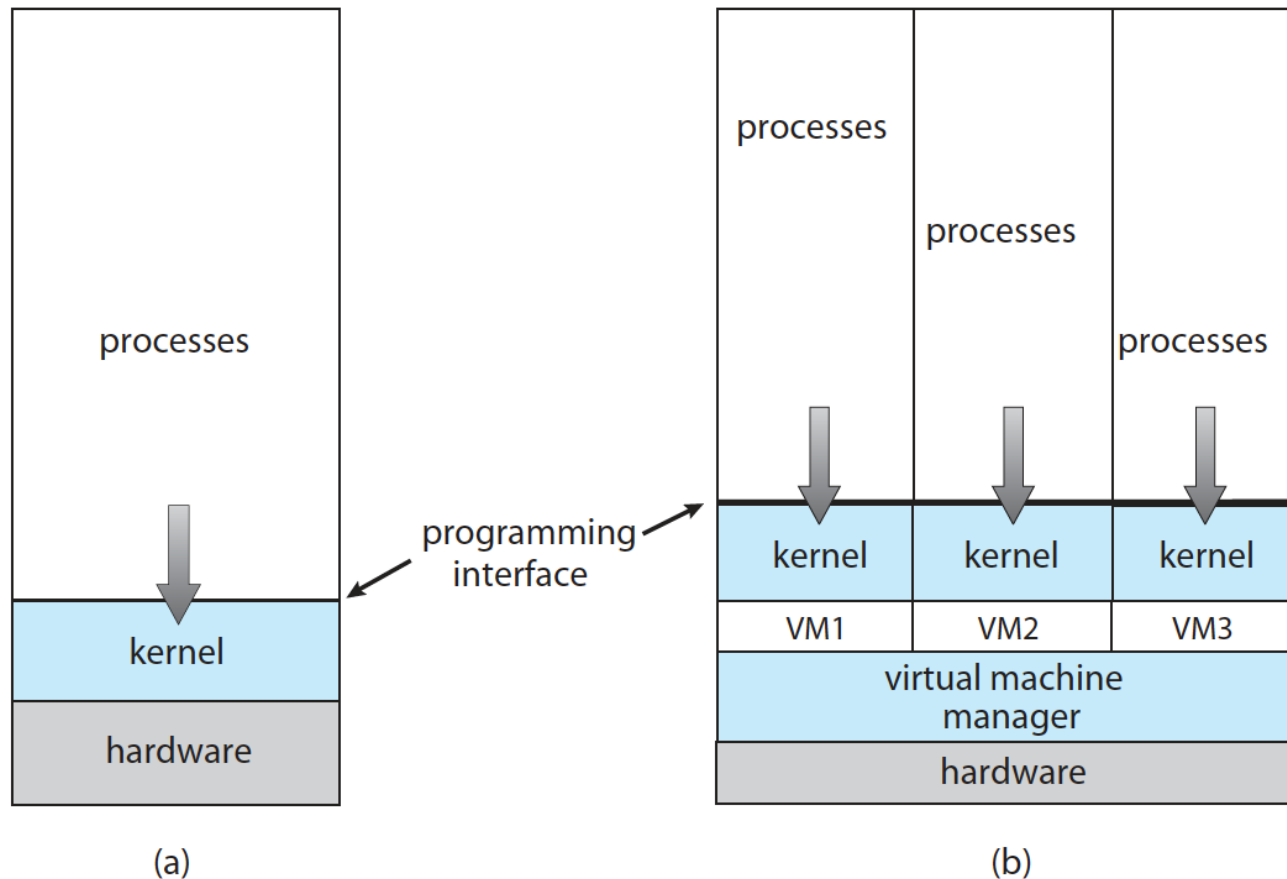
- Hardware is abstracted into **several different execution environments**.
  - Virtual machines.
- Each virtual machine provides an interface that is **identical to the bare hardware**.
- A *guest* process/kernel can run on top of a virtual machine.
  - We can run several operating systems on the same *host*.
  - Each virtual machine will run another operating system (guest).

# Virtual Machines



- 1) VMM is a *hosted virtual machine* (running on a host OS) – Type-2 Hypervisor  
Example: VMware Workstation

# Virtual Machines



- 2) VMM is *running directly on hardware* like an OS: **Type-1 Hypervisor**  
Example: Wmware ESX Hypervisor

# Virtualization

---

- Vmware (Virtual Machine Manager– VMM)
  - Abstracts Intel x86 hardware
- Java virtual machine (JVM)
  - Specification of an **abstract computer**
- .NET Framework
  - Specification of an abstract computer again

# Operating System Debugging

---

- Failure analysis
  - Log files
  - Core dump
  - Crash dump
- Performance tuning
  - Monitor system performance
    - Add code to kernel
    - Use system tools like “top”
- DTrace
  - Facility for dynamically adding probes to a running system (both to processes and to the kernel)
  - Probes can be queried using D programming language to obtain info

# Operating System Generation

---

- Configure the kernel
- Compile the kernel

# System Boot

---

- Bootstrap program (loader) locates the kernel, loads it and starts the kernel.
- This can be a two-step procedure.
  - **Bootstrap** program loads another more **complex boot** program
  - That boot program loads the **kernel**
- Then control is given to the kernel.
- Kernel starts the environment and makes the computer ready to interact with the user (via a GUI or command shell).
- Details depend on the system.

# Performance Monitoring and Tuning

- We may need to **monitor system performance**.
- There are **tools** for that
- In Linux, for example,
  - ps
  - top
  - vmstat (memory usage statistics)
  - netstat (network usage statistics)
  - iostat (I/O device usage statistics)
- In Linux, **/proc file system** gives kernel information (statistics)
  - cd /proc
    - we see lots of directories
    - one **directory per process** (named with process id – **pid** )

1	111	122	137	156	165	22582	3062	360	41	527	62	857	asound	meminfo
10	1115	123	138	1562	166	22833	3063	361	4134	529	63	858	buddyinfo	misc
100	11183	1233	1389	1564	167	23	3064	362	425	53	64	859	bus	modules
10028	112	1237	1394	1567	168	23456	3065	364	43	54	6426	86	cgroups	mounts
101	1128	124	14	157	169	23716	3066	365	434	547	65	87	cmdline	mtrr
102	113	125	140	1576	17	23963	3067	366	44	556	66	88	consoles	net
1035	11317	1253	1400	1580	170	245	3068	367	445	559	6742	8819	cpuinfo	pagetypeinfo
104	11328	1254	1406	1581	171	24510	3069	368	446	56	68	89	crypto	partitions
105	114	126	141	1584	172	246	3070	369	45	580	6869	9	devices	sched_debug
10547	11447	1261	1418	1587	173	248	3071	37	458	563	69	90	diskstats	schedstat
10552	11451	12663	142	1589	174	249	3072	370	454	5648	691	9044	dma	scsi
106	1152	128	1422	1590	18238	25	3073	371	457	566	6963	92	driver	self
1064	116	129	143	1592	18245	259	3074	372	46	567	7	9253	execdomains	slabinfo
1065	11670	13	144	1594	183	251	3075	373	466	57	70	9260	fb	softirqs
1069	117	130	146	1595	18381	25126	3076	374	467	571	7002	9262	filesystems	stat
107	11731	13047	147	1599	186	257	3077	375	469	572	71	9263	fs	swaps
1071	11739	13049	148	16	19	26	3078	376	47	573	7129	9264	interrupts	sys
1072	11741	131	149	160	191	267	3079	377	477	575	7173	9265	ionem	sysrq-trigger
1074	11742	13105	1490	1600	196	27	3080	378	485	576	72	9266	ioparts	sysvipc
1078	118	13108	15	1601	198	27688	3081	38	49	577	74	9276	irq	thread-self
108	11839	13113	150	1602	199	27695	3082	387	498	58	75	9277	kallsyms	timer_list
1080	11840	1312	1511	1604	2	27718	3083	388	499	580	76	9278	kcrc	tty
1082	119	13133	1516	1606	20	28	31	389	5	581	77	93	keys	uptime
1088	12	1314	1518	161	200	284	31672	39	51	582	78	9356	key-users	version
1090	120	13149	152	1611	201	29	32	393	510	584	8	9367	kmsg	version_signature
1092	1204	132	1522	1614	202	3046	33	396	513	585	80	94	kpagecgroup	vmlallocinfo
1095	1205	1325	1523	1616	203	3047	34	397	515	59	81	95	kpagecount	vmstat
1097	12054	134	153	162	21	3058	35	4	516	590	82	96	kpageflags	zoneinfo
1099	12169	1344	154	163	21786	3059	357	40	518	597	8248	98	loadavg	
11	12170	135	155	1635	22	3060	358	4006	52	60	83	99	locks	
110	12189	136	1559	164	22368	3061	359	4032	526	603	84	acpi	mdstat	



# Summary

---

- **System calls** are the programming interface to kernel (used by programs)
- **System programs** are useful for users (a nice environment for program development and running)
- Different approaches to **structure kernel**
  - monolithic
  - layered
  - microkernel
- **Virtualization** is important
  - Has various forms.

# References

---

- Operating System Concepts, Silberschatz et al.
- Modern Operating Systems, Andrew S. Tanenbaum et al.
- Computer Systems, A Programmer's Perspective, 3<sup>rd</sup> edition, R. E. Bryant and D. R. O'Hallaron, Pearson, 2016.

---

# **Additional Study Material**