



Bilkent University
Department of Computer Engineering
CS342 Operating Systems

Threads

Last Update: Oct 3, 2022

Outline and Objectives

Outline

- Overview
- Multithreading Models
- Thread Libraries
- Threading Issues
- Operating System Examples
- Windows XP Threads
- Linux Threads
- Re-entrancy
- Thread specific data

Objectives

- To introduce the notion of a thread — a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- To discuss the APIs for the Pthreads, Win32, and Java thread libraries
- To examine issues related to multithreaded programming

Threading Concept


Threads and thread usage

- A **process** has normally a single thread of control; that means a single stream (flow) of instructions executed sequentially.
- A process always has **at least one thread** (default thread or main thread).
- In a single-threaded process, if the thread (process) **blocks**, no activity can be done as part of the process.
 - Better would be being able to **run multiple concurrent threads** (tasks) as part of the same process.
- With **multithreading** we have that.
 - A process can have multiple threads running **concurrently**.
 - Each thread runs in a *sequential* manner.
 - Threads **share code** (instructions) and **data** (global variables) of the process.


Benefits of using multiple threads in an application (process)

- **Responsiveness**
 - If one thread blocks, another can run.
 - One thread may always wait for the user, for example.
- **Easy resource/data sharing**
 - Threads can easily share resources, global variables.
- **Economy**
 - Creating a thread is faster than creating a process.
 - Context switching among threads may be faster.
- **Scalability**
 - Multiple CPUs in a computer can be utilized by a process. Each thread running on another available CPU.


Threads and thread usage



CPU



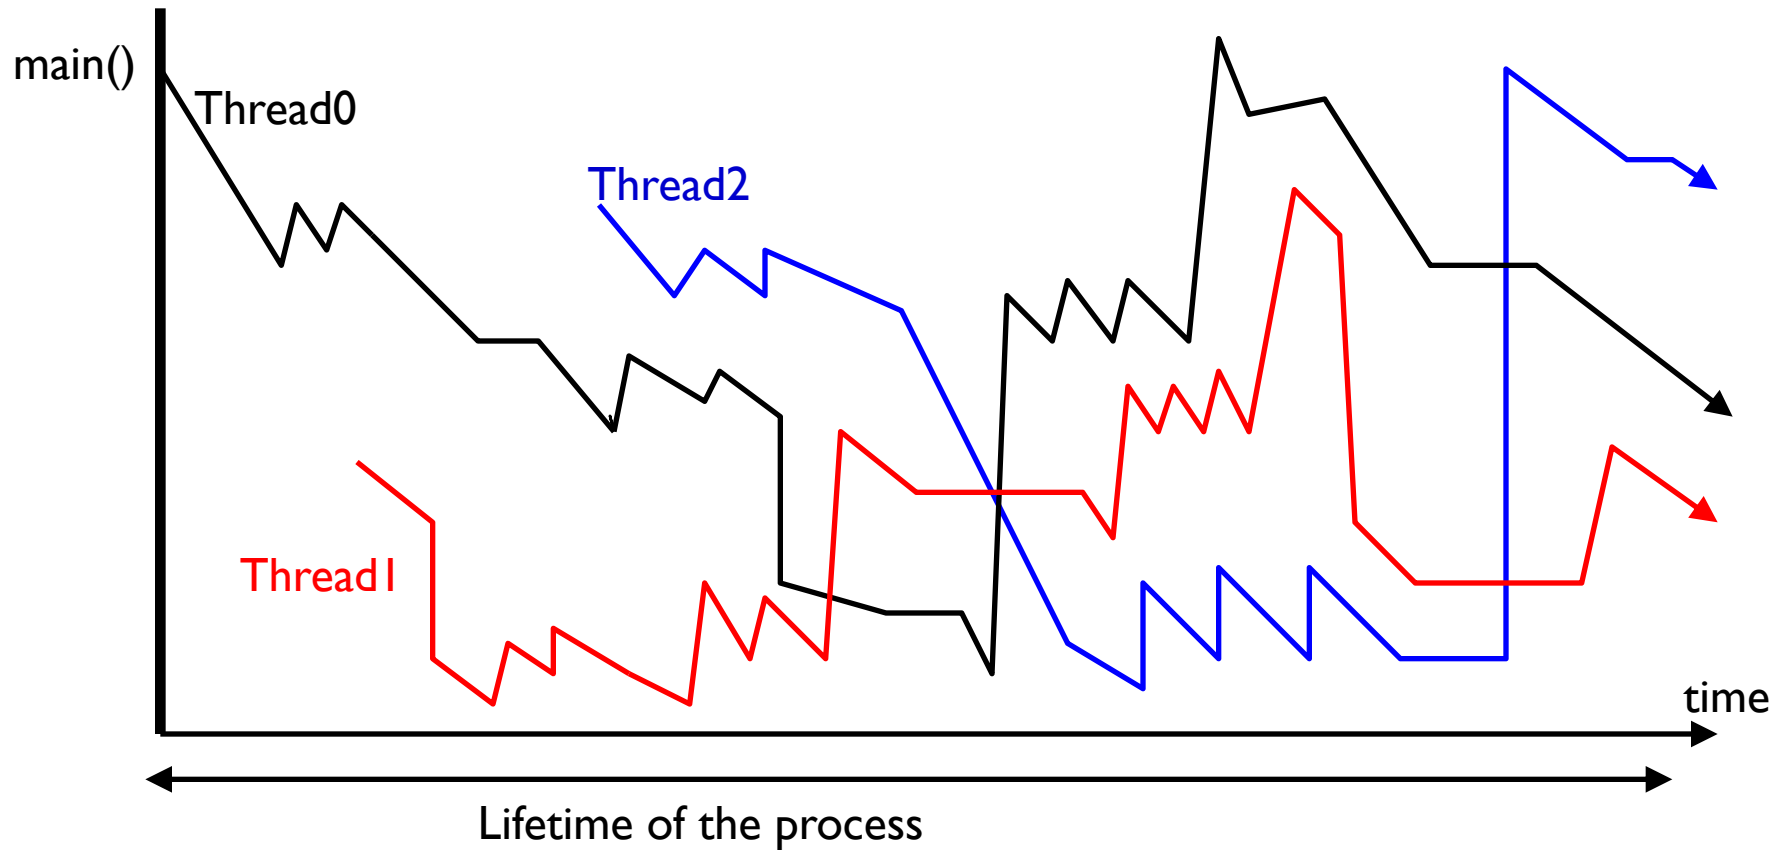
single-threaded process




multi-threaded process

a multithreaded process' execution flows: threads


Instructions of the Program




Multithreading concept



CPU




single-threaded process



multi-threaded process

Multithreading concept



Schedulable Entities


We can select one of them and run

Multithreading concept

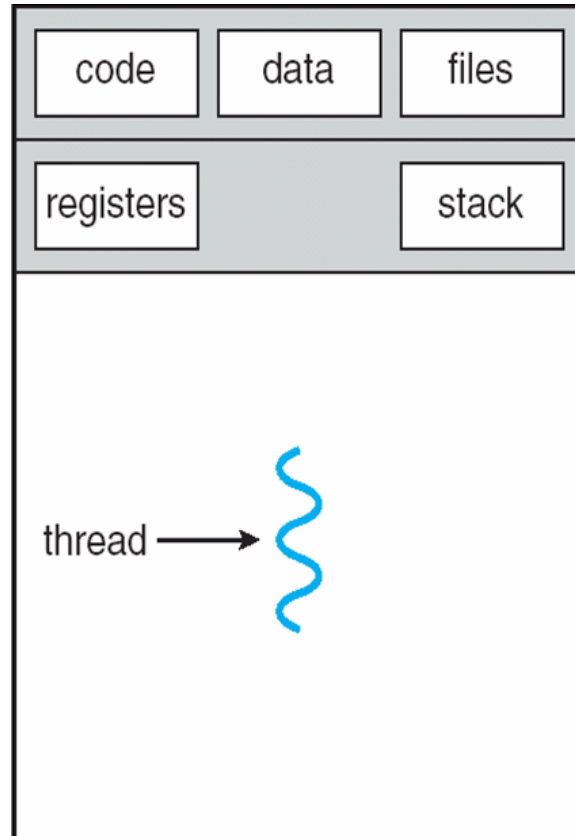
```
function1(...)
{
    ....
}

function2(...)
{
    ....
}

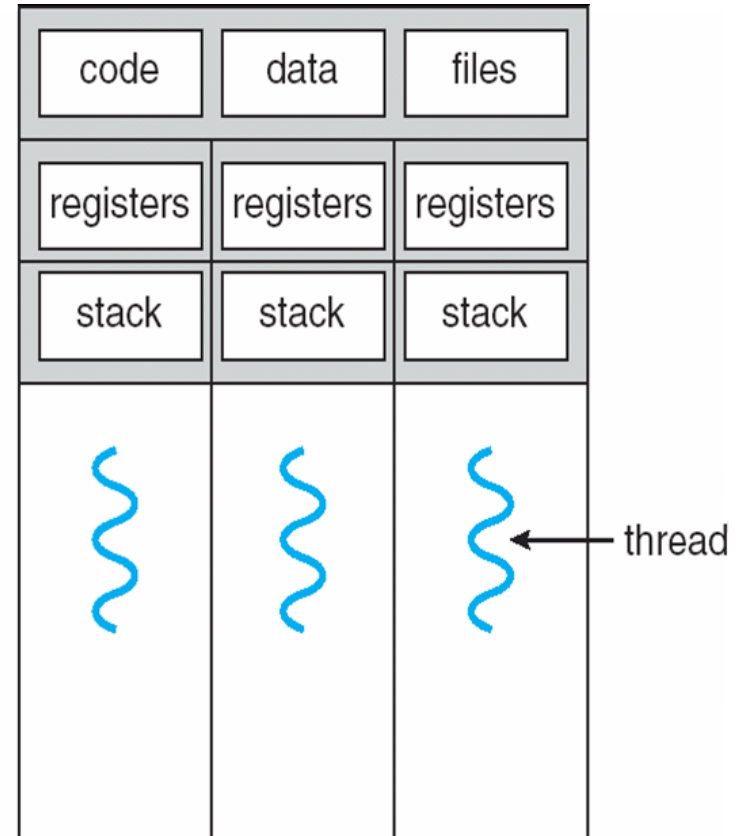
main()
{
    ....
    thread_create (function1 ,...);
    ....
    thread_create (function2, ...);
    ....
    thread_create (function1, ...);
    ....
}
```



Single and multithreaded processes



single-threaded process



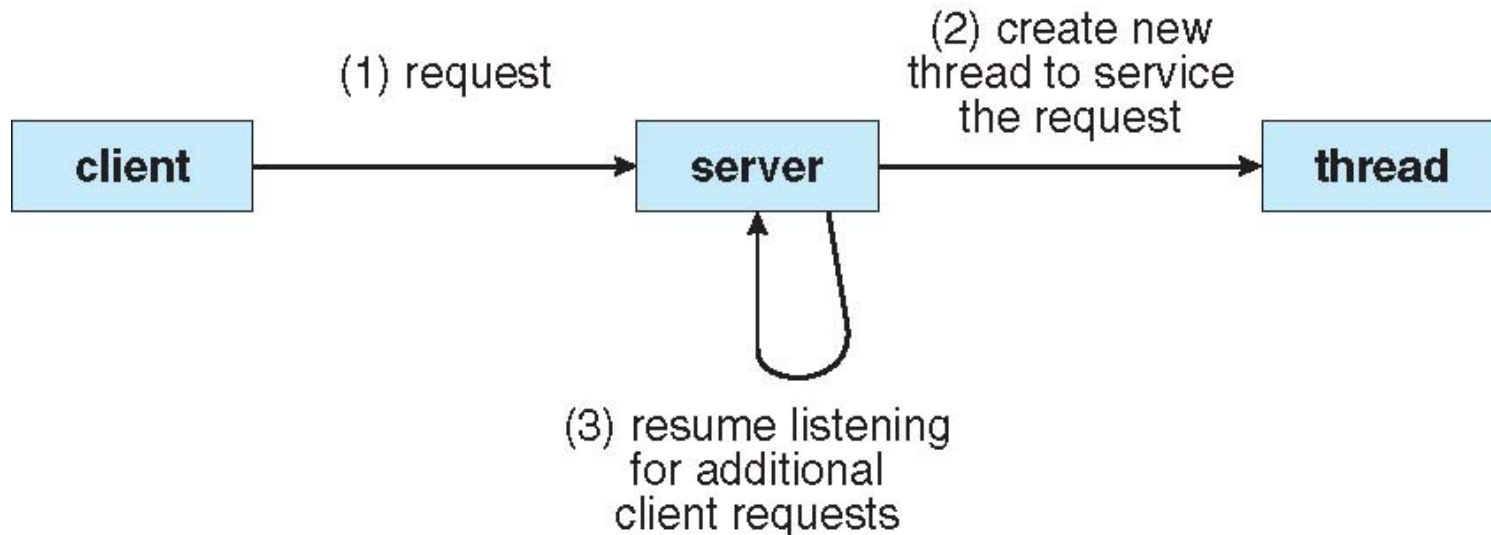
multithreaded process

Multicore programming and multithreading challenges

Threading **challenges** include

- **Dividing activities**: come up with concurrent tasks
- **Balance**: tasks should be of similar importance and load
- **Data splitting**: data may need to be split as well
- **Data dependency**: data dependencies should be considered; need synchronization of activities.
- **Testing and debugging**: Debugging a multithreaded program is more difficult.
 - Programmer should explicitly care about coordinated access to **shared** variables.


Multithreaded Server Architecture



Using threads to implement a server is common.


- Provides better performance (multicore machines).
- Natural structuring of the server.
- Ease of programming. A thread executes sequentially.

Concurrent Execution on a Single-core System



- In a single CPU computer, multiple threads will **share** the same CPU.
- CPU switched among the threads frequently.
- We don't have physical parallelism, but we have **concurrency** (pseudo-parallelism).

Parallel Execution on a Multicore System



- If there are multiple CPUs, some threads can run in parallel.
- All threads will run concurrently.

Speedup: Amdahl's Law

- What is the potential performance gain (**speed up**) with multiple CPUs in system?
- A program may have both **serial** and **parallelizable** parts.
 - Serial part may not be parallelized.
- Assume the program executes and finishes in 1 time unit in a single CPU system. Then:
 - S : fraction of time to execute the serial portion ($0 \leq S \leq 1$) in a single-CPU system.
 - $1 - S$: is the fraction of time to execute the parallelizable portion in a single-CPU system.

Speedup: Amdahl's Law

- N : number of processing cores.
 - N tasks/threads can run in parallel.
 - Parallelizable portion of the program can finish N times faster (at most).
- We define **speed-up** as the ratio of program completion time in a single-CPU system to the completion time in an N -CPU system.
- Then we have the following formula for the speed-up (called Amdahl's law).

$$\text{speedup} \leq \frac{1}{S + \frac{1-S}{N}}$$

Threading Support

Threading Support


- Multithreading can be supported by:
 - **User level libraries** (without Kernel being aware of it): library creates and manages threads (user-space implementation)
 - **Kernel itself**: kernel creates and manages threads (kernel-space implementation).
- No matter which method is used, threads can be created, used, and terminated by a program via a set of functions that are part of a **Thread API** (a thread library)
- Three are primary thread libraries: POSIX threads, Java threads, Win32 threads (for Windows 32 bit systems).

Multithreading Models

- A user process creates one or more threads (N threads) by making thread API calls.
- Kernel can create one or more threads (M threads) for the user process. Kernel creates at least one thread (the process itself).
- A relationship must exist between user threads and kernel thread(s).
 - Mapping user-level threads to kernel-level threads.
- Three common ways of establishing such a relationship.
 - Many-to-One model.
 - One-to-One model.
 - Many-to-Many model.

Many-to-One Model: Implementing Threads in User Space


- Many user-level threads mapped to a single kernel thread
- Examples:
 - Solaris Green Threads.
 - GNU Portable Threads.
- Thread management done at user space, by a thread library.



Kernel supports process concept;
not threading concept

Many-to-One Model: Implementing Threads in User Space


- No need for kernel support for multithreading (+)
- Thread creation is fast (+)
- Switching between threads is fast (+)
- A blocking system call blocks the whole process (-)
- A thread has to explicitly call a function to give the CPU to some other thread. Example: `thread_yield()`.
- All threads will run in a single CPU. Will not utilize multiple CPUs. (-)



One-to-One Model:

Implementing Threads in Kernel Space


- Kernel implements threading and can manage threads, schedule threads.
- Examples (nearly all modern OSs): Windows, Linux, ...
- All these kernels have threading support. They schedule processes and also threads.
- Each user-level thread maps to a single kernel thread



One-to-One Model:

Implementing Threads in Kernel Space

- Provides more concurrency; when a thread blocks, another can run. Blocking system calls are not a problem anymore. (+)
- Multiple processors can be utilized (+).
- Kernel can stop a long running thread and run another thread. No need for explicit request from a thread to be suspended. (+)
- Need system calls to create and control threads. Making system call takes more time than a library call. (-)
- Thread switching is costly; (-)



Threading API

Thread Libraries

- Thread library provides programmers with an API for creating and controlling threads
- Programmer just have to know the thread library interface (API).
- Threads may be implemented in user space or kernel space.
- Thread API library may be entirely in user space or may get kernel support for threading.

Pthreads Library

- **Pthreads**: POSIX (Portable Operating System Interface) threads
- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization.
- **API** specifies **behavior** of the thread library, implementation is up to development of the library.
- Common in UNIX-like operating systems
 - Solaris, Linux, Mac OS X.

Pthreads Example

- We will show a program that creates a new thread, and in this way will have two threads:
 - the initial/main thread (i.e, the process)
 - the new thread (both threads have equal power)
- The program will just create the new thread to do a simple computation: will sum all integers up to a given parameter value N

$$sum = 1 + 2 + \dots + N = \sum_{i=1}^N i$$

- The main thread will wait until sum is computed into a global variable.
- Then the main thread will print the result.

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum;          /* shared sum by threads - global variable */
void *runner (void *param); /* thread start function */
```

Pthreads Example

```
int main(int argc, char *argv[]){
    pthread_t tid;          /* id of the created thread */
    pthread_attr_t attr;    /* set of thread attributes */

    if (argc != 2) {
        fprintf (stderr, "usage: a.out <value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf (stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }

    pthread_attr_init (&attr);
    pthread_create (&tid, &attr, runner, argv[1]);
    pthread_join (tid, NULL);
    printf ("sum = %d\n", sum);
}
```

Pthreads Example

```
void *runner (void *param)
{
    int i;

    int upper;

    upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; ++i)
        sum += i;

    pthread_exit(0);
}
```


Pthreads Example

```
int main(...)
{
    ...
    ...
    pthread_create(&tid, ..., runner, ...);

    pthread_join(tid);

    printf (... , sum, ...);
}

runner (...)
{
    ...
    sum = ...
    pthread_exit();
}
```



Compiling and running the program

- We can put the above code into a `.c` file, say `mysum.c`
- To use Pthreads functions, we need to include the **header file** `pthread.h` in our program.
- We also need to link with the **pthread library** (the Pthreads API functions are not implemented in the standard C library). The way to do that is using the `-l` option of the C compiler. After `-l` you can provide a library name like `pthread`.
- We can compile+link our program as follows:
`gcc -Wall -o mysum -lpthread mysum.c`
- Then we run it as:
`./mysum 6`
- It will print out 21

Windows Threads

```
/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}
```

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */
```

Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Java threads may be created by
 - Extending the Thread class
 - Implementing the Runnable interface
 - Example given in the book

From Single-threaded to Multithreaded

- Many programs are written as a single threaded process.
- If we want to convert a single-threaded program to a multi-threaded program, we need to be careful about the following:
 - the global variables
 - the library functions we use

From Singlethread to Multithreaded

```
int status;

func1(...) {
    ....
    status = ...
    do_something_based_on(status);
}

func2(...) {
    ...
    status = ...
    do_something_based_on(status);
}

main() {
    ....
    thread_create(..., func1, ...);
    thread_create(..., func2, ...);
}
```

- We can have problem here.
- Just after func1 of thread 1 updated status, a thread switch may occur and 2nd thread can run and update status.
- Then thread 1 will run again, but will work with a different status value.
Wrong result!

Thread-safe / Reentrant libraries

- Many library procedures may not be **reentrant**.
 - They are not designed to have a second call to itself from the same process before it is completed (not re-entrant).
 - They may be using global variables. Hence may not be thread-safe.
- We have to be sure that we use **thread-safe** (reentrant) library routines in multi-threaded programs we are developing.

From Single-threaded to Multithreaded

- Scope of variables:
 - Normally we have: global, local
 - With threads, we want: global, local, and thread-specific (thread wide)
- **thread-specific**: global inside the thread (**thread-wide global**), but not global for the whole process. Other threads can not access it. *But all functions of the thread can access.*
- But we may not have language support to define such variables.
 - C can not do that. Normally in C we have just global variables and local variables. We do not have thread-wide (thread-local or thread-global) variables.
- Therefore thread API has special functions that can be used to create such variables – data. This is called thread specific data.

Thread Specific Data (TSD)

- Allows each thread to have its **own copy of data**.
- Each thread refers to the data with the **same name (key)**.

Thread Specific Data

- In POSIX Pthreads Library we have the following functions to create thread specific data that has an associated name (key).

`pthread_key_create()`

`pthread_key_delete()`

`pthread_setspecific()`

`pthread_getspecific()`

- key is seen by all threads.
- in each thread, key is associated with thread specific data.

Example:

```
pthread_key_t key1;  
pthread_key_create (&key1, ...);
```

do this once


// In each thread we can do the following:

```
char *buf = malloc(1024);
```

```
pthread_set_specific (key1, buf); // associate buf (buffer pointer) with key1
```

```
buf = pthread_get_specific (key1); // get pointer to access data
```

Thread Specific Data



Each thread is using the same variable name (i.e., key) to get access to *its own specific data*.

Examples from Operating Systems

Operating System Examples

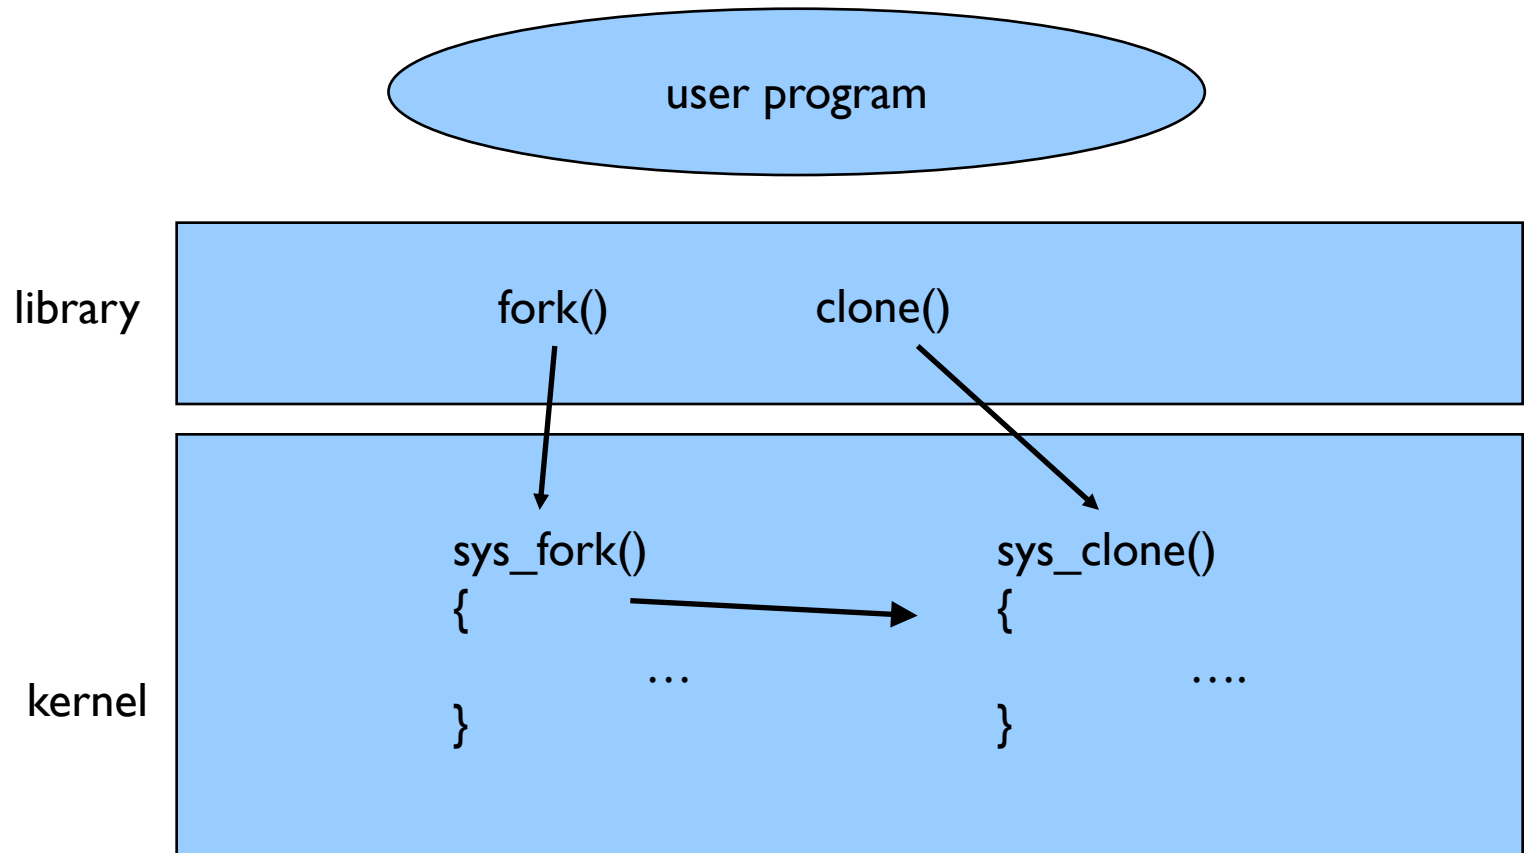
- Windows XP Threads
- Linux Threads

Linux Threads

- Linux refers to them as *tasks* rather than *threads*.
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process).

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.


Clone() and fork()



Many-to-Many Model & Two-level Model


Many-to-Many Model

- Allows **many user level threads** to be mapped to **many kernel threads**
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the *ThreadFiber* package



Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to a kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier




Implicit Threading

Implicit Threading

- Applications can create hundreds or thousands of threads.
 - Can be creating *explicitly*. (App programmer worries)
 - Can be created by compiler, library, platform, etc. (easier for App programmer) - called *Implicit* Threading.
- Some alternatives for implicit threading
 - Thread Pools
 - OpenMP (Compiler Supported)
 - Grand Central Dispatch (in Mac OS X)

Thread Pools

- Create a number of threads in a pool where they wait for work.
- Advantages: 1) faster to start a thread. 2) enables limiting the number of threads application is using – pool size)
- Incoming task requests are queued.
- An available thread takes a request and processes it.
- When done, the busy thread becomes available again (waits for work).



OpenMP

- Is a set of compiler directives as well as an API for programs in C and C++.
- Parallel programming in shared memory environments.
- Parallel regions identified by programmer.
- Examples:

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```

→ Number of threads ==
Number of cores

```
#pragma omp parallel for
for (i = 0; i < N; i++) {
    c[i] = a[i] + b[i];
}
```

Grand Central Dispatch (GCD)

- A combination of extensions to the C language, and API, and a run-time library that allows app developers to identify sections of code to run in parallel.
- Programmer identifies the region to be parallel:

```
{ printf("I am a block"); }
```

- GCD schedules such blocks for run-time execution by placing them in a dispatch queue (serial or concurrent).
- Blocks can run in parallel in different cores.

Other Threading Issues

Threading Issues

- Semantics of `fork()` and `exec()` system calls
- Thread cancellation of target thread
 - Asynchronous or deferred
- Signal handling
- Thread pools
- Thread-specific data
- Scheduler activations

Semantics of fork() and exec()

- Does fork() duplicate only the calling thread or all threads?
- How should we implement fork?
- logical thing to do is:
 - 1) If `exec()` will **be called** after fork(), there is no need to duplicate the threads. They will be replaced anyway.
 - 2) If `exec()` will **not be called**, then it is logical to **duplicate** the threads as well; so that the child will have as many threads as the parent has.
- So we may implement two system calls: like fork1 and fork2!

Thread Cancellation

- Terminating a thread before it has finished
 - Can be asynchronous or deferred.
- **Asynchronous cancellation** terminates the target thread immediately
- **Deferred cancellation** allows the *target thread* to periodically check if it should be cancelled
 - Canceller thread indicates a target thread to be cancelled.
 - Target threads performs checks at cancellation points and if it is safe gets terminated.

Thread Cancellation

```
pthread_t tid;                                Cancelling thread

pthread_create (&tid, 0, worker, NULL);
...
/* cancel the thread */
pthread_cancel(tid); /* requesting cancellation */
```

cancellation modes

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

Deferred cancellation

```
while (1) {                                    Target thread
    /* do some work */
    ...
    /* check if there is cancel request */
    pthread_testcancel();
}
```

mode can be set with
`pthread_setcancelstate()`

Signal Handling

- If a **signal** is sent to a Multithreaded Process, **who will receive** and handle that?
- In a single threaded process, it is obvious.
- In a multi-threaded process, the options are:
 - deliver signal thread to which signal applies
 - deliver signal to all threads of the process
 - deliver signal a specific thread responsible for handling signals

Signal Handling

- **Signals** are used in **Unix** systems to notify a process that a particular event has occurred
 - a signal is generated by a particular event or a process
 - A signal is **delivered** to a process (same or different process)
 - a signal is handled
- A **signal handler** is used to process signals
 - In kernel (default handler)
 - In process (if process specified a handler)
- Handled **asynchronously**

a C program using signals

```
#include <stdio.h>
#include <signal.h>
#include <stdlib.h>

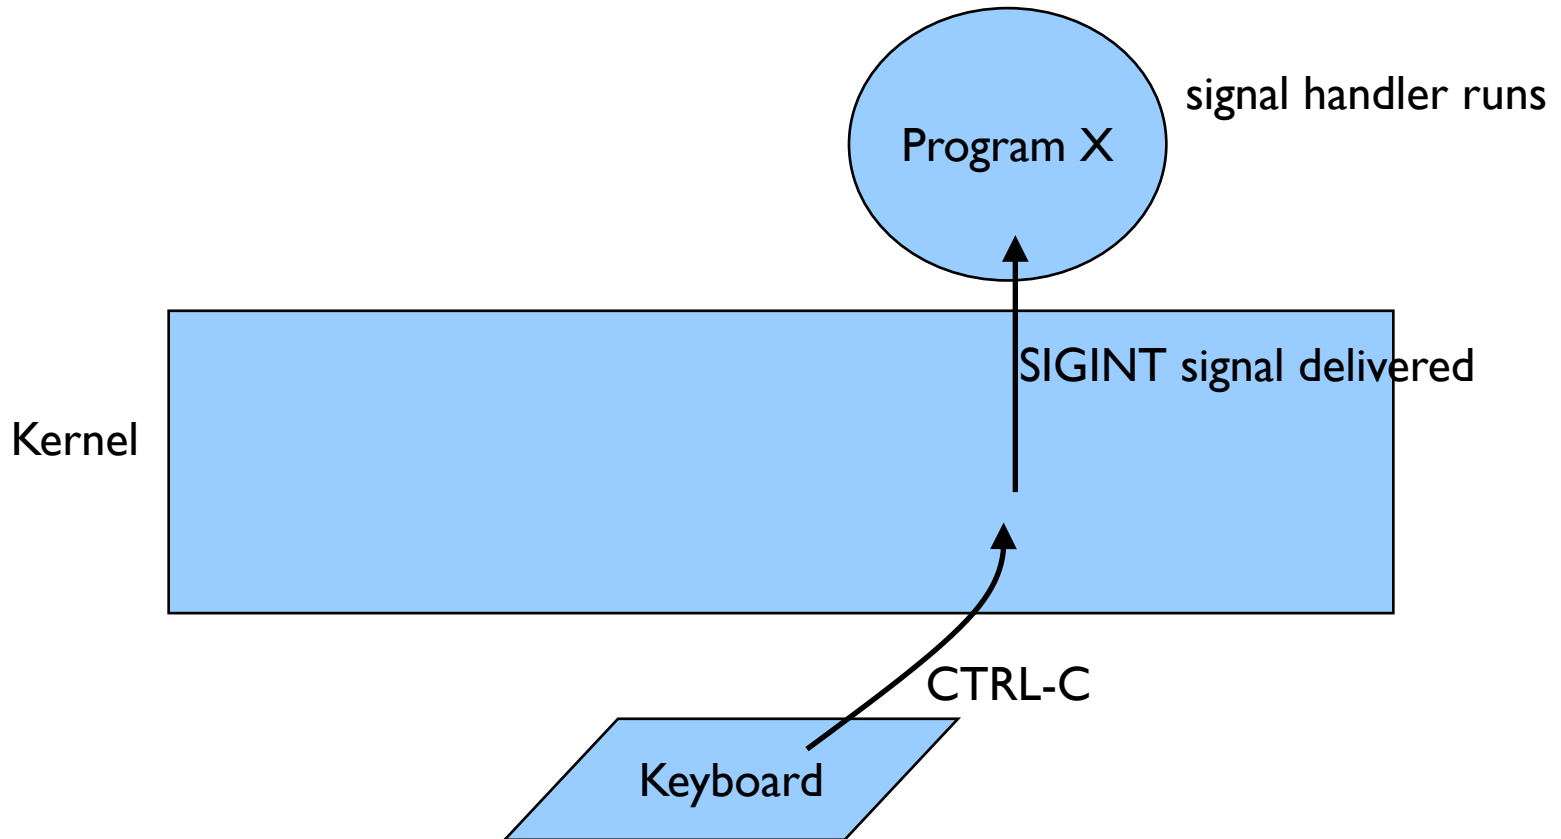
static void sig_int_handler() {
    printf("received SIGINT signal.\n");
    fflush(stdout);
    exit(0);
}

int main() {
    signal(SIGINT, sig_int_handler);

    while (1)
        ;
}
```

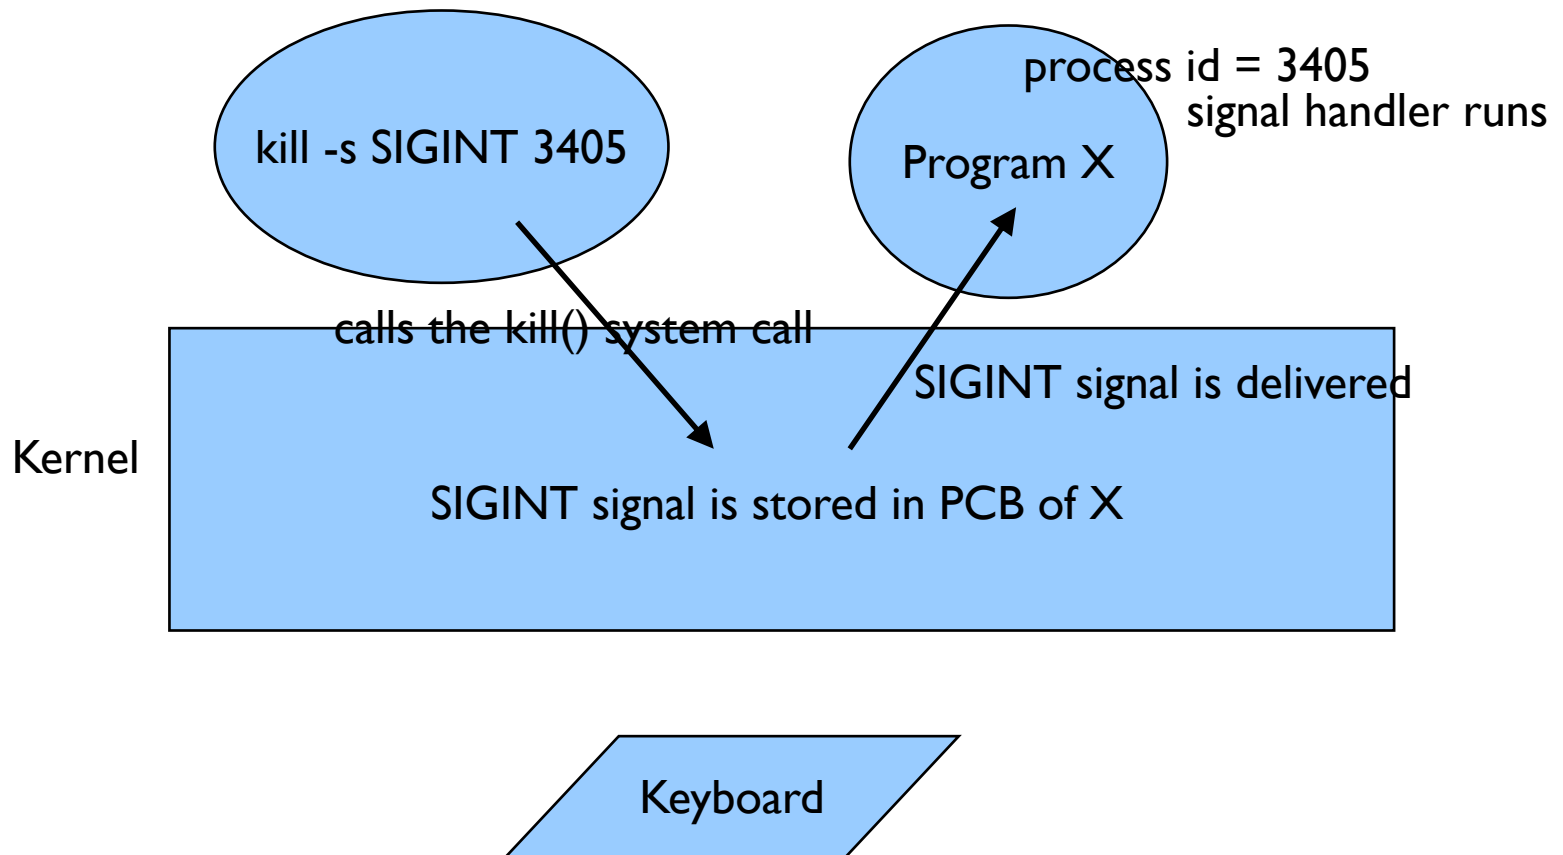
- While a program is running, if we press **CTRL-C**, we are sending a **SIGINT** signal to the program.
- By default, SIGINT is handled by kernel. Kernel terminates the program.
- But if we specify a handler function (shown left), then our program can handle it.
- Kernel will notify our process with a signal when user presses CTRL-C.

delivering signal (notifying)



Sending a signal

In Unix, we have the **kill program** that can be used to send a signal to a process (uses the `kill()` system call).



Some Signals

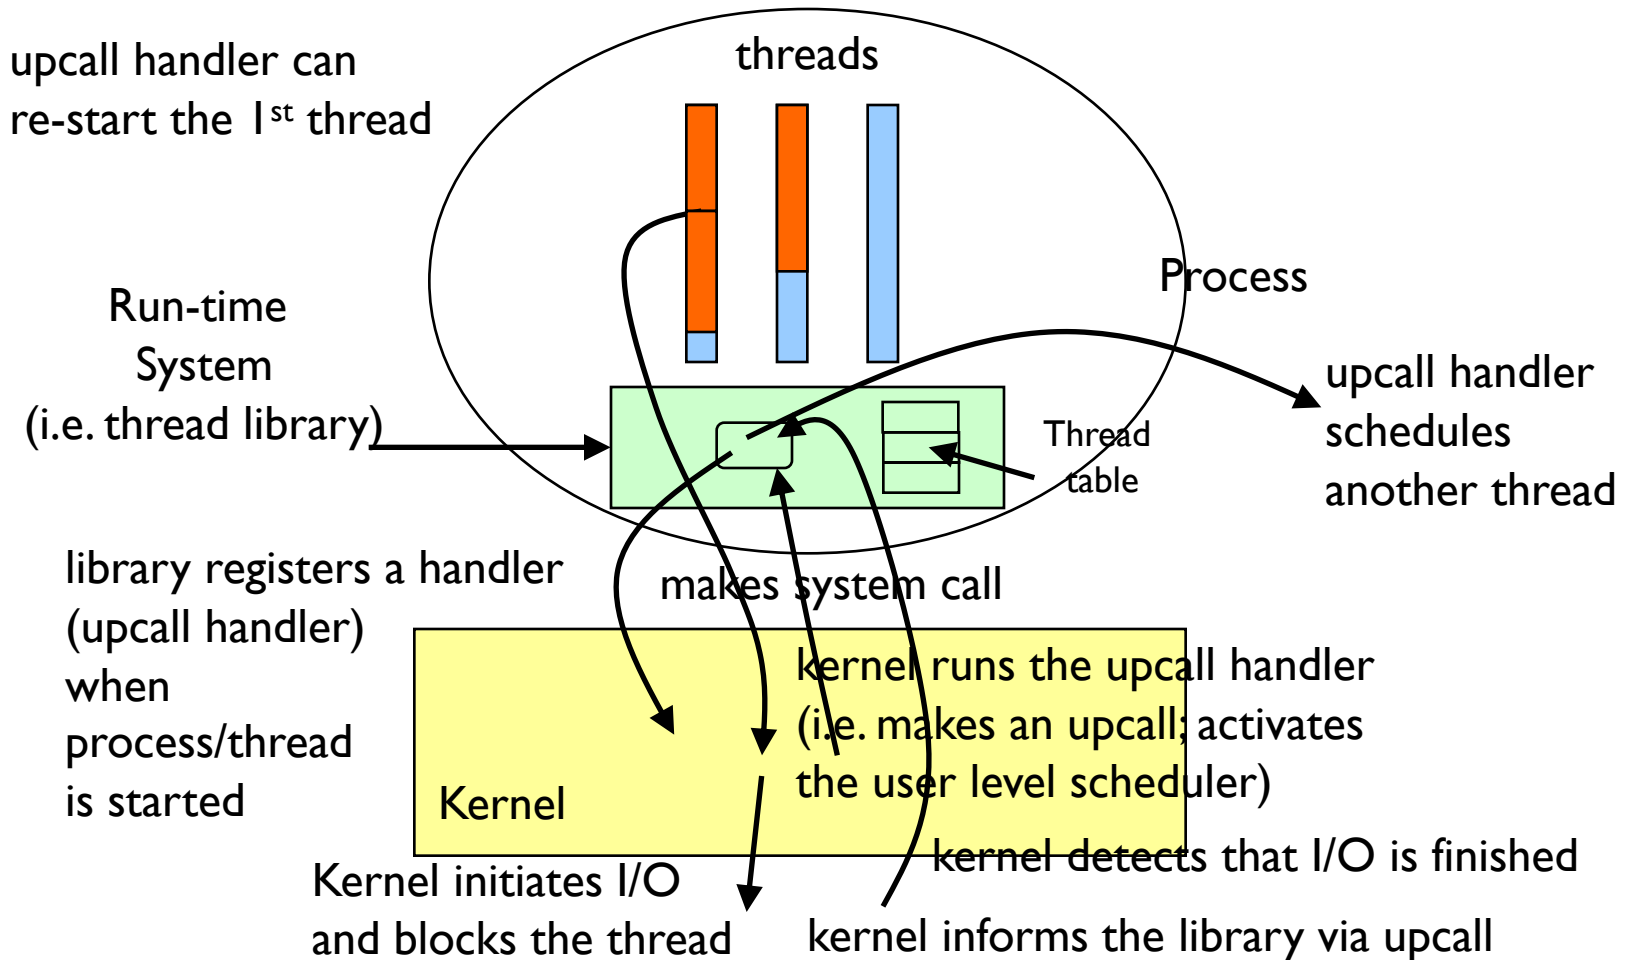
SIGABRT	Process abort signal.
SIGALRM	Alarm clock.
SIGBUS	Access to an undefined portion of a memory object.
SIGCHLD	Child process terminated, stopped, or continued.
SIGCONT	Continue executing, if stopped.
SIGFPE	Erroneous arithmetic operation.
SIGHUP	Hangup.
SIGILL	Illegal instruction.
SIGINT	Terminal interrupt signal.
SIGKILL	Kill (cannot be caught or ignored).
SIGPIPE	Write on a pipe with no one to read it.
SIGQUIT	Terminal quit signal.
SIGSEGV	Invalid memory reference.
SIGSTOP	Stop executing (cannot be caught or ignored).
SIGTERM	Termination signal.

Additional Study Material

Scheduler Activations

- Kernel threads are good, but they are slower if we create short threads too frequently, or threads wait for each other too frequently.
- Is there a middle way?
 - Schedule Activation
- Goal is to mimic kernel threads at user level with some more kernel support. But kernel will not create another thread for each user thread (M:1 or M:M model).
- Avoid unnecessary transitions between user and kernel space.

Scheduler Activations: Upcall mechanism



References

- Operating System Concepts, Silberschatz et al. Modern Operating Systems, Andrew S. Tanenbaum et al.