



Bilkent University
Department of Computer Engineering
CS342 Operating Systems

CPU Scheduling

Last Update: Mar 17, 2023

Objectives and Outline

Outline

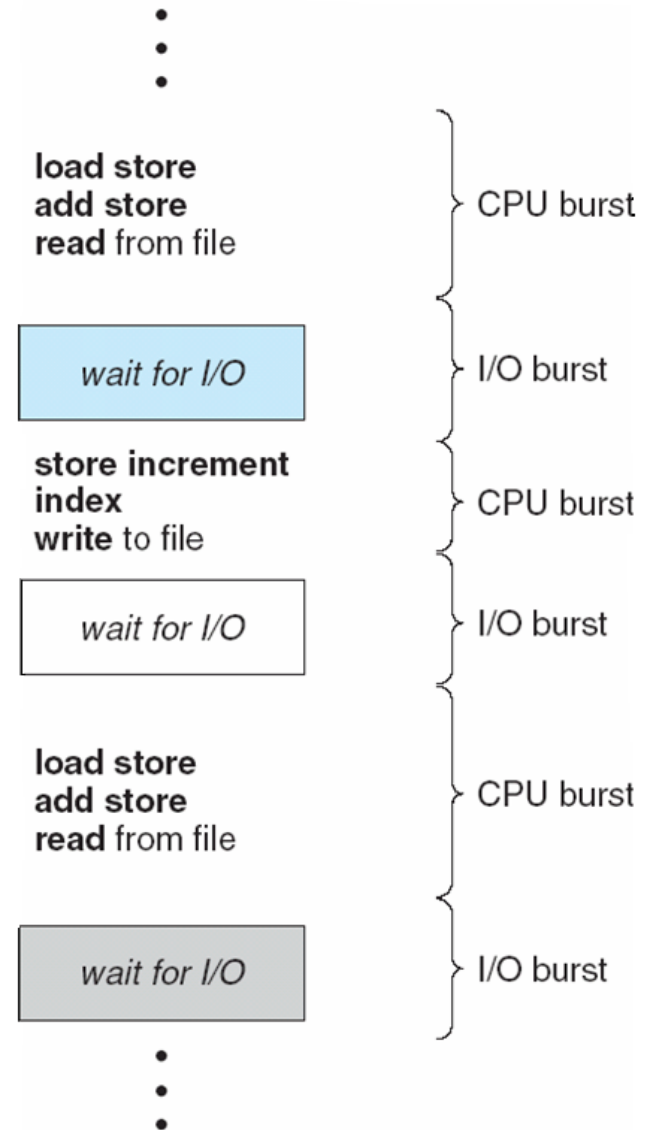
- Basic Concepts
- Scheduling **Criteria**
- Scheduling **Algorithms**
- Thread Scheduling
- Multiple-Processor Scheduling
- Operating Systems **Examples**
- Algorithm **Evaluation**

Objective


- To introduce **CPU scheduling**, which is the basis for multi-programmed (multi-tasking) operating systems
- To describe various CPU-scheduling algorithms
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system

Basic Concepts

- Maximum CPU utilization obtained with **multiprogramming**.
- Process execution usually consists of a cycle of CPU execution and I/O wait.
- CPU–I/O Burst Cycle repeated.



Histogram of CPU-burst Times



CPU burst distribution

CPU Scheduler


- **Selects** from among the **processes** (**threads**) in memory that are **ready** to execute, and allocates the CPU to one of them
- CPU **scheduling decisions** may take place when a process:
 1. Switches from running to waiting state (e.g., system call)
 2. Switches from running to ready state (upon interrupts)
 3. Switches from waiting to ready (e.g., I/O completed)
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**

CPU Scheduler: kernel design

- Virtually all modern operating systems are preemptive.
 - A process is **not allowed** to run as long as it wishes.
- A **preemptive** system is **more responsive** and suitable for real-time tasks.
- But preemptive scheduling may cause **race conditions** among processes/threads sharing data.
 - A context switch can happen at any moment (not only when waiting for I/O)
 - A process may be in middle of updating a shared data.
- Preemption effects **kernel design**.
 - A process running in kernel mode (as a result of a system call or interrupt) may be in middle of updating a shared kernel structure.
 - Kernel should use locks, etc., to avoid race conditions.

Dispatcher

- **Dispatcher module** gives control of the CPU to the process selected by the CPU scheduler; this involves:
 - switching context from one process to the other
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- **Dispatch latency**: time it takes for the dispatcher to stop one process and start another running.



Linux tools

- Assume you have a process started with pid x.
- Type at command shell:
`cat /proc/x/status`
- At the bottom of the output you will see the **number of context switches** that the process experiences. Like the following:

```
voluntary_ctxt_switches: 9  
nonvoluntary_ctxt_switches: 18
```

- The **vmstat** command can give statistics about context switches as well.

Scheduling Criteria

Many criteria exist for comparing scheduling algorithms.

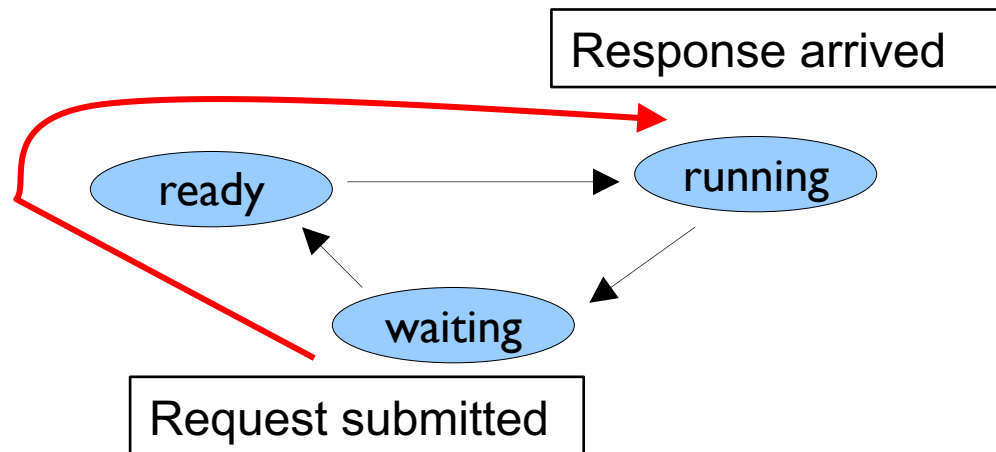
- **CPU utilization:** keep the CPU as busy as possible
 - **Throughput:** number of processes that complete their execution per time unit
 - **Turnaround time:** amount of time to execute a particular process
 - **Waiting time:** *total* amount of time a process has waited in the ready queue (in ready state)
- Maximize CPU utilization
 - Maximize throughput
 - Minimize turnaround time
 - Minimize waiting time

$$\text{Turn. Time} = \text{Finish time (completion time)} - \text{Start time (arrival time)}$$

Scheduling Criteria

- **Response time**: amount of time it takes from when a request was submitted **until the first response** is produced (for interactive systems)
 - In other words, elapsed time between the arrival to ready queue and the first execution in the CPU.

Important for Interactive Processes
(interacting with a user)



We want to minimize response time

Scheduling Algorithms

Scheduling algorithms

- We will describe several algorithms.
- We will illustrate their operation and calculate metrics.
 - Several processes are considered.
- An **accurate illustration** should consider **many** CPU and I/O bursts for each process.
- For **simplicity**, for each process, **a single CPU burst** is considered.
- Basic metric used: waiting time
- Algorithms are discussed assuming system has **one core** (CPU).
 - Later, we will discuss scheduling issues in multicore (**multiprocessor**) systems.


First-Come, First-Served (FCFS) Scheduling

- Simple scheduling algorithm.
- Process that requests the CPU first, is allocated the CPU first.
- FIFO queue can be used in its implementation.
- The PCB of the process that becomes ready is added to the **tail** of the queue.
- When a running process finishes its CPU burst (by terminating or by doing a I/O call and going into waiting state), it is removed from the CPU and ready queue. Then the **head** of the queue is allocated the CPU.
- Code of FCFS is simple.
- Time to wait till next execution can be quite long.

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time (ms)</u>
P ₁	24
P ₂	3
P ₃	3

- Suppose that the processes arrive at time 0 in the **order**: P₁, P₂, P₃. The **Gantt Chart** for the schedule (illustrating the scheduling) is:



- Waiting time** for P₁ = 0; P₂ = 24; P₃ = 27
- Average waiting time**: $(0 + 24 + 27)/3 = 17$ ms

FCFS Scheduling (Cont)

Suppose that the same processes arrive at time 0 in the following order

P_2, P_3, P_1

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3 \text{ ms}$
- Much better than previous case

FCFS Scheduling (Cont)

- The **average waiting time** in FCFS is usually **not minimal**.
- The average waiting time can be vary **substantially** if CPU burst times of processes **vary greatly**.
- FCFS is **nonpreemptive**. A process is allowed to run until it finishes its CPU burst (which will happen when **process requests an I/O operation** or **terminates**).
- Causes **convoy effect**: short processes behind long process
 - Results in low CPU and device utilization.
 - Short CPU burst processes are I/O bound.
 - Long CPU burst processes are CPU-bound.
- Not good for interactive systems.


Shortest-Job-First (SJF) Scheduling

- Associate with **each process** the **length of its next CPU burst**.
- Use these lengths to schedule the process with the **shortest** time.
- When CPU becomes available, the process that has the **smallest next CPU burst** is allocated the CPU.
 - A better name would be *shortest-next-cpu-burst* algorithm.
 - Since scheduling depends on next cpu burst of a process, rather than its total length.
- If there is a tie (two processes having the same next shortest CPU burst), the tie can be broken by applying FCFS.

Example of SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0.0	6
P ₂	0.0	8
P ₃	0.0	7
P ₄	0.0	3

- SJF scheduling chart



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$ ms

Shortest-Job-First (SJF) Scheduling

- SJF is **optimal** for **waiting time**: gives **minimum average waiting time** for a given set of processes.
- **Moving** a short process before a long process (**exchanging** these processes in the ready list) **decreases** the waiting time of the short process more than it **increases** the waiting time of the long process.
 - Therefore the average waiting will decrease with each such exchange.
- The difficulty in **implementing** SJF is **knowing** the length of the next CPU burst for each process.
- But we can **approximate** the SJF scheduling.
 - **Predict** the length of the next CPU burst.
 - We **expect** next CPU burst to be similar in length to the previous CPU burst (or bursts).

Determining Length of Next CPU Burst

- We can **predict** (**estimate**) the length of the next CPU burst.
- Prediction can utilize the the lengths of the previous CPU bursts of a process.
- We can use **exponential averaging** of previous burst lengths as the prediction method.
- We can keep a **running exponential average**:
 - Updated after each CPU execution (CPU burst) completed.
- We need to do this for each process.
 - When the current **CPU burst finishes (for a process)**, we need to update the **prediction** (running exp avg) (of **that process**).
- When scheduling is needed, we **select** the process that **has smallest predicted value** (smallest running exp avg).

Predicting the Length of next CPU Burst: Exponential Averaging

- Let t_n denotes the length of the n^{th} CPU burst.
 - This is the **actual** length (**known** after burst finishes)
- Let τ_{n+1} denote the **predicted value** for the next CPU burst (i.e., the **new avg**).
- Assume the first CPU burst is Burst₀ and its length is t_0 .
- τ_n denotes the previous estimate (i.e., the **old avg**) (before t_n)
- Define α (a weight factor) to be: $0 < \alpha < 1$.
- Then we **define** the new avg:

$$\tau_{n+1} = \alpha \cdot t_n + (1 - \alpha) \cdot \tau_n$$

Exponential Averaging

- We have **CPU bursts** as: $\text{Burst}_0, \text{Burst}_1, \text{Burst}_2, \dots, \text{Burst}_n, \text{Burst}_{n+1}$.
- The actual lengths of these bursts are denoted by: $t_0, t_1, t_2, t_3, \dots, t_n, t_{n+1}$. Let τ_0 be initial estimate (i.e., estimate for Burst_0) and let it be a constant value like 10 ms . Then

$$\tau_1 = \alpha t_0 + (1 - \alpha)\tau_0$$

- If we **expand** the formula, we get:


$$\begin{aligned} \tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + (1 - \alpha)^2 \alpha t_{n-2} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^n \alpha t_0 + (1 - \alpha)^{n+1} \tau_0 \end{aligned}$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1, **each successive term** has **less weight** than its predecessor.

Exponential Averaging

- If $\alpha = 0$:
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- If $\alpha = 1$:
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts.
- Usually we have α between 0 and 1, for example 0.5.

Exponential Averaging



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

An exponential average with $\alpha = 1/2$ and $\tau = 10$ ms.

Example

- $\tau_0 = 10$ ms. Measured CPU bursts: $t_0 = 8$ ms, $t_1 = 16$ ms, $t_2 = 20$ ms, $t_3 = 10$ ms
- Assume $\alpha = \frac{1}{2}$. Then we have:
 - $\tau_1 = \frac{1}{2} \times 8 + \frac{1}{2} \times 10 = 9$
 - $\tau_2 = \frac{1}{2} \times 16 + \frac{1}{2} \times 9 = 12.5$
 - $\tau_3 = \frac{1}{2} \times 20 + \frac{1}{2} \times 12.5 = 16.25$
 - $\tau_4 = \frac{1}{2} \times 10 + \frac{1}{2} \times 16.25 = 13.125$.
- The next CPU burst length is estimated to be 13.125 ms. After that burst has completed, it is measured as t_4 .
- We **update τ** for a process, when its **burst has completed**.
- Hence we will maintain up-to-date τ value for each process. We will select the process with minimum τ .


Shortest Remaining Time First (SRTF)

- **Preemptive** version of SJF.
- While a job A is running, if a new job B comes whose length is shorter than the remaining time of job A, then B preempts A and B runs.

Shortest Remaining Time First (SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0.0	8
P ₂	1.0	4
P ₃	2.0	9
P ₄	3.0	5

- SRJF scheduling chart



- Average waiting time = $(9 + 0 + 2 + 15) / 4 = 6.5$ ms

Example

- Assume we have the following processes. Find out the finish time, waiting time and turnaround time of each process for the following scheduling algorithms: FCFS, SJF, SRTF.

Process	Arv time	CPU Burst
A	0	30
B	5	20
C	10	12
D	15	10

Example

FCFS: Processes will run in the order they arrive.

The following is the finish, turnaround, waiting time of each process.

	Arv	Burst	Finish	Turnaround	Waiting
A	0	30	30	30	0
B	5	20	50	45	25
C	10	12	62	52	40
D	15	10	72	57	47

Example

SJF: running order will be: A(30) D(10) C(12) B(20)

	Arv	Burst	Finish	Turnaround	Waiting
A	0	30	30	30	0
B	5	20	72	67	47
C	10	12	52	42	30
D	15	10	40	25	15

Example

SRTF: running order will be: A(5) B(5) C(12) D(10) B(15) A(25)

	Arv	Burst	Finish	Turnaround	Waiting
A	0	30	72	72	42
B	5	20	47	42	22
C	10	12	22	12	0
D	15	10	32	17	7

Round Robin (RR) Algorithm

- Each process is allowed to run a small unit of time (called *time quantum* or *time slice*) in CPU.
- Time quantum (time slice) is usually a value between 10 and 100 ms.
- Then another process is run.
- That means a process is not allowed to execute in CPU as long as it wishes (i.e., until its CPU burst completes).
 - The running process can be forcefully taken from CPU (preempted) if its time quantum (time slice) finishes, even though the CPU burst of the process did not finish.

Round Robin (RR) Algorithm

- Can be implemented by using a **queue**
 - The ready processes are put in a queue.
- The **head** of the queue is selected to run next (when scheduling needed).
- A process that has run and finished its time quantum is put to the **tail** of the queue. **Preemptive** scheduling.
- A new arriving process (a new process or a process completing an I/O wait) is added to the **tail** of the queue.
- A process that finishes its CPU burst will no longer be in the queue (has terminated or is waiting for something).
- When a process finishes its CPU burst **before** its time quantum expires, again another process is **scheduled** immediately.

Round Robin (RR) Algorithm

- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
- No process waits more than $(n - 1)q$ time units (response time)
- Performance
 - q large \Rightarrow behaves like FIFO
 - q small \Rightarrow good CPU sharing, good response time


q must be large with respect to context switch time, otherwise overhead is too high.

Example of RR with time quantum = 4 time units

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	3

All processes
arrived at time 0.

- The Gantt chart is:



Waiting times:

P₁: 10-3 = 6 ms

P₂: 4-0 = 4 ms

P₃: 7-0 = 7 ms

Avg = 17/3 = 5.66 ms

- Typically, RR has **higher average turnaround time** than SJF and FCFS, but it has **better response time**.

Example

Process	Arrival Time (ms)	CPU Burst Length (ms)
A	0	40
B	15	25
C	25	30
D	35	45
E	55	25

Finish time of each process?

a) Round Robin $q = 30$

b) Round Robin $q = 10$

Example


Solution

	RR $q=30$	RR $q=10$
A	95	80
B	55	105
C	85	125
D	165	165
E	150	150

RR vs FCFS

- Round Robin is good for fast response (small response time), but not for small turnaround time.

Assume 3 jobs all arrived at time 0. Each has a CPU burst = 10 ms



Turnaround times


A: 20 ms

B: 25 ms

C: 30 ms

Avg TT = 25 ms

Worst response time = 10 ms



Turnaround times

A: 10 ms


B: 20 ms

C: 30 ms

Avg TT = 20 ms


Worst response time = 20 ms

Time Quantum and Context Switches



If time quantum (q) is small, a process **experiences** more **context switches** during a CPU **burst** execution. Hence we will have more overhead.

Average turnaround time varies with time quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

When time quantum (q) gets **larger**, average **turnaround time** gets **smaller**. Exceptions can occur for particular values of burst lengths and time quantum.

Priority Scheduling

- A **priority number** (integer) is associated with **each process**
- The CPU is **allocated** to the process with the **highest priority**.
 - Equal priority processes are scheduled in FCFS order.
- **Convention**: the *smaller* the priority number, the *higher* the priority, or vice versa.
- Has two versions:
 - **Preemptive** (higher priority process preempts the running one)
 - **Non-preemptive**
- SJF is a priority scheduling (non-preemptive) where priority is the *predicted next CPU burst time*.
- SRTF is preemptive priority scheduling.

Example

	Arv	CPU burst	Priority
A	0	20	3
B	5	15	2
C	10	20	0
D	25	15	1
E	30	20	1

Nonpreemptive priority scheduling:

AAAACCCCDDEEEEBBB

assuming each letter is 5 time units

Finish times: A: 20, B: 90, C: 40, D: 55, E: 75

Preemptive priority scheduling:

ABCCCCDDDEEEEBBAAA

Finish times: A: 90, B: 75, C: 30, D: 45, E: 65

In case of a **tie**, we can choose the one that is early in the queue (FCFS).

Priority Scheduling

- Problem: **Starvation (indefinite blocking)**.
 - Low priority processes may never have a chance to execute.
- Solution: **Aging**.
 - As time progresses, increase the priority of the process.

Priority Scheduling and Round Robin together.

- Priority scheduling can be combined with RR.
 - Highest priority process is run first.
 - If there are multiple processes with the same priority, they can be served with RR. We have RR queue used for each priority.

All bursts arrived at time 0.

Time quantum $q = 2$


P1 4 (3)

P2 5 (2)

P3 8 (2)

P4 7 (1)

P5 3 (3)



(priority)

In this example: smaller priority number indicates higher priority.

burst time

Multilevel Queue Scheduling

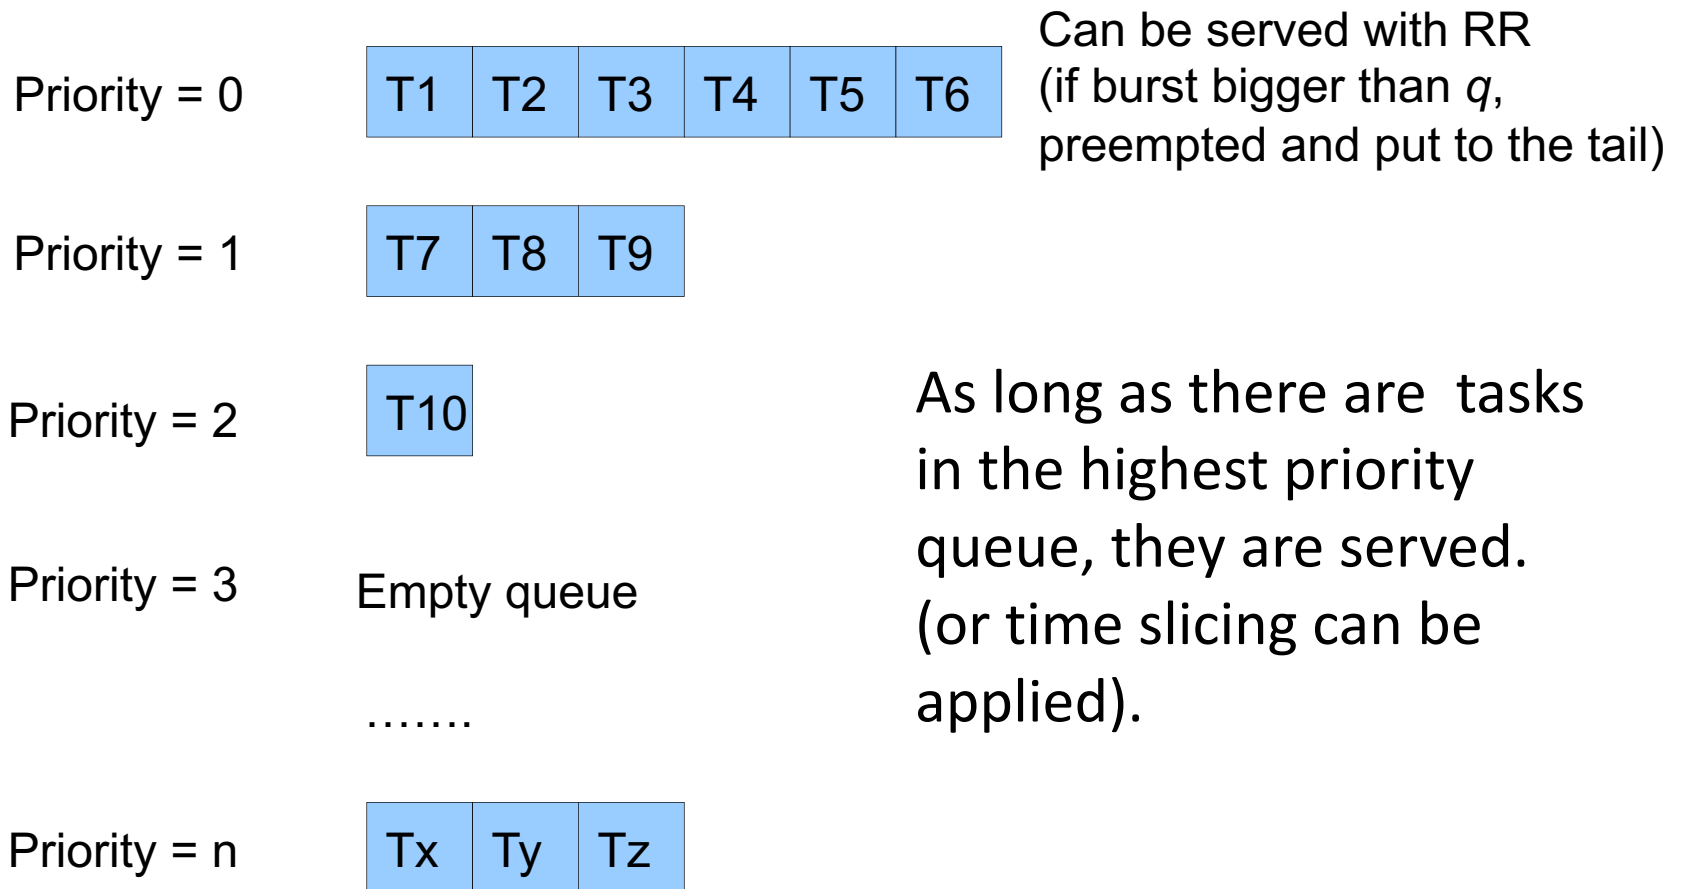
- Ready queue can be **partitioned** into separate (**multiple**) **queues**.
- For example: We may have 2 queues
 - foreground (interactive – interacting with a user), and
 - background (non-interactive) processes.
- **Foreground** processes have **higher priority** (to have small response time, since they are interactive)
- Each queue has its **own scheduling algorithm**.
 - foreground – RR (we can not wait a process too long)
 - background – FCFS (response time is not important)

Multilevel Queue Scheduling

Scheduling between queues

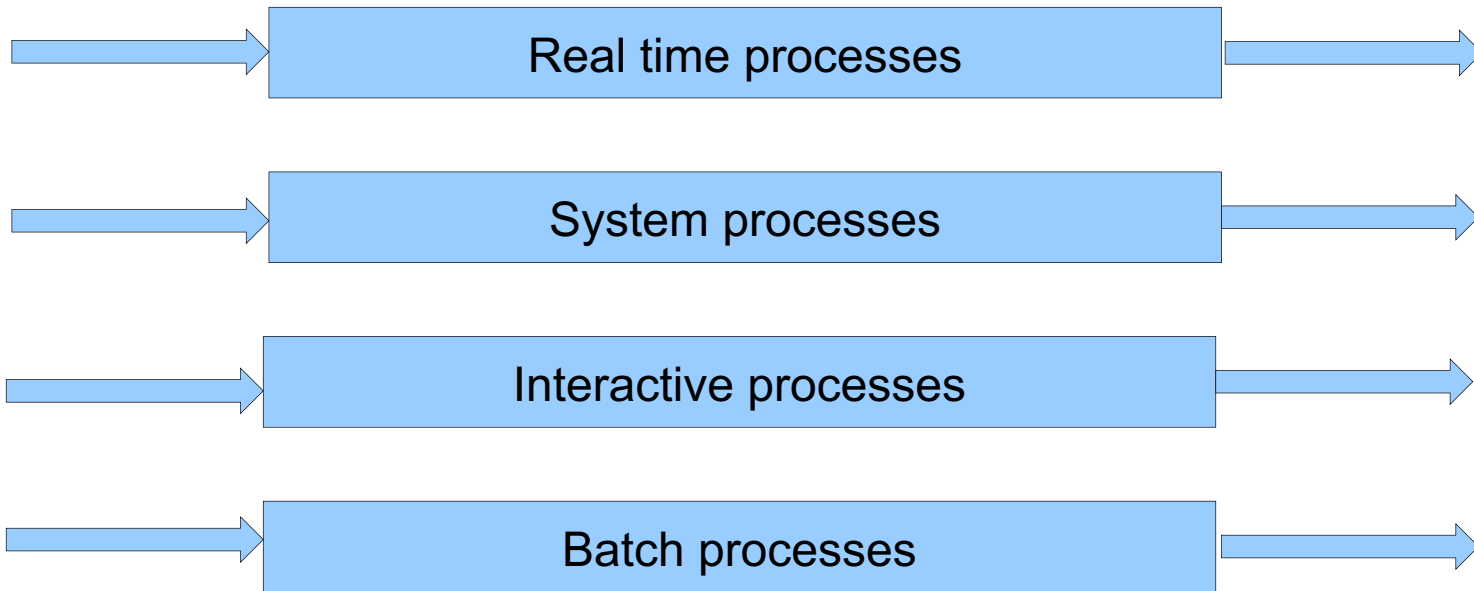
- Scheduling must be done **between the queues**. We can have the following alternatives:
 - Fixed (**strict**) **priority scheduling**. Serve foreground processes first. If no foreground process, then start serving background processes.
 - Possibility of **starvation**.
 - **Time slicing**. Each queue gets a certain amount of CPU time, which it can schedule among its processes. For example:
 - **80% of time** serve foreground processes with RR.
 - **20% of time** serve serve background processes with FCFS.

Multilevel Queue Scheduling



Multilevel Queue Scheduling

Highest priority



Lowest priority

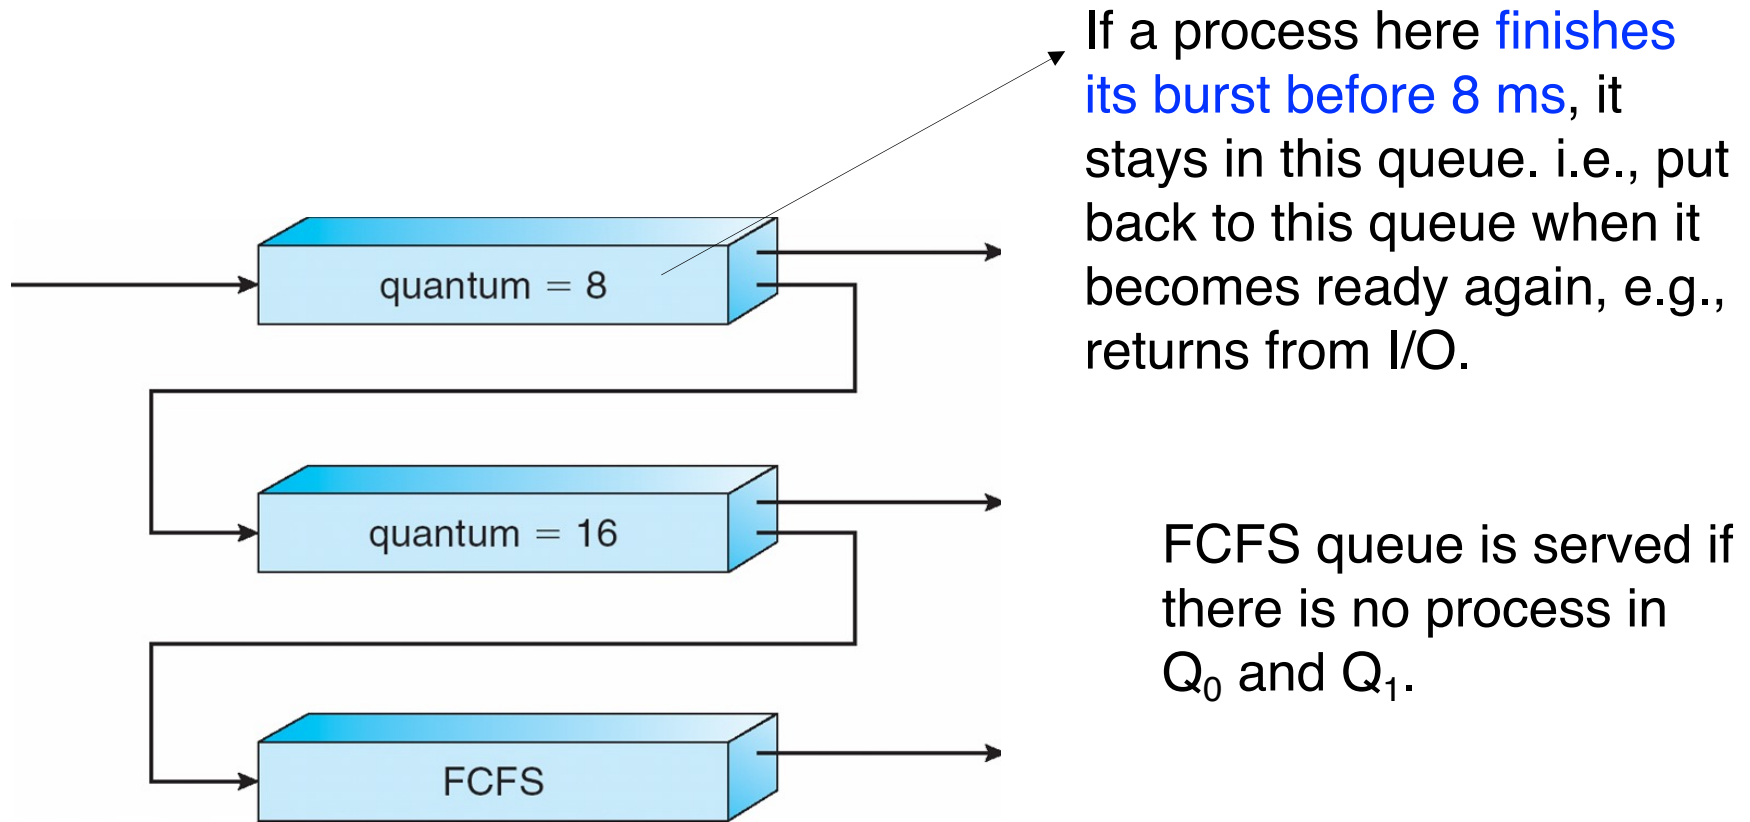
Multilevel Feedback Queue (MLFQ) scheduling

- A **general** scheduling algorithm.
- A process can **move between the various queues**. Aging can be implemented this way.
- A **specific** Multilevel-Feedback-Queue scheduler is **defined** by the following **parameters**:
 - **number** of queues
 - scheduling **algorithm** for each queue.
 - method to decide when to **upgrade** a process.
 - method to decide when to **downgrade** a process.
 - method to decide which queue **an arriving** process will **enter**.

Example of Multilevel Feedback Queue

- For **example**, we can have the parameters as follows.
- Three queues:
 - Q_0 – RR with $q = 8$ ms
 - Q_1 – RR time $q = 16$ ms
 - Q_2 – FCFS
- Scheduling rules:
 - A new job enters queue Q_0 .
 - A job in Q_0 has 8 ms time quantum. If it does **not finish its burst** in 8 ms, it is moved to queue Q_1 (seems it is non-interactive)
 - At job in Q_1 16 ms time quantum. If it **does not finish its burst** in 16 ms, it is moved to queue Q_2 (seems it is really non-interactive)
 - At Q_2 , jobs are served with FCFS.

Multilevel Feedback Queues



To prevent starvation, a process that waits too long in a lower-priority queue, may gradually be moved to a higher priority queue.

Thread Scheduling

Thread Scheduling Issues

- Thread API may give some amount of control to the application programmers about how threads will be scheduled.
- Some systems implementing M:1 (many-to-one) and/or M:M models, may allow specifying at which level new threads will be scheduled:
 - at user-level (by thread library), or
 - at kernel-level (by kernel).

Thread Scheduling Issues

- Scheduling at **user** level **means process-contention scope** (PCS), since scheduling competition is happening within the process
 - In Pthreads: PTHREAD_SCOPE_PROCESS option
- Scheduling at **kernel** level **means system-contention scope** (SCS), since scheduling competition is happening within the kernel with other threads seen by the kernel
 - In Pthreads: PTHREAD_SCOPE_SYSTEM option
- Since Linux uses 1:1 model, it schedules with SCS.
- In Linux, we can also specify the scheduling class/policy. This effects the scheduling algorithm used by the kernel.
 - FIFO: FCFS scheduling
 - RT: real time scheduling
 - OTHER: schdeluling algorithm for ordinary processes (default)

Example: Pthread Scheduling API

```
int main(int argc, char *argv[])
{
    int i; pthread_t tid[5]; pthread_attr_t attr;

    pthread_attr_t attr; // get the default attributes
    pthread_attr_setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    /* set the scheduling policy - FIFO, RT, or OTHER */
    pthread_attr_setschedpolicy(&attr, SCHED_OTHER);

    for (i = 0; i < 5; i++) // create five threads
        pthread_create(&tid[i], &attr, runner, NULL);
    for (i = 0; i < 5; i++)
        pthread_join(tid[i], NULL);
}

void *runner(void *param)
{
    printf("I am a thread\n");
    pthread_exit(0);
}
```

Multiprocessor Scheduling


Multiple-Processor Scheduling

- CPU scheduling more complex when multiple CPUs are available in a computer.
- Algorithms we have seen can be applied. Additionally, there are **more issues** to consider.
- Different multiprocessor architectures:
 - Multicore CPUs
 - Multithreaded cores (Hyperthreading – HW threads)
 - NUMA systems (non-uniform memory access)
 - Heterogenous multiprocessing (CPUs are different)

Multiprocessor scheduling



- Asymmetric multiprocessing:
 - Only one processor accesses the kernel data structures, alleviating the need for data sharing.
- Symmetric multiprocessing (SMP):
 - Each processor is self-scheduling.
 - Kernel can run on any processor when a hardware or software interrupt occurs.
 - Common in our desktop and laptop computers.

Multiple-Processor Scheduling: SMP



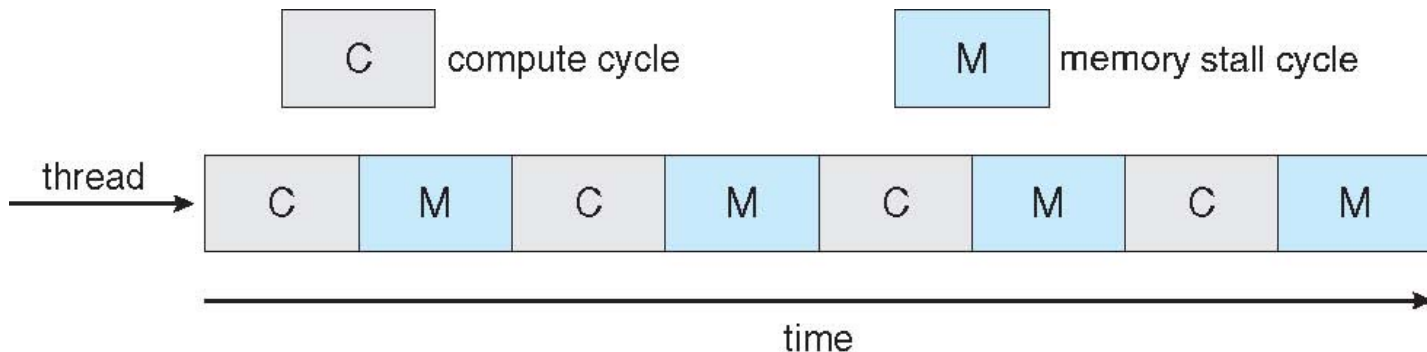
SMP

- We have two alternatives for enqueueing ready processes (threads):
 - a) All processes can be put into a **common** ready **queue**
 - b) Each processor can have its own **private queue** of ready processes.



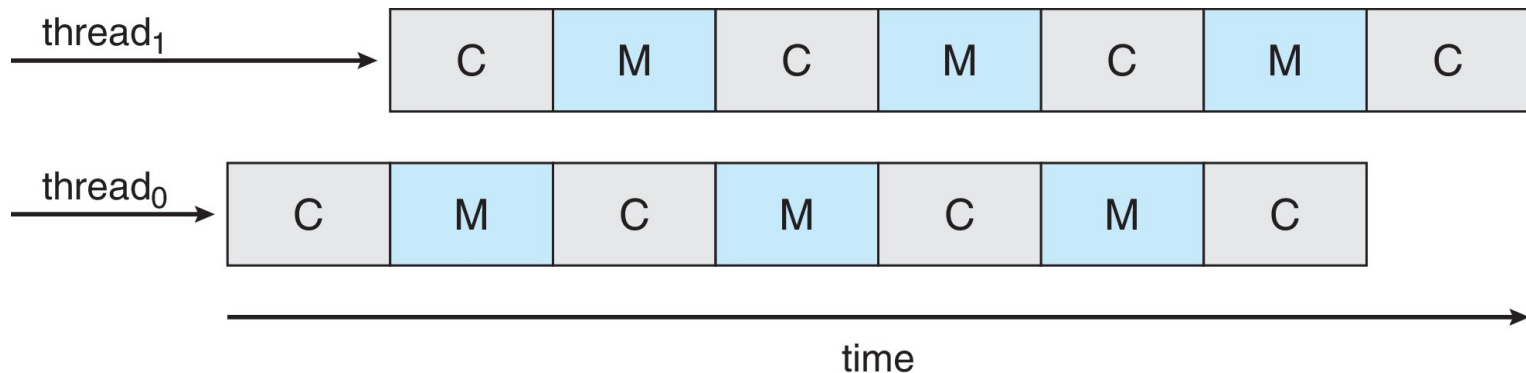
Multicore Processors

- Recent trend is to place multiple processor **cores** (CPUs) on **same physical chip** (die)
- Faster and consumes less power
- Multiple (hw) threads per core also growing (**hardware threading**)
 - Takes advantage of *memory stall* to make progress on another thread while memory retrieve is happening.



Multithreaded Multicore System


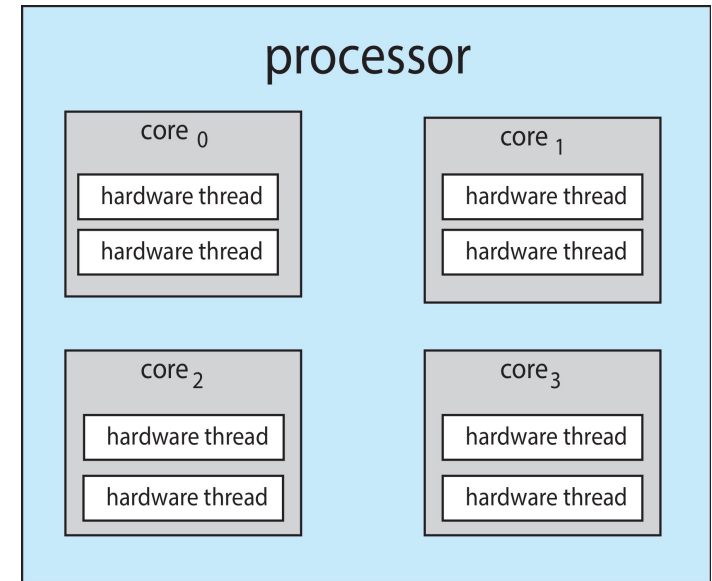
- Each core has more than 1 hardware threads (for example, 2).
- If one thread has a memory stall, hw switches to another thread. This is done at hardware level.




Appears to OS as two CPUs

Multithreaded Multicore System

- **Chip-multithreading (CMT)** assigns **each core multiple hardware threads** (Intel refers to this as **hyperthreading**.)
- On a quad-core system with 2 hardware threads per core, the **operating system sees 8 logical processors**.



NUMA and CPU Scheduling



NUMA: non-uniform memory access. Multiple nodes (chips) in a computer. Each node has CPU and Memory. Memory in a chip is accessible to other CPUs on other chips, but with more delay.

Processor Affinity

- A thread (process) has *affinity* to the CPU it is running (due to *per CPU cache*).
- If run in the same CPU, no need to re-populate the cache.
- *Soft affinity*: system *tries* to schedule to the same CPU
- *Hard affinity*: system *always schedules* to the same CPU
- Linux has `sched_setaffinity()` system call for hard affinity.

Load Balancing

- If load is unbalanced (in a NUMA system):
 - Load balancing among the cores on the same chip can be performed.
 - Load balancing between different NUMA processors (chips) is not desirable.
 - Cache needs to be populated
 - Access to memory in other chip is slower.

Heterogenous Multiprocessing (HMP)

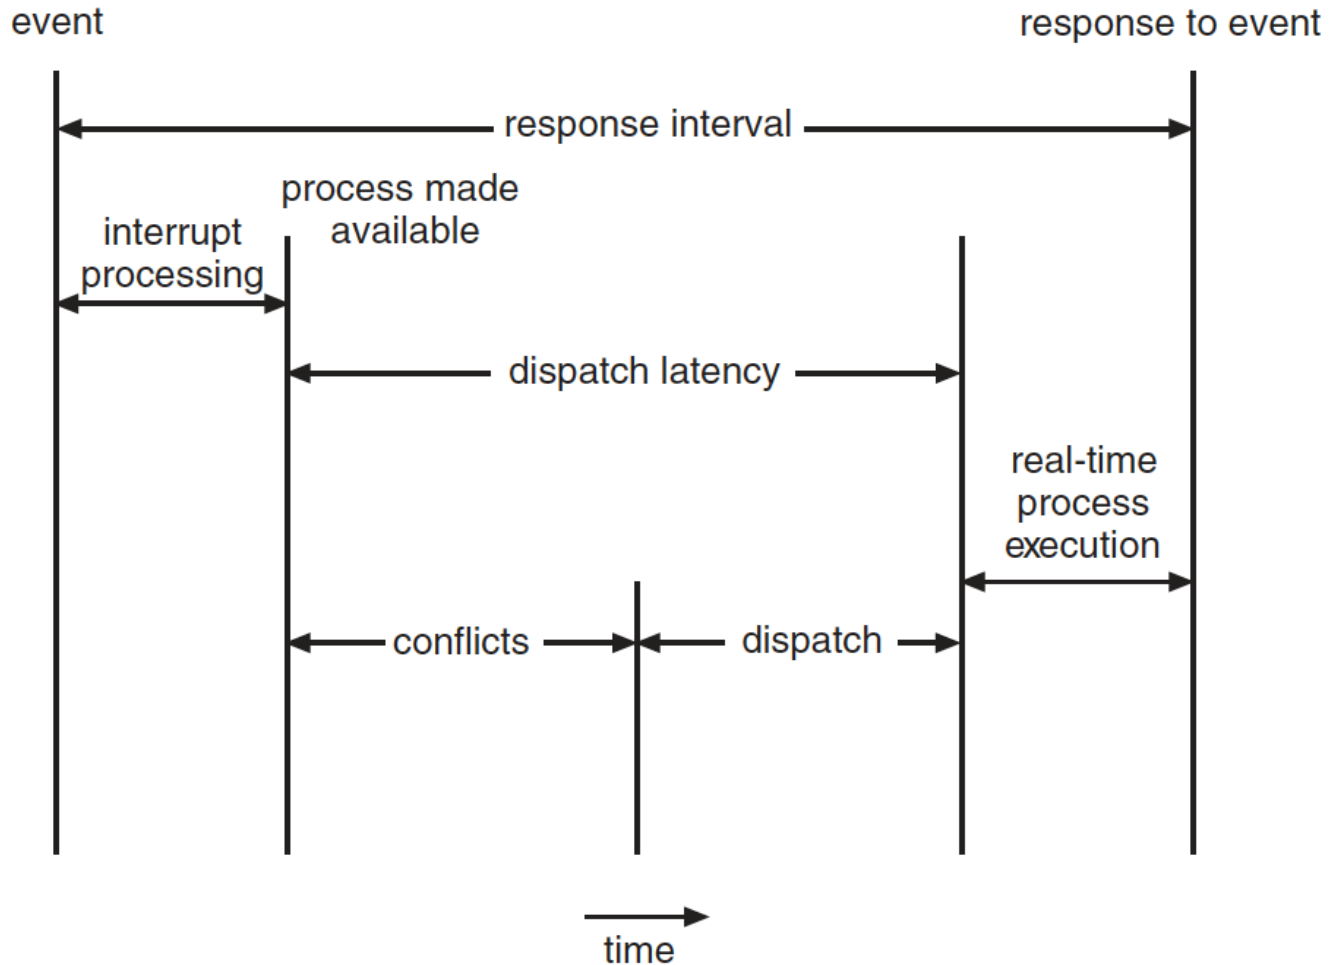
- In mobile devices, there are
 - Little cores (slow CPUs) (consuming less power)
 - Big cores (fast CPUs) (consuming more power)
- Tasks are assigned to cores based on their processing needs.
- Tasks that need high performance are scheduling to big cores. But they should not run too long.
- Tasks in little cores can run too long.

Real-Time CPU Scheduling

Real time systems

- Real-time processes (**tasks**) running in a system
 - There are hard and soft real-time systems
- Commonly, RT systems are **event-driven** systems
 - When an event occurs: need to respond (run a task) as quickly as possible
 - Minimizing **latency** is important
- Latency
 - **Interrupt latency**: time from interrupt till running the respective ISR, which marks the related task available (ready)
 - **Dispatch latency**: time between task becomes ready until it gets running

Latency



Scheduling

- **Priority based** scheduling
 - Priority based algorithm used
 - Preemption may be provided as well
 - **Soft real-time systems** apply this.
 - No hard guarantees
 - **Example**: video player in a computer
- **Hard real-time** guarantees
 - There are **deadlines** for processes (tasks) to execute
 - Processes may be periodic or not
 - **Admission control** required

Hard guarantees

- Periodic tasks:
 - Has period (p), i.e., has rate $= \frac{1}{p}$
 - Requires some certain amount of cpu-time (t) periodically.
 - Has deadline (d)
 - We must have: $0 \leq t \leq d \leq p$
 - Example scheduling algorithm: Rate monotonic scheduling
- Non periodic tasks:
 - Example scheduling algorithm: Earliest deadline first (EDF)

Rate monotonic scheduling


- **Static priority** used
 - Priority inversely proportional with period.
- **Preemptive**: high-priority task preempts low-priority task
- CPU burst time (processing time) equal in all periods.
- CPU burst must be executed before the next period.
 - *Implicit Deadline = beginning of next period.*
- Can underutilize CPU.
- Is optimal for static priority case.
 - That means if any static-priority scheduling algorithm can meet all the deadlines, then the rate-monotonic algorithm can too.

Rate monotonic scheduling

Example: P1: cputime:20; period 50

P2: cputime:35; period 100


$$20/50 + 35/100 = 75 / 100 \text{ less than } 1$$



Example: P1: cputime:25; period 50

P2: cputime:35; period 80

$$25/50 + 35/80 = 375/400 \text{ less than } 1$$



deadline missed

Rate monotonic scheduling

- A feasible schedule (not missing deadlines) for N processes is guaranteed to exist if combined CPU utilization of N processes is less than a **bound** (worst-case CPU utilization)

$$\text{Worst-case CPU utilization} = N(2^{1/N} - 1)$$
$$\text{Combined utilization} \leq \text{Worst-case utilization}$$


- With $N = 2$, worst-case utilization is 83%. That means:
 - When combined utilization (for $N = 2$) is 75%, we are guaranteed to schedule 2 periodic processes.
 - When combined utilization (for $N = 2$) is above 83%, we are not guaranteed to schedule 2 periodic processes.
- As N approaches **infinity**, utilization bound is **69%**.

Earliest Deadline First (EDF)

- Dynamically assigns priorities depending on deadlines.
 - When a process becomes runnable (ready), it announces its deadline.
 - The process that has the earliest deadline is scheduled.
- Tasks do not need to be periodic
 - CPU bursts may be different.
- Theoretically optimal: can meet deadlines and can utilize CPU 100%.

Earliest Deadline First (EDF)

- Example: Two periodic processes (tasks):
P1: cputime: 25; period 50
P2: cputime: 35; period 80



Proportional Share Scheduling


- Allocate T shares among all applications (processes)
- If an application has N shares, it can get N/T of CPU time.
- System ensures that each process gets CPU time proportional to its shares.
- Admission control policy required.

Assume $T=200$ (total 200 shares)

Task 1 has 50 shares

Task 2 has 100 shares

Task 3 has 50 shares



25% of cpu time

50% of cpu time

25% of cpu time

Operating System Examples

- Linux scheduling

Linux Scheduling: CFS Algorithm

- To schedule ordinary processes (not real-time), Linux uses an algorithm called CFS: **Completely Fair Scheduling**.
- Linux has two **scheduling classes** at the moment. 1) real-time, and 2) time-sharing.
- Each **class** uses a **fixed priority range** (set of priority values that can be assigned to processes internally)
 - 1) real-time class (0-99).
 - 2) time-sharing (default scheduling) class **(100-140)**. **CFS** algorithm is used. Our ordinary processes are in this class.

Linux Scheduling: CFS Algorithm

- Fair share scheduling. If all processes are of equal priority, they can get the same share of the cpu time (approximately).
- Efficient and scalable. Selects a task to run next very quickly.
- Processes have priorities. 40 different priority values (called nice values).
 - $-20, -19, \dots, -1, 0, 1, 2, \dots, 18, 19$
 - -20 is highest priority.
 - Can be set using the nice command. Default value is 0.
- CFS assigns a proportion of CPU time based on nice value indirectly.

CFS

- CFS uses virtual runtime (**vruntime**) for each process (granularity: **nanoseconds**)
- vruntime indicates **how long a process has used** the cpu so far (**virtually**).
- When scheduling decision
- n is to be made, task with the **smallest vruntime is picked** to run next (*since it could not use the cpu as much as others so far*).
- If a process with **nice value 0** runs **x ms** in cpu (actual time), its vruntime is increased by **x ms**.
- The vruntime is **advanced more** (than x) for a **lower** priority (higher nice value) process, and **less** (than x) for a **higher** priority process, when such a process really runs **x ms** in cpu.

CFS

- What is the **time slice** for a process?
 - Time slice is set **dynamically**.
- CFS has a parameter called **targeted latency** (or, **sched_latency**).
 - Typical sched_latency is 48 ms.
- CFS uses sched_latency to **determine** how long (i.e., timeslice) a process will run in cpu before considering a context switch.
 - If there are n ready processes of **equal** priority, then the timeslice of each will be $48/n$. Over the latency period, each process will have a **chance to run**.
 - For example: if there 4 processes of equal priority, each will get 12 ms timeslice, i.e., each process is allowed to run 12 ms, then a context switch is considered.

CFS

- If **too many processes** (say 100), timeslice will be too small, and overhead will be too high.
- Therefore there is **minimum timeslice value**. It is called **min_granularity**.
 - It can be **6 ms**, for example. Then no process will run less than 6 ms, even though there are many processes.
 - For example: if 10 **processes**, $48/12 = 4.8$, but timeslice will be 6. That means it will take **60 ms** (instead of 48) to run these 10 processes; but that is fine.
- **Timer** can tick every **1 ms**. Then CFS checks if current task has reached its time slice. If not reached, process continues, otherwise context switch happens.

CFS

- CFS considers process **priority** as well (nice value).
- CFS maps **nice values** to **weights** using the following **table**.
- **Default nice value is 0**. smaller value ==> higher priority
- **Higher priority** process has **more weight**.

```
static const int prio_to_weight[40] = {  
    /* -20 */      88761,      71755,      56483,      46273,      36291,  
    /* -15 */      29154,      23254,      18705,      14949,      11916,  
    /* -10 */      9548,       7620,       6100,       4904,       3906,  
    /*  -5 */      3121,       2501,       1991,       1586,       1277,  
    /*   0 */      1024,        820,        655,        526,        423,  
    /*   5 */       335,        272,        215,        172,        137,  
    /*  10 */       110,         87,         70,         56,         45,  
    /*  15 */        36,         29,         23,         18,         15,  
};
```

CFS

- A process k gets its cpu share (**timeslice**) depending on its **weight** as:

$$timeslice_k = \frac{weight_k}{\sum_{i=0}^{n-1} weight_i} \cdot sched_latency$$

- Example: assume we just have 2 ready processes, A and B, with nice values as -5 and 0 .
 - Then $weight(A) = 3121$, and $weight(B)$ is 1024 .
 - Then A's timeslice will be $\cong \frac{3}{4} \times sched_latency = 36$ ms.
 - B's time slice will be $\cong 12$ ms.

CFS

- A process's **vruntime** is **advanced** considering its **weight** as well.
- When a process k runs t ms (timeslice) in cpu, t is its **actual runtime**. Then, its **vruntime** is **updated** as below:

$$vruntime_k = vruntime_k + \frac{weight_0}{weight_k} \cdot actualruntime_k$$

- The vruntime of a **higher priority** process advances **slower**.
 - *Then it will have a better chance for **quick** scheduling.*
- The vruntime of a **lower priority** process advances **faster**.
 - *We say, decay for a lower priority process is higher.*

CFS


- While **executing**, a process k may **start I/O** (system call). Then it will sleep/block in waiting state (no **longer** in **ready** state).
- When k returns from waiting (I/O completed), its **vruntime** can be **very behind**.
- To **prevent monopolizing** the cpu by k , such process k 's vruntime is set to be equal to the **minimum vruntime** (of all **ready processes** at that moment) **minus** some **constant**, after it returns from I/O.
- In this way process k gets scheduled **very quickly** but does not monopolize the cpu.
- A process returning from I/O wait (or just created) has its timeslice reset according to its weight.

CFS

- CFS does not have different queues for different priorities.
- CFS keeps all processes in ready state in a **red-black tree** (a **balanced binary search tree**) with respect to their **vruntimes**.
- The process on the **left lower side** is the one to run next (has the smallest vruntime).
 - **Selection:** $O(1)$
 - A pointer is kept to the minimum vruntime node.
 - **Insertion** of a process into tree: $O(\log N)$
- No Round Robin applied (no queue structure). Instead CFS just **picks** the **process k** with **minimum vruntime** from the tree and **runs it** for **timeslice(k)** ms.

CFS

- **Red black tree** example. There are 10 runnable processes. T7 is picked up by the scheduler to run next.



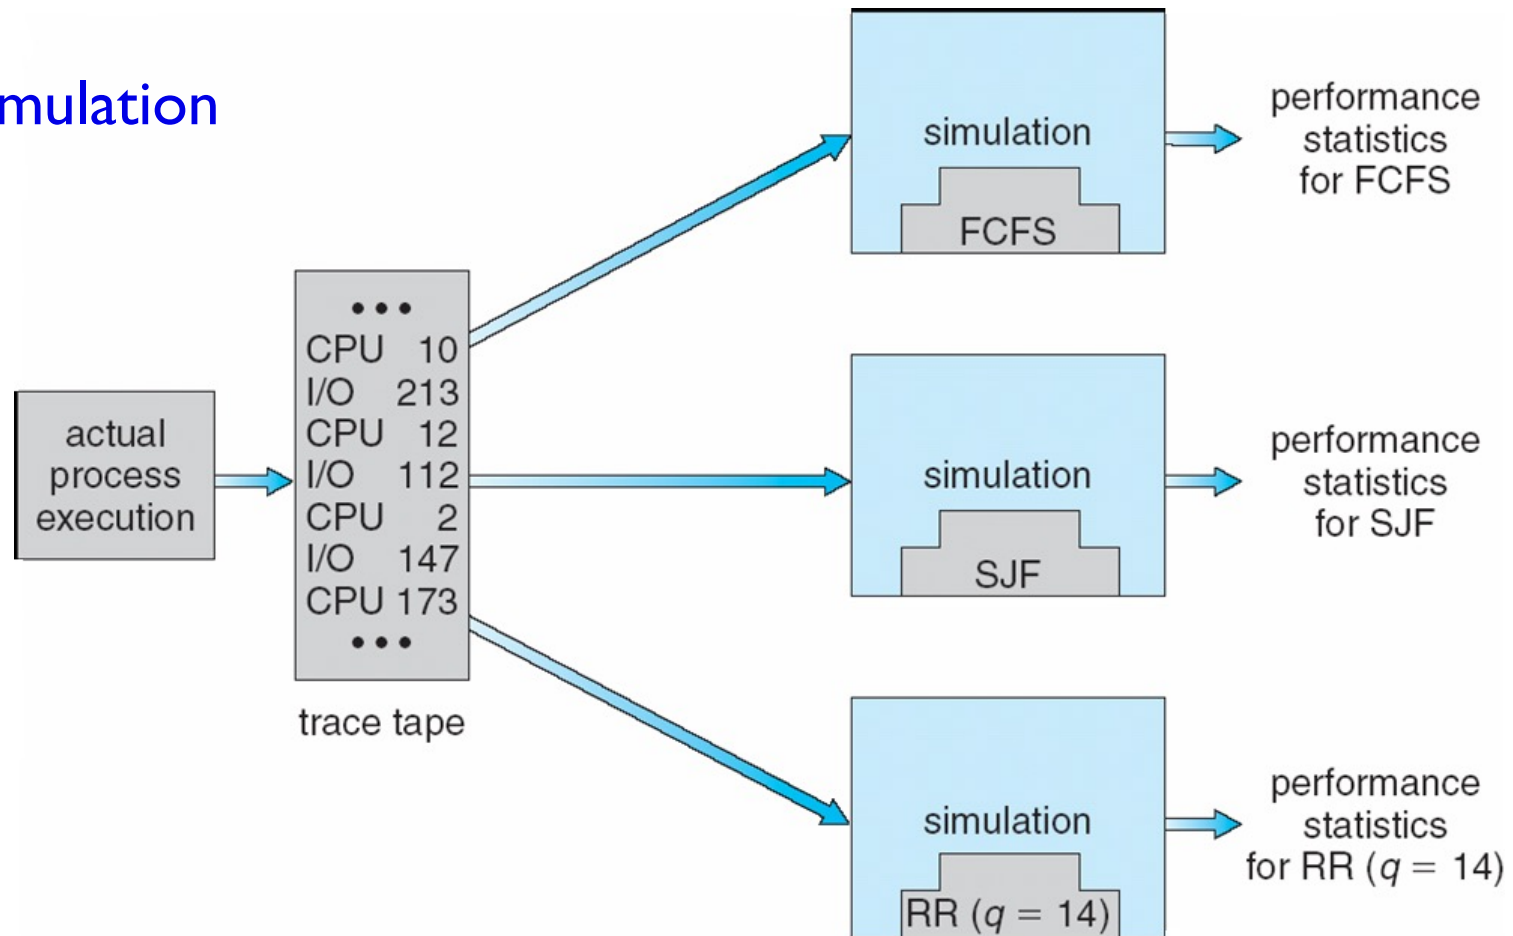
Algorithm Evaluation

Algorithm Evaluation

- Analytic evaluation
 - **Deterministic modeling**: takes a particular predetermined workload and defines the performance of each algorithm for that workload. Valid for a **particular** scenario and input.
 - **Queuing models**.
- Simulation
- Implementation
 - Implement in OS and see how it is performing.

Evaluation of CPU schedulers by Simulation

Simulation



Queueing Models


- Average **throughput**, **utilization**, and **waiting time**, **queue length** can be computed by using **queueing models** for a system. There is a branch of mathematics called queuing theory.
 - One or more CPUs and one or more ready queues make the system.
- **Arrival times of tasks** modeled by a **distribution** (inter-arrival times)
 - *Exponential distribution* is a good approximation for stochastic arrivals (as opposed to deterministic arrivals)
- **Service times** (required cpu times, i.e., burst lengths) modeled by a **distribution**.
 - *Exponential distribution* is a good approximation for required CPU times as well.

Queueing Models

- For a queueing system, which does not stay idle when there are tasks in the system, the following formula, called **Little's formula** holds:

$$N = W \times \lambda$$

- **N** : number of tasks (cpu bursts) in the queue (excluding the one in the cpu)
 - **W** : average waiting time in the queue for a task (not including the time spent in CPU, i.e., the service time)
 - **λ** : arrival rate (number of tasks arriving per second)
- Assuming a **stable** system (service rate $>$ arrival rate)
- Valid** for any arrival and service time distribution.



Queueing Models

- **Example:** Assume arrival rate is 30 tasks/sec. What is the average queue length (excluding the task executing), if a task waits on the average 200 ms in the queue (not including the cpu execution time)?
- **Answer:**
 - $N = W \times \lambda$
 - Therefore, $N = 200 \text{ ms} \times 30 \frac{\text{tasks}}{\text{sec}} = 6$. There are 6 tasks in the system, on the average, that are waiting in the queue.
- **Note** that if *arrivals* and service times are *deterministic*, no queuing and therefore no waiting in the queue will happen. Waiting in the queue happens due to *random* (stochastic) arrivals, i.e., *bursty* arrivals.

References

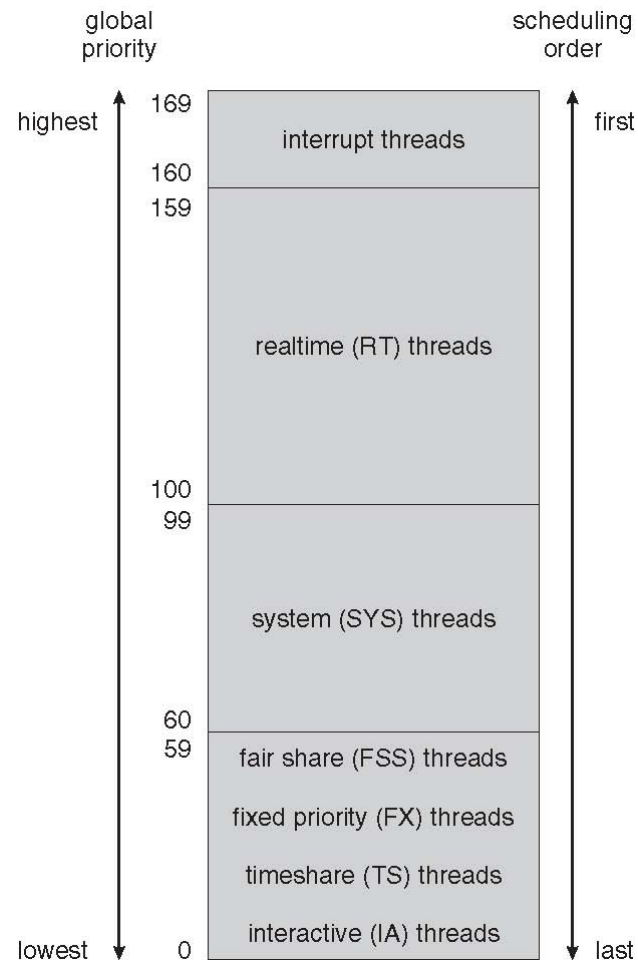
- Operating System Concepts, Silberschatz et al., Wiley.
- Modern Operating Systems, Andrew S. Tanenbaum et al.
- Operating Systems: Principles and Practice, Anderson et al., 2014.
- OSTEP, Arpaci-Dusseau et al.

Additional material (optional)

Solaris Dispatch Table

priority	time quantum	time quantum expired	return from sleep
0	200	0	50
5	200	0	50
10	160	0	51
15	160	5	51
20	120	10	52
25	120	15	52
30	80	20	53
35	80	25	54
40	40	30	55
45	40	35	56
50	40	40	58
55	40	45	58
59	20	49	59

Solaris Scheduling



Windows Priorities

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1