



Bilkent University
Department of Computer Engineering
CS342 Operating Systems

Memory Management Strategies

Last Update: April 24, 2023

Objectives and Outline

Objectives

- Describe ways of **organizing** memory hardware
- discuss various **memory-management techniques**, including paging and segmentation
- Description of the **Intel x86 architecture**, which supports both pure segmentation and segmentation with paging

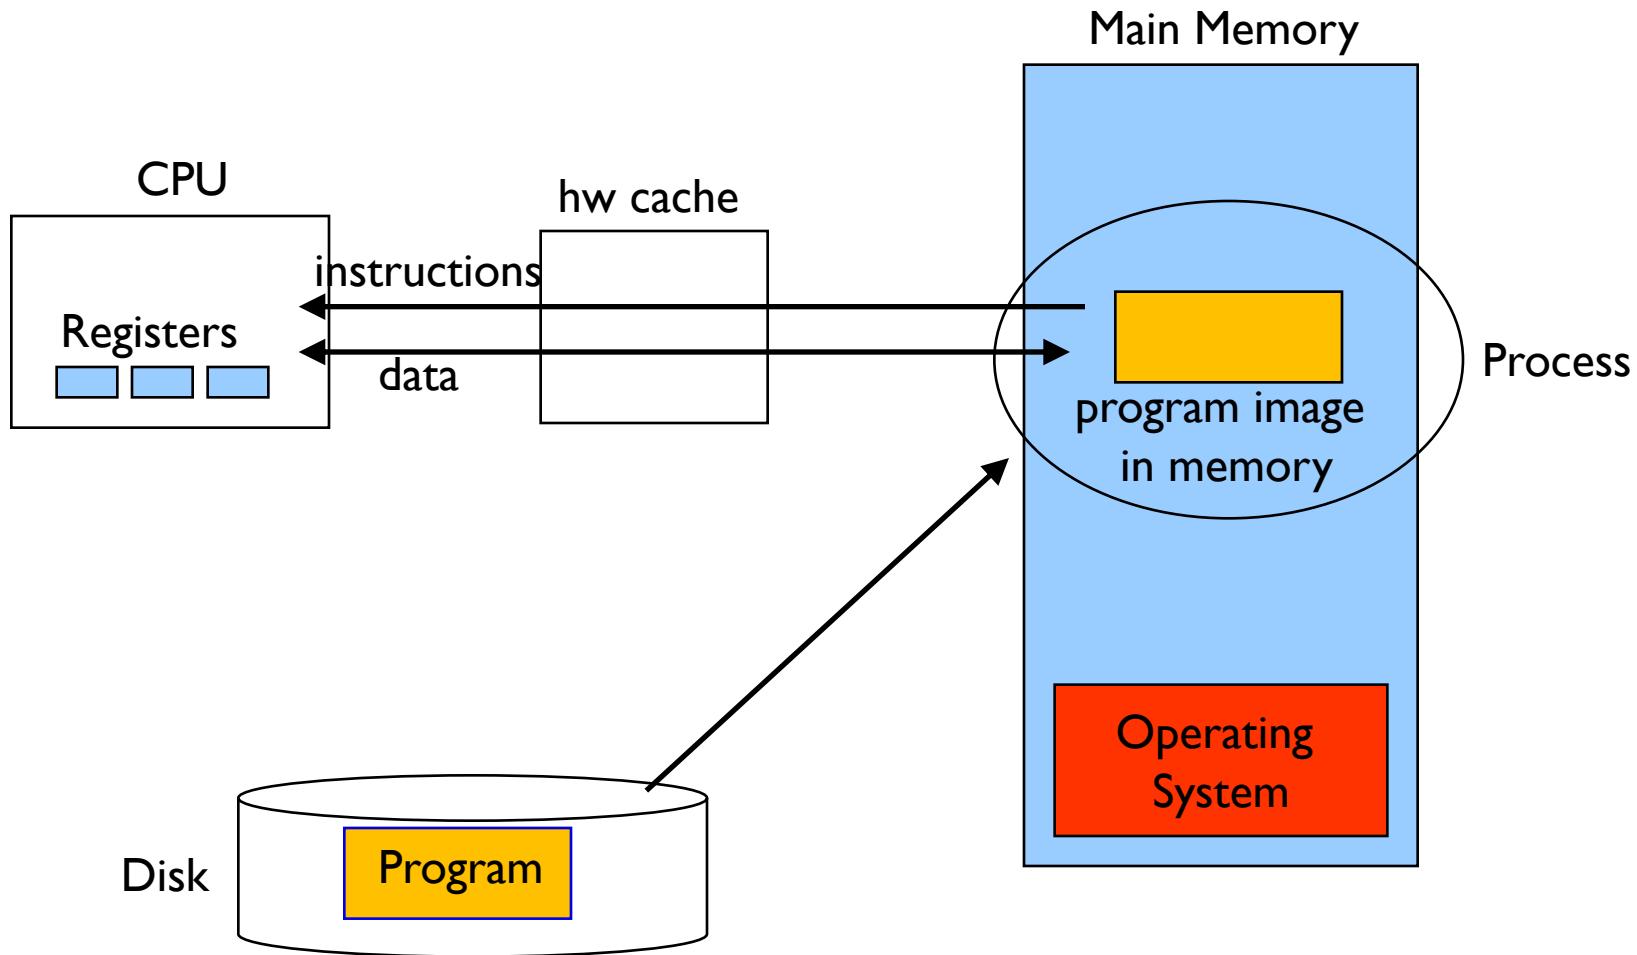
Outline

- **Background**
 - Address space
 - Logical address space
 - MMU
- **Contiguous Memory Allocation**
- **Paging**
 - Structure of the Page Table
- **Segmentation**

Background

- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory (Physical Memory) and registers are only storage CPU can access directly
 - Register access in one CPU clock cycle (or less)
 - Main memory can take many cycles.
 - Cache sits between main memory and CPU registers
- Physical Memory (RAM): a sequence of bytes. Each byte has an address.
- Protection of memory required to ensure correct operation

Background

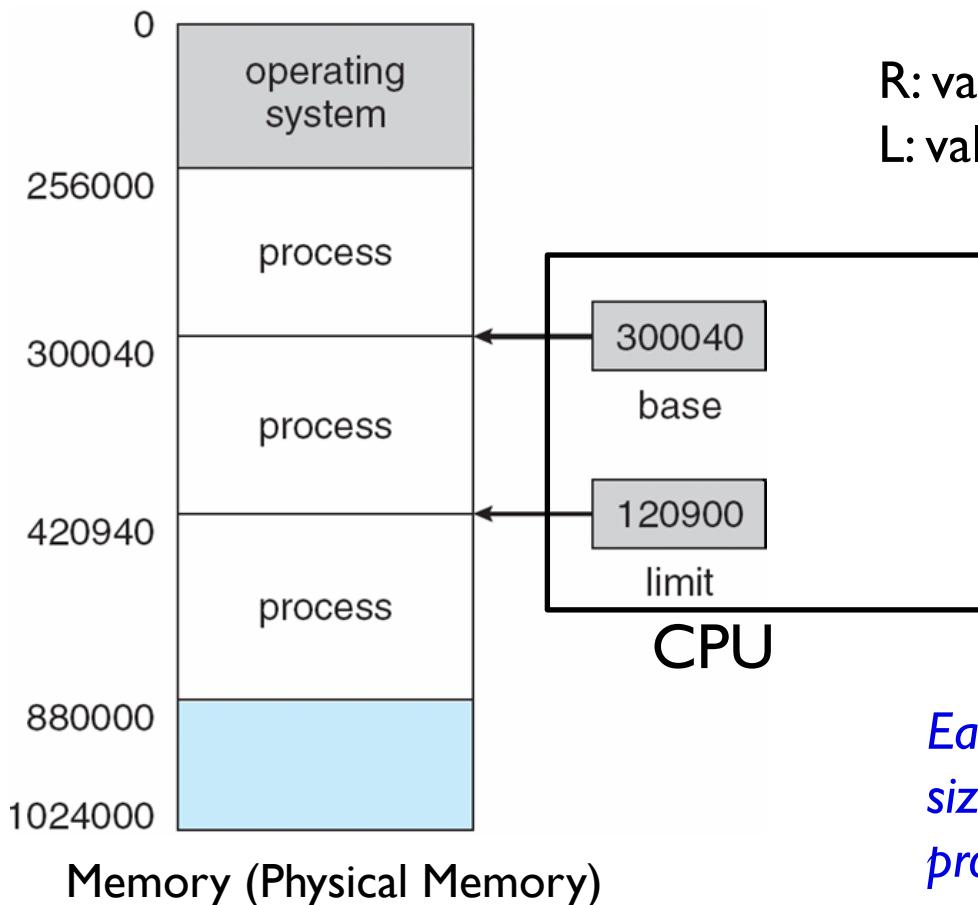


Multiple programs and Protection

- Multiple programs started **share** main memory.
- Programs should not access each other's memory and OS memory. **Protection** needed.
 - OS can and should access everything.
 - A program is allowed to access the physical memory allocated to itself only.
- **Hardware support** needed for protection.
- Having **base** and **limit** registers at CPU can be used for a basic protection mechanism.
 - **Base register**: contains the **start physical address** (**start PA**) of the program to run.
 - **Limit (bound) register**: contains the **size** of the program.

Protection

A pair of **base** and **limit** registers define the **physical address space** of a process running in the CPU.



R: value in base reg (start PA process)
L: value in limit reg (size of process)

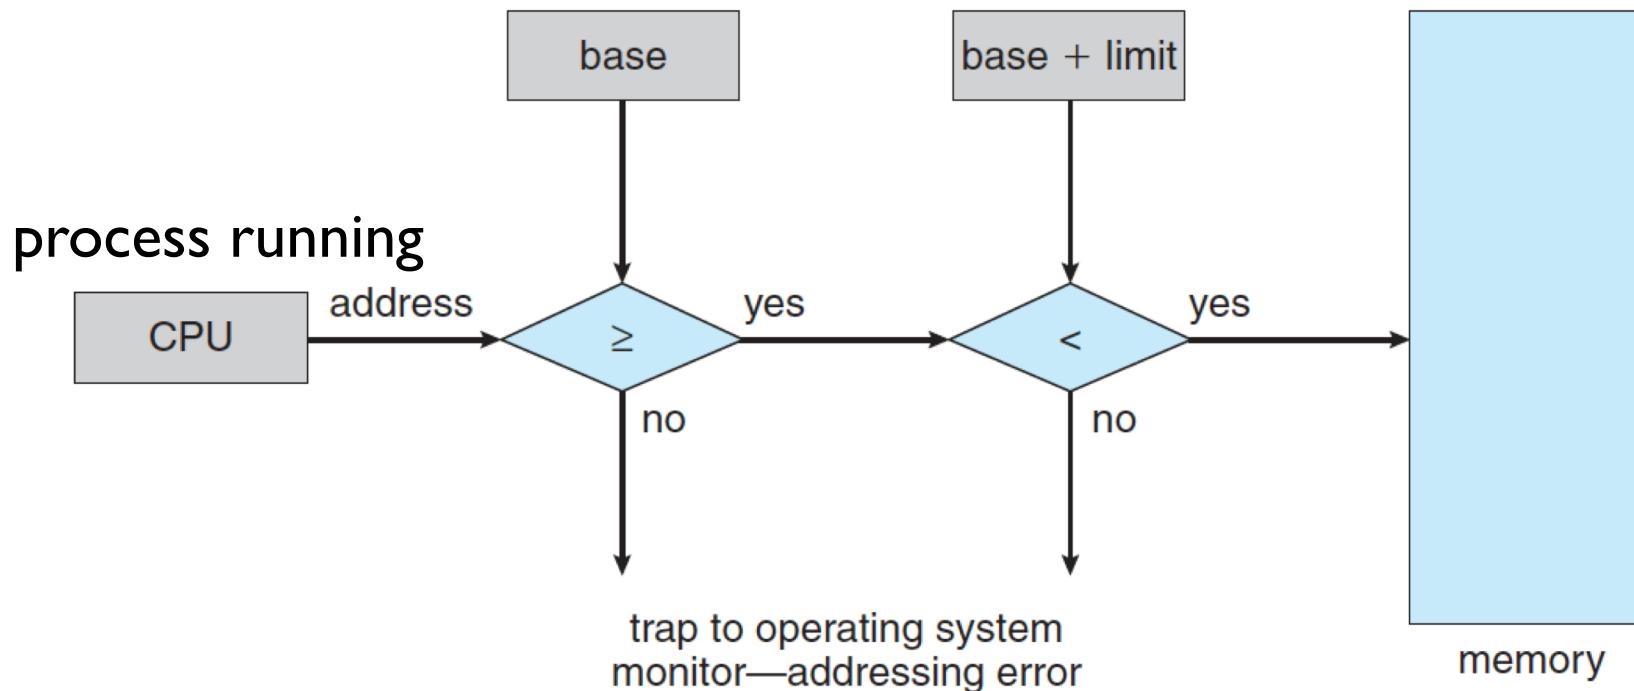
PA range: [R, R+L]

Process can only access
that range.
Access to outside the
range: memory error.

Each process has its own start PA and
size. Can be stored in the PCB of the
process.

Protection

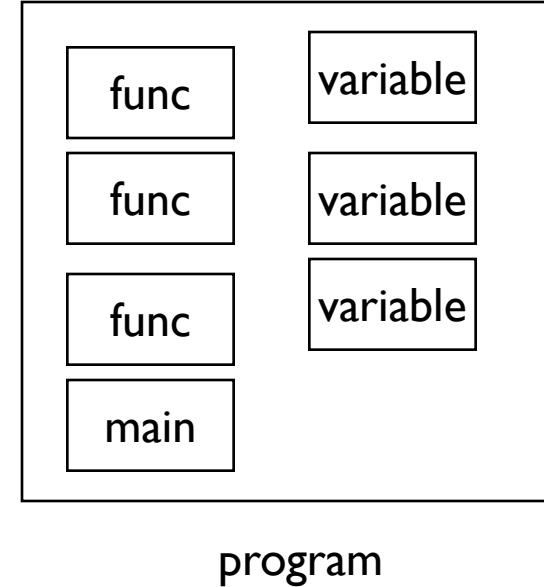
Base and limit registers are set by OS before selected process (by scheduler) starts execution in CPU.



Hardware checks every address.
Access to memory is protected.

Program addresses and memory

- Source code uses symbolic names for program elements (variables, functions, etc.)
- When machine code is generated (via compiler or assembler) numeric **memory addresses are used** for instructions, data, and functions.
- Those addresses can be
 - physical memory addresses (**PA**),
 - or, **logical memory addresses (LA)**.

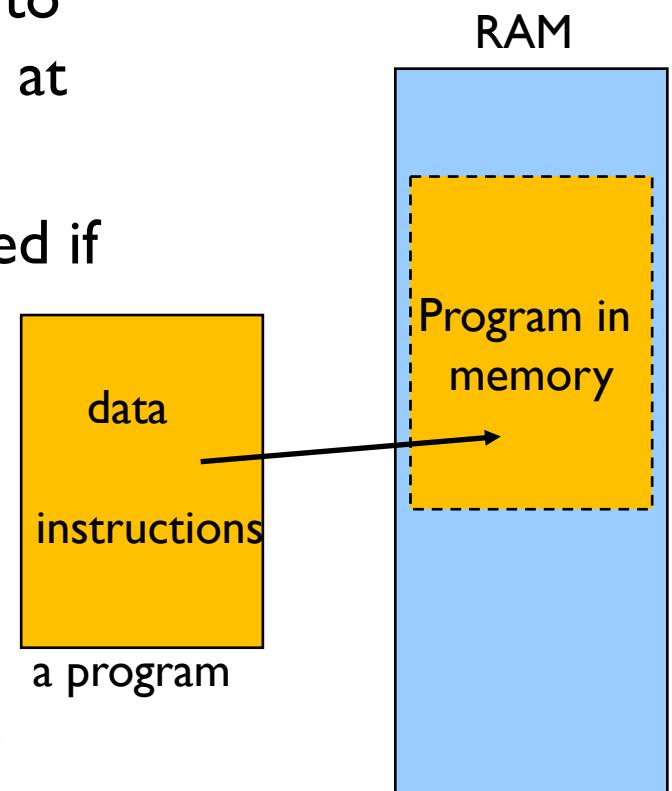


Finally, program elements (instructions, data, etc.) should be loaded to physical locations and will have physical addresses.

Binding of Instructions and Data to Memory

- **Address binding** of instructions and data to (physical) memory addresses can happen at three different stages.

- **Compile time**: absolute code generated if start PA address known.
- **Load time**: compiler generates relocatable code. Loader modifies addresses wrt start PA address.
- **Execution time**: Binding at run time. Program uses logical addresses *all* the time. Mapping from LA to PA done by HW at runtime. Program can be moved at run time.



Program addresses and memory

Assume they are physical addresses

Program

Add	12
Mov	8
...	4
Jump 8	0

RAM

physical addresses
of RAM

	44
	40
	36
	32
	28
	24
	20
	16
Add	12
Mov	8
...	4
Jump 8	0

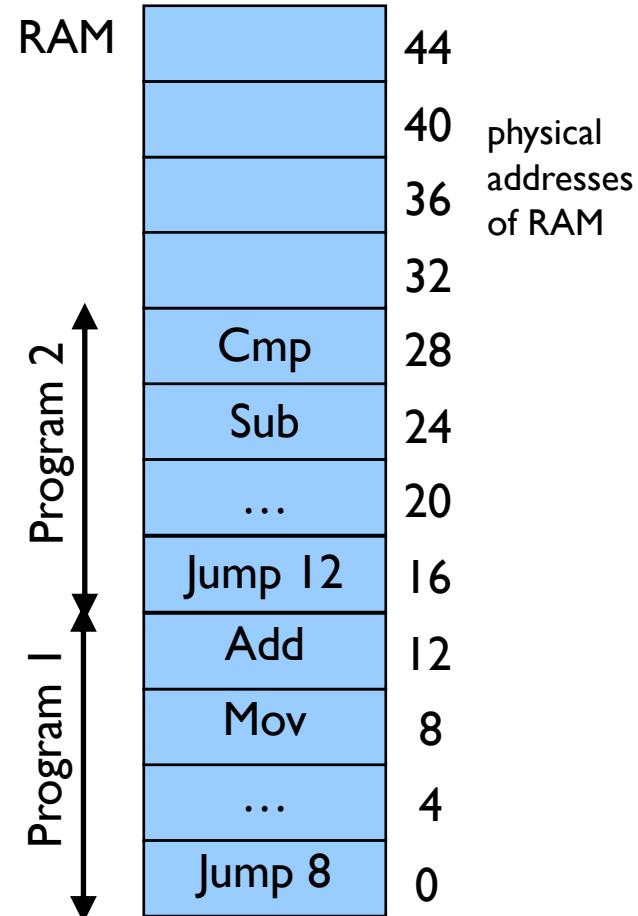
In very early systems they were just **physical memory** addresses. A program has to be loaded to that address to run. **No relocation.**

Program addresses and memory

Program 2 will not work correctly.
Since it is not loaded at address 0.

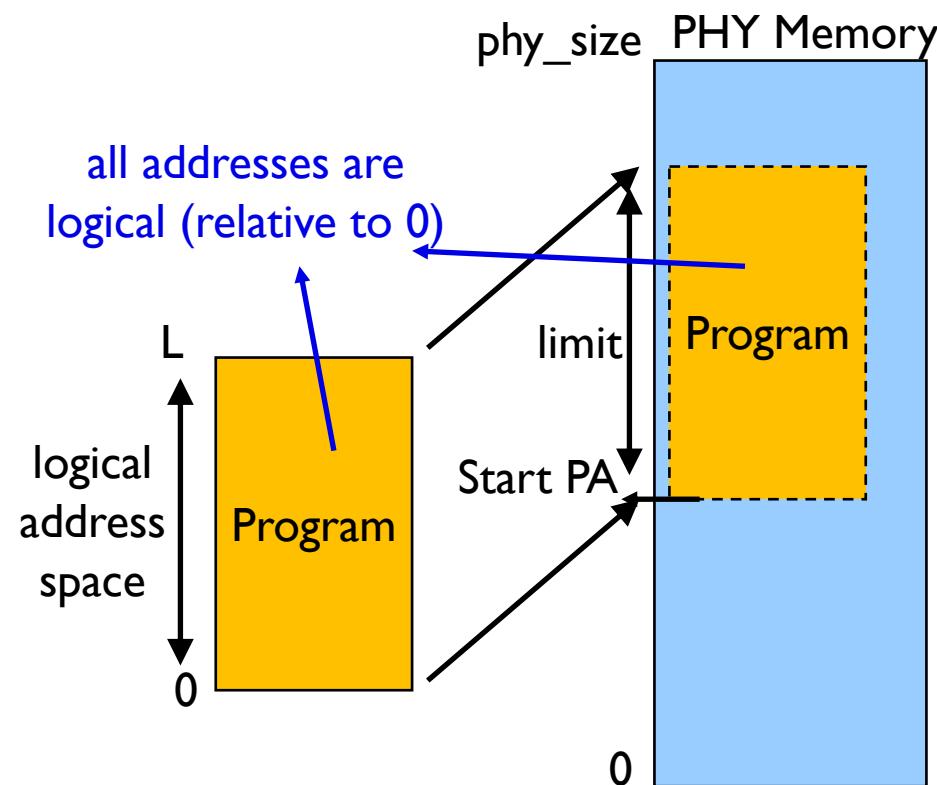
Program 1		Program 2	
Add	12	Cmp	12
Mov	8	Sub	8
...	4	...	4
Jump 8	0	Jump 12	0

Use of physical address in the program is not flexible and complicates the job of the programmer/compiler and also complicates the management of memory (OS).



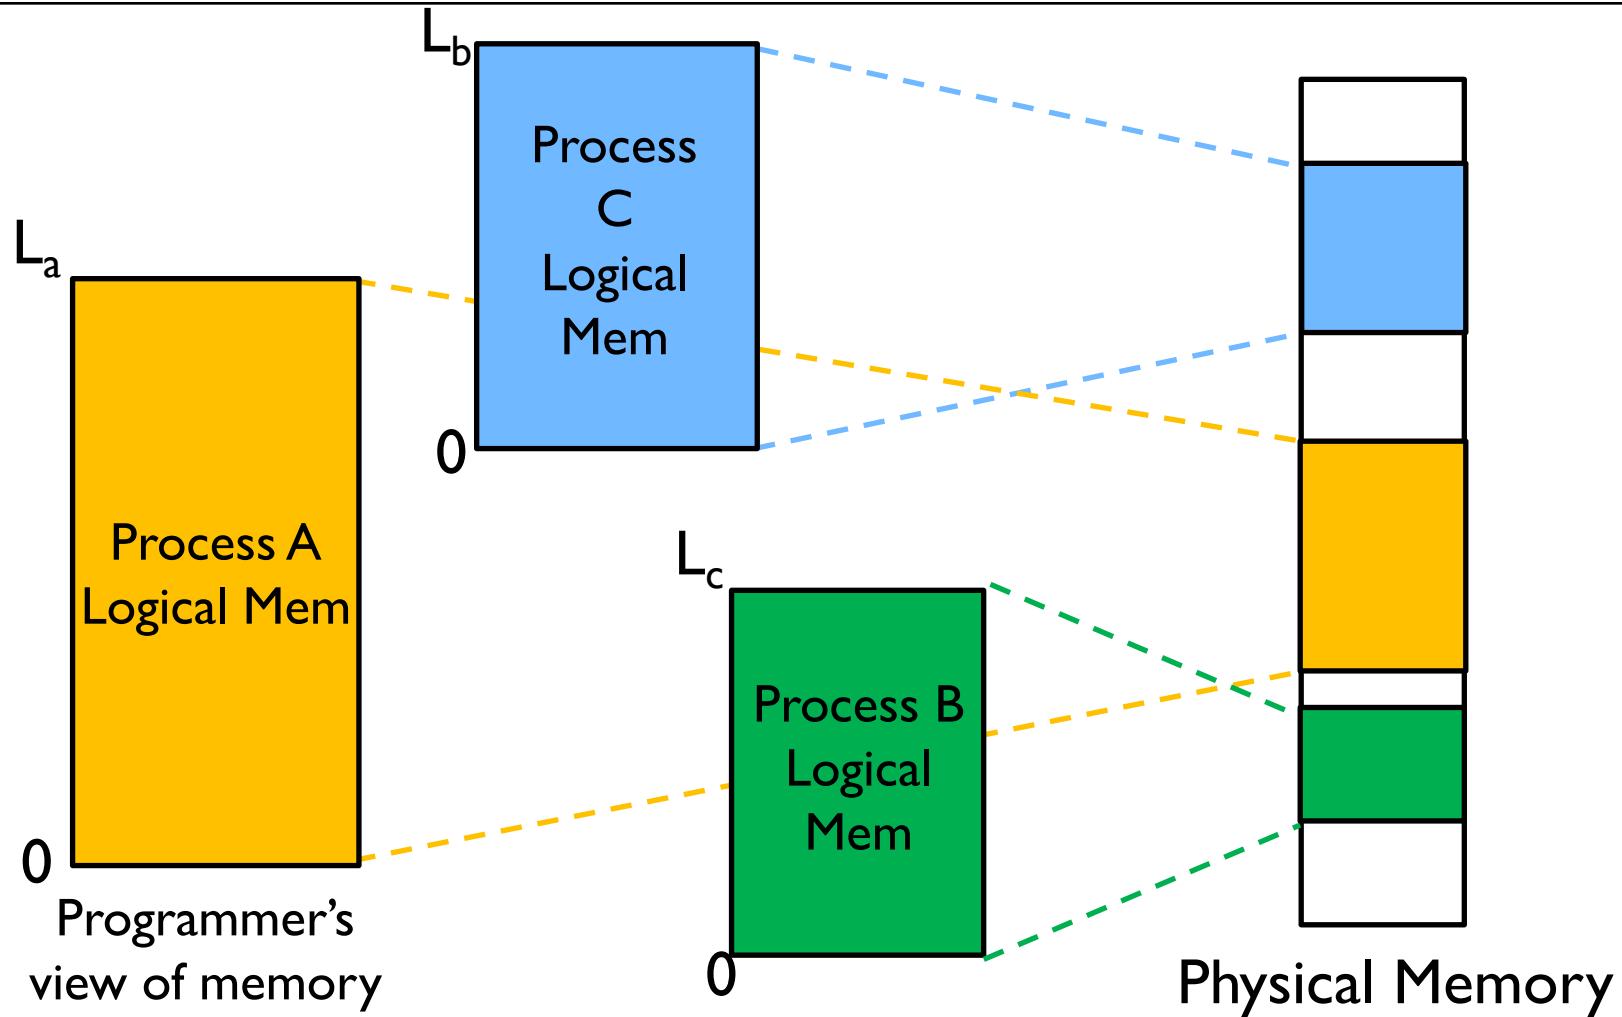
Logical address space concept for a process

- We need **logical address space** (**logical memory**) concept, that is different than the physical RAM (main memory) addresses.
- A program uses logical addresses. They start at 0.
- Set of logical addresses used by the program is its logical address space.
 - Logical address space can be, for example, $[0, L)$.
 - L is size of logical memory.



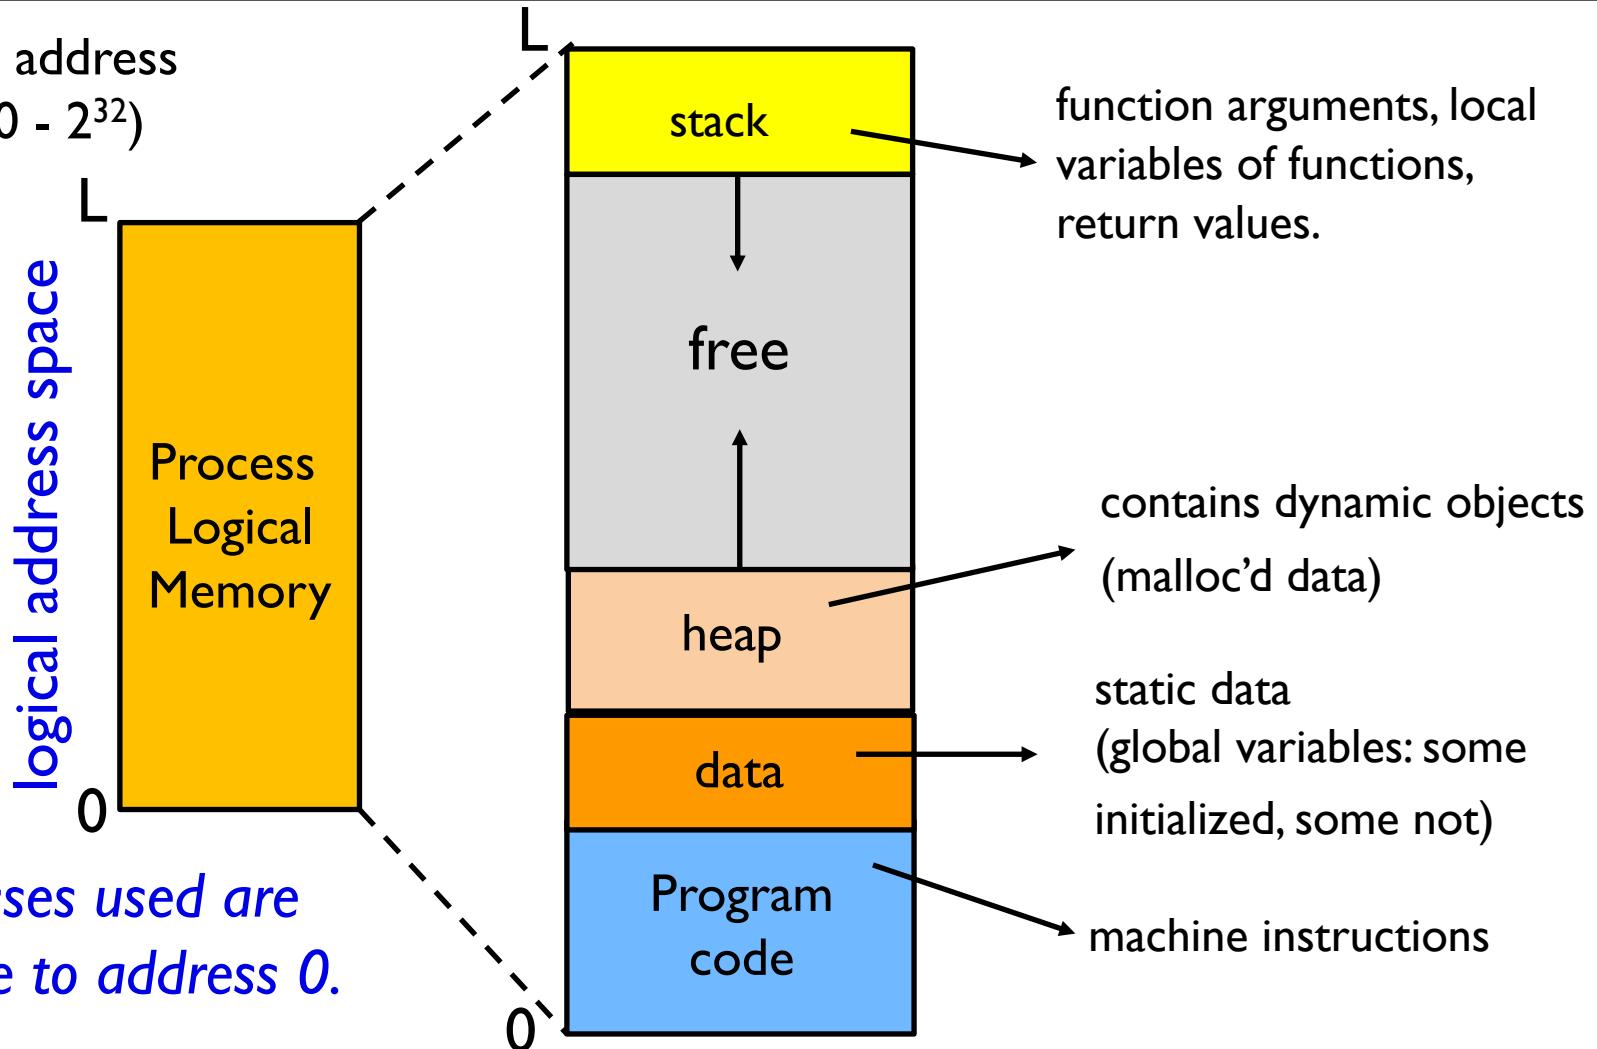
*In this way, memory is **virtualized**. Logical address space has to be **mapped** somewhere in physical memory when program started.*

Logical address space concept



Logical (virtual) memory layout of a process

example address
space: $[0 - 2^{32}]$

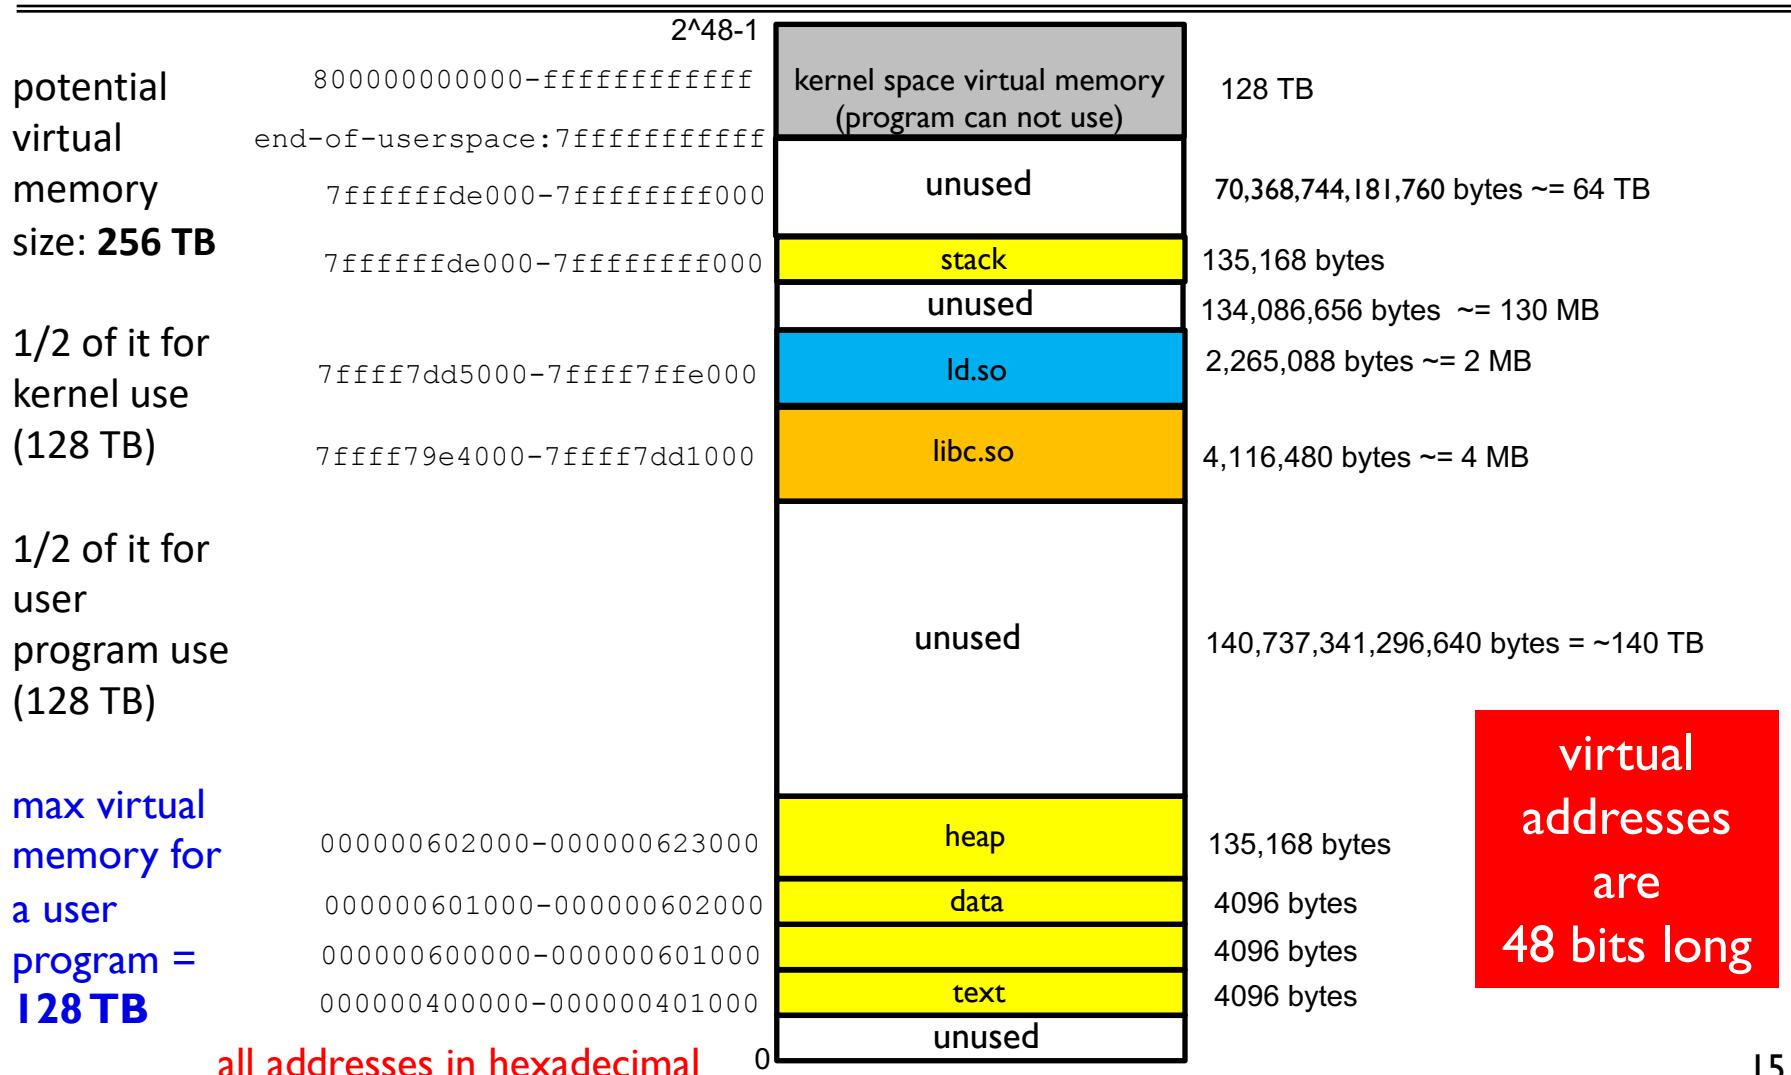


*addresses used are
relative to address 0.*

a real-life example: a **64-bit machine** – **x86-64 / Linux OS**

VM layout for a process (small process)

very large logical address space (potential VM is too big)



VM layout of a process

64-bit machine – x86-64 / Linux OS

output of:

`cat /proc/pid/maps`

```
00400000-00401000 r-xp 00000000 08:01 50081283 [code]
00600000-00601000 r--p 00000000 08:01 50081283 [code-system]
00601000-00602000 rw-p 00001000 08:01 50081283 [data]
00602000-00623000 rw-p 00000000 00:00 0 [heap]
7fffff79e4000-7fffff7bcb000 r-xp 00000000 08:01 67109118 /lib/x86_64-linux-gnu/libc-2.27.so
7fffff7bcb000-7fffff7dcb000 ---p 001e7000 08:01 67109118 /lib/x86_64-linux-gnu/libc-2.27.so
7fffff7dcb000-7fffff7dcf000 r--p 001e7000 08:01 67109118 /lib/x86_64-linux-gnu/libc-2.27.so
7fffff7dcf000-7fffff7dd1000 rw-p 001eb000 08:01 67109118 /lib/x86_64-linux-gnu/libc-2.27.so
7fffff7dd1000-7fffff7dd5000 rw-p 00000000 00:00 0
7fffff7dd5000-7fffff7dfc000 r-xp 00000000 08:01 67108949 /lib/x86_64-linux-gnu/ld-2.27.so
7fffff7fdf000-7fffff7fe1000 rw-p 00000000 00:00 0
7fffff7ff7000-7fffff7ffa000 r--p 00000000 00:00 0 [vvar]
7fffff7ffa000-7fffff7ffc000 r-xp 00000000 00:00 0 [vdso]
7fffff7ffc000-7fffff7ffd000 r--p 00027000 08:01 67108949 /lib/x86_64-linux-gnu/ld-2.27.so
7fffff7ffd000-7fffff7ffe000 rw-p 00028000 08:01 67108949 /lib/x86_64-linux-gnu/ld-2.27.so
7fffff7ffe000-7fffff7fff000 rw-p 00000000 00:00 0
7fffffffde000-7fffffff000 rw-p 00000000 00:00 0 [stack]
```

virtual memory regions of a process with id=pid

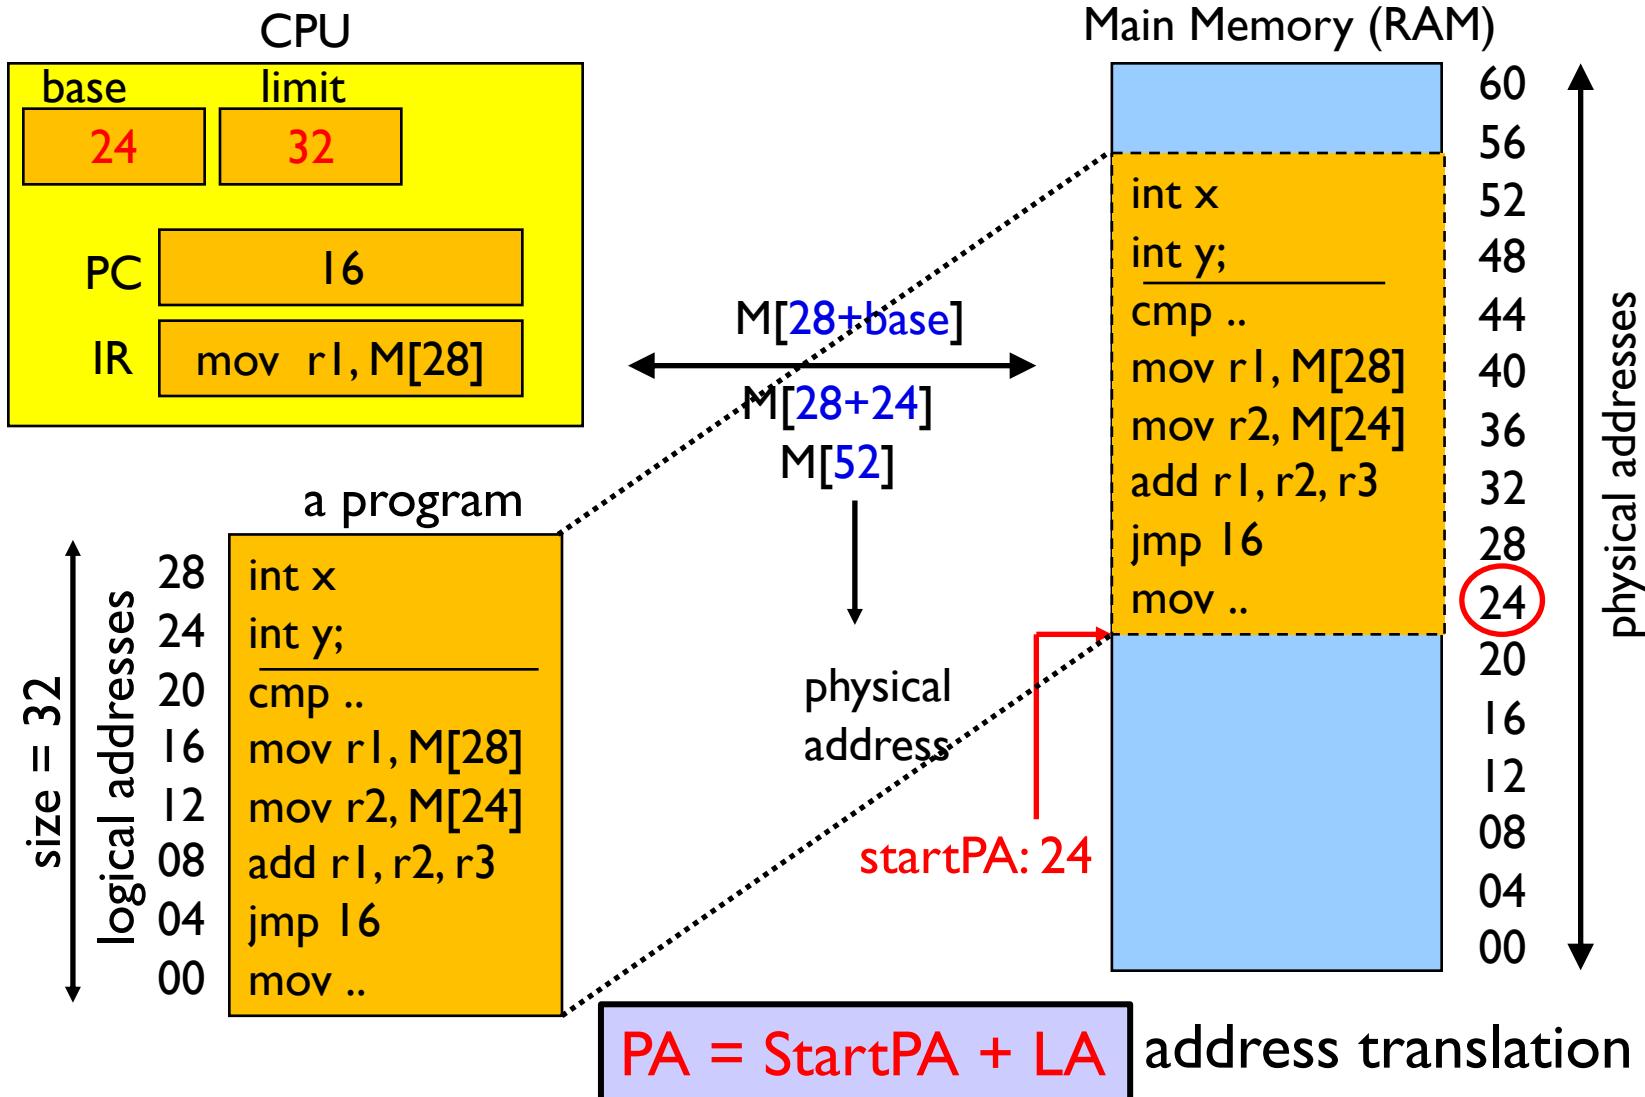
We can also use `pmap` command to get VM layout information for a process.

Logical vs. Physical Address Space

- The **concept of a logical address** space that is bound to a separate physical address space is **central** to proper memory management
 - **Logical address (LA)** – generated by the CPU; also referred to as **virtual address**.
 - **Physical address (PA)** – address seen by the memory unit (RAM)
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes (program in memory has absolute addresses).
- Logical (virtual) and physical addresses differ in execution-time address-binding scheme. Program in memory has logical addresses.

Logical and physical addresses

a small example (stack, etc., not shown)



Memory-Management Unit (MMU)

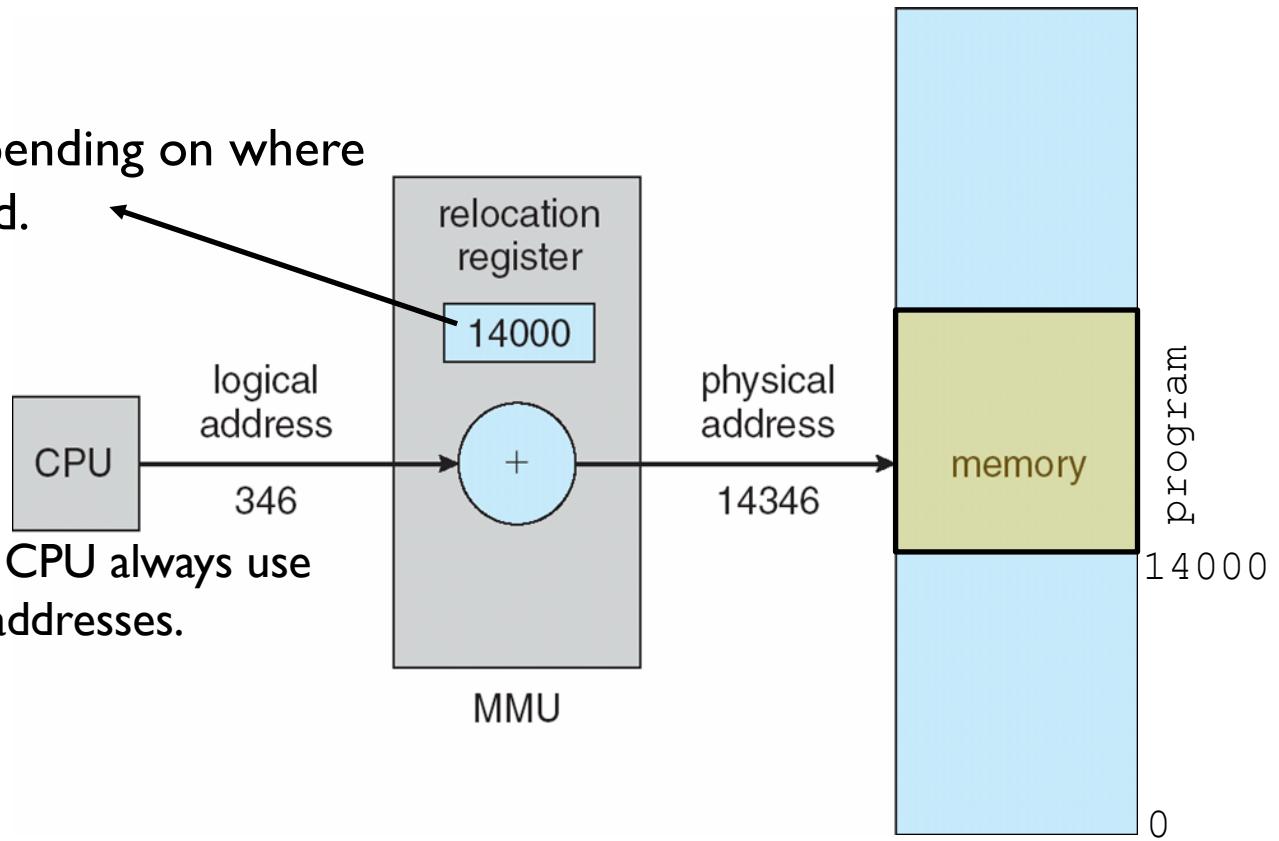
- Hardware device that translates logical (virtual) addresses to physical addresses.
- In MMU, the start PA address of the program is loaded into **base register** (which is also called **relocation register**).
- The value in the **base register** is **added** to every logical address to obtain physical address.
- The user program deals with *logical* addresses; it never sees the *real* physical addresses.
 - *Pointer values in a C program are all virtual addresses.*
Programmer does not see physical addresses.
- **Base (relocation) register** together with **limit (bound) register** is used for both **protection** and **relocation**.

Address translation by HW

HW translates each *logical address* to a *physical address* on-the-fly at runtime.

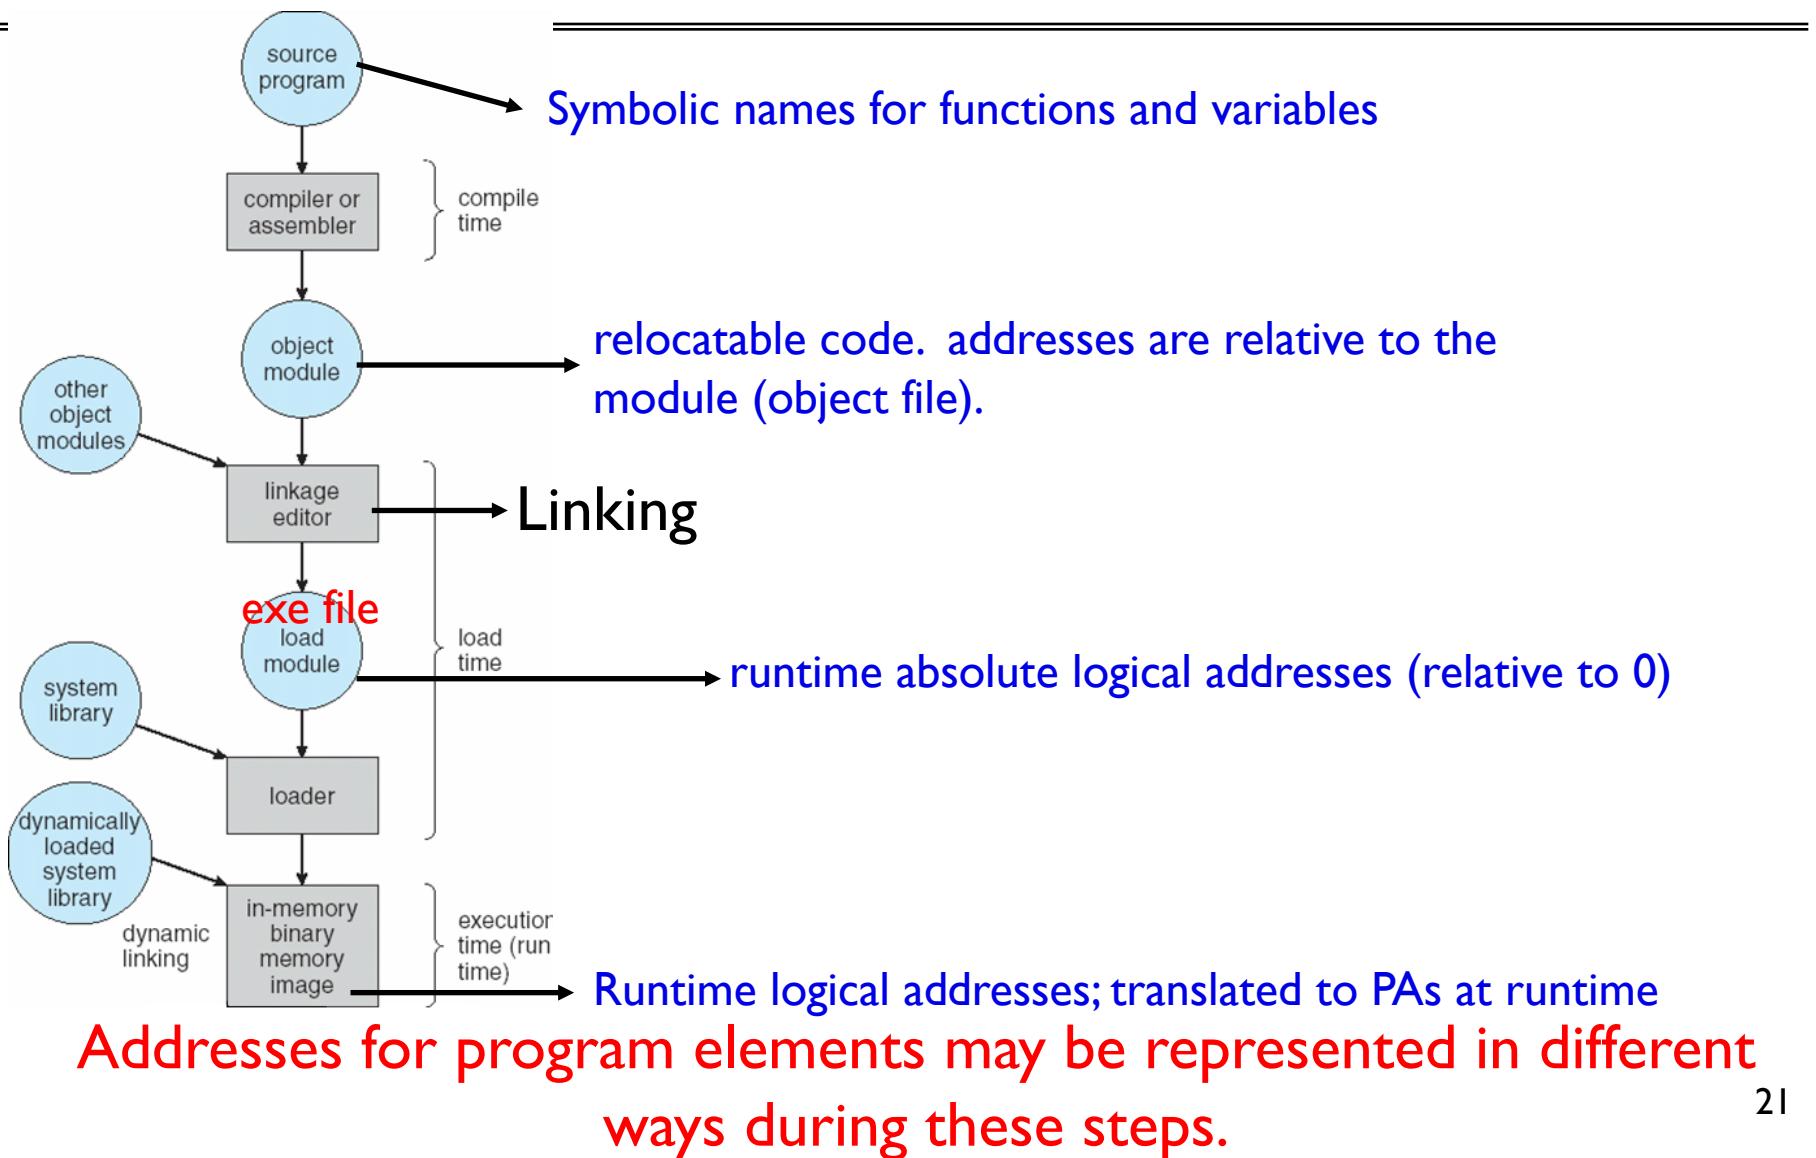
set up by OS depending on where program is loaded.

Programs and CPU always use logical addresses.



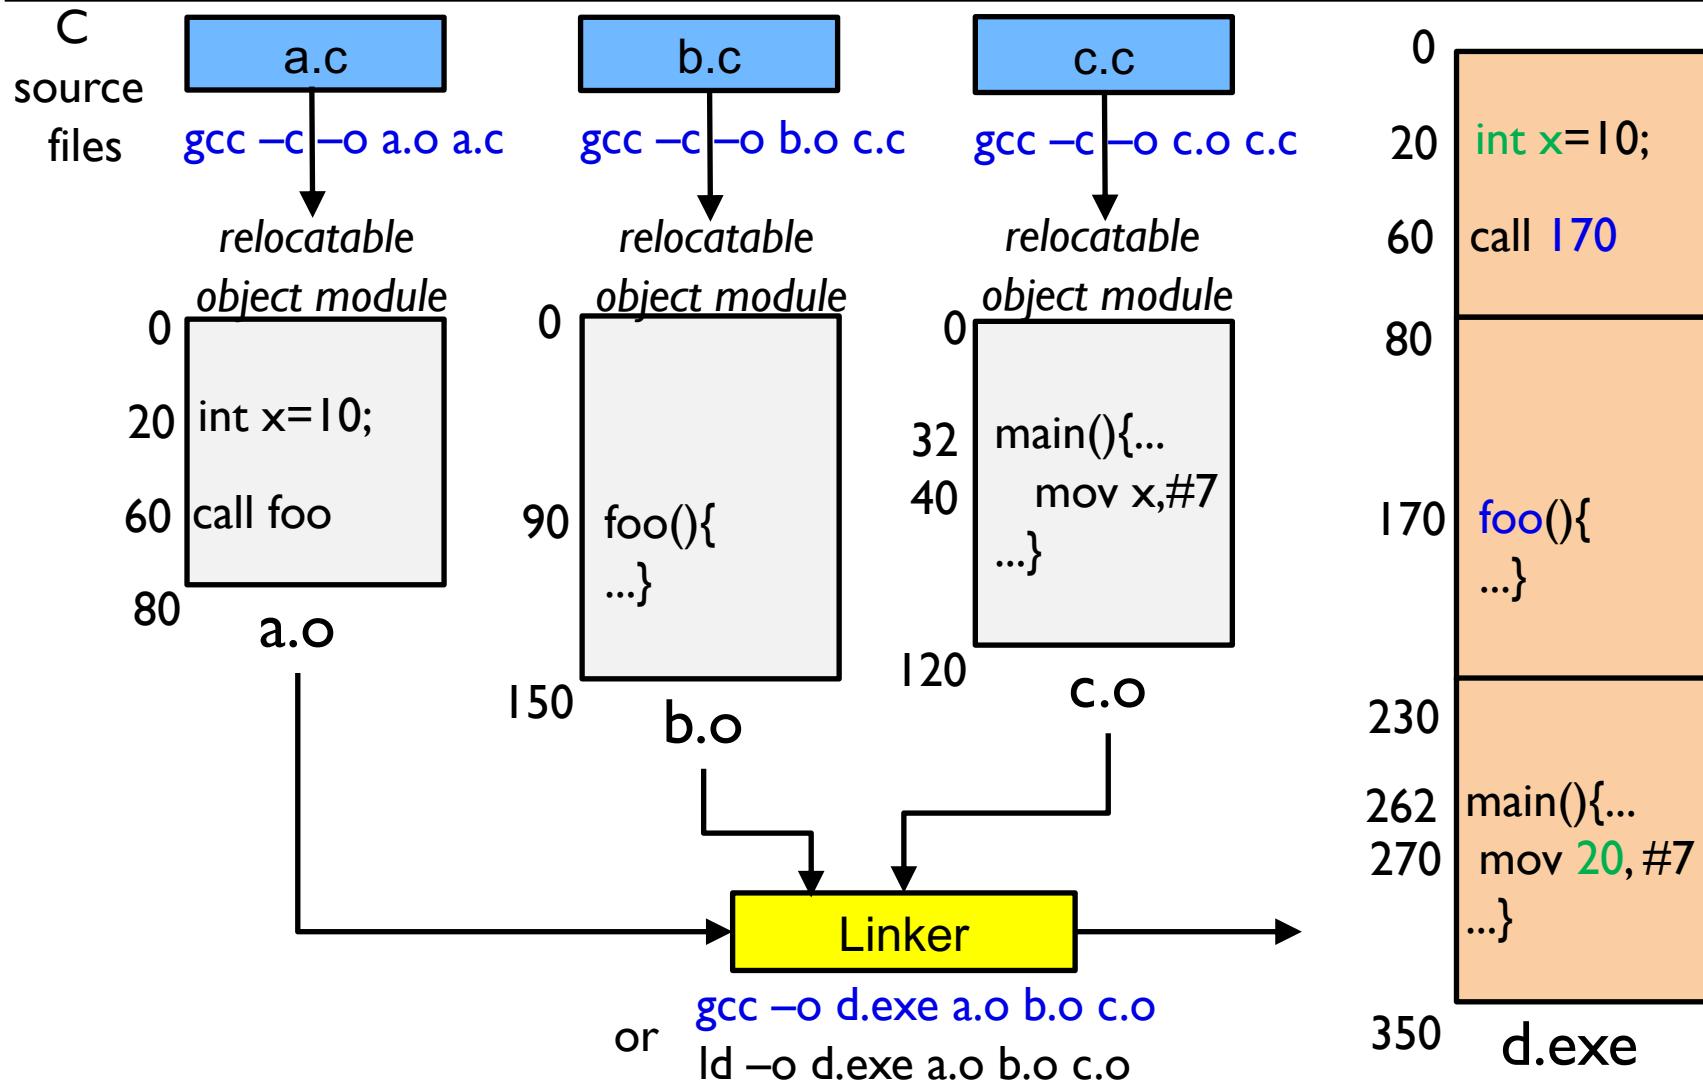
relocation register enables program to be loaded at any location in main memory. Program can be moved as well, if needed.

Multistep Processing of a User Program



a tiny example for linking of multiple modules of a program

(not real code; pseudo-code - to illustrate the idea)



Dynamic Loading

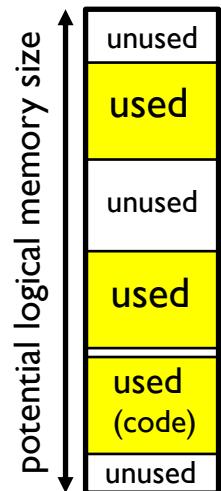
- Routine is not loaded until it is called.
- Better memory-space utilization; unused routine is never loaded
 - Useful when large amounts of code are needed to handle infrequently occurring cases.
- No special support from the operating system is required,
 - implemented through program design.

Dynamic Linking

- Linking postponed until execution time
- Small piece of code, *stub*, used to locate the appropriate memory-resident **library routine**.
- Stub replaces itself with the address of the routine, and executes the routine
- **Dynamic linking** is particularly useful for *libraries*.
 - Standard C library is shared library that is dynamically linked, not statically linked.
 - You can link statically if you want. (`gcc -static`)
- Also known as **shared libraries**.

Basic Memory Allocation Strategies

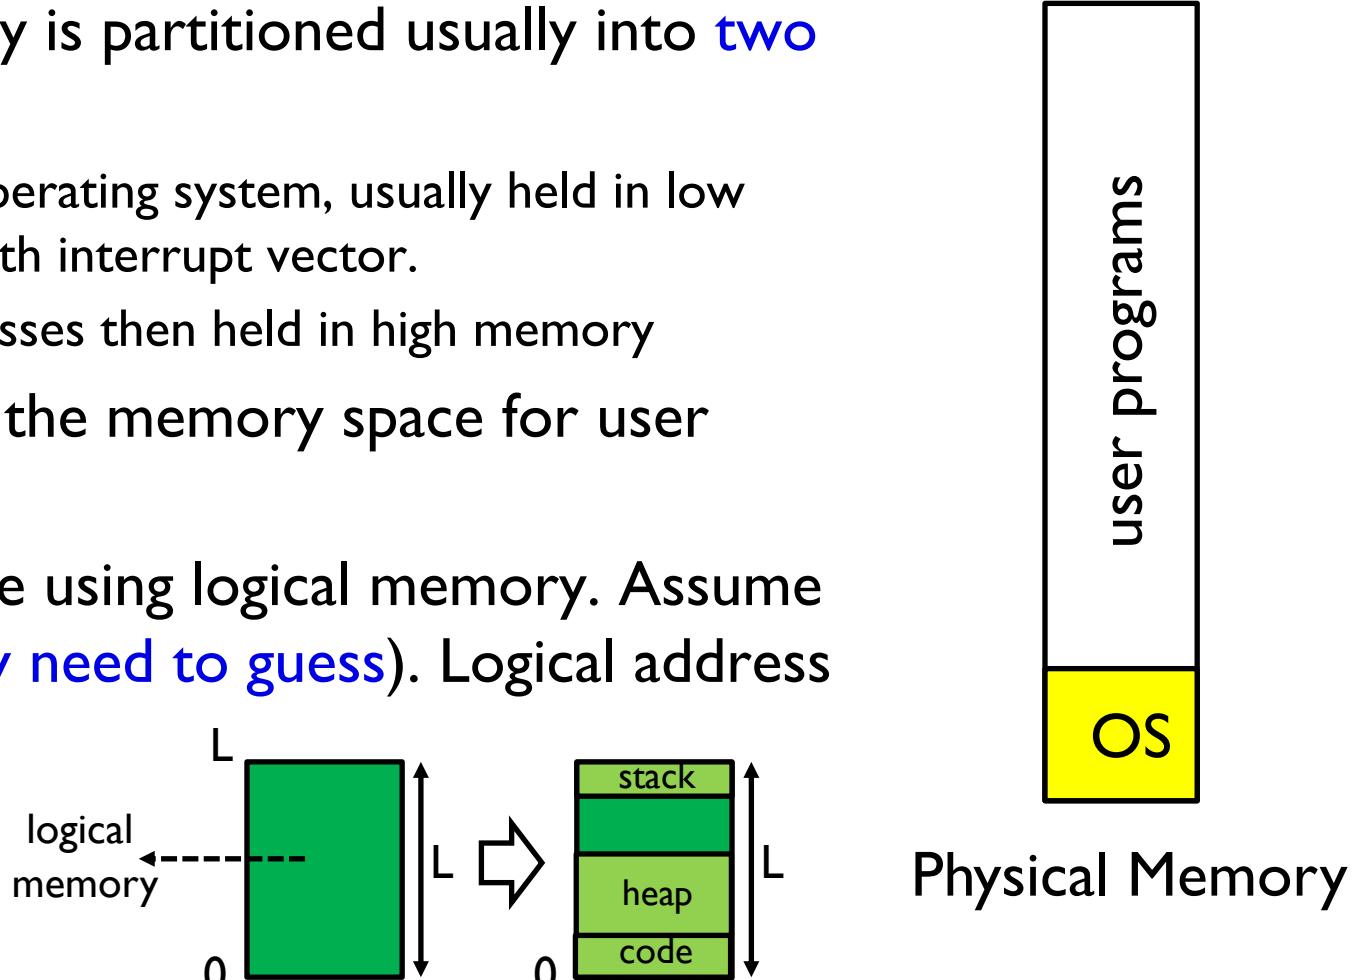
- We will see three basic main memory allocation strategies
 - 1) Contiguous allocation (not much used)
 - 2) Segmentation (used in some system)
 - 3) Paging (heavily used in modern systems)
- In this chapter we will assume that all the **used** logical memory content of a process **must be** in main memory for the process to run (**no partial** loading of **used** regions).
 - note that a program may not use all its potential logical memory. a portion of it may be used. For example, logical memory in a 32 bit system can be as big as 4 GB.
- *In the next chapter (virtual memory), we will relax this assumption (partial loading will be possible to run).*



Contiguous Memory Allocation

Contiguous Allocation

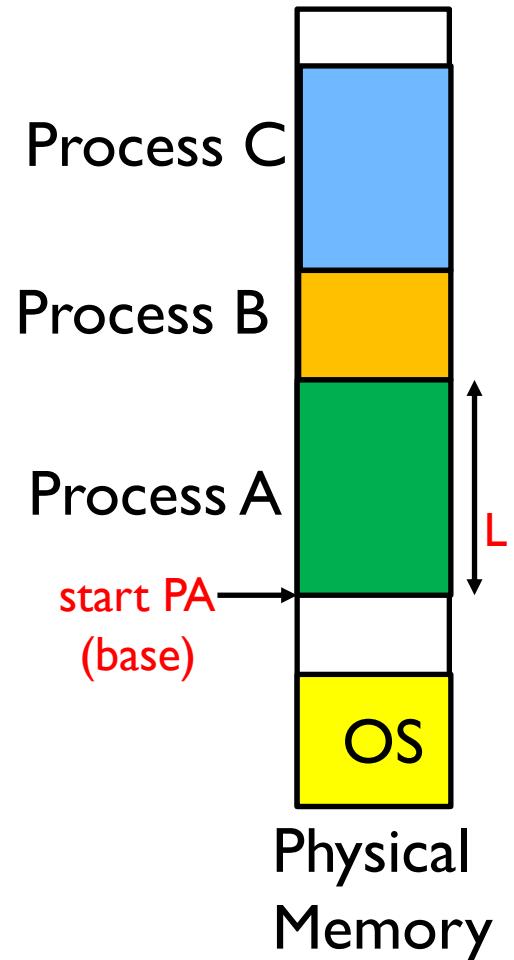
- Main memory is partitioned usually into two partitions:
 - Resident operating system, usually held in low memory with interrupt vector.
 - User processes then held in high memory
- OS manages the memory space for user programs.
- Processes are using logical memory. Assume size = L (may need to guess). Logical address range: [0,L).



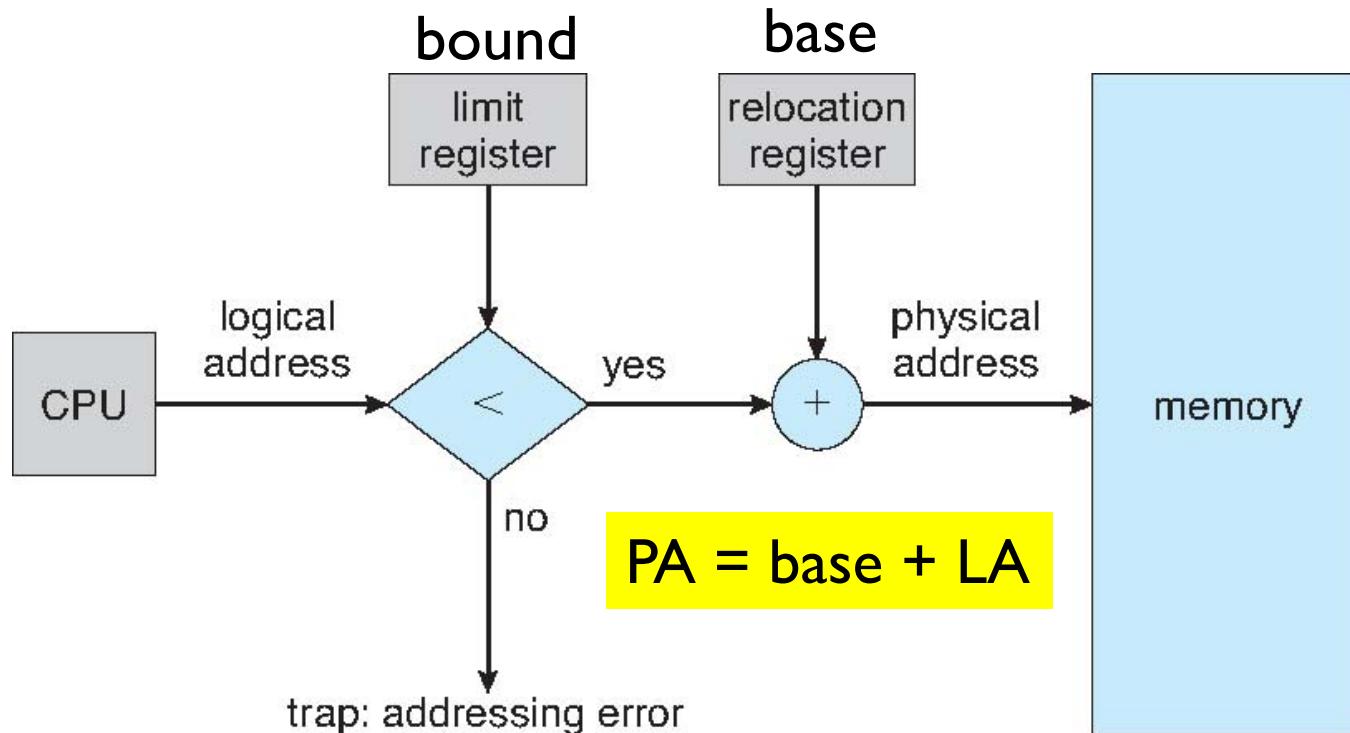
We assume logical memory of a process fits into physical memory

Contiguous Allocation

- A process is allocated a **contiguous region** of physical memory (of size L) by OS.
- OS loads the program from executable file into allocated space.
- **PCB** of the process includes the **start PA** and **size (L)**.
- When process is scheduled, **start PA** and **size loaded** into relocation (**base**) and limit registers of CPU.
 - **relocation (base) register** contains smallest PA.
 - **limit (bound) register** contains range of logical addresses.



HW support for contiguous allocation in terms of relocation and limit registers

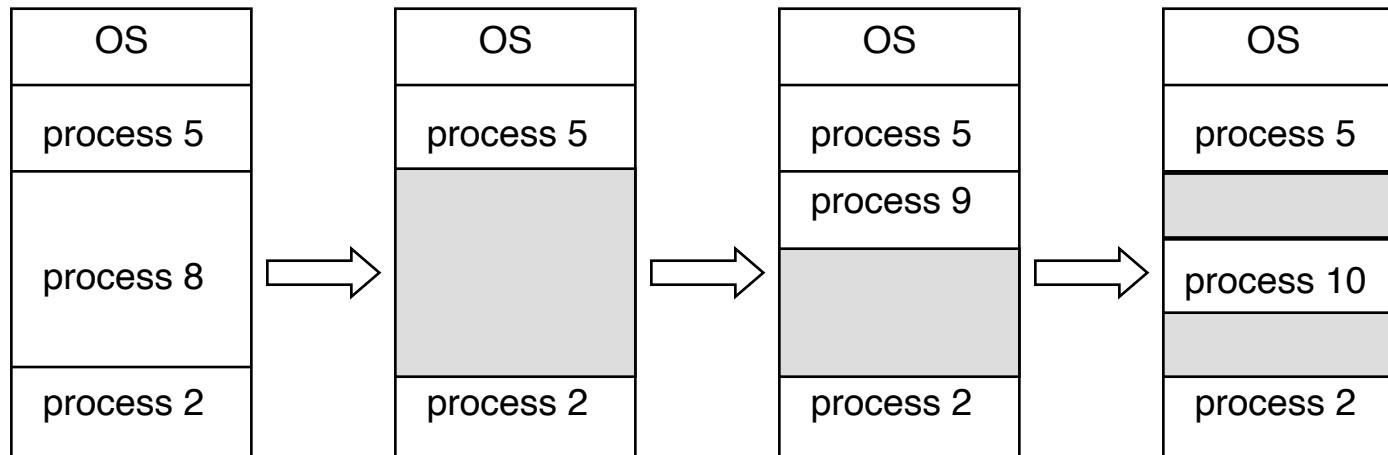


Hardware (MMU) translates logical addresses to physical addresses
dynamically at run time.
Also checks the addresses (protection).

In this way, **protection** and **relocation** provided.

Contiguous Allocation

- Memory is partitioned to programs, in a **contiguous** manner.
 - Each partition contains a **program** or is **free (hole)**.
 - **Hole** – block of available (free) memory; holes of various size are scattered throughout memory after some time.



Contiguous Allocation

- Operating system maintains information about:
 - a) allocated partitions (in PCBs)
 - b) free partitions (holes)
- When a process arrives, it is allocated memory from a hole large enough to accommodate it.
- How to satisfy a request of size n from a list of free holes?
 - This is **dynamic storage allocation problem**.
 - Various strategies exists.

Dynamic Storage-Allocation Problem

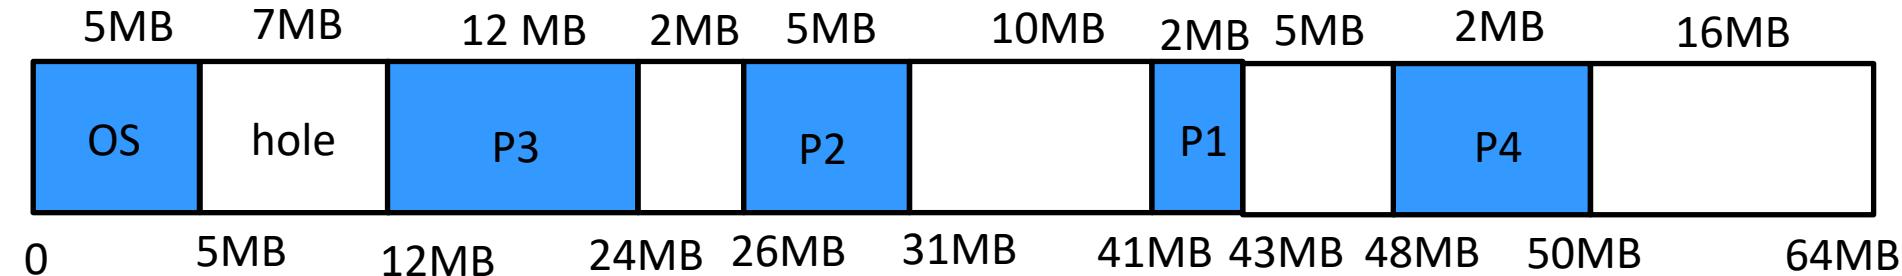
Searching and selecting a free block (hole) for request of size n

- **First-fit:** allocate the *first* hole that is big enough.
- **Next-fit:** use first-fit, but next time continue from where the previous search stopped.
- **Best-fit:** allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size.
 - Produces the smallest leftover hole.
- **Worst-fit:** allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole

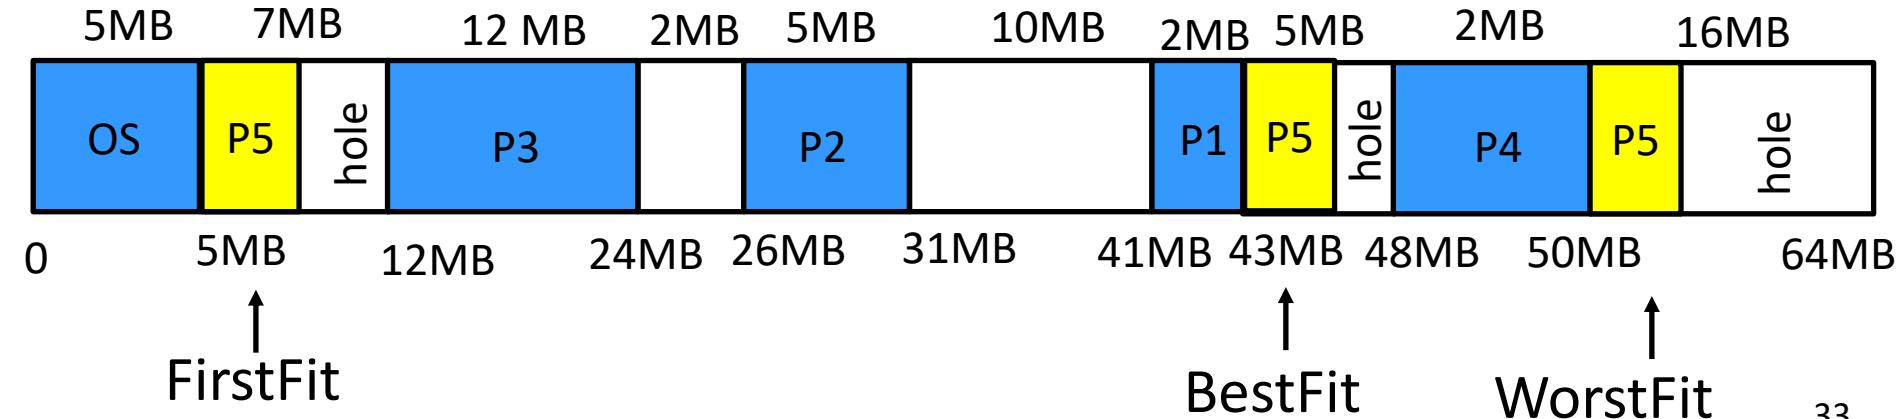
First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Small example: where a new program will be loaded with FirstFit, BestFit, WorstFit?

memory state at t1

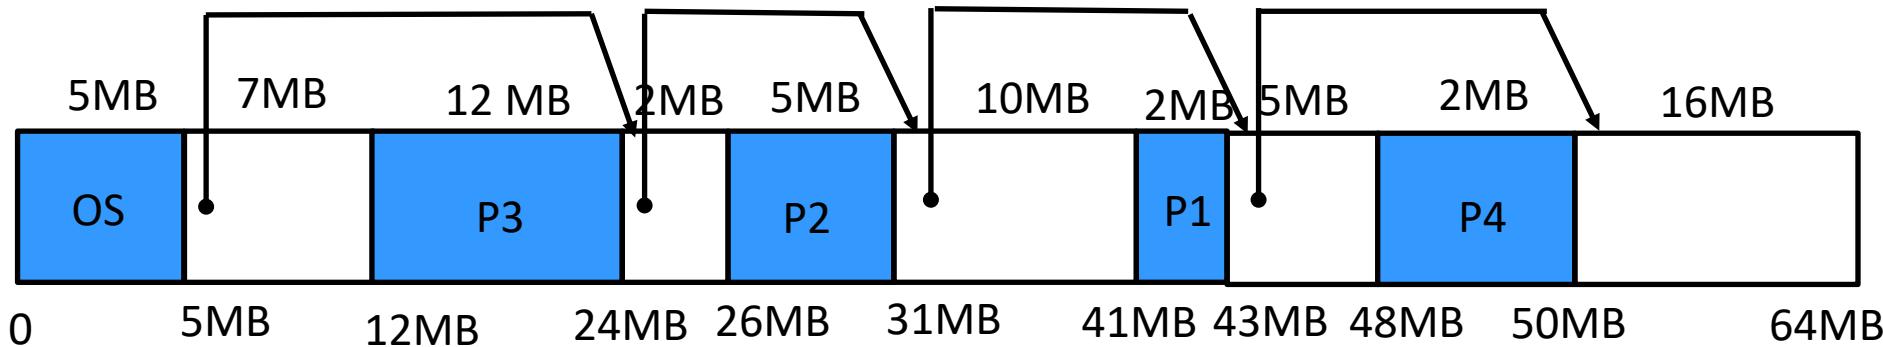


New program P5 will be loaded at t2. size of program P5=3MB



Small example: keeping track of holes.

Use holes (free blocks) themselves. No need for another structure.



each hole will have such a header at the beginning.

```
struct hole_hdr {  
    struct hole_hdr *next_hole; // pointer to the next hole  
    struct hole_hdr *prev_hole; // pointer to the prev hole  
    int size; // size of this hole  
};
```

we can have a singly or doubly linked list of holes.

Fragmentation

- Contiguous allocation causes external fragmentation.
- External Fragmentation – total memory space exists to satisfy a request, but it is not contiguous.
- Internal Fragmentation – allocated memory may be slightly larger than the requested memory; this size difference is memory internal to a partition (allocation), but not being used.
- Reduce external fragmentation by compaction
 - Shuffle memory contents to place all free memory together in one large block

Segmentation

Segmentation

- Memory-management scheme that supports user's/programmer's view of memory
- A program is a collection of segments (logical contiguous units)
 - A segment is a logical program unit such as code part, or heap part of the program.
 - Course-grained segments: data segment, code segment, stack segment, heap segment. A few segments.
 - Fine grained segments: main program, procedure, function, method, object, local variables, global variables, common block, stack, symbol table, arrays, ... In this case, hundreds of segments may exist in a program.
- Segments are of variable size.

User's View of a Program

Which segments will a program have is *decided by the programmer or compiler.*

instruction address consists of:

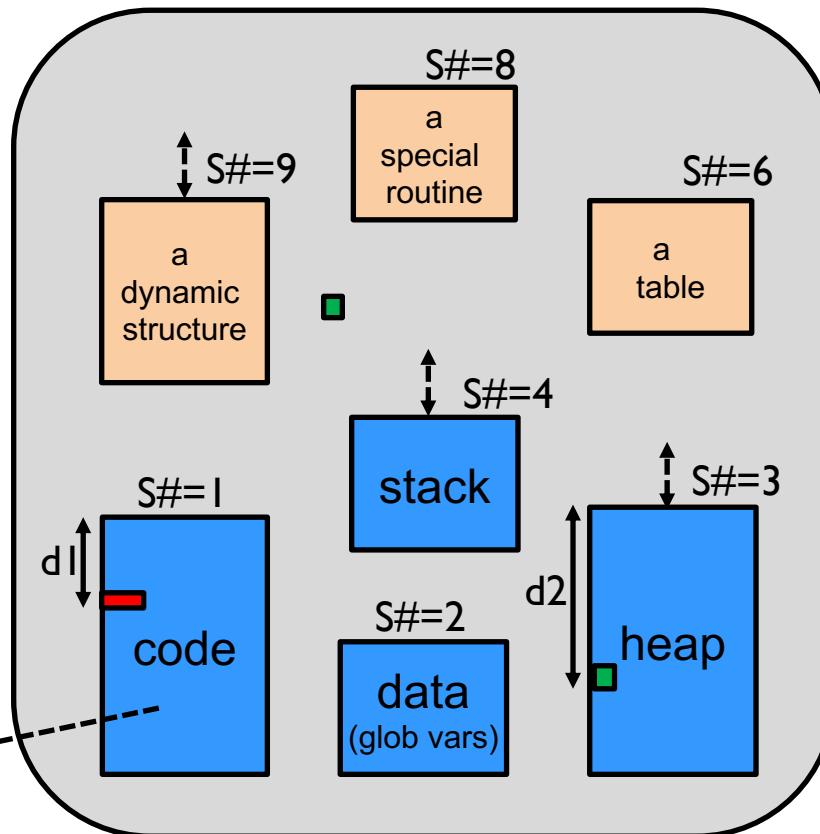
- 1) segment number (#1)
- 2) displacement inside segment (d1)
 $LA = <1, d1>$



object address consists of:

- 1) segment number (#3)
- 2) displacement inside segment (d2)
 $LA = <3, d2>$

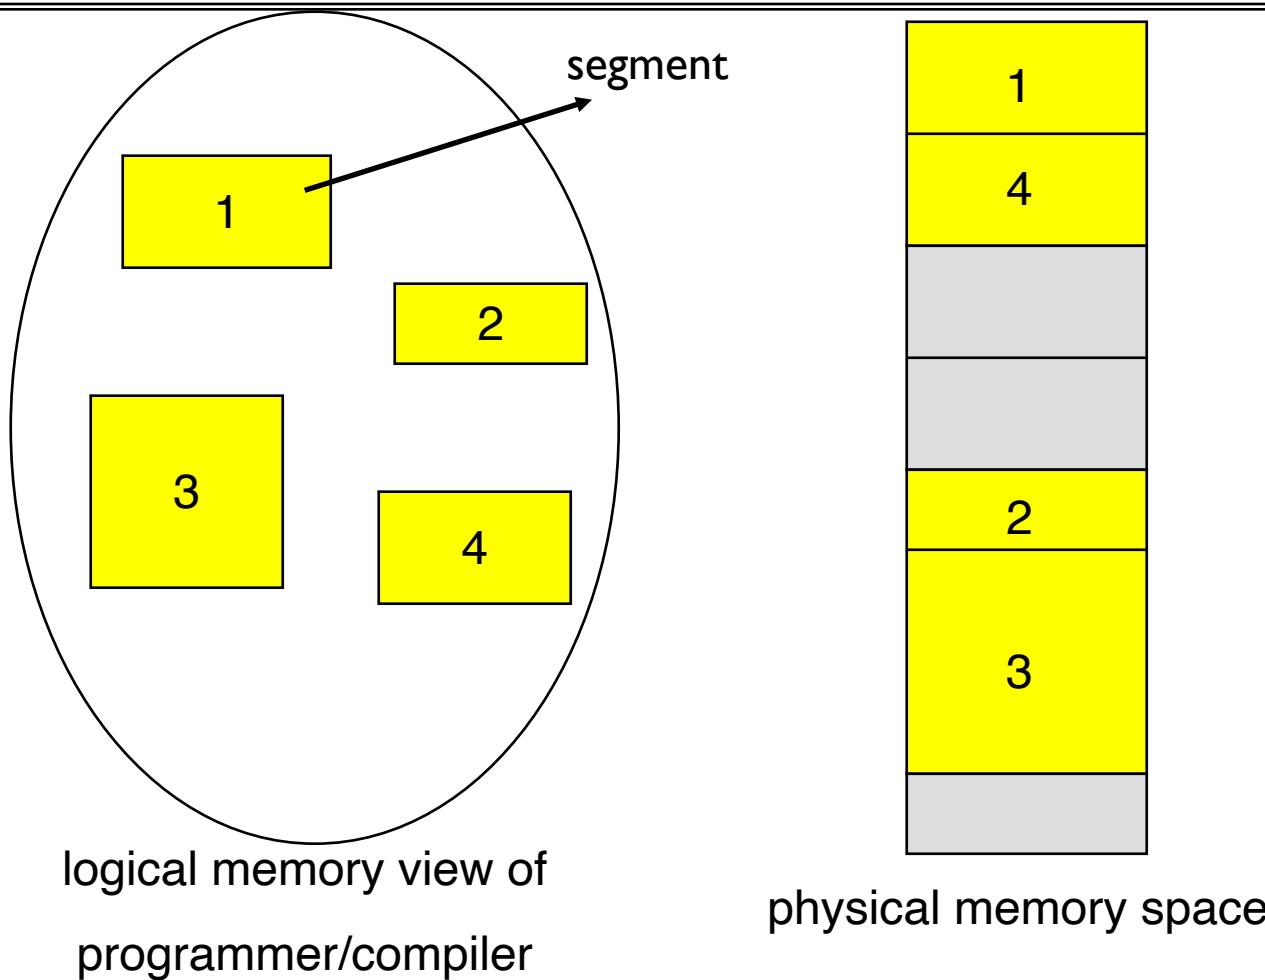
a segment



logical address space (logical memory)

segments
(of variable size).
some static.
some may be
dynamic.

Logical View of Segmentation

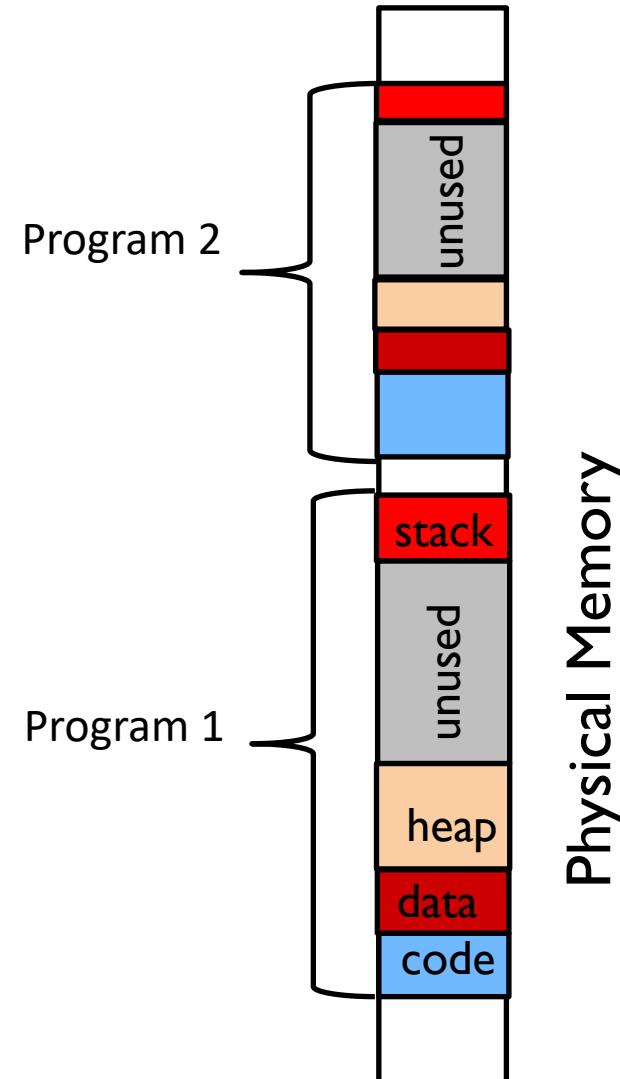


OS needs to allocate contiguous memory to each segment.

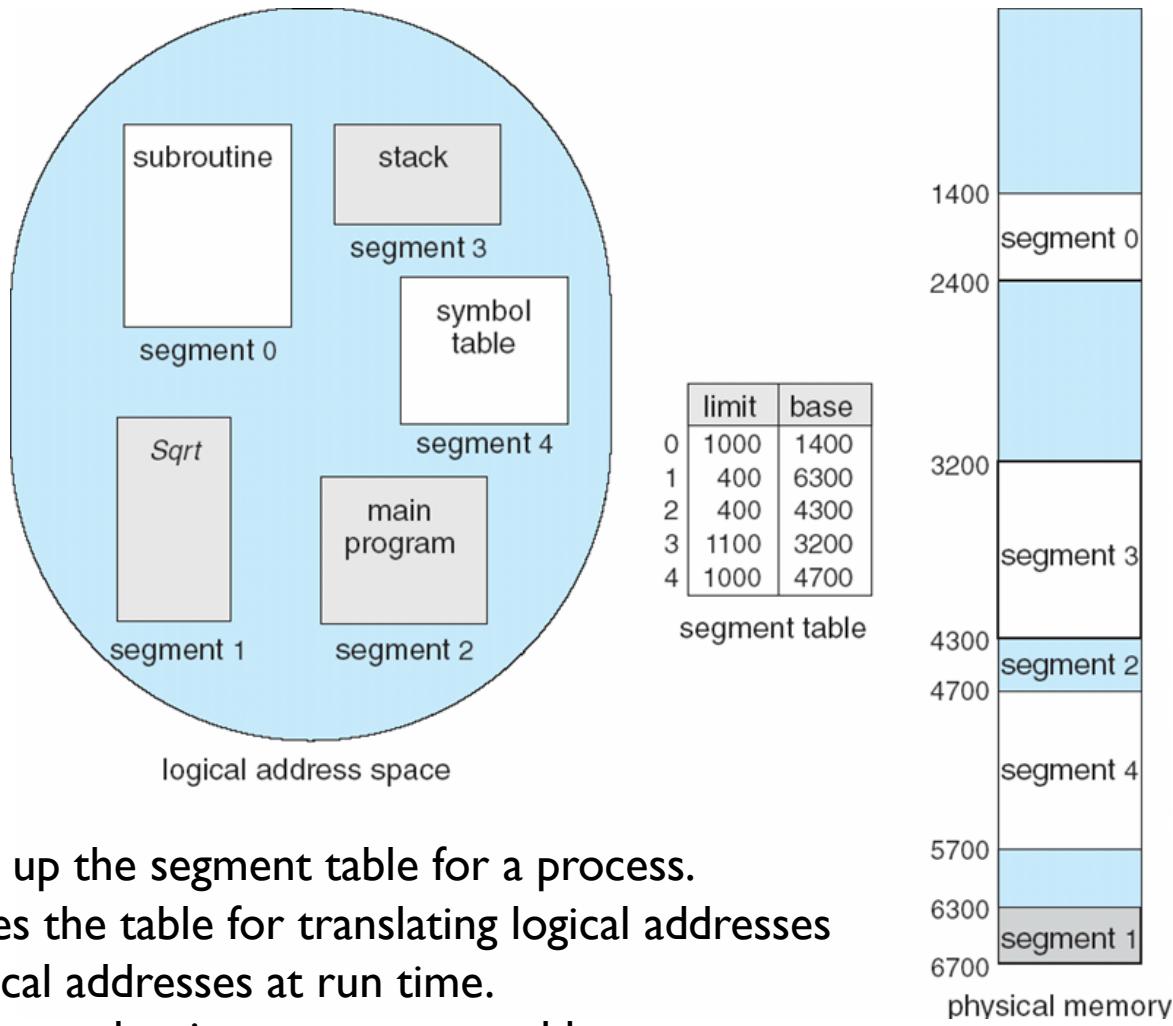
Segments do not have to be next to each other.

Contiguous Allocation versus Segmentation

- With contiguous allocation, large amount of physical memory may stay **unused** due **sparse logical address space**.
- With segmentation, we don't have this problem.
- Additionally, segments are of smaller in size; easier to find an empty hole.



Example of Segmentation



OS sets up the segment table for a process.

HW uses the table for translating logical addresses to physical addresses at run time.

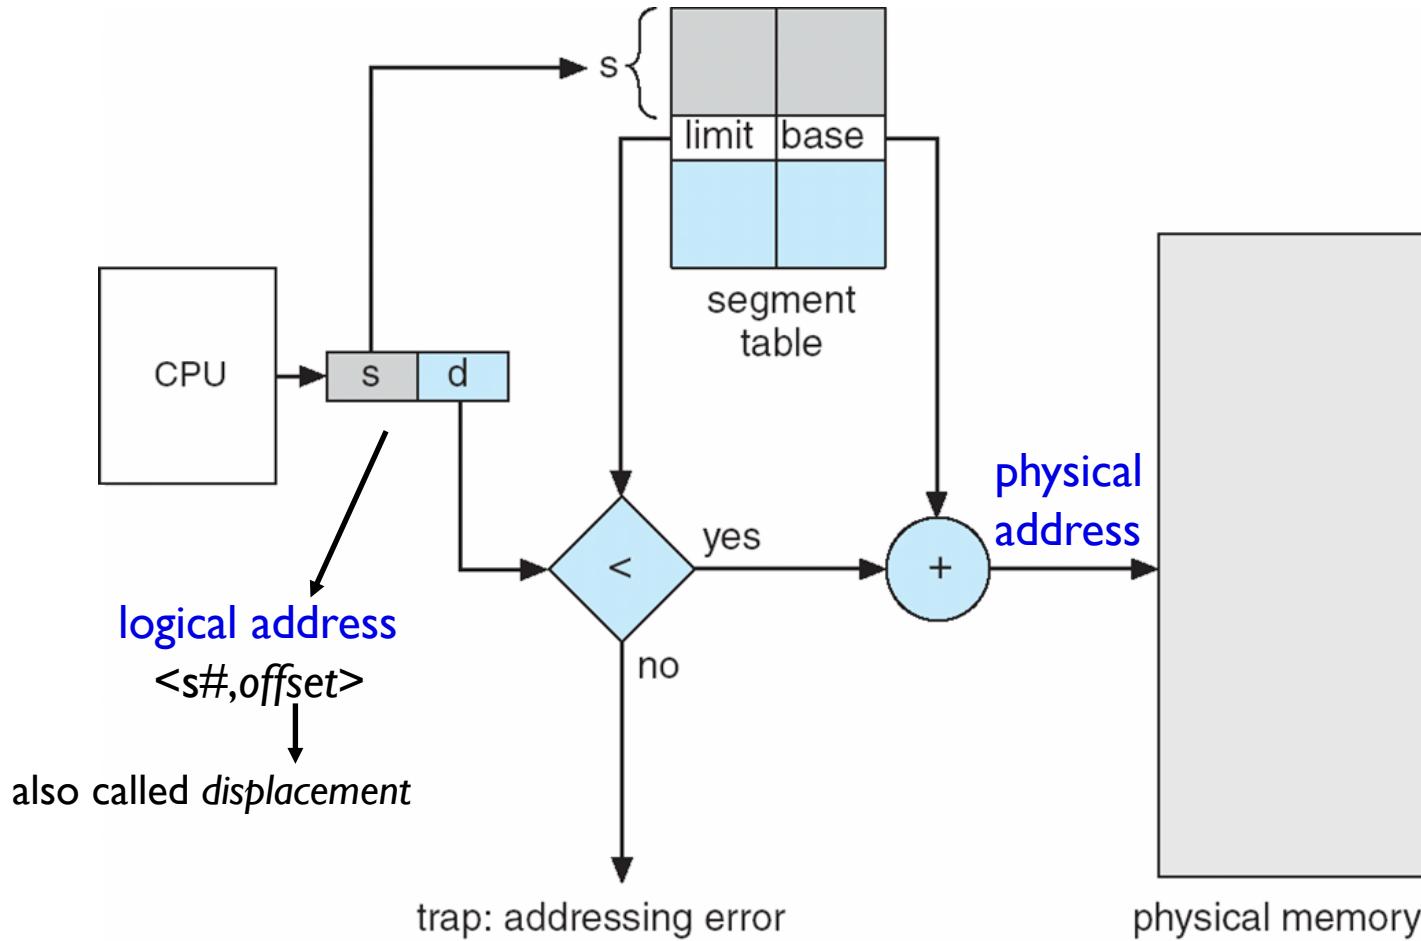
Each process has its own segment table

Segmentation Architecture

- Logical address is a two tuple (it is 2D):
 $\langle \text{segment-number}, \text{offset} \rangle$
- Segment table: each entry has **base** and **limit** register values. Using them, table maps two-dimensional LAs to one-dimensional PAs.
 - **base** – contains the start PA of segment in memory.
 - **limit** – specifies the length of the segment.
- Segment-table base register (STBR) points to the segment table's location in memory.
- Segment-table length register (STLR) indicates number of segments used by a program.

segment number s is legal if s < STLR

Segmentation Hardware address translation



HW maps 2D logical addresses to 1D physical addresses.

Segmentation Architecture

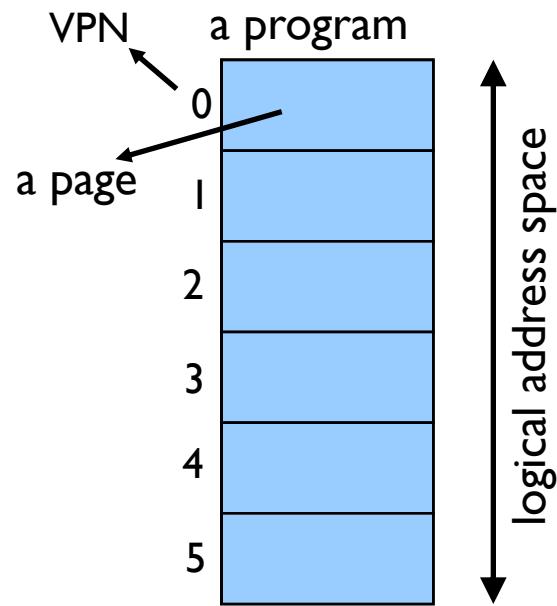
- Protection
 - With each entry in segment table associate:
 - validation bit = 0 \Rightarrow illegal segment
 - read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
 - Code segment: READ-ONLY; sharable; ...
 - Data segment: READ-WRITE; not-sharable
- Since segments vary in length, memory allocation is a **dynamic storage-allocation problem**
 - (similar to contiguous allocation)
- A

Paging

Paging

- Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
- Divide **physical memory** into fixed-sized blocks called **frames** (size is power of 2, a typical size: 4096)
- Divide **logical memory** into blocks of same size called **pages**
- Keep track of all free frames
- To run a **program of size n pages**, need to find **n free frames** and load program
- Set up a **page table** to translate logical to physical addresses
- Internal fragmentation

Paging



program: set of pages

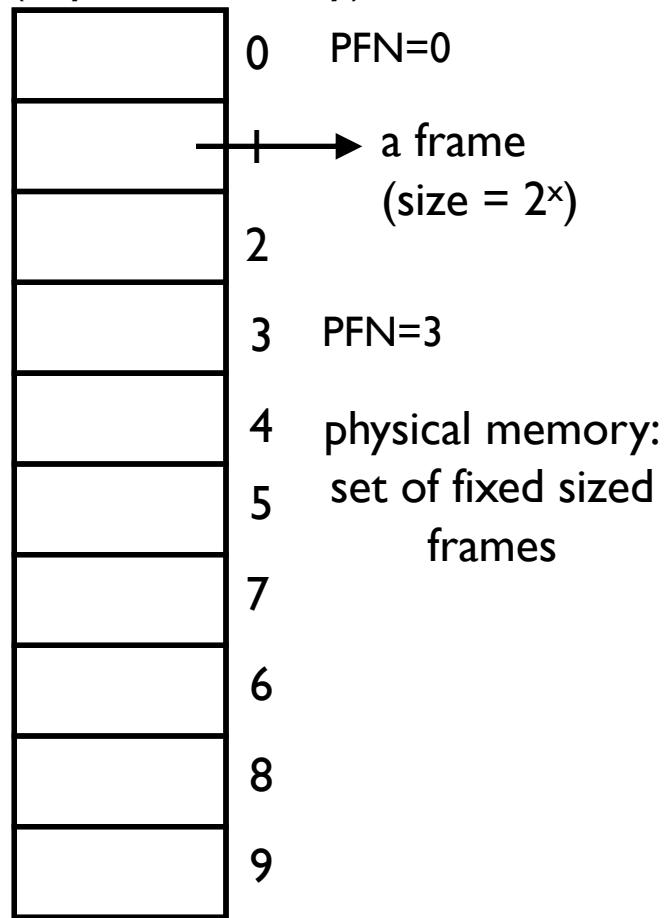
a **page** contains program **content**.

a **frame** is physical memory that can contain/store a page (it is a **container**).

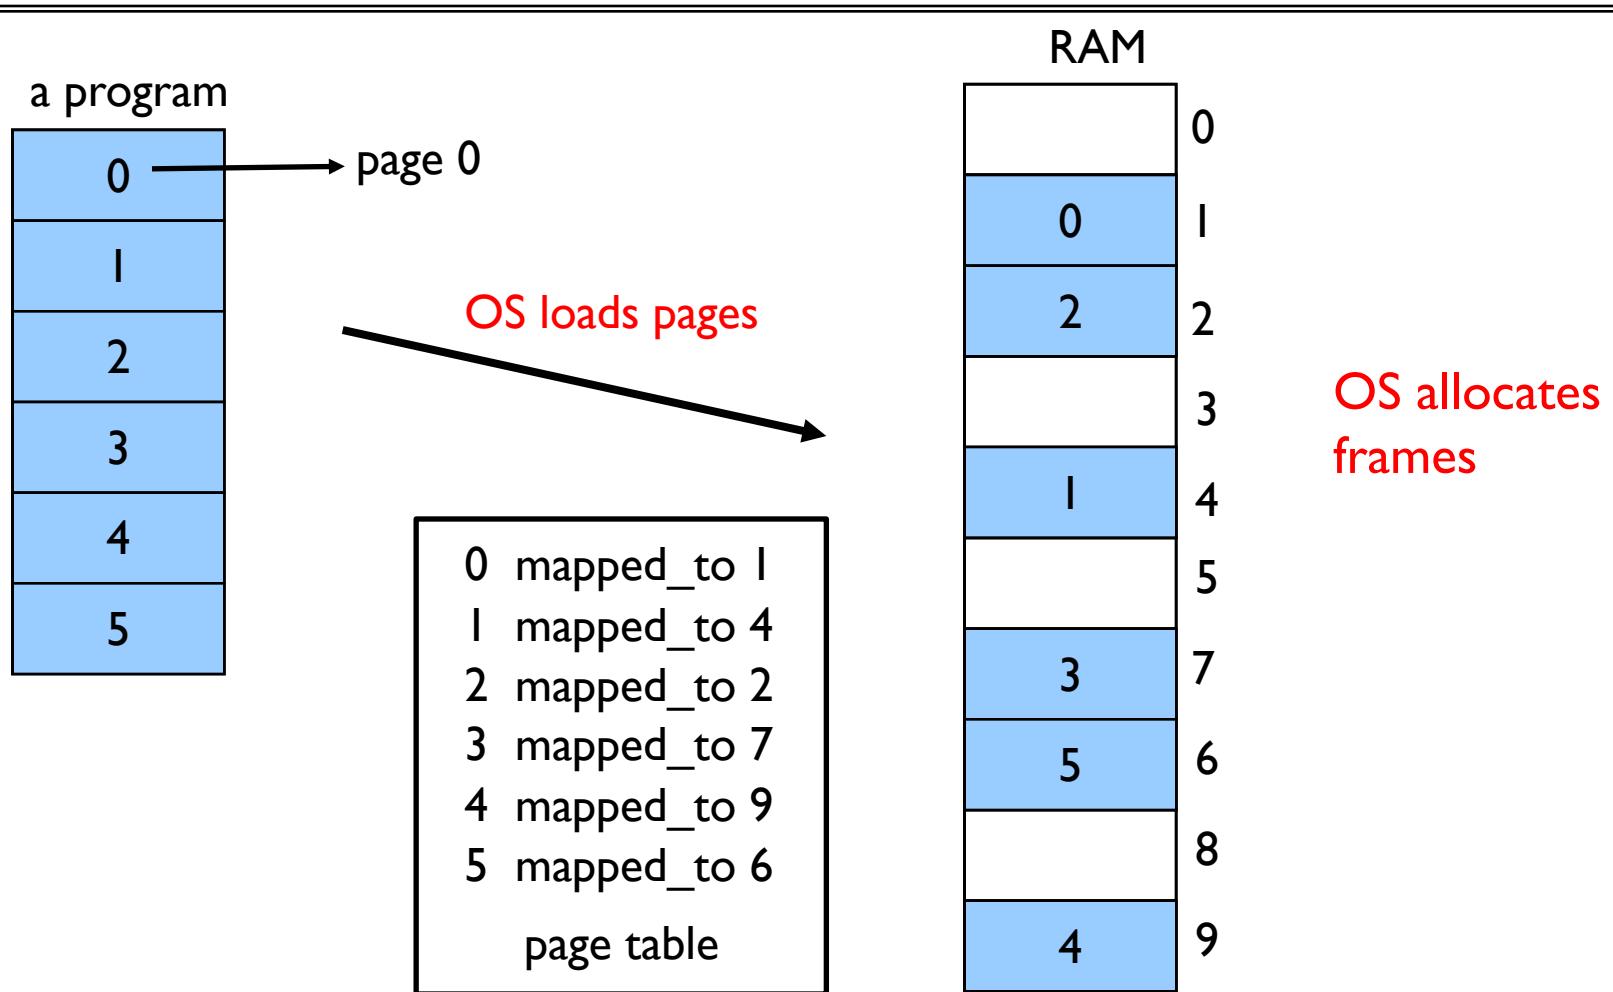
VPN: virtual page number
PFN: physical frame number

Page size = Frame size
typical size:
4 KB

RAM (Physical Memory)

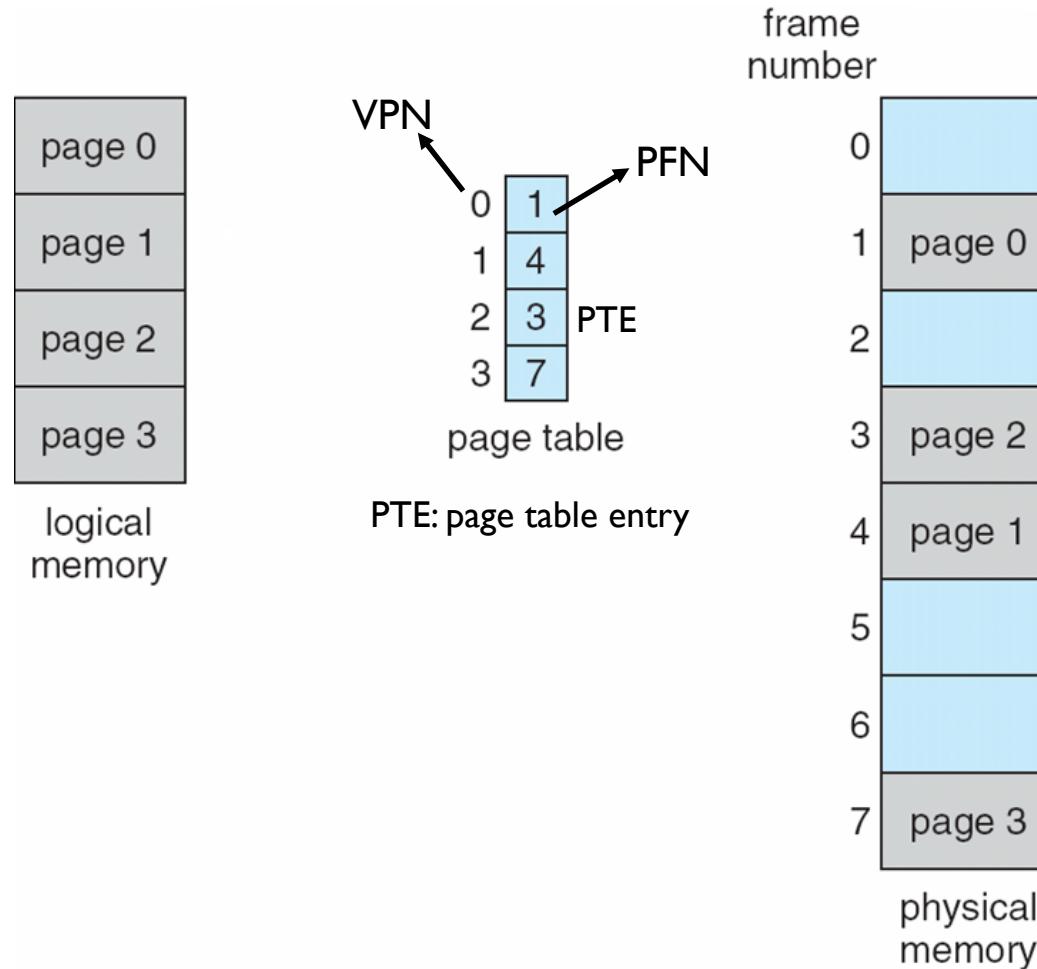


Paging



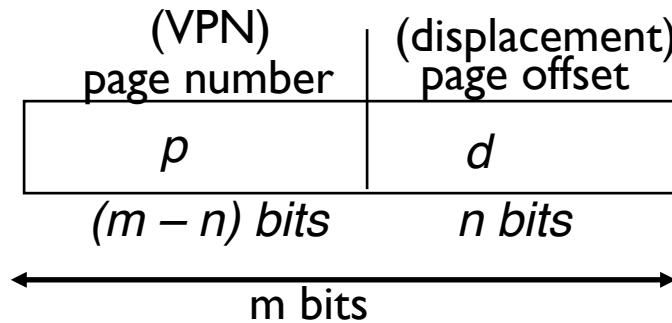
OS sets up the page table

Example



Address Translation Scheme

- Assume logical addresses are m bits. Then logical address space is 2^m bytes.
- Assume page size is 2^n bytes.
- Logical Address generated by CPU is divided into:
 - **Page number (p)** – (also called **VPN**) used as an index into a page table which contains base address (**PFN**) of each page in physical memory
 - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit



Simple example

Assume m is 3 and n is 2

Logical addresses

000

001

010

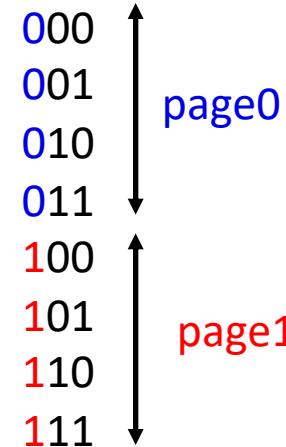
011

100

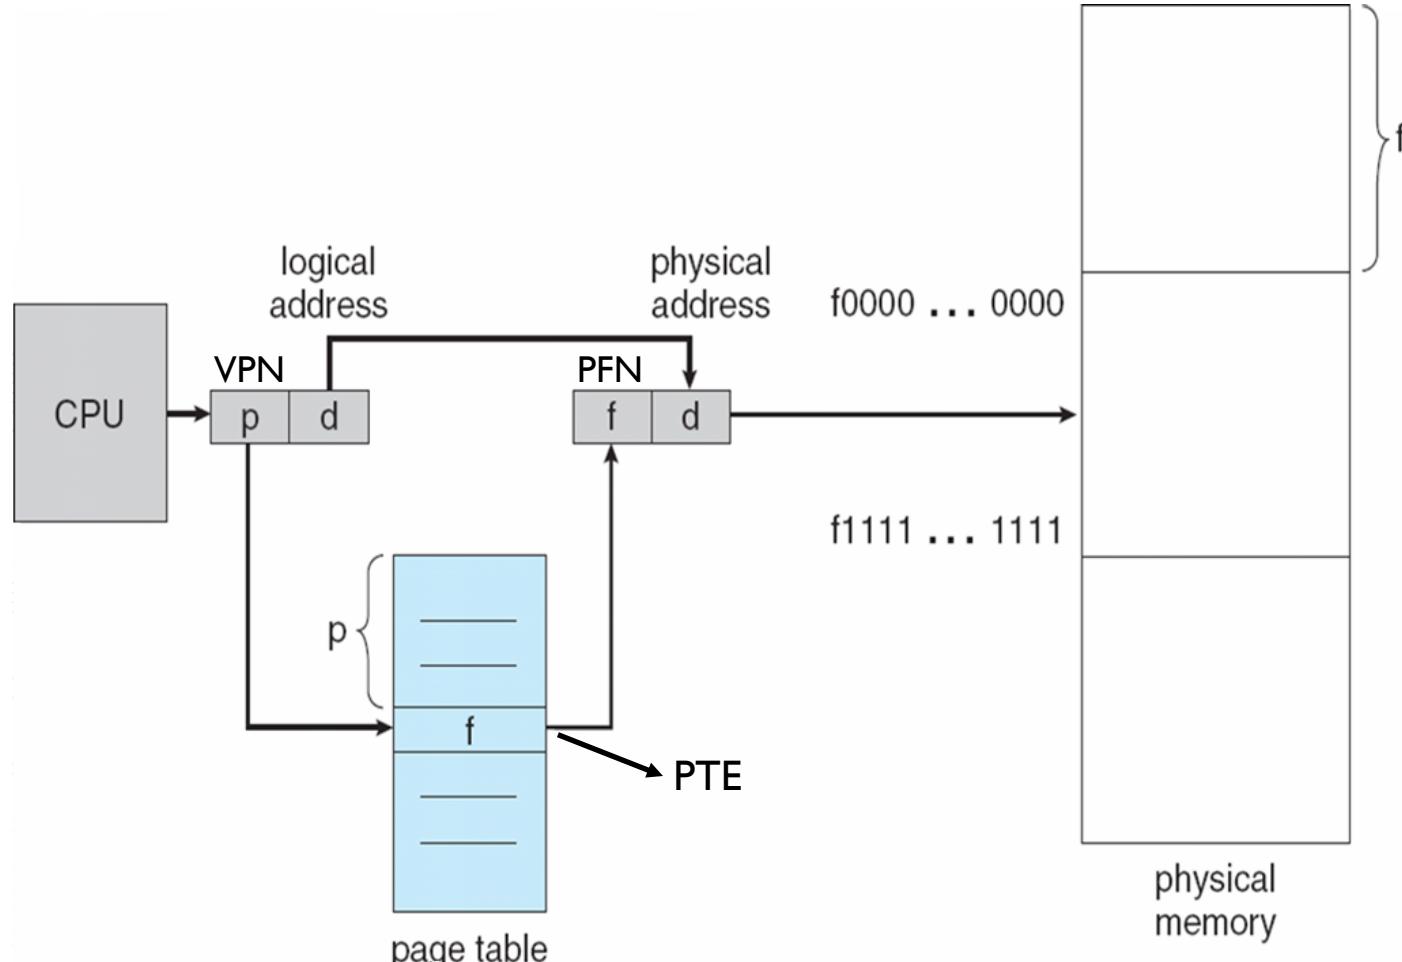
101

110

111



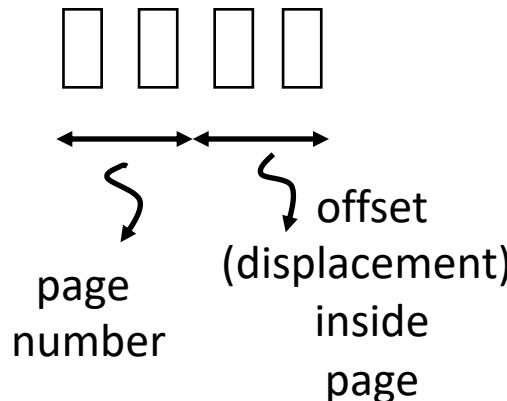
Paging Hardware: address translation



Paging and Address Translation Example

page size = 4 bytes
 $= 2^2$

4 bit logical address



0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

32 byte memory

LA = 5
PA = ?

5 is 0101
PA = 11001

LA = 11
PA = ?

11 is 1011
PA = 00111

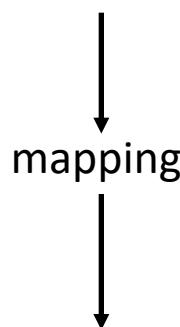
LA = 13
PA = ?

13 is 1101
PA = 01001

Address translation example 2

16 bit logical address

0010 0000000000100
p# offset



f# offset
110 0000000000100
15 bit physical address

15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

page size = 4096 bytes
(offset is 12 bits)

frame number
valid/invalid bit

page table

Address translation example 3

$m=3; 2^3 = 8$ logical addresses

$n=2$; page size = $2^2 = 4$

1 bit for page#

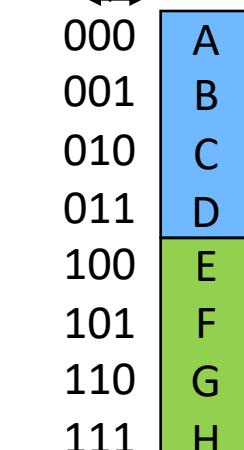
page 0

page 1

page table

0	11
1	10

2 bits for offset



Logical Memory

0000

0001

0010

0011

0100

0101

0110

0111

1000

1001

1010

1011

1100

1101

1110

1111

frame 00

frame 01

frame 10

frame 11

2 bits for frame#

Physical Memory

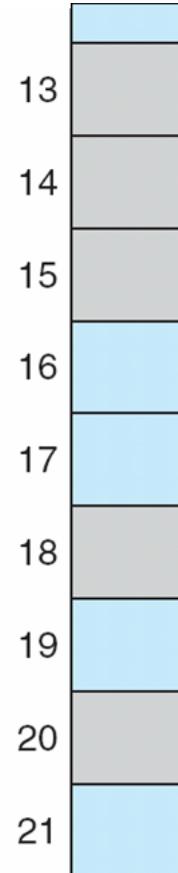
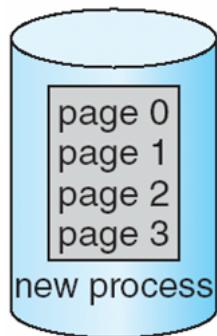


Free Frames

OS also
keeps info
about the
frames
in its frame
table

free-frame list

14
13
18
20
15

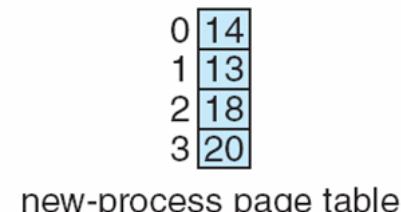
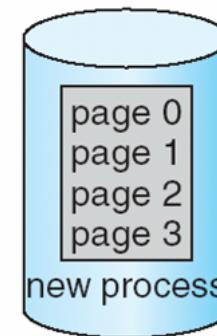


(a)

Before allocation

free-frame list

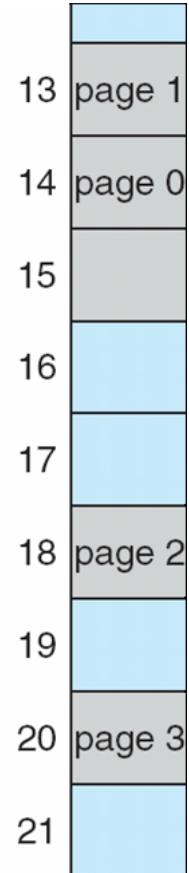
15



new-process page table

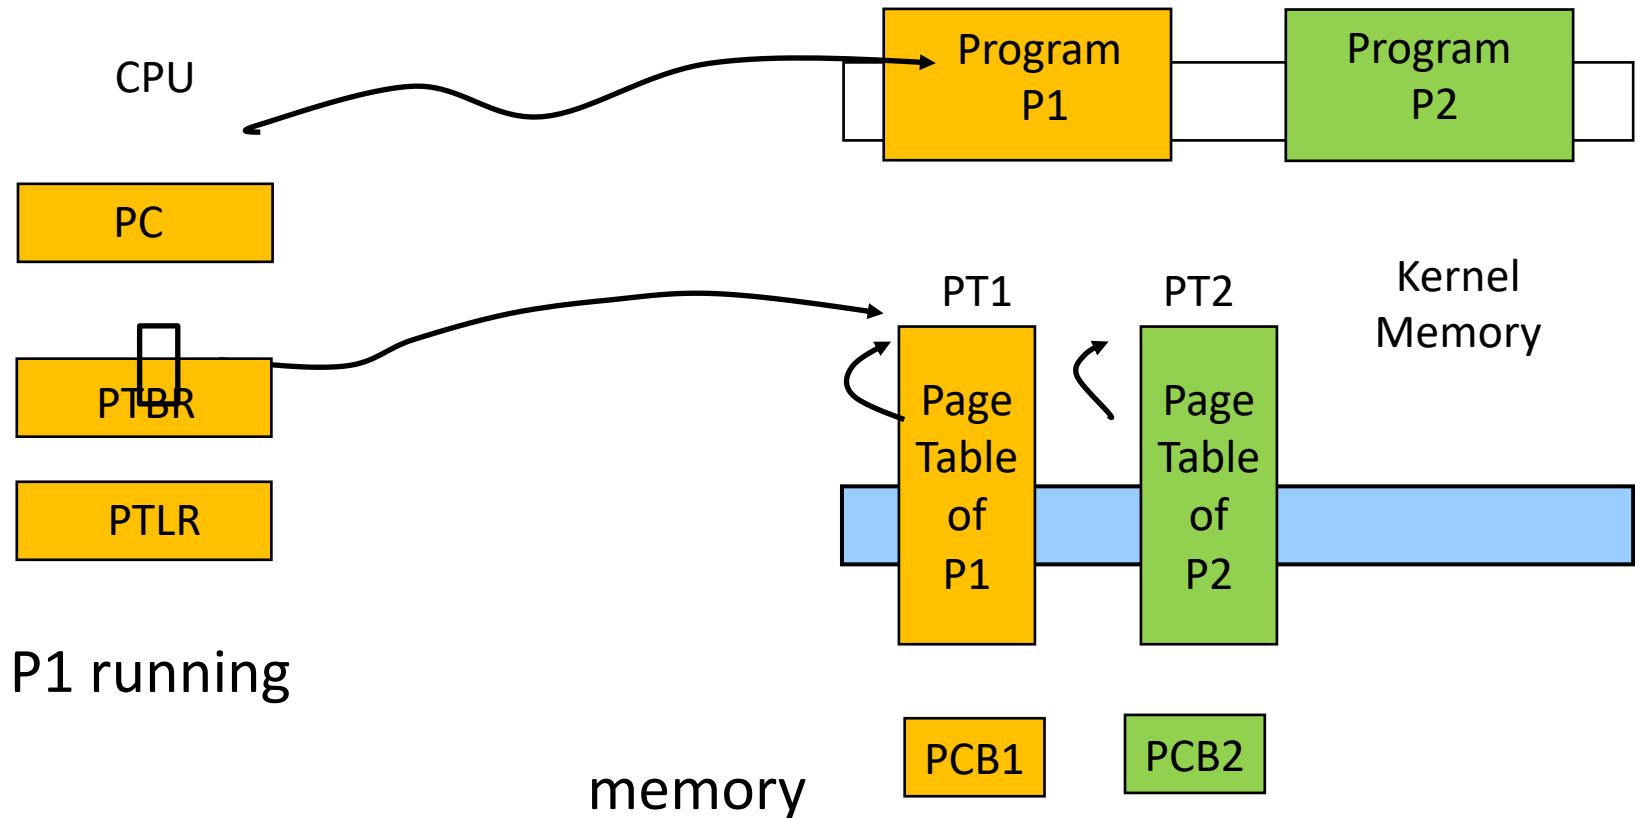
(b)

After allocation



Implementation of Page Table

- Page table is kept in main memory.
- Page-table base register (PTBR) points to the page table.
- Page-table length register (PTLR) indicates size of the page table.

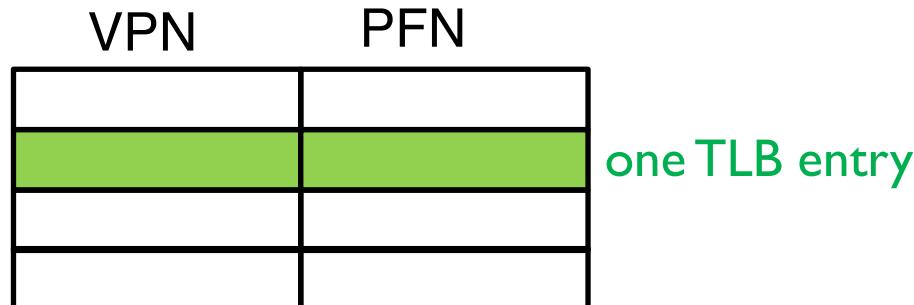


TLB

- With paging, every data/instruction access requires **two memory accesses**.
 - One for the page table
 - One for the data/instruction.
- Solve it by using of a special **fast-lookup hardware** cache called associative memory or **translation look-aside buffers (TLBs)**
- **Learned (p#, f#) mappings** can be **cached in TLB**.
 - TLB: a set of entries each entry storing a different mapping.
 - Number of entries limited (~ 1024 , for example) (TLB size)
 - Can be accessed very fast.

TLB Associative Memory

- Associative memory – parallel search



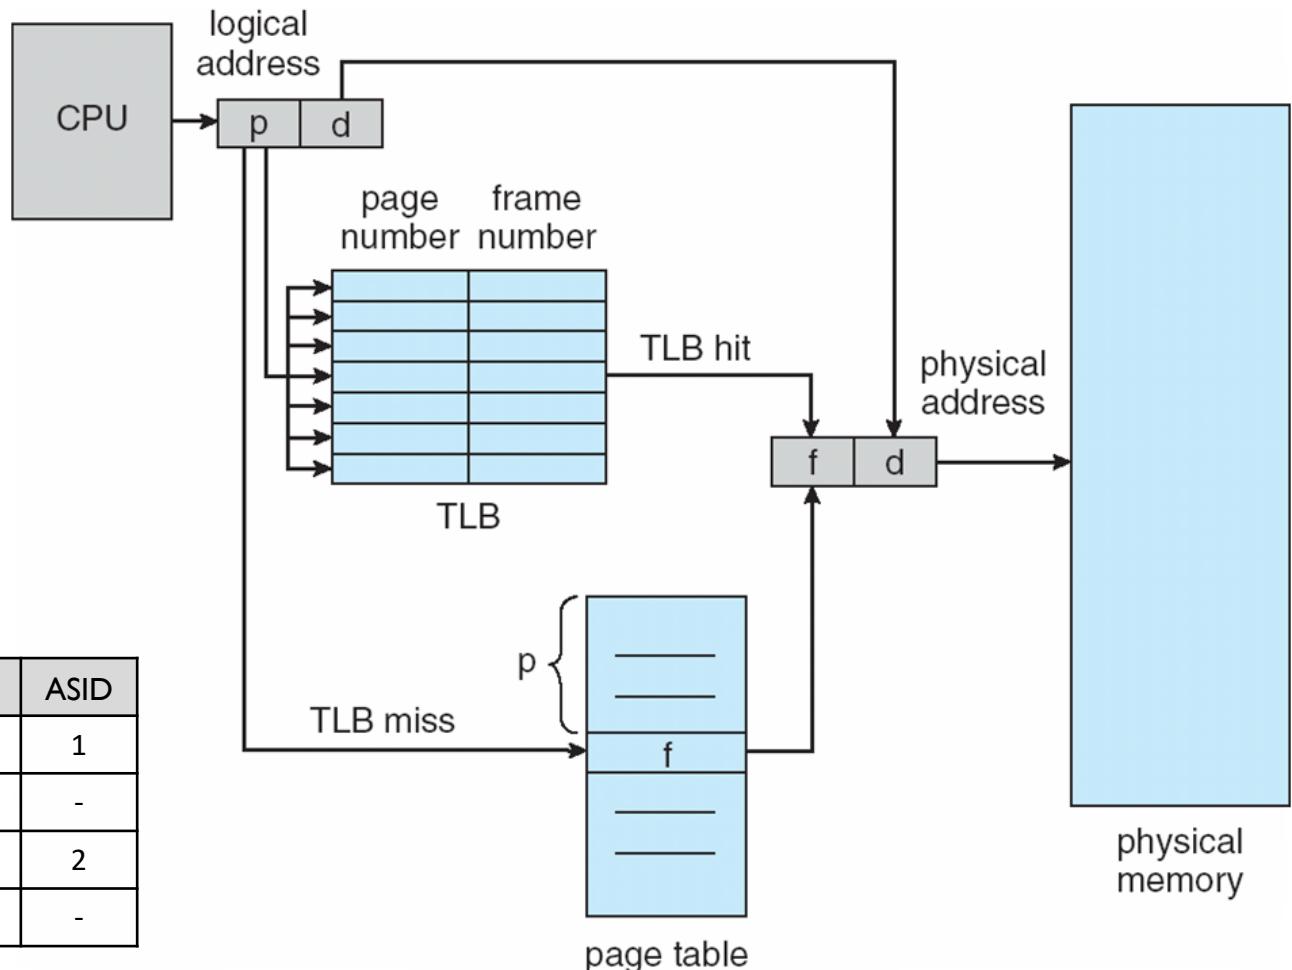
Address translation (p, d)

- If p is in TLB, get **frame #** out
- Otherwise get **frame #** from page table in memory

Some TLBs store *address-space identifiers (ASIDs)* in each **TLB entry** – uniquely identifies each process to provide address-space protection for that process.

ASID	VPN	VFN

Paging Hardware With TLB



Effective Memory Access Time

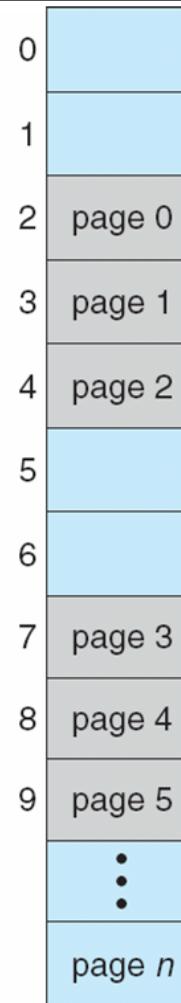
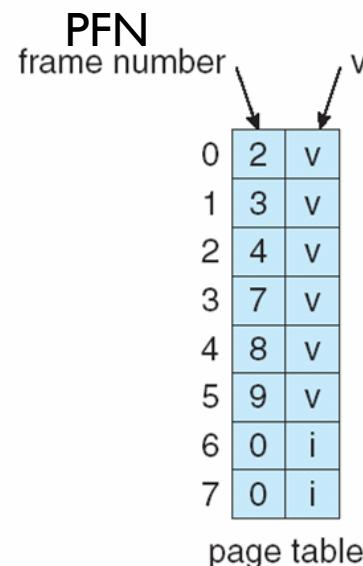
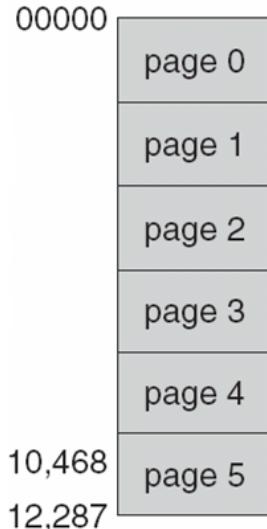
- TLB Lookup = ε time unit
- Assume memory access time is **1 time unit**.
- **Hit ratio** – percentage of times that a page number is found in the TLB; ratio related to the TLB size.
- Hit ratio = α
- **Effective (average) Access Time (EAT) to Memory**

$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$

Memory Protection

- Memory protection implemented by associating a protection bits with each page
 - Read only page
 - Executable page
 - Read-write page
- **Valid-invalid** bit attached to **each entry** in the page table:
 - “valid” indicates that the page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space

Valid (v) or Invalid (i) Bit In A Page Table

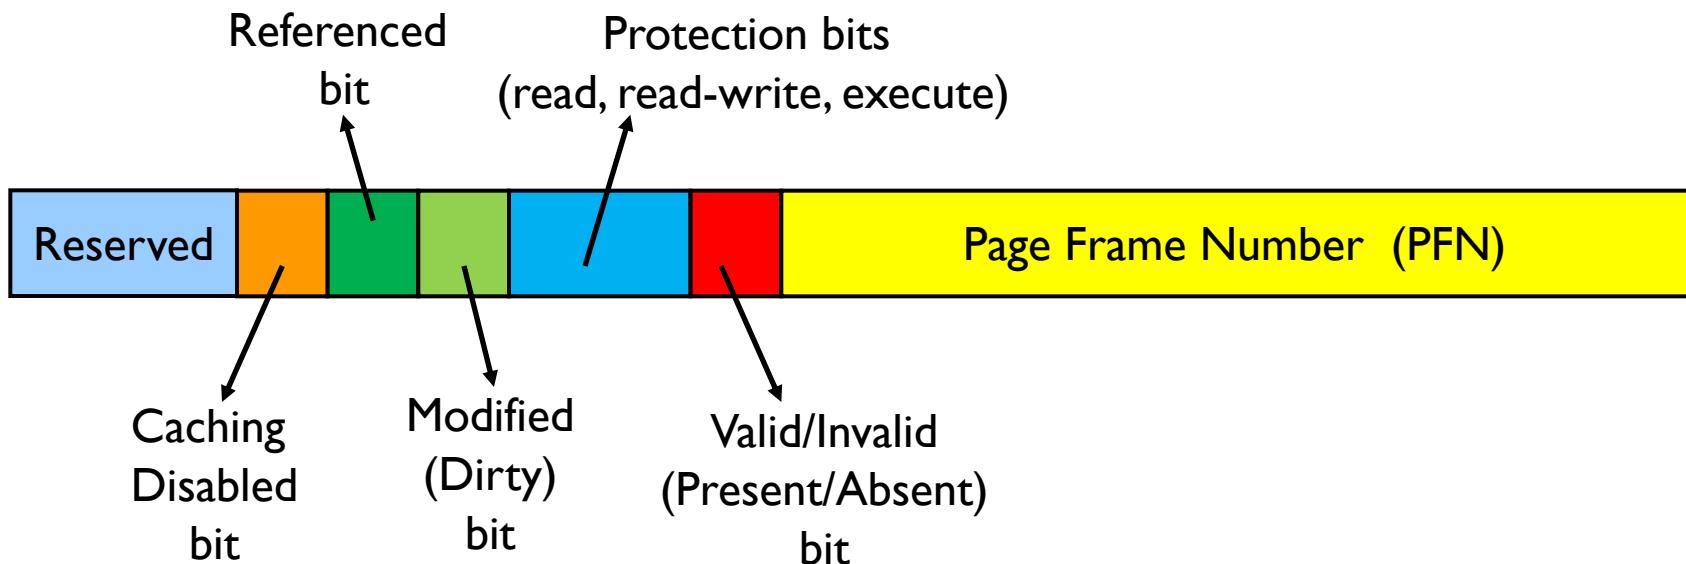


page 6 and 7 is not used by the program

accessing page 6 or 7 generates memory access error (segmentation fault).

Page Table Entry (PTE) Structure

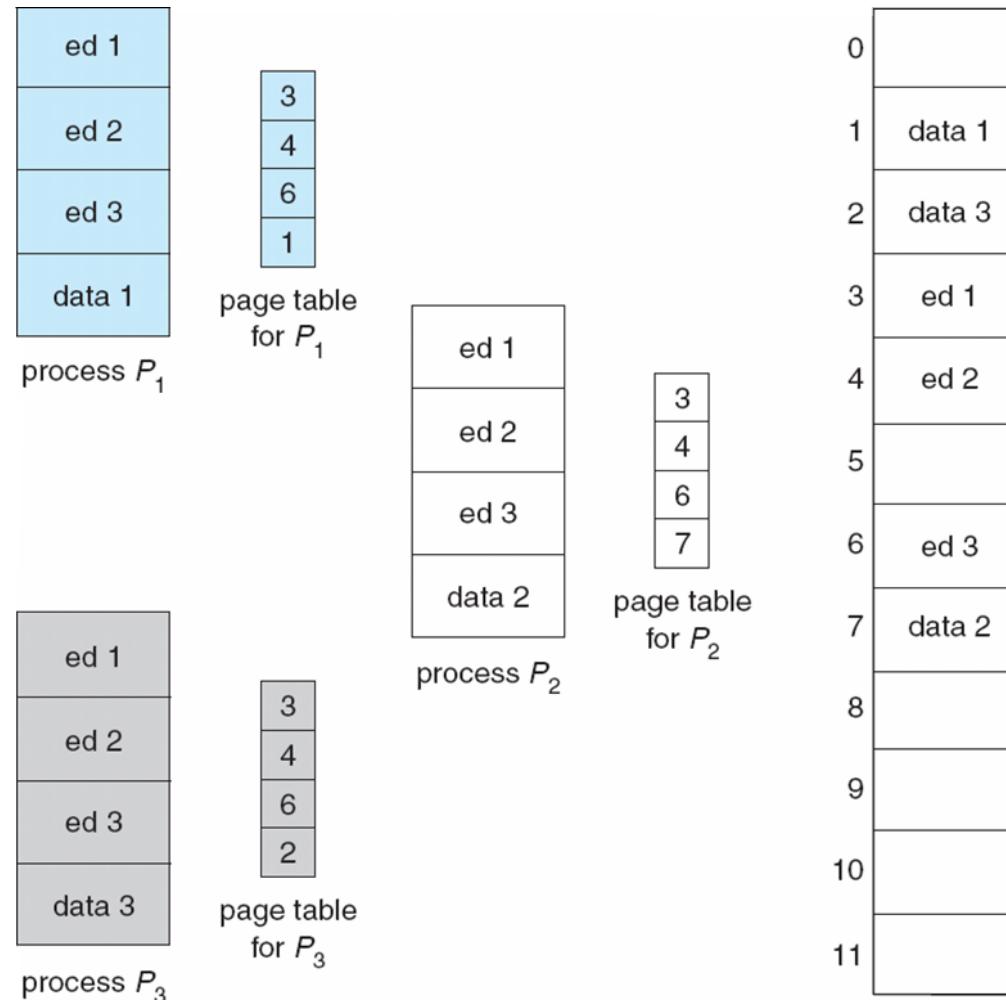
- A typical size of a page table entry can be **32 bits** (in a 32-bit system). Depends on the architecture.
- Typically we have the following fields in a page table entry.



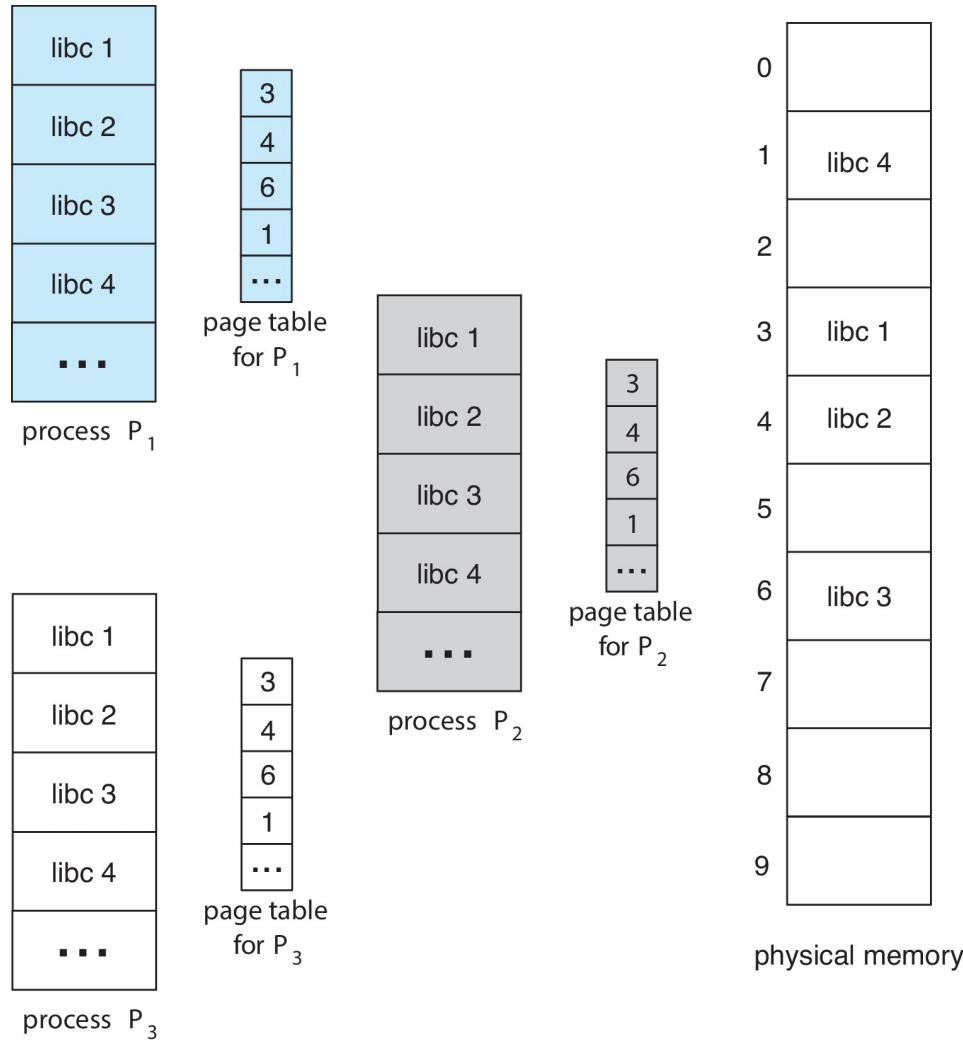
Shared Pages

- A program can be started multiple times (same program executed by multiple process)
 - Code pages can be shared
 - One copy of read-only (reentrant) code shared among processes (can be used for sharing run-time libraries, etc., as well).
 - Shared code
 - Private data
 - Each process keeps a separate copy of the data

Shared Pages Example

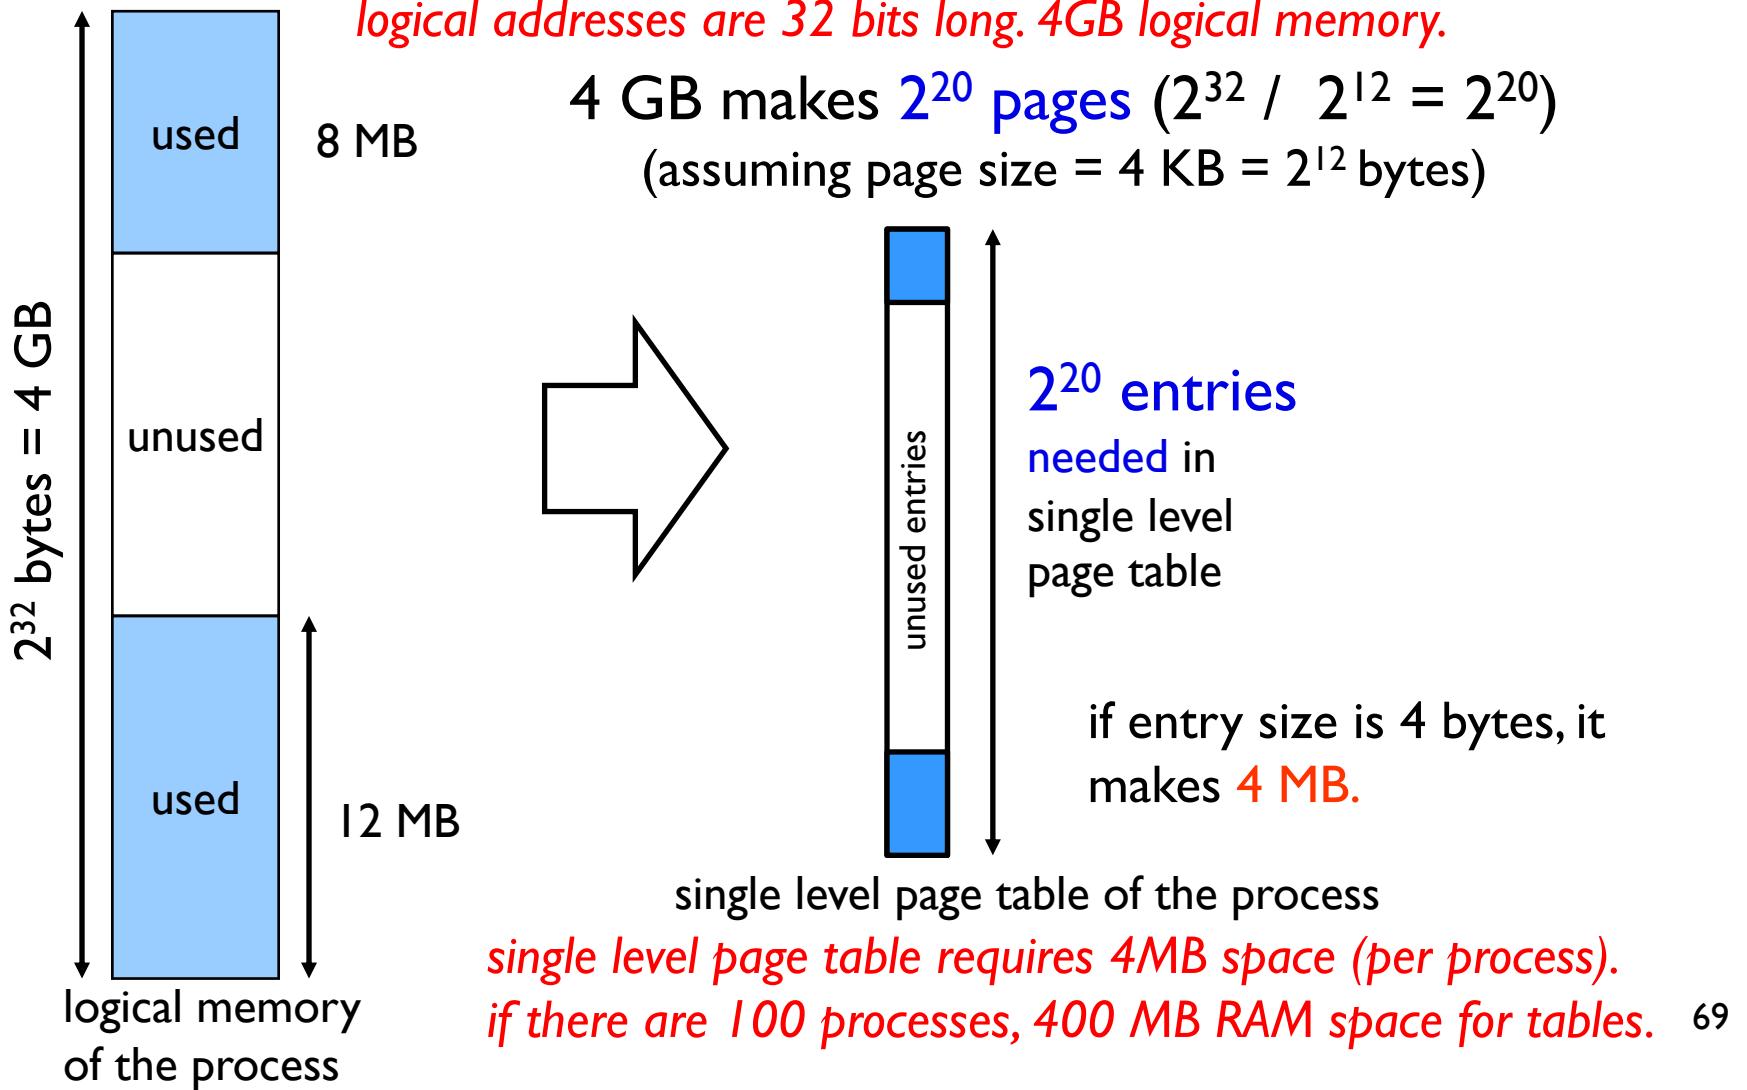


Shared Pages Example: sharing a library: libc



Structure of the Page Table

single level paging may not work for large address spaces. example: 32 bit system.

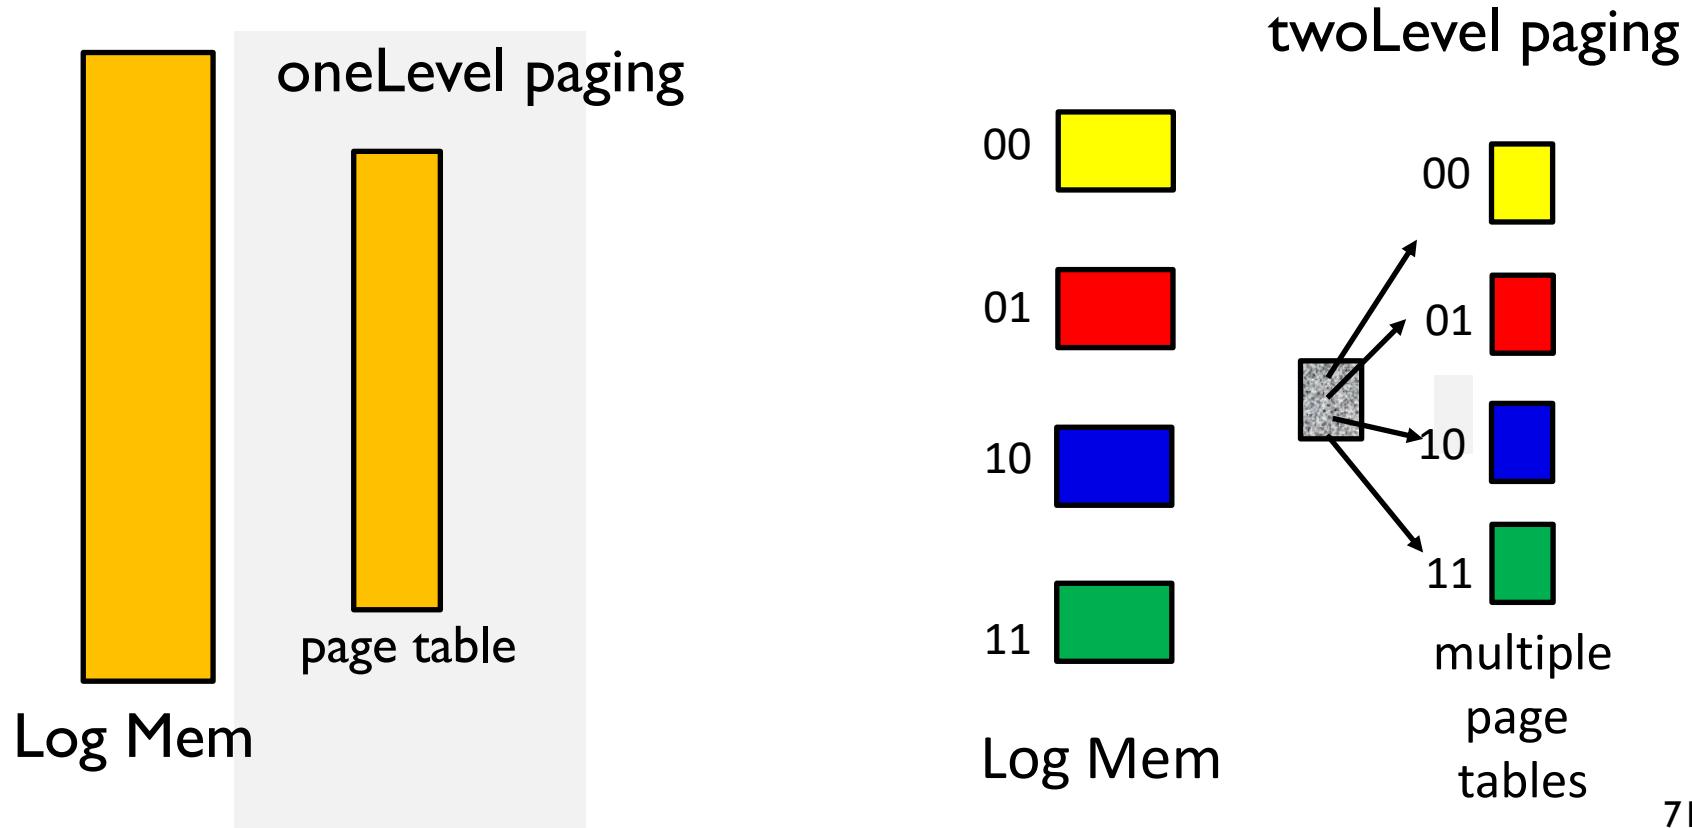


Structure of the Page Table

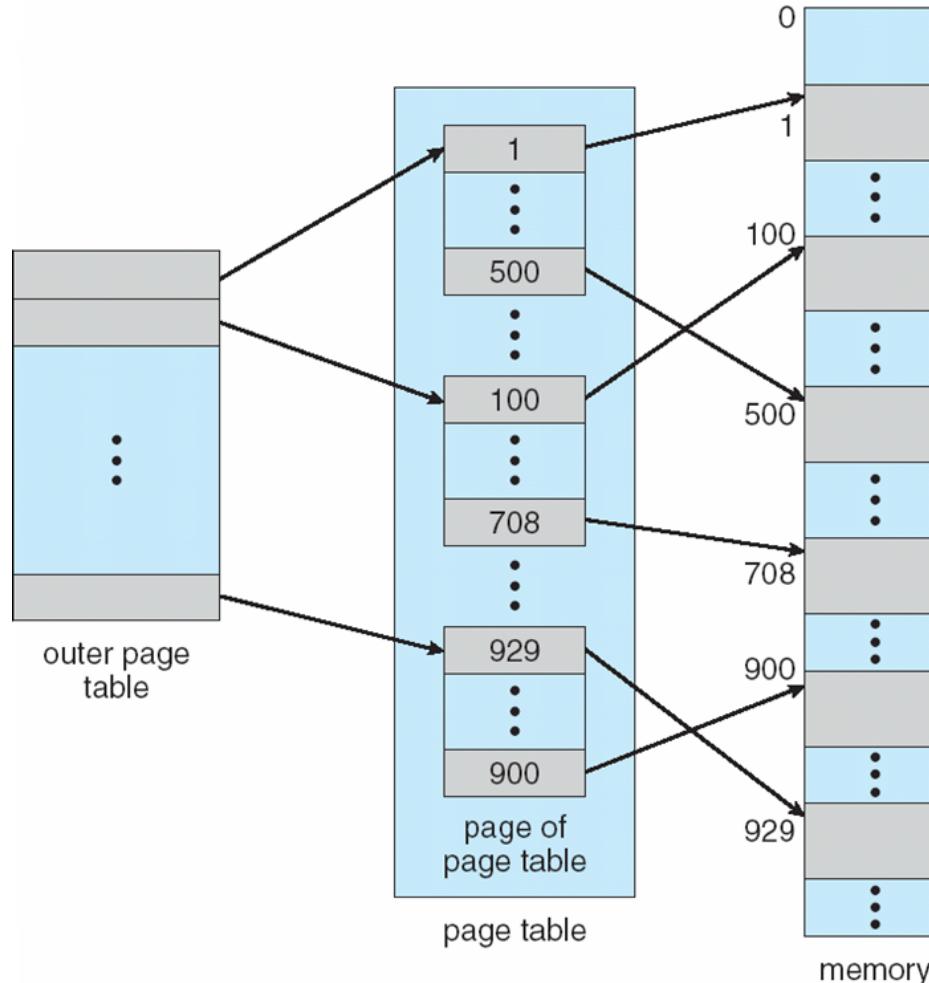
- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

Hierarchical Page Tables

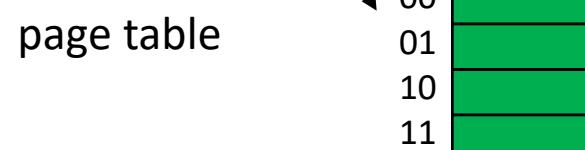
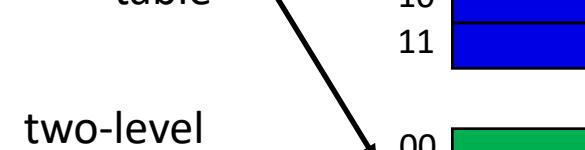
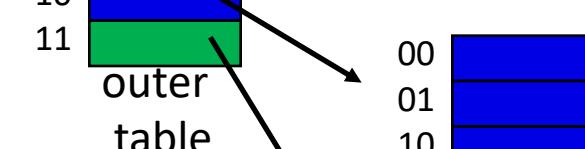
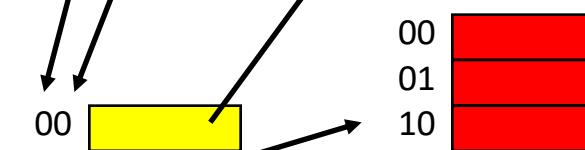
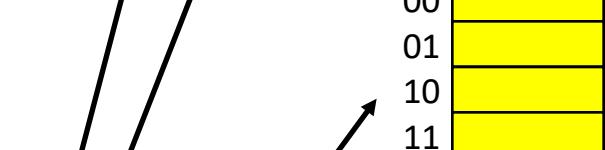
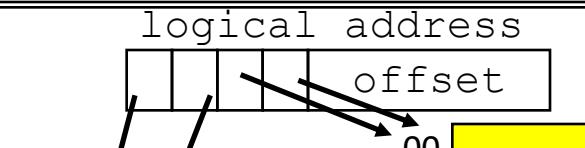
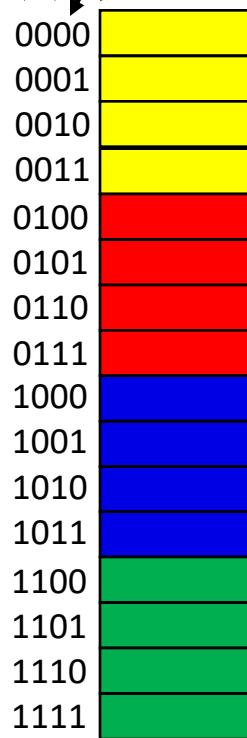
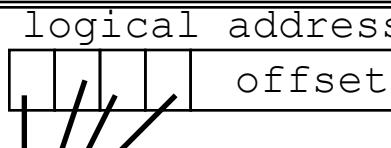
- Break up the **logical address space** into **subspaces**. Break up the page table into **sub-tables**. A sub-table for a sub-space.
- Two-level paging, is the simplest.



Two-Level Paging Scheme

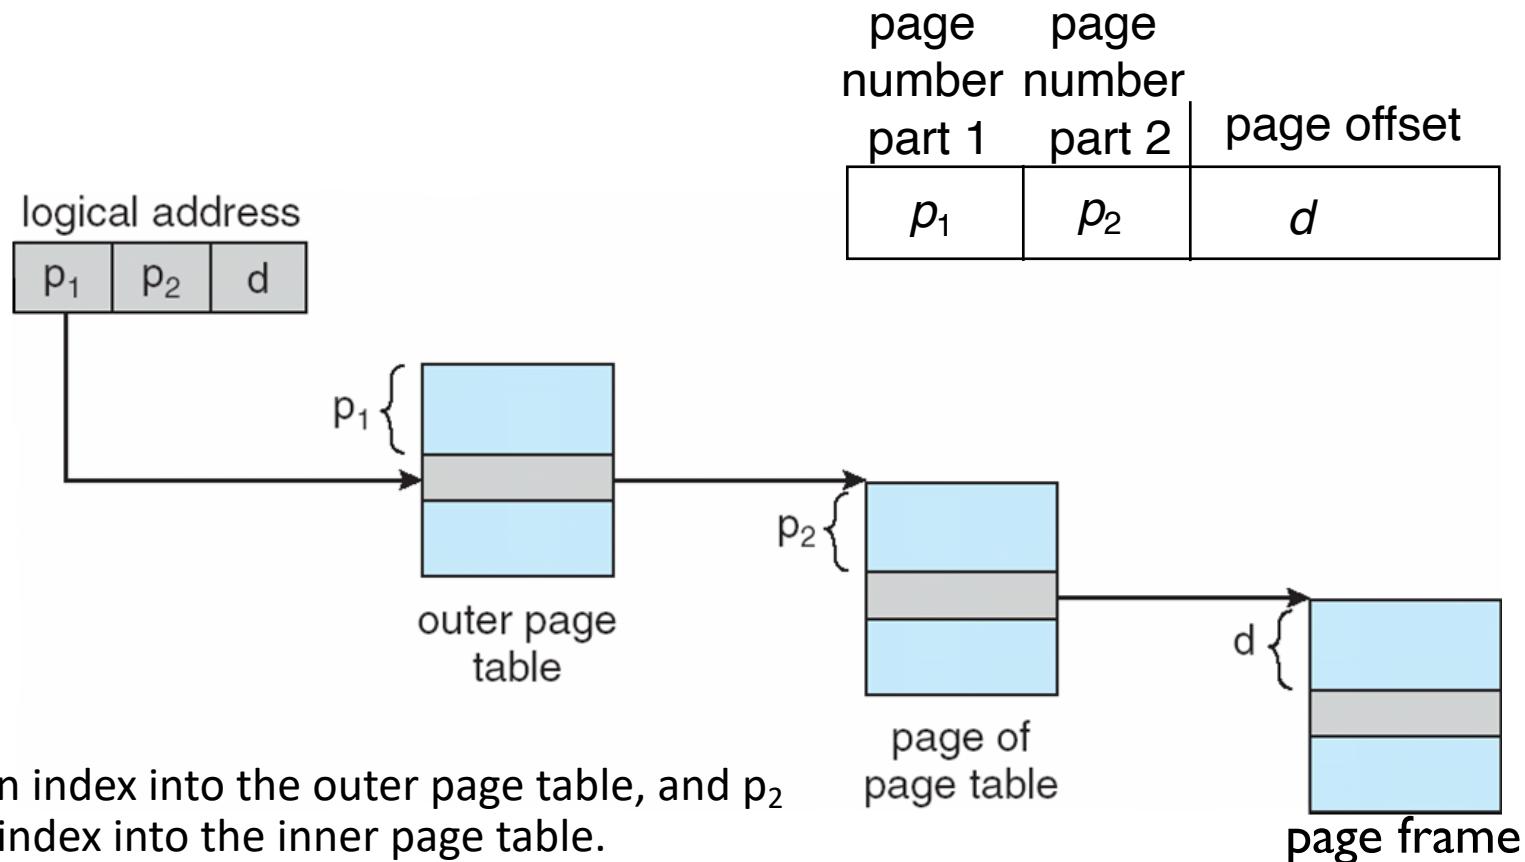


Two-Level Paging Scheme



Address-Translation Scheme

A logical address is divided into 3 parts:



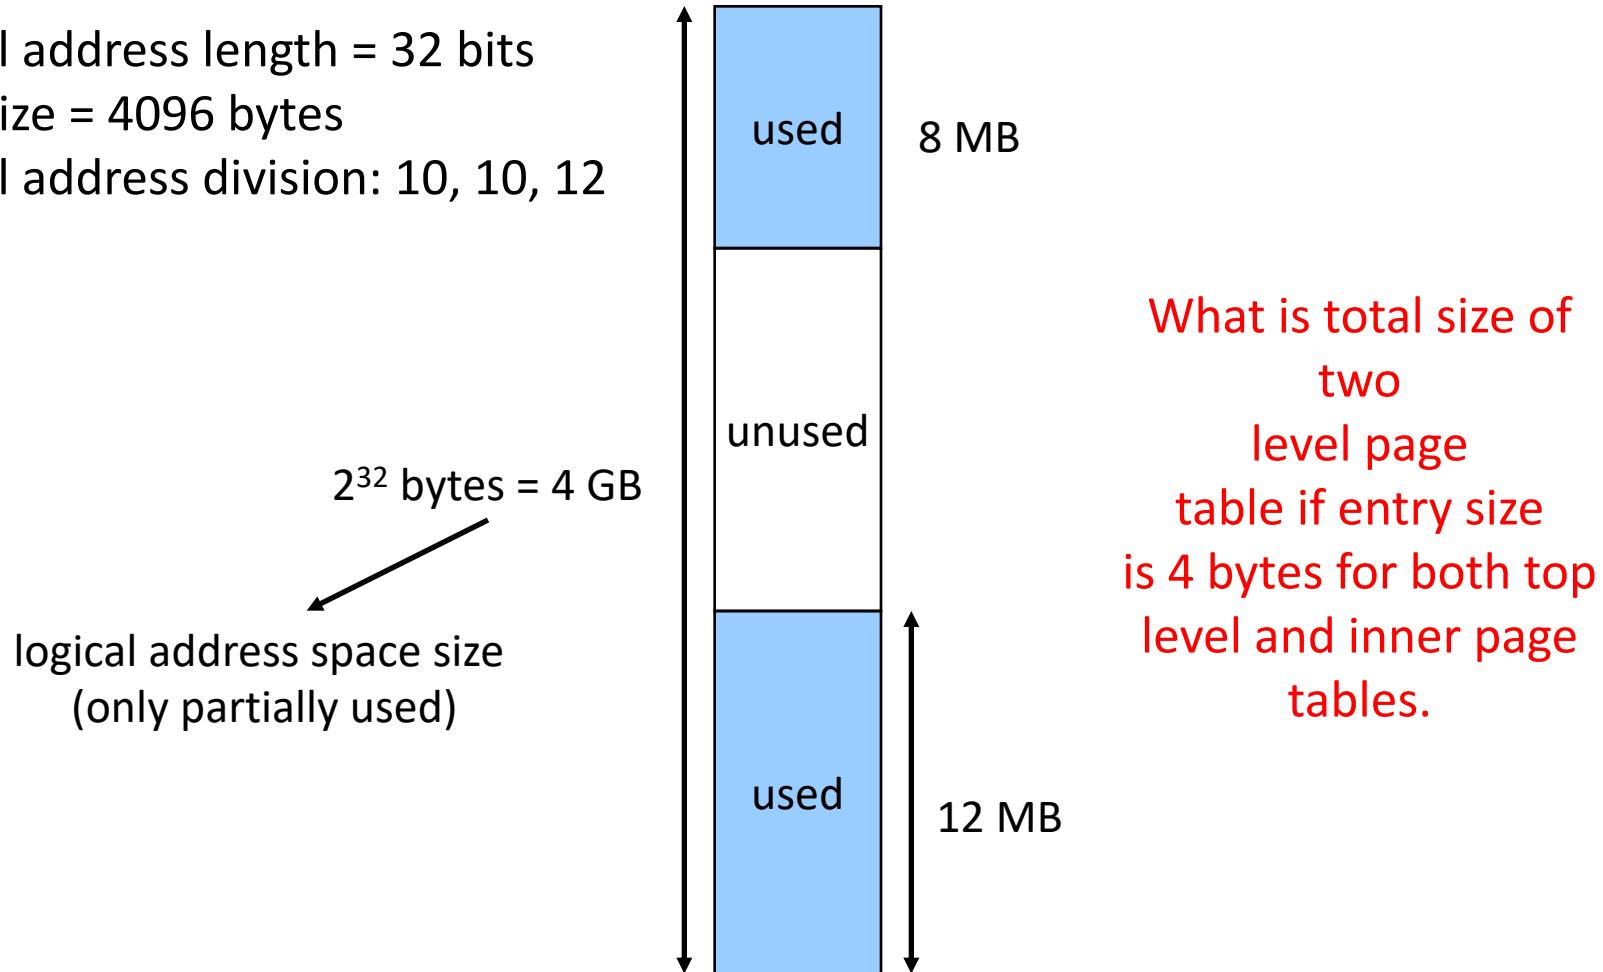
p_1 is an index into the outer page table, and p_2 is the index into the inner page table.

Example: two level page table

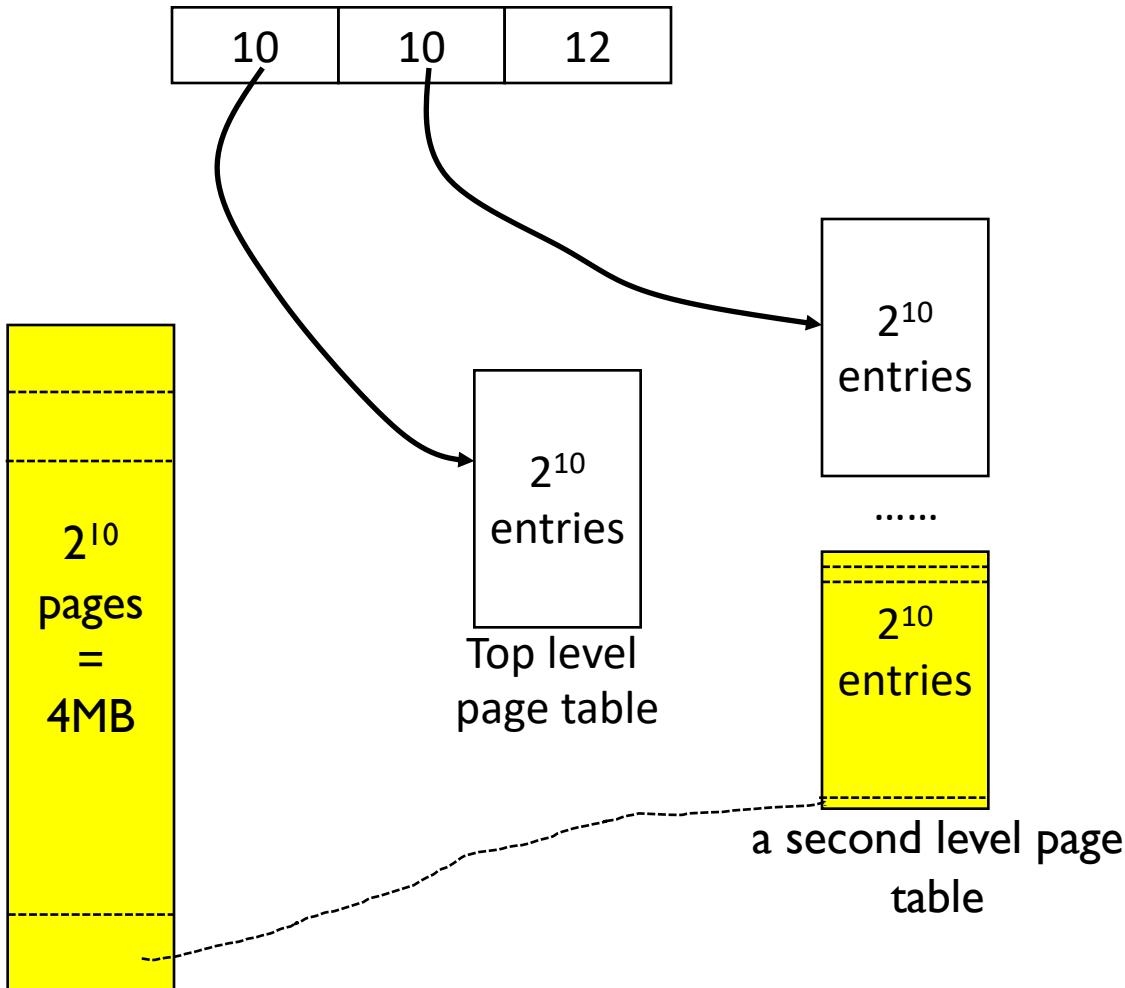
- Assume:
 - 32 bit virtual (logical addresses) used
 - 2 level paging
 - Address split scheme: [10, 10, 12]
 - 10 bits page number part 1
 - 10 bits for page number part 2
 - 12 bits for offset
 - **Page size = 2^{12} = 4096 bytes = 4 KB**
 - An outer (top level) table will have 2^{10} entries
 - An inner (second level) table will have 2^{10} entries

Example: two level page table space requirement

logical address length = 32 bits
pagesize = 4096 bytes
logical address division: 10, 10, 12



Example: two level page table space requirement

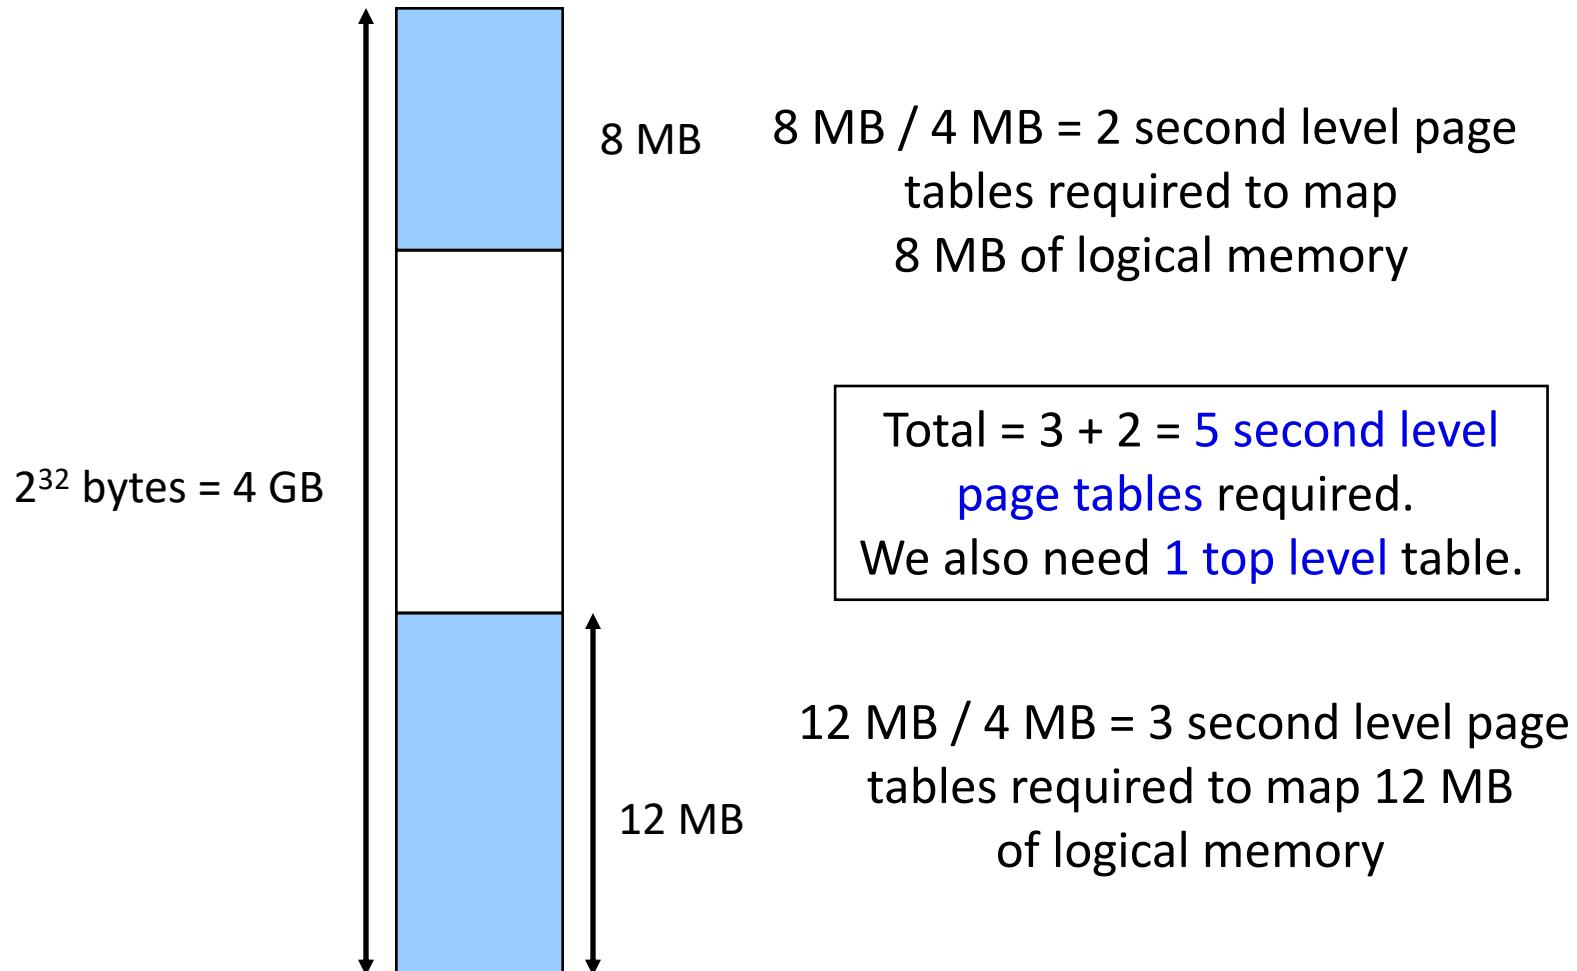


Each entry of a second level page table translates a page# to a frame#;
i.e., each **entry maps a page** which is 4096 bytes

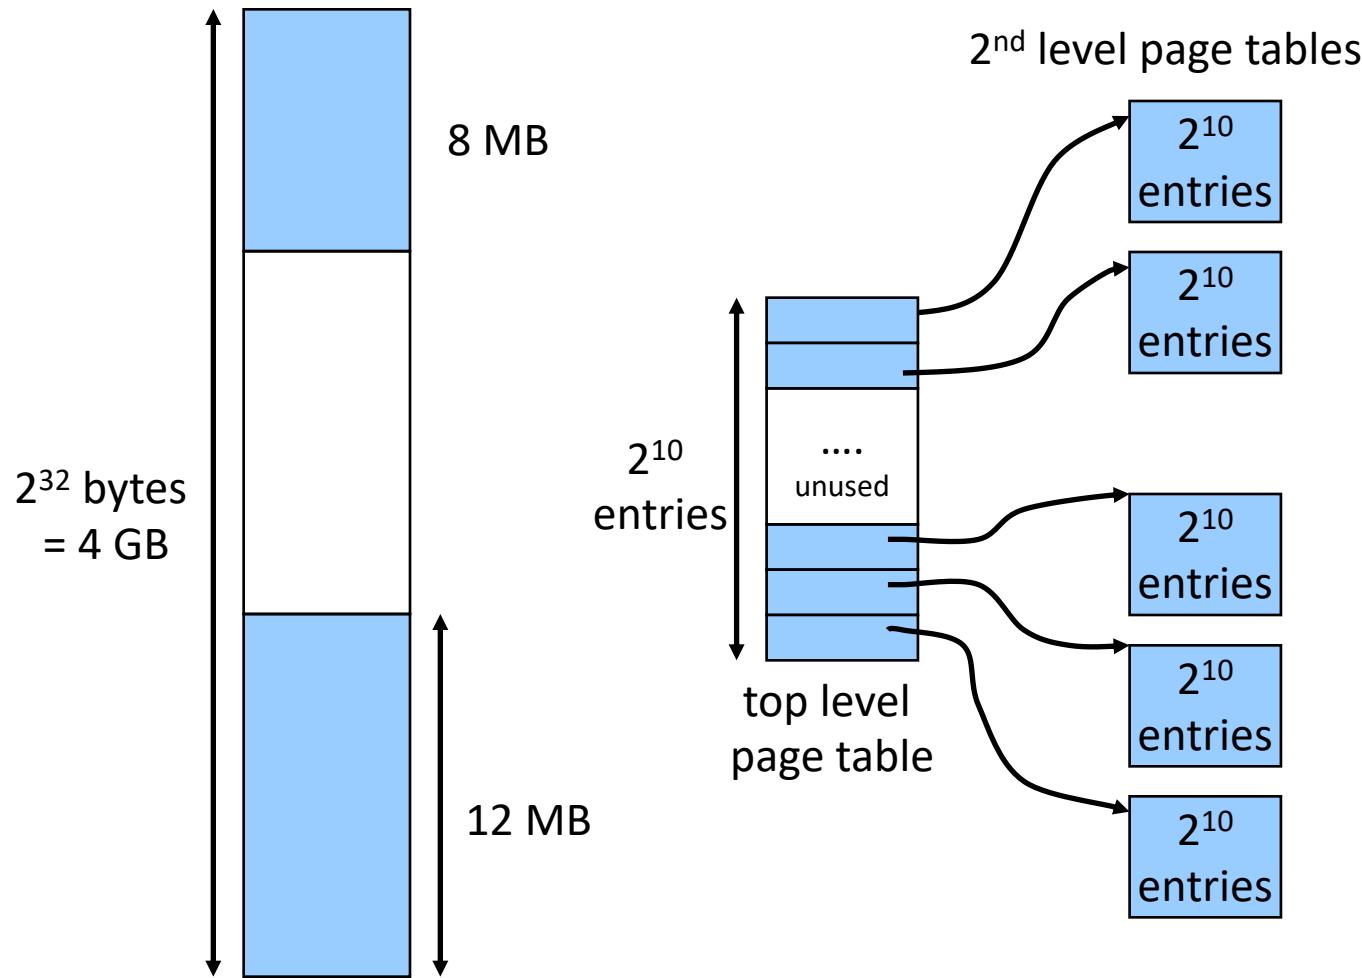
There are 1024 entries
In a second level page table

Hence, a second level page table can map $2^{10} * 2^{12} = 2^{22} = 4$ MB of logical address space

Example: two level page table space requirement



Example: two level page table space requirement



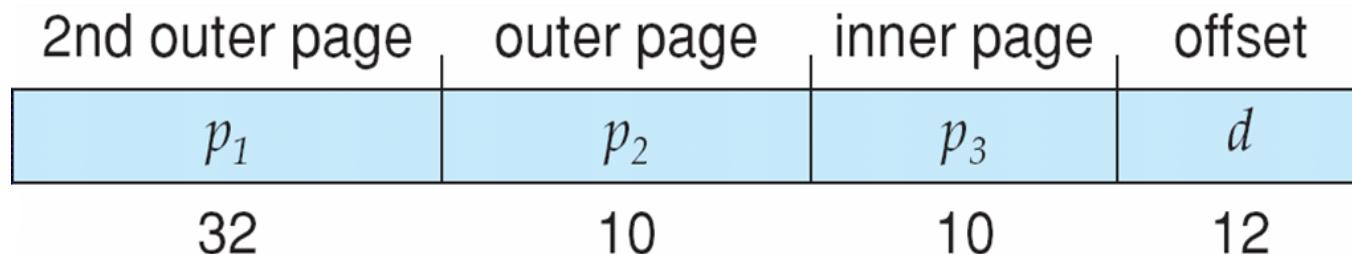
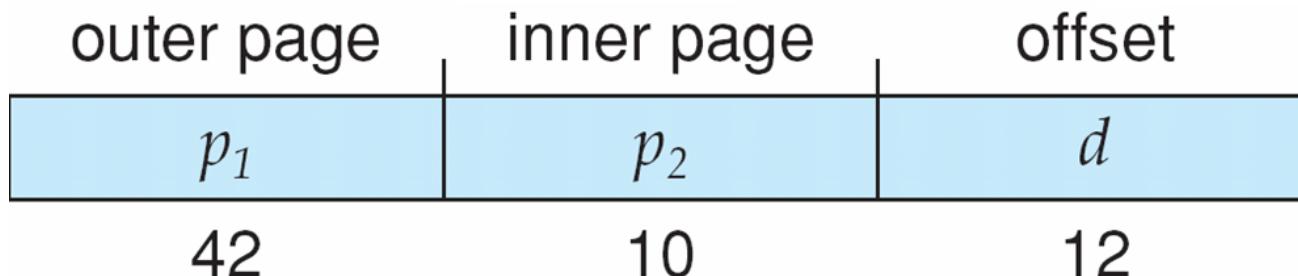
In total 6
tables
required.
 $1+5 = 6$.
Each table
has
size=4KB.

Then we
need 6×4
KB = 24 KB
for page
table
information

We assume entry size=4 bytes for both top level and second level tables. 79

Three-level Paging Scheme

64 bit addresses



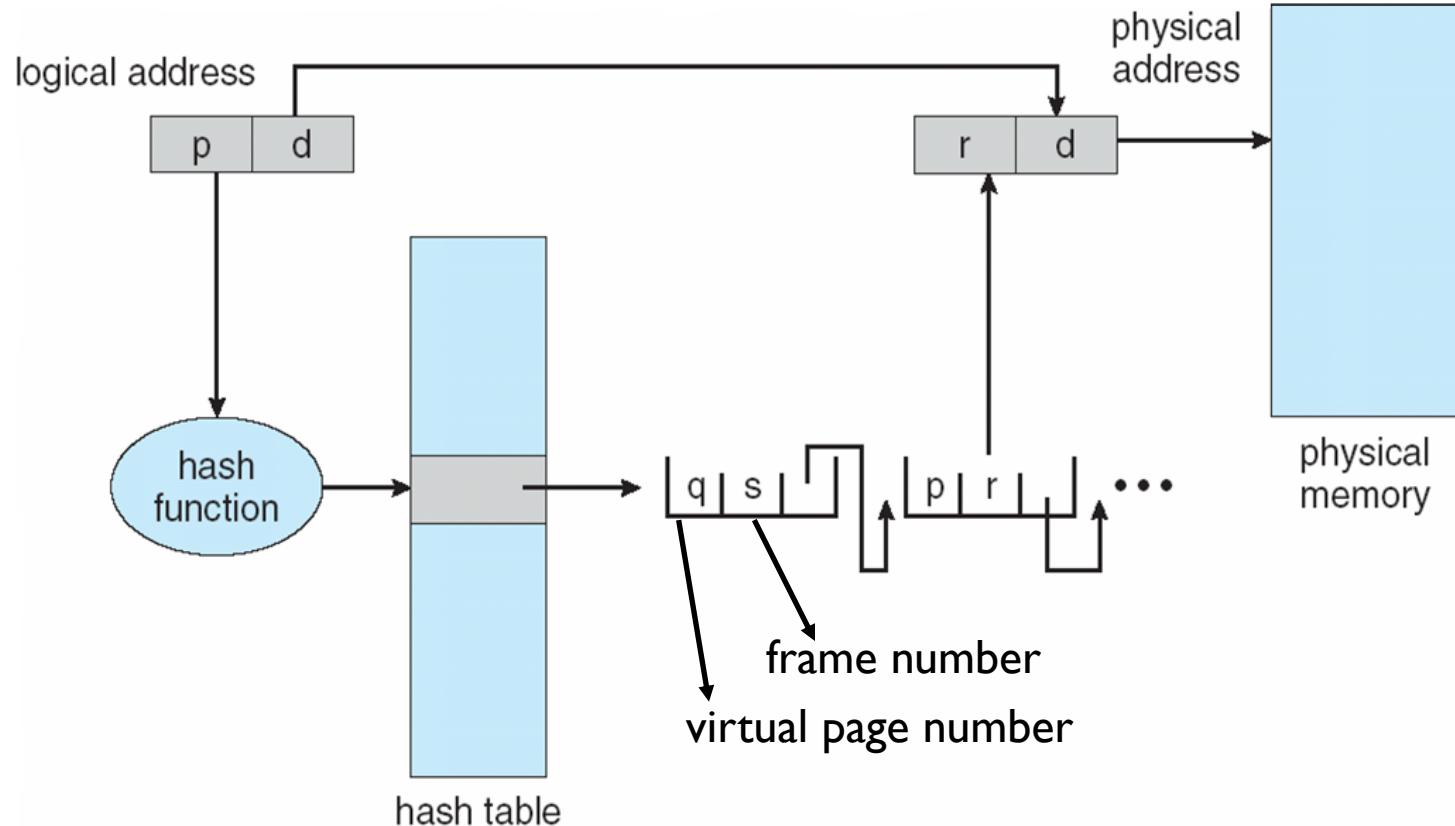
Top level table too big: 2^{32} entries

Even 3-level paging is not enough

Hashed Page Tables

- Common in address spaces > 32 bits
 - Large address spaces
- Page table is a hash table
- A virtual page number is hashed into a page table entry
- A page table entry contains a chain of elements hashing to the same location
 - each element = <a virtual page number, frame number>
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted.

Hashed Page Table



We can limit the hash table size per process, no matter how large the address space is (smaller the hash table, more collisions we have)

Inverted Page Table

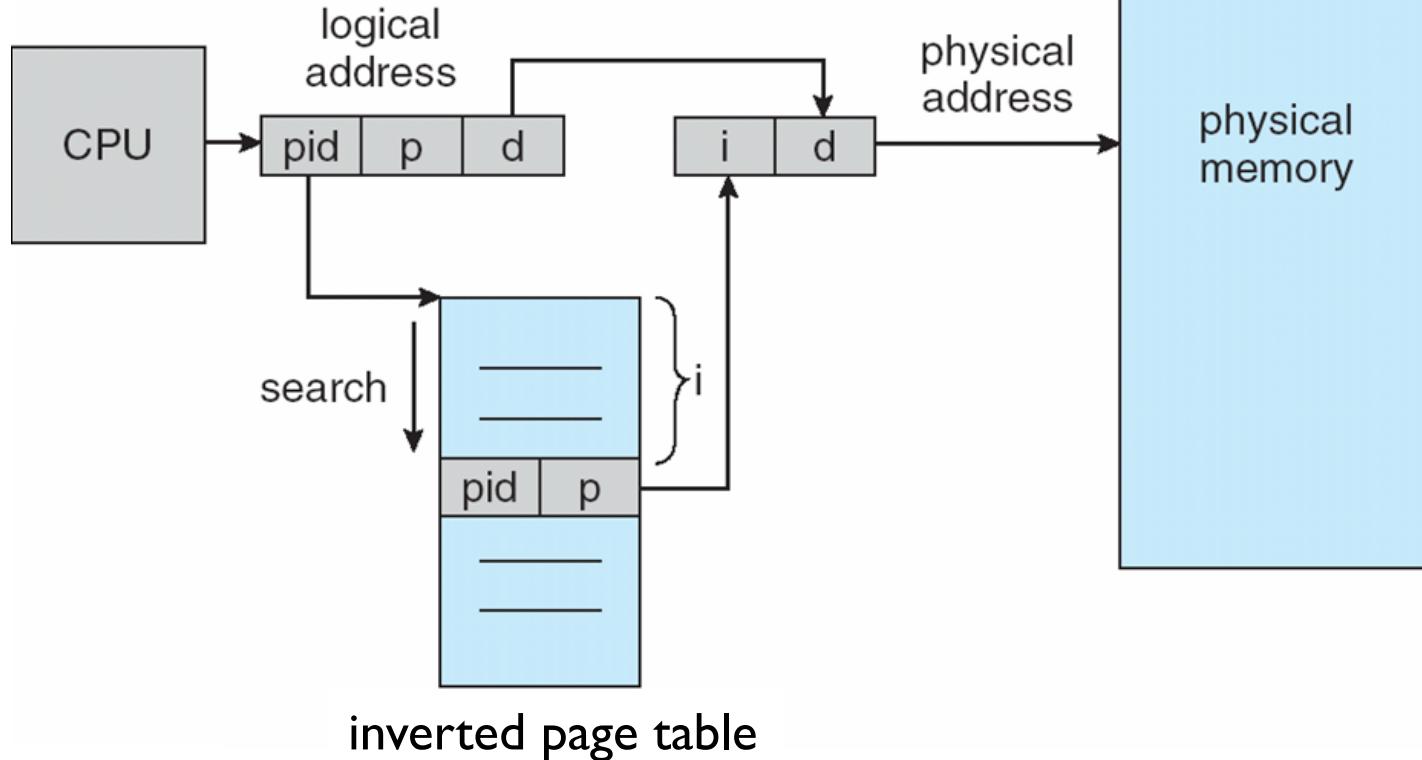
- One entry for each real **page frame of physical memory**
- Entry consists of the page number of the virtual page stored in that real memory location (frame), with information about the process that owns that page.
 - Entry content: **<pid, virtual page number>**
- Decreases memory needed to store page table
- All processes share the same inverted table
 - **One inverted table in the system.**
- **Searching** is problematic (search the p# in the table).
 - Given a p#, we need to find f# (therefore search needed)
- Search time can be reduced by using a **hash table** as well.

Inverted Page Table Architecture

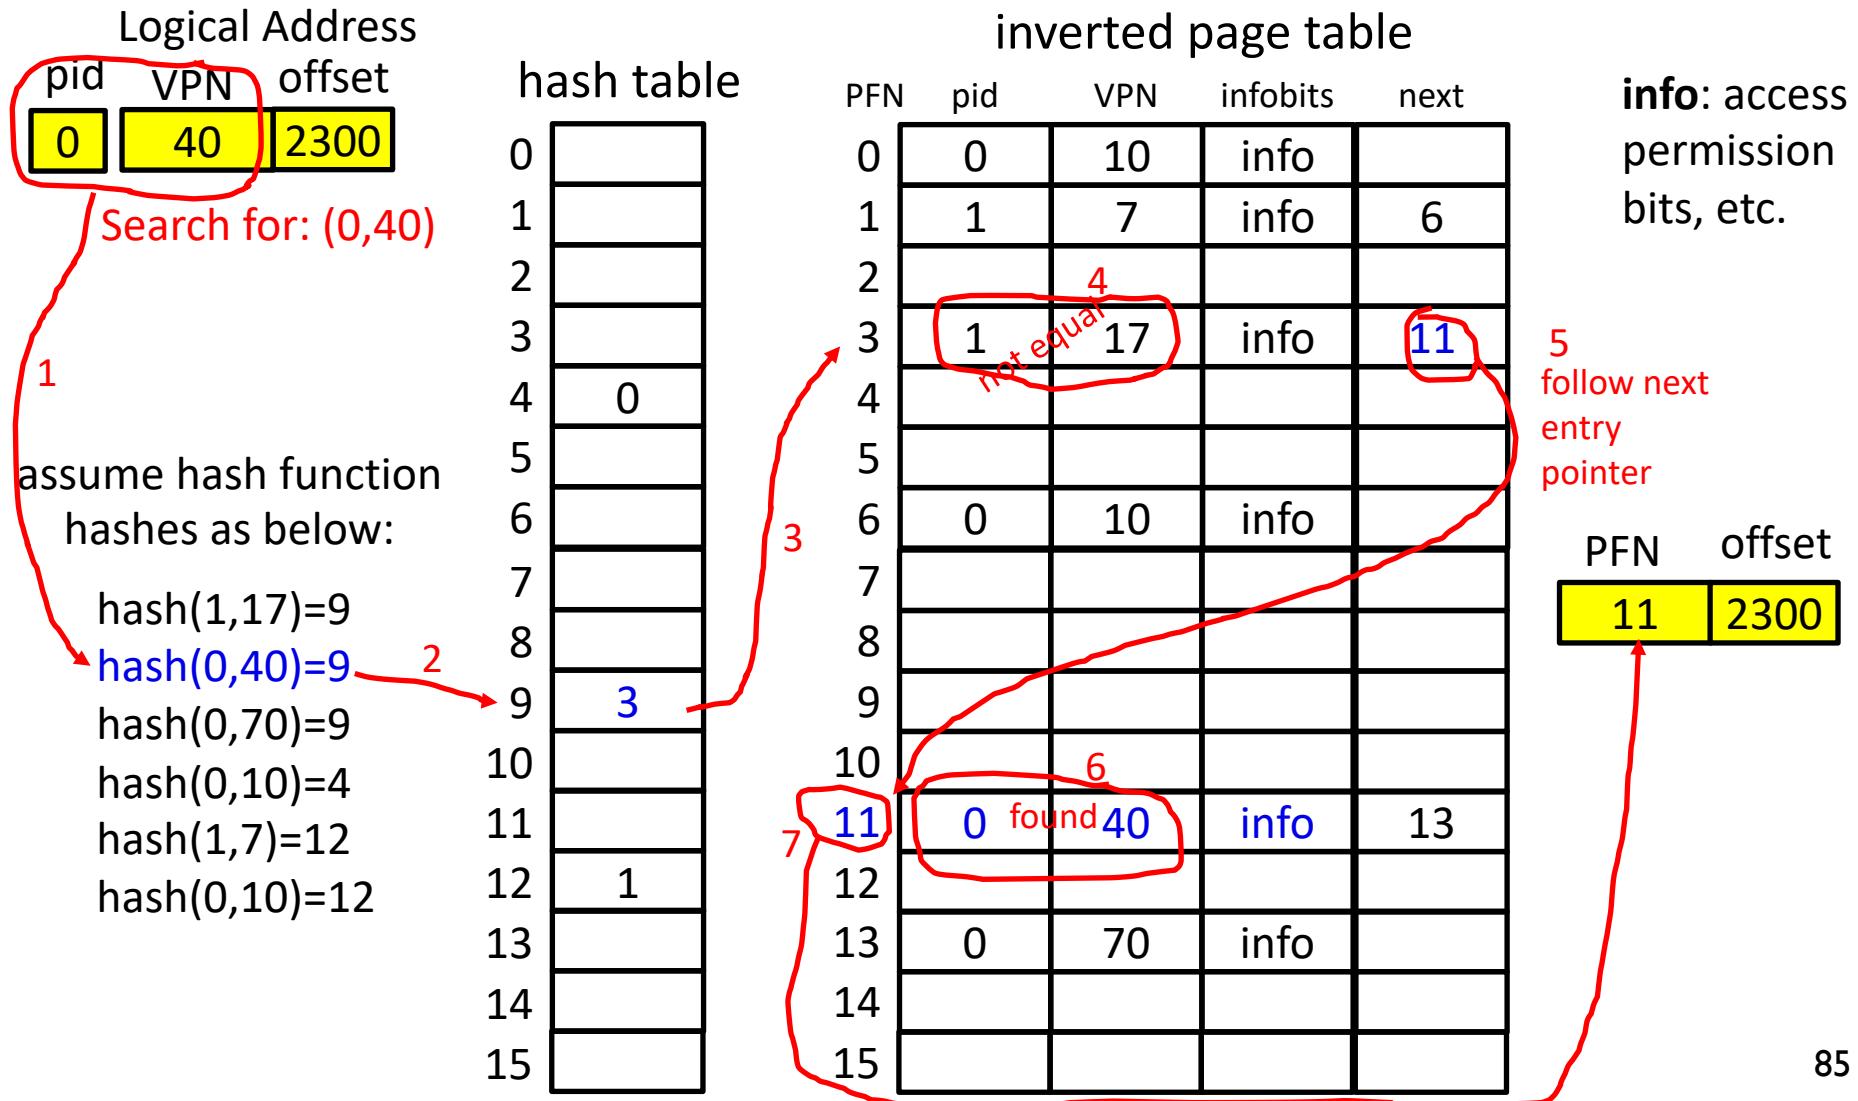
A logical address can be 64 bits long.

If offset is 12 bits (4 KB pages): VPN is 52 bits.

PFN can be n bits (if physical memory is 2^{n+12} bytes).



Inverted Table Example with a hash table in front

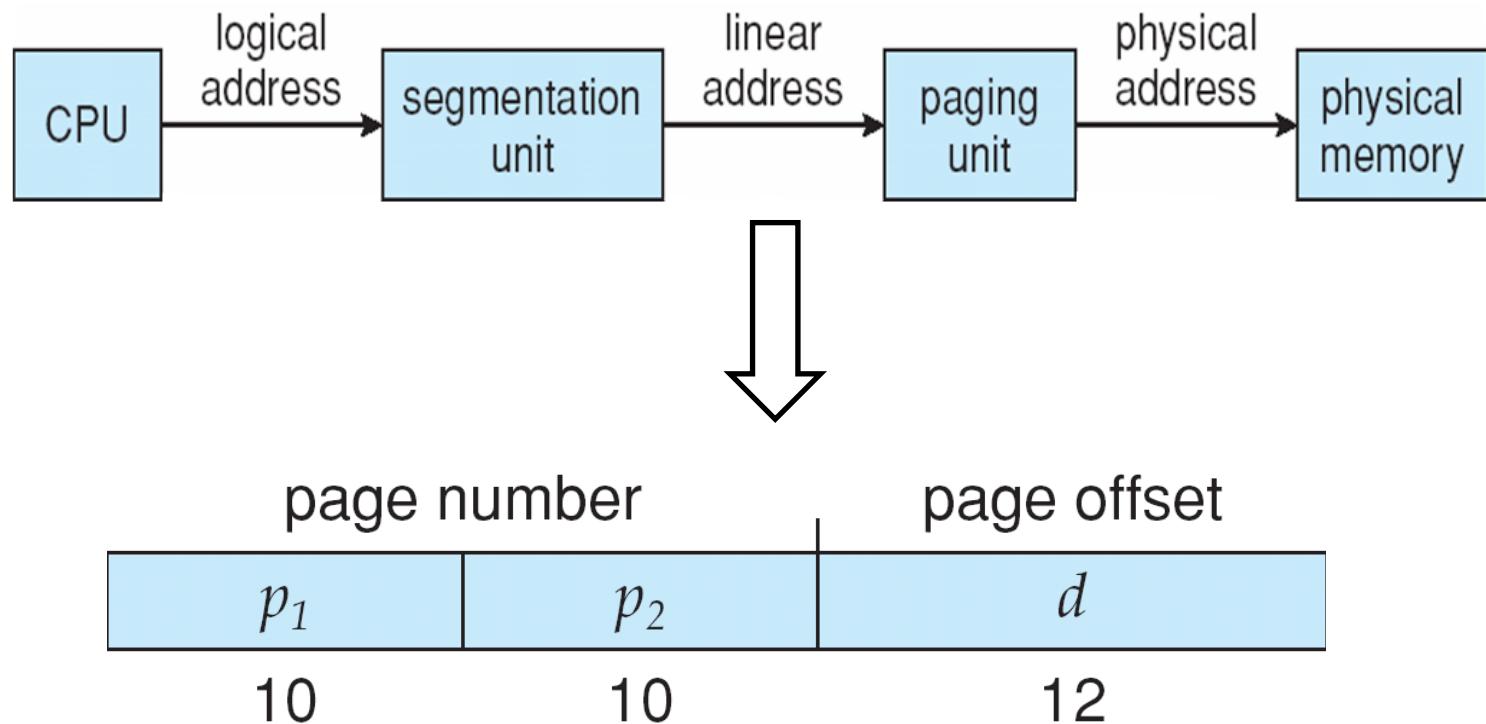


Example from a real system.
Intel x86 architecture
support for memory management

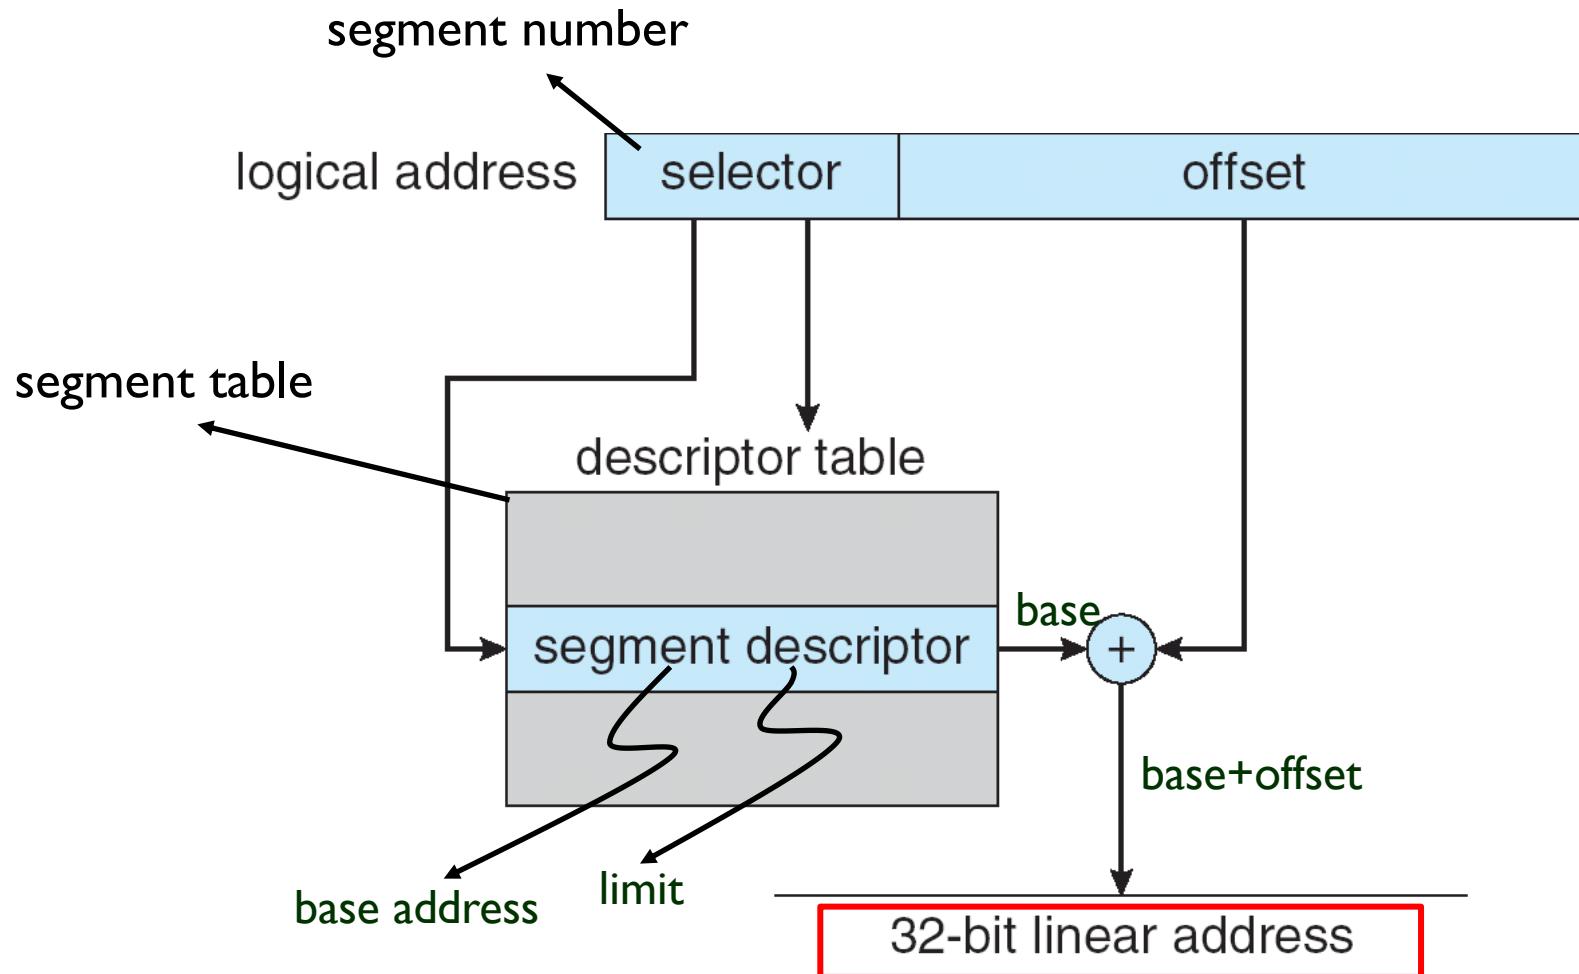
Example: The Intel 32 bit (IA32) architecture (x86 architecture)

- Supports both segmentation and segmentation with paging
- CPU generates **logical address** (<segment#, offset> pair)
(this is a **two-dimensional address**)
 - Given to segmentation unit
 - Which produces **linear addresses**
(this is a **one dimensional address**)
 - Linear address given to paging unit
 - Which generates **physical address** in main memory
 - Paging unit forms equivalent of MMU

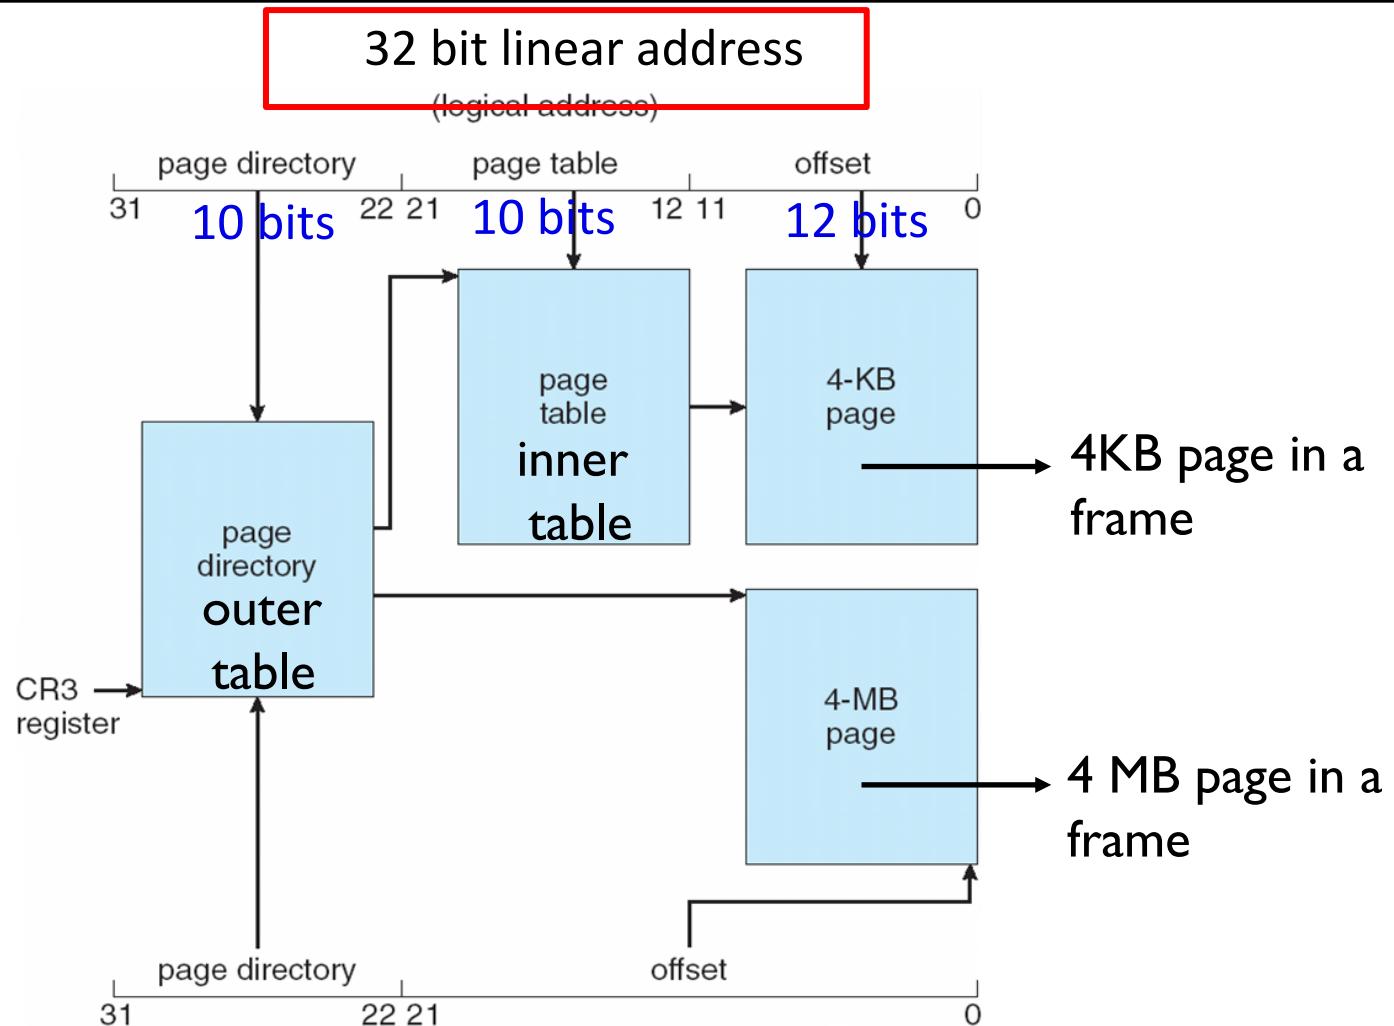
Logical to Physical Address Translation in Intel IA32



Intel IA32 Segmentation



IA32 Paging Architecture



IA32 Paging Architecture

- IA32 architecture allows a page size of either 4 KB or 4 MB.
- For 4 KB pages, two-level paging scheme is used in a 32 bit machine
 - Address division: <10, 10, 12> bits
- For 4 MB pages, we can skip the inner page tables (secondary page tables). A top level page table entry will point directly to a 4 MB page.

Linux on IA32: Segmentation

- Linux is designed to run on different hardware platforms:
 - x86,
 - Arm,
 - Motorola,
 - Sparc,
 - MIPS,
 - ...
- Therefore it does **not rely on segmentation** and makes minimal use of segmentation in Intel IA32 architecture.

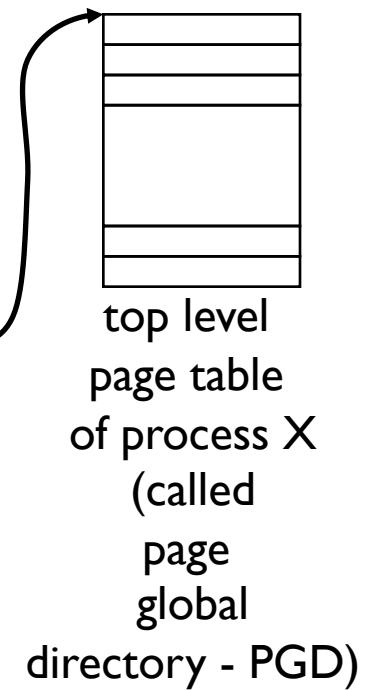
Linux Paging: some kernel structures (32 bit)

```
struct task_struct  
{  
    ...  
    ...  
    struct mm_struct *mm;  
}
```

the PCB object
of a process X

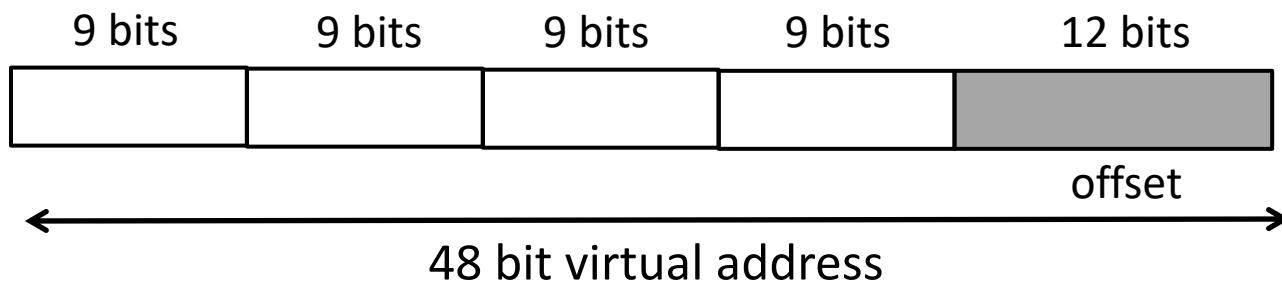
```
struct mm_struct  
{  
    ...  
    pgd_t *pgd;  
    ....  
}
```

mm object
of process X
(keeps memory
management
related
information)



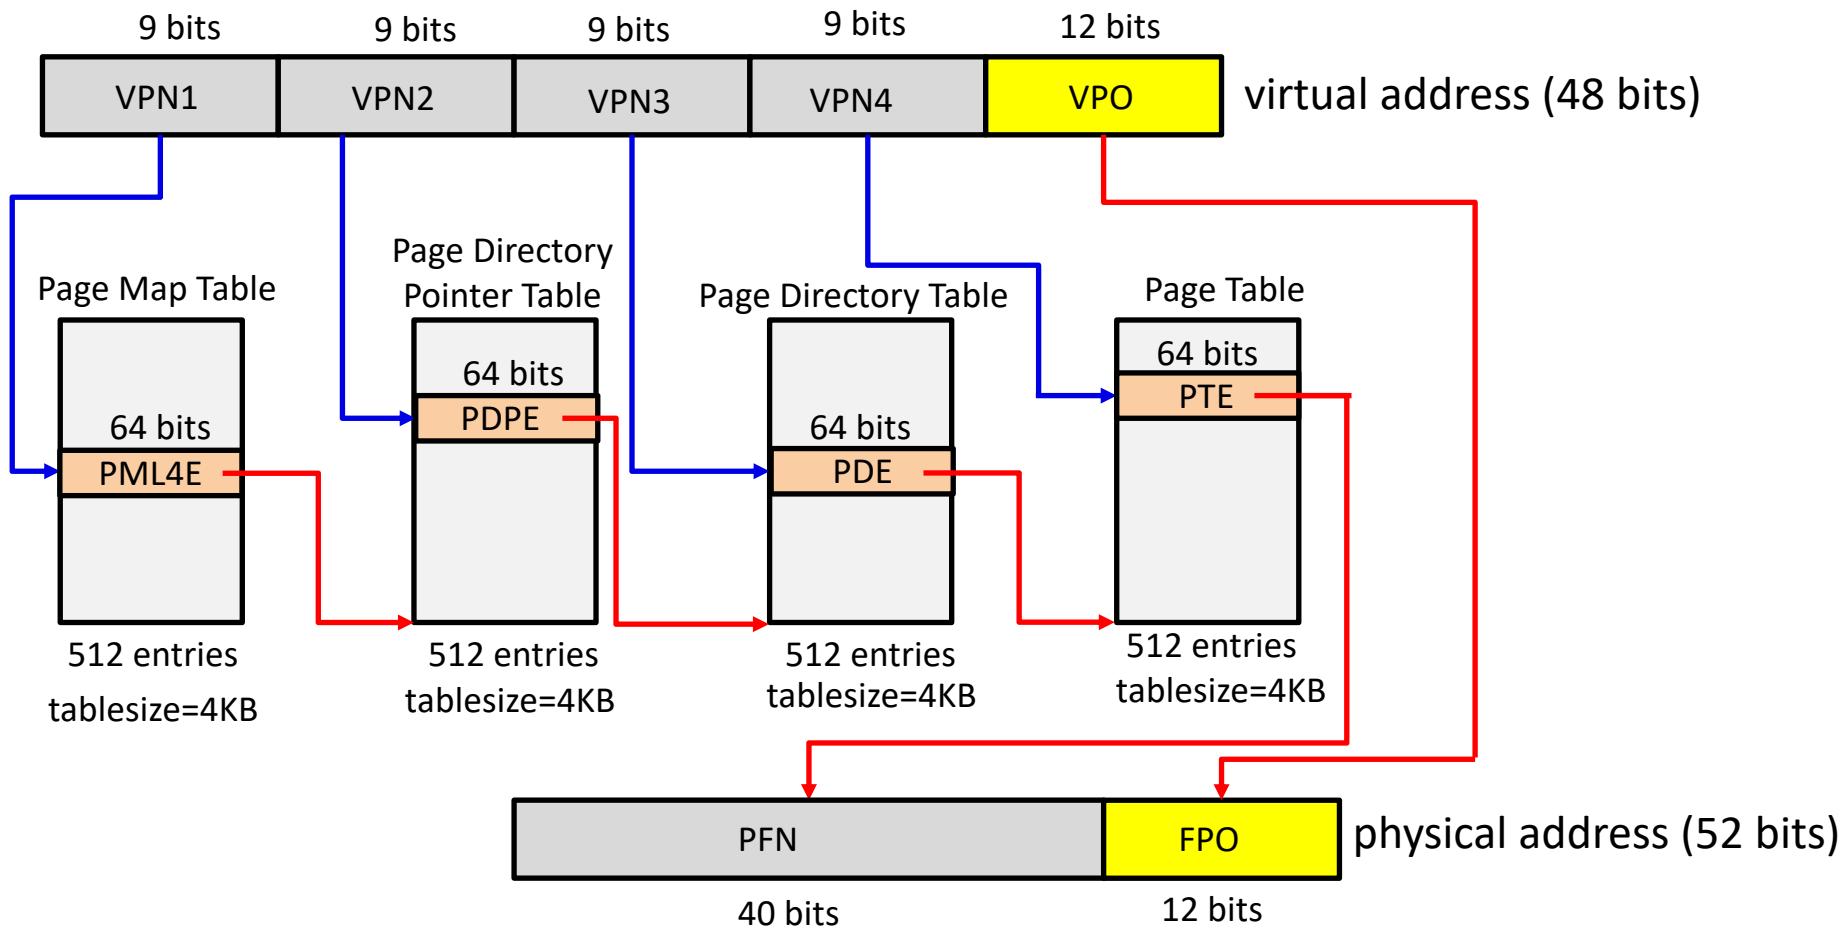
4 level paging in x86-64 architecture

- Intel 64 bit architecture (x86-64) uses **four level paging**. Linux supports x86-64.
- Virtual addresses are 48 bits long. Address division scheme is: [9,9,9,9,12]



x86-64 - translating virtual addresses to physical addresses

pagesize=4KB=framesize



References

- Operating System Concepts, Silberschatz et al, Wiley.
- Modern Operating Systems, Andrew S. Tanenbaum.
- Operating Systems: Three Easy Pieces, Arpaci-Dusseau et al.
- Intel Architecture Manuals.

Additional Study Material

Swapping

- A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk

Schematic View of Swapping

